

SECOND EDITION

EFFECTIVE C

AN INTRODUCTION TO
PROFESSIONAL C PROGRAMMING

ROBERT C. SEACORD



UPDATED TO
COVER C23



CONTENTS IN DETAIL

[PRAISE FOR EFFECTIVE C](#)

[TITLE PAGE](#)

[COPYRIGHT](#)

[DEDICATION](#)

[ABOUT THE AUTHOR, CONTRIBUTOR AND TECHNICAL REVIEWER](#)

[FOREWORD TO THE SECOND EDITION](#)

[FOREWORD TO THE FIRST EDITION](#)

[ACKNOWLEDGMENTS](#)

[INTRODUCTION](#)

[A Brief History of C](#)

[The C Standard](#)

[The CERT C Coding Standard](#)

[Common Weakness Enumeration](#)

[Who This Book Is For](#)

[What's in This Book](#)

1

[GETTING STARTED WITH C](#)

[Developing Your First C Program](#)

[Compiling and Running a Program](#)

[Function Return Values](#)

[Formatted Output](#)

[Editors and Integrated Development Environments](#)

[Compilers](#)

[GNU Compiler Collection](#)

[Clang](#)

[Microsoft Visual Studio](#)

[Portability](#)

[Implementation-Defined Behavior](#)
[Unspecified Behavior](#)
[Undefined Behavior](#)
[Locale-Specific Behavior and Common Extensions](#)

[Summary](#)

2

OBJECTS, FUNCTIONS, AND TYPES

[Entities](#)

[Declaring Variables](#)

[Swapping Values, First Attempt](#)

[Swapping Values, Second Attempt](#)

[Object Types](#)

[Boolean](#)

[Character](#)

[Arithmetic](#)

[void](#)

[Derived Types](#)

[Function](#)

[Pointer](#)

[Array](#)

[Structure](#)

[Union](#)

[Tags](#)

[Type Qualifiers](#)

[const](#)

[volatile](#)

[restrict](#)

[Scope](#)

[Storage Duration](#)

[Storage Class](#)

[static](#)

[extern](#)

[thread_local](#)

[constexpr](#)

[register](#)

[typedef](#)

[auto](#)

[typeof Operators](#)

[Alignment](#)

[Variably Modified Types](#)

[Attributes](#)

[Summary](#)

3

ARITHMETIC TYPES

[Integers](#)

[Padding, Width, and Precision](#)

- [Integer Ranges](#)
- [Integer Declarations](#)
- [Unsigned Integers](#)
- [Signed Integers](#)
- [Bit-Precise Integer Types](#)
- [Integer Constants](#)
- [Floating-Point Representation](#)
 - [Floating Types and Encodings](#)
 - [C Floating-Point Model](#)
 - [Floating-Point Arithmetic](#)
 - [Floating-Point Values](#)
 - [Floating Constants](#)
- [Arithmetic Conversion](#)
 - [Integer Conversion Rank](#)
 - [Integer Promotions](#)
 - [Usual Arithmetic Conversions](#)
 - [An Example of Implicit Conversion](#)
 - [Safe Conversions](#)
- [Summary](#)

4

EXPRESSIONS AND OPERATORS

- [Simple Assignment](#)
- [Evaluations](#)
- [Function Invocation](#)
- [Increment and Decrement Operators](#)
- [Operator Precedence and Associativity](#)
- [Order of Evaluation](#)
 - [Unsequenced and Indeterminately Sequenced Evaluations](#)
 - [Sequence Points](#)
- [sizeof Operator](#)
- [Arithmetic Operators](#)
 - [Unary + and –](#)
 - [Logical Negation](#)
 - [Additive](#)
 - [Multiplicative](#)
- [Bitwise Operators](#)
 - [Complement](#)
 - [Shift](#)
 - [Bitwise AND](#)
 - [Bitwise Exclusive OR](#)
 - [Bitwise Inclusive OR](#)
- [Logical Operators](#)
- [Cast Operators](#)
- [Conditional Operator](#)
- [alignof Operator](#)
- [Relational Operators](#)
- [Compound Assignment Operators](#)
- [Comma Operator](#)

[Pointer Arithmetic](#)
[Summary](#)

5

CONTROL FLOW

[Expression Statements](#)
[Compound Statements](#)
[Selection Statements](#)

[if](#)
[switch](#)

[Iteration Statements](#)

[while](#)
[do...while](#)
[for](#)

[Jump Statements](#)

[goto](#)
[continue](#)
[break](#)
[return](#)

[Summary](#)

6

DYNAMICALLY ALLOCATED MEMORY

[Storage Duration](#)

[The Heap and Memory Managers](#)
[When to Use Dynamically Allocated Memory](#)

[Memory Management](#)

[malloc](#)
[aligned_alloc](#)
[calloc](#)
[realloc](#)
[reallocarray](#)
[free](#)
[free_sized](#)
[free_aligned_sized](#)

[Memory States](#)

[Flexible Array Members](#)

[Other Dynamically Allocated Storage](#)

[alloca](#)
[Variable-Length Arrays](#)

[Debugging Allocated Storage Problems](#)

[dmalloc](#)
[Safety-Critical Systems](#)

[Summary](#)

7

CHARACTERS AND STRINGS

[Characters](#)

[ASCII](#)

[Unicode](#)

[Source and Execution Character Sets](#)

[Data Types](#)

[Character Constants](#)

[Escape Sequences](#)

[Linux](#)

[Windows](#)

[Character Conversion](#)

[Strings](#)

[String Literals](#)

[String-Handling Functions](#)

[<string.h> and <wchar.h>](#)

[Annex K Bounds-Checking Interfaces](#)

[POSIX](#)

[Microsoft](#)

[Summary](#)

8

INPUT/OUTPUT

[Standard I/O Streams](#)

[Error and End-of-File Indicators](#)

[Stream Buffering](#)

[Predefined Streams](#)

[Stream Orientation](#)

[Text and Binary Streams](#)

[Opening and Creating Files](#)

[fopen](#)

[open](#)

[Closing Files](#)

[fclose](#)

[close](#)

[Reading and Writing Characters and Lines](#)

[Stream Flushing](#)

[Setting the Position in a File](#)

[Removing and Renaming Files](#)

[Using Temporary Files](#)

[Reading Formatted Text Streams](#)

[Reading from and Writing to Binary Streams](#)

[Endian](#)

[Summary](#)

9

PREPROCESSOR

[The Compilation Process](#)

[File Inclusion](#)

[Conditional Inclusion](#)

- [Generating Diagnostics](#)
- [Using Header Guards](#)
- [Macro Definitions](#)
 - [Macro Replacement](#)
 - [Type-Generic Macros](#)
 - [Embedded Binary Resources](#)
 - [Predefined Macros](#)
- [Summary](#)

10

PROGRAM STRUCTURE

- [Principles of Componentization](#)
 - [Coupling and Cohesion](#)
 - [Code Reuse](#)
 - [Data Abstractions](#)
 - [Opaque Types](#)
- [Executables](#)
- [Linkage](#)
- [Structuring a Simple Program](#)
- [Building the Code](#)
- [Summary](#)

11

DEBUGGING, TESTING, AND ANALYSIS

- [Assertions](#)
 - [Static Assertions](#)
 - [Runtime Assertions](#)
- [Compiler Settings and Flags](#)
 - [GCC and Clang Flags](#)
 - [Visual C++ Options](#)
- [Debugging](#)
- [Unit Testing](#)
- [Static Analysis](#)
- [Dynamic Analysis](#)
- [AddressSanitizer](#)
 - [Running the Tests](#)
 - [Instrumenting the Code](#)
 - [Running the Instrumented Tests](#)
- [Summary](#)
- [Future Directions](#)

APPENDIX: THE FIFTH EDITION OF THE C STANDARD (C23)

- [Attributes](#)
- [Keywords](#)
- [Integer Constant Expressions](#)
- [Enumeration Types](#)
- [Type Inference](#)
- [typeof Operators](#)

[K&R C Functions](#)
[Preprocessor](#)
[Integer Types and Representations](#)
[unreachable Function-Like Macro](#)
[Bit and Byte Utilities](#)
[IEEE Floating-Point Support](#)

REFERENCES

INDEX

PRAISE FOR *EFFECTIVE C*

“*Effective C* will teach you C programming for the modern era. . . . This book’s emphasis on the security aspects of C programming is unmatched. My personal recommendation is that, after reading it, you use all of the available tools it presents to avoid undefined behavior in the C programs you write.”

—PASCAL CUOQ, CHIEF SCIENTIST, TRUSTINSOFT

“An excellent introduction to modern C.”

—FRANCIS GLASSBOROW, ACCU

“A good introduction to modern C, including chapters on dynamic memory allocation, on program structure, and on debugging, testing, and analysis.”

—STACK OVERFLOW, THE DEFINITIVE C BOOK LIST

“A worthwhile addition to a C programmer’s bookshelf.”

—IAN BRUNTLETT, ACCU

“This is why you should program in C. Because other languages don’t open portals to hell.”

—MICHAŁ ZALEWSKI, FORMER CISO, SNAP INC.

EFFECTIVE C

2nd Edition

An Introduction to Professional C Programming

by Robert C. Seacord



**no starch
press®**

San Francisco

EFFECTIVE C, 2ND EDITION. Copyright © 2025 by Robert C. Seacord.

All rights reserved. No part of this work may be reproduced or transmitted in any form or by any means, electronic or mechanical, including photocopying, recording, or by any information storage or retrieval system, without the prior written permission of the copyright owner and the publisher.

First printing

28 27 26 25 24 1 2 3 4 5

ISBN-13: 978-1-7185-0412-7 (print)

ISBN-13: 978-1-7185-0413-4 (ebook)



Published by No Starch Press®, Inc.

245 8th Street, San Francisco, CA 94103

phone: +1.415.863.9900

www.nostarch.com; info@nostarch.com

Publisher: William Pollock

Managing Editor: Jill Franklin

Production Manager: Sabrina Plomitallo-González

Production Editor: Jennifer Kepler

Developmental Editor: Jill Franklin

Cover Illustrator: Gina Redman

Interior Design: Octopod Studios

Technical Reviewers: Vincent Mailhol and Martin Sebor

Copyeditor: Lisa McCoy

Proofreader: Dan Foster

Indexer: Michael Goldstein

The Library of Congress has catalogued the first edition as follows:

Names: Seacord, Robert C., author.

Title: Effective C : an introduction to professional C programming / Robert C. Seacord.

Description: San Francisco : No Starch Press, Inc., 2020. | Includes bibliographical references and index.

Identifiers: LCCN 2020017146 (print) | LCCN 2020017147 (ebook) | ISBN 9781718501041 (paperback) | ISBN 1718501048 (paperback) | ISBN 9781718501058 (ebook)

Subjects: LCSH: C (Computer program language)

Classification: LCC QA76.73.C15 S417 2020 (print) | LCC QA76.73.C15 (ebook) | DDC 005.13/3--dc23

LC record available at <https://lccn.loc.gov/2020017146>

LC ebook record available at <https://lccn.loc.gov/2020017147>

For customer service inquiries, please contact info@nostarch.com. For information on distribution, bulk sales, corporate sales, or translations: sales@nostarch.com. For permission to translate this work: rights@nostarch.com. To report counterfeit copies or piracy: counterfeit@nostarch.com.

No Starch Press and the No Starch Press iron logo are registered trademarks of No Starch Press, Inc. Other product and company names mentioned herein may be the trademarks of their respective owners. Rather than use a trademark symbol with every occurrence of a trademarked name, we are using the names only in an editorial fashion and to the benefit of the trademark owner, with no intention of infringement of the trademark.

The information in this book is distributed on an “As Is” basis, without warranty. While every precaution

has been taken in the preparation of this work, neither the author nor No Starch Press, Inc. shall have any liability to any person or entity with respect to any loss or damage caused or alleged to be caused directly or indirectly by the information contained in it.

To my granddaughters, Olivia and Isabella, and to all the young women who will
grow up to be scientists and engineers

About the Author

Robert C. Seacord (rscs@robertseacord.com) is the standardization lead at Woven by Toyota, where he works on the software craft. Robert was previously a technical director at NCC Group, the manager of the Secure Coding Initiative at Carnegie Mellon University's Software Engineering Institute, and an adjunct professor in the School of Computer Science and the Information Networking Institute at Carnegie Mellon. Robert is the convener of the ISO/IEC JTC1/SC22/WG14, the international standardization working group for the C programming language. He is the author of other books, including *The CERT® C Coding Standard*, 2nd edition (Addison-Wesley, 2014); *Secure Coding in C and C++*, 2nd edition (Addison-Wesley, 2013); and *Java Coding Guidelines: 75 Recommendations for Reliable and Secure Programs* (Addison-Wesley, 2014). He has also published more than 50 papers on software security, component-based software engineering, web-based system design, legacy-system modernization, component repositories and search engines, and user interface design and development.

About the Contributor

Aaron Ballman (aaron@aaronballman.com) is the lead maintainer for Clang, a popular open source compiler for C, C++, and other languages. Aaron is an expert for the JTC1/SC22/WG14 C programming language and JTC1/SC22/WG21 C++ programming language standards committees. His primary professional focus has been in helping programmers recognize mistakes in their code through better language design, diagnostics, and tooling. When not thinking about programming, he enjoys spending quiet moments in the woods of rural Maine with his family.

About the Technical Reviewer to the First Edition

Martin Sebor is a principal software engineer with the GNU Toolchain Team at Red Hat. His focus is the GCC compiler and the areas of detecting, diagnosing, and preventing security-related issues in C and C++ programs, as well as implementing optimizations of string-based algorithms. Prior to joining Red Hat in 2015, he worked as a compiler toolchain engineer at Cisco. Martin has been a

member of the C++ standardization committee since 1999 and a member of the C language committee since 2010. He lives with his wife near the small town of Lyons, Colorado.

About the Technical Reviewer to the Second Edition

Vincent Mailhol is a senior product security engineer at Woven by Toyota, where he architects and develops C software for the trusted execution environment of automotive electronic control units. He is a member of the company's cryptographic board and, together with Robert, established the company's secure and safety coding guidelines. Prior to joining Woven in 2020, he was the lead of the pentesting team at ESCRYPT Japan, a subsidiary of the Bosch Group, specializing in automotive security. In his free time, Vincent is an active contributor and maintainer of the Linux kernel's CAN subsystem, also known as SocketCAN.

FOREWORD TO THE SECOND EDITION

When I started in cybersecurity over 27 years ago, I learned my trade primarily by finding and exploiting unsafe memory handling in C programs—a class of vulnerability that, even at the time, was over 20 years old. In my career at BlackBerry, as I waded through torrents of code review, I saw firsthand how dangerous C could be to the improperly initiated. Now, as a chief technology officer for the UK’s National Cyber Security Centre, I see the consequences of poorly written C code on our connected society every day at a national level.

Today we still face numerous challenges in writing secure and professional C. The many innovations in compiler- and operating system-level mitigations can be and are regularly undermined. And even while we see advanced innovations in other modern languages and hardware, there is still a growing demand for C, particularly in Internet of Things (IoT) or other highly resource-constrained environments, while it also sustains what we have. Professional C coupled with hardware-level architectures such as CHERI is how we secure the environments that will never migrate to other languages.

Robert is the authority on how to program professionally and securely in C. For over a decade, I have recommended his material to customers and internal teams alike. There is no better person to teach how to code C in a professional and, among other things, secure manner.

Writing professional C today means writing code that is performant, safe, and secure. By doing so, you will be able to contribute to our connected society without increasing its technical debt.

This book will help those with little or no C experience quickly develop the knowledge and skills to become professional C programmers and will provide a strong foundation for developing systems that are performant, safe, and secure.

Ollie Whitehouse
CTO, National Cyber Security Centre, United Kingdom

FOREWORD TO THE FIRST EDITION

The first time I came across Robert Seacord's name was in 2008. Robert was already well known in the C programming galaxy for his work on *The CERT® C Coding Standard* and Annex K of the C standard. But in 2008, it had been only a few years since—young and foolish—I had embarked on the Frama-C project to guarantee the absence of undefined behavior in C programs. At some point, a CERT Vulnerability Note about how C compilers (and GCC in particular) removed certain pointer-arithmetic overflow checks piqued my interest. The compilers had reason to eliminate the checks; naively written, they invoked undefined behavior when the overflow was present.

The C compilers were also allowed to tell the programmer nothing about what they had done wrong, even at the maximum warning level. Undefined behavior in C can be harsh. I had set out to solve this exact problem. Robert was one of the authors of that note.

Effective C will teach you C programming for the modern era. It will help you establish good habits to keep you from using undefined behavior, whether voluntarily or through negligence. Let the reader be warned: in large C programs, avoiding ordinary programming errors alone may not suffice for dodging undefined behavior caused by arbitrary inputs!

This book's emphasis on the security aspects of C programming is unmatched. My personal recommendation is that, after reading it, you use all of the available tools it presents to avoid undefined behavior in the C programs you write.

Pascal Cuoq
Chief scientist, TrustInSoft

ACKNOWLEDGMENTS

I would like to acknowledge the contributions of all the folks who made this possible. I'll start with Bill Pollock at No Starch Press, who doggedly pursued me to write a C book.

I would like to thank Ollie Whitehouse and Pascal Cuoq for the excellent forewords.

Aaron Ballman has been a valuable partner for the duration of this effort. In addition to contributing two chapters, he reviewed everything (often multiple times) and helped me solve problems from the profound to the mundane.

Douglas Gwyn, emeritus at the US Army Research Laboratory and an emeritus member of the C standards committee, helped review all the chapters. When my writing was not up to his standards, he guided me in the right direction.

Martin Sebor served as the official technical reviewer for the first edition and Vincent Mailhol for the second edition of this book.

In addition to Aaron, Doug, and Martin, several other distinguished members of the C and C++ standards committee reviewed chapters, including Jim Thomas, Thomas Köppe, Niall Douglas, Tom Honermann, and JeanHeyd Meneide. Technical reviewers among my colleagues at Woven by Toyota include J.F. Bastien, Carlos Ramirez, and Jørgen Kvalsvik. Technical reviewers among my former colleagues at NCC Group include Nick Dunn, Jonathan Lindsay, Tomasz Kramkowski, Alex Donisthorpe, Joshua Dow, Catalin Visinescu, Aaron Adams, and Simon Harraghy. Technical reviewers outside these organizations include Geoff Clare, David LeBlanc, David Goldblatt, Nicholas Winter, John McFarlane, and Scott Aloisio.

I would also like to thank the following professionals at No Starch who ensured a quality product: Jill Franklin, Jennifer Kepler, Elizabeth Chadwick, Frances Saux, Zach Lebowsky, Annie Choi, Barbara Yien, Katrina Taylor, Natalie Gleason, Derek Yee, Laurel Chun, Gina Redman, Sharon Wilkey, Emelie Battaglia, Dapinder Dosanjh, Lisa McCoy, Morgan Vega Gomez, and Sabrina

Plomitallo-González. Finally, I would like to thank Michael Goldstein for preparing the index.

INTRODUCTION



C was developed as a system programming language in the 1970s, and even after all this time, it remains incredibly popular. System languages are designed for performance and ease of access to the underlying hardware while providing high-level programming features. While other languages may offer newer language features, their compilers and libraries are typically written in C.

Carl Sagan once said, “If you wish to make an apple pie from scratch, you must first invent the universe.” The inventors of C did not invent the universe; they designed C to work with a variety of computing hardware and architectures that, in turn, were constrained by physics and mathematics. C is layered directly on top of computing hardware, making it more sensitive to evolving hardware features, such as vectorized instructions, than higher-level languages that typically rely on C for their efficiency.

According to the TIOBE index (<https://www.tiobe.com/tiobe-index/>)—whose rankings are based on the number of skilled engineers, courses, and third-party vendors for each language—C has been either the most popular programming language or second most popular since 2001. The popularity of the C programming language can most likely be attributed to several tenets of the language referred to as the *spirit of C*:

- Trust the programmer. The C language assumes you know what you’re doing and lets you. This isn’t always a good thing (for example, if you don’t know what you’re doing).
- Don’t prevent the programmer from doing what needs to be done. Because

C is a system programming language, it needs to handle a variety of low-level tasks.

- Keep the language small and simple. The language is designed to be close to the hardware and to have a small footprint.
- Provide only one way to do an operation. Also known as *conservation of mechanism*, the C language tries to limit the introduction of duplicate mechanisms.
- Make it fast, even if it isn't guaranteed to be portable. Allowing you to write optimally efficient code is the top priority. The responsibility of ensuring that code is portable, safe, and secure is delegated to you, the programmer.

C is used as a target language for compilers to build operating systems, to teach fundamentals of computing, and for embedded and general-purpose programming.

There is a large amount of legacy code written in C. The C standards committee is extremely careful not to break existing code, providing a smooth pass for modernizing this code to take advantage of modern language features.

C is often used in embedded systems because it is a small and efficient language. Embedded systems are small computers that are embedded in other devices, such as cars, appliances, and medical devices.

Your favorite programming language and library are written in C (or were at one time). There are many libraries available for C. This makes it easy to find libraries that can be used to perform common tasks.

Overall, C is a powerful and versatile language that is still widely used today. It is a good choice for programmers who need a fast, efficient, and portable language.

A Brief History of C

The C programming language was developed in the early 1970s at Bell Labs as a system implementation language for the nascent Unix operating system and remains incredibly popular today (Ritchie 1993). System languages are designed for performance and ease of access to the underlying hardware while providing high-level programming features. While other languages may offer newer language features, their compilers and libraries are typically written in C. It serves as a lingua franca for translating between various systems and languages.

C was first described in 1978 by Kernighan and Ritchie in the book *The C Programming Language* (Kernighan and Ritchie 1988). It is now defined by

revisions of the ISO/IEC 9899 standard (ISO/IEC 2024) and other technical specifications. The C standards committee is the steward of the C programming language, working with the broader community to maintain and evolve the C language. In 1983, the American National Standards Institute (ANSI) formed the X3J11 committee to establish a standard C specification, and in 1989, the C standard was ratified as ANSI X3.159-1989, “Programming Language C.” This 1989 version of the language is referred to as *ANSI C* or *C89*.

In 1990, the ANSI C standard was adopted (unchanged) by a joint technical committee of the International Organization for Standardization (ISO) and the International Electrotechnical Commission (IEC) and published as the first edition of the C standard, C90 (ISO/IEC 9899:1990). The second edition of the C standard, C99, was published in 1999 (ISO/IEC 9899:1999), and a third edition, C11, in 2011 (ISO/IEC 9899:2011). The fourth version, published in 2018 as C17 (ISO/IEC 9899:2018), repairs defects in C11. The latest version of the C standard (as of this writing) is the fifth version, published in 2024 as C23 (ISO/IEC 9899:2024). As of September 2023, I am the convenor of ISO/IEC JTC1/SC22/WG14, the international standardization working group for the programming language C.

In the 20 years the TIOBE Programming Community index has tracked programming language popularity, C has remained in first or second place (TIOBE Index 2022).

The C Standard

The C standard (ISO/IEC 9899:2024) defines the language and is the final authority on language behavior. While the standard can be obscure to impenetrable, you need to understand it if you intend to write code that’s portable, safe, and secure. The C standard provides a substantial degree of latitude to implementations to allow them to be optimally efficient on various hardware platforms. *Implementations* is the term the C standard uses to refer to compilers and is defined as follows:

A particular set of software, running in a particular translation environment under particular control options, that performs translation of programs for, and supports execution of functions in, a particular execution environment.

This definition indicates that each compiler with a particular set of command line flags, along with the C standard library, is considered a separate implementation, and different implementations can have significantly different *implementation-defined behavior*. This is noticeable in GNU Compiler

Collection (GCC), which uses the `-std=` flag to determine the language standard. Possible values for this option include `c89`, `c90`, `c99`, `c11`, `c17`, and `c23`. The default depends on the version of the compiler. If no C language dialect options are given, the default for GCC 13 is `-std=gnu17`, which provides extensions to the C language. For portability, specify the standard you're using. For access to new language features, specify a recent standard. C23 features have been available since GCC 11. To enable C23 support, add the compiler option `-std=c23` (or possibly `-std=c2x`). All the examples in this book are written for C23.

Because implementations have such a range of behaviors, and because some of these behaviors are undefined, you can't understand the C language by just writing simple test programs to examine the behavior. (If you want to try this, the Compiler Explorer is an excellent tool; see <https://godbolt.org>.) The behavior of the code can vary when compiled by a different implementation on different platforms or even the same implementation using a different set of flags or a different C standard library implementation. Code behavior can even vary between *versions* of a compiler. The C standard specifies which behaviors are guaranteed for all implementations and where you need to plan for variability. This is mostly a concern when developing portable code but can also affect the security and safety of your code.

The CERT C Coding Standard

The CERT® C Coding Standard: 98 Rules for Developing Safe, Reliable, and Secure Systems, 2nd edition (Addison-Wesley Professional, 2014), is a reference book I wrote while managing the secure coding team at the Software Engineering Institute at Carnegie Mellon University. The book contains examples of common C programming mistakes and how to correct them. Throughout this book, we reference some of those rules as a source for detailed information on specific C language programming topics.

Common Weakness Enumeration

MITRE's Common Weakness Enumeration (CWE) is a list of common hardware and software weaknesses that can be used to identify weaknesses in source code and operational systems. The CWE list is maintained by a community project with the goals of understanding flaws in software and hardware and creating automated tools that can be used to identify, fix, and prevent those flaws.

Occasionally, we will reference specific CWEs in this book when discussing classes of defects that can lead to security vulnerabilities. For more information on CWE, see <https://cwe.mitre.org>.

Who This Book Is For

This book is an introduction to the C language. It is written to be as accessible as possible to anyone who wants to learn C programming, without dumbing it down. In other words, we didn't overly simplify C programming in the way many other introductory books and courses might. These overly simplified references will teach you how to compile and run code, but the code might still be wrong. Developers who learn how to program C from such sources will typically develop substandard, flawed, insecure code that will eventually need to be rewritten (often sooner than later). Hopefully, these developers will eventually benefit from senior developers in their organizations who will help them unlearn these harmful misconceptions about programming in C and help them start developing professional-quality C code. On the other hand, this book will quickly teach you how to develop correct, portable, professional-quality code; build a foundation for developing security-critical and safety-critical systems; and perhaps teach you some things that even the senior developers at your organization don't know.

Effective C: An Introduction to Professional C Programming, 2nd edition, is a concise introduction to essential C language programming that will soon have you writing programs, solving problems, and building working systems. The code examples are idiomatic and straightforward. You'll also learn about good software engineering practices for developing correct, secure C code.

In this book, you'll learn about essential programming concepts in C and practice writing high-quality code with exercises for each topic. Code listings from this book and additional materials can be found on GitHub at <https://github.com/rcseacord/effective-c>. Go to this book's page at <https://nostarch.com/effective-c-2nd-edition> or to <http://www.robertseacord.com> to check for updates and additional material, or contact me if you have additional questions or are interested in training.

What's in This Book

This book starts with an introductory chapter that covers just enough material to get you programming right from the start. After that, we circle back and examine the basic building blocks of the language. The book culminates with two chapters that will show you how to compose real-world systems from these basic building blocks and how to debug, test, and analyze the code you've written. The

chapters are as follows:

Chapter 1: Getting Started with C You'll write a simple C program to become familiar with using the main function. You'll also look at a few options for editors and compilers.

Chapter 2: Objects, Functions, and Types This chapter explores basics like declaring variables and functions. You'll also investigate the principles of using basic types.

Chapter 3: Arithmetic Types You'll learn about the integer and floating-point arithmetic data types.

Chapter 4: Expressions and Operators You'll learn about operators and how to write simple expressions to perform operations on various object types.

Chapter 5: Control Flow You'll learn how to control the order in which individual statements are evaluated. We'll introduce expression and compound statements that define the work to be performed. We'll then cover the control statements that determine which code blocks are executed and in what order: selection, iteration, and jump statements.

Chapter 6: Dynamically Allocated Memory You'll learn about dynamically allocated memory, which is allocated from the heap at runtime. Dynamically allocated memory is useful when the exact storage requirements for a program are unknown before runtime.

Chapter 7: Characters and Strings This chapter covers the various character sets, including ASCII and Unicode, that can be used to compose strings. You'll learn how strings are represented and manipulated using the legacy functions from the C standard library, the bounds-checking interfaces, and POSIX and Windows application programming interfaces (APIs).

Chapter 8: Input/Output This chapter will teach you how to perform input/output (I/O) operations to read data from, or write data to, terminals and filesystems. I/O involves all the ways information enters or exits a program. We'll cover techniques that make use of C standard streams and POSIX file descriptors.

Chapter 9: Preprocessor You'll learn how to use the preprocessor to include files, define object- and function-like macros, and conditionally include code based on implementation-specific features.

Chapter 10: Program Structure You'll learn how to structure your program into multiple translation units consisting of both source and include files. You'll also learn how to link multiple object files together to create libraries and executable files.

Chapter 11: Debugging, Testing, and Analysis This chapter describes tools and techniques for producing error-free programs, including compile-time and runtime assertions, debugging, testing, static analysis, and dynamic analysis. The chapter also discusses which compiler flags are recommended for use in different phases of the software development process.

Appendix: The Fifth Edition of the C Standard (C23) This appendix enumerates some of the additions and changes in C23. It's a convenient way to learn what's new in C and to identify changes from the previous C standard (C17). This book is updated from the previous edition to cover the features and behaviors of C23. According to 2022 polling data from JetBrains (<https://www.jetbrains.com/lp/devecosystem-2022/c/>), 44 percent of C programmers use C99, 33 percent use C11, 16 percent use C17, and 15 percent use an embedded version of C.

You're about to embark on a journey from which you will emerge a newly minted but professional C developer.

1

GETTING STARTED WITH C



In this chapter, you'll develop your first C program: the traditional “Hello, world!” program. We'll examine the various aspects of this simple C program, compile it, and run it. Then I'll review some editor and compiler options and lay out common portability issues you'll quickly become familiar with as you code in C.

Developing Your First C Program

The most effective way to learn C programming is to start writing C programs, and the traditional program to start with is “Hello, world!” Open your favorite text editor and enter the program in [Listing 1-1](#).

hello.c

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    puts("Hello, world!");
    return EXIT_SUCCESS;
}
```

Listing 1-1: The “Hello, world!” program

The first two lines use the `#include` preprocessor directive, which behaves as if you replaced it with the contents of the specified file at the exact same location. In this program, `<stdio.h>` and `<stdlib.h>` are both headers. A *header* is a source file that, by convention, contains the definitions, function declarations, and constant definitions required by the users of the corresponding source file. As the filenames suggest, `<stdio.h>` defines the interface for C standard input/output (I/O) functions, and `<stdlib.h>` declares several general utility types and functions and defines several macros. You need to include the declarations for any library functions that you use in your program. (You'll learn more about the appropriate use of headers in [Chapter 9](#).)

Here, we include `<stdio.h>` to access the declaration of the `puts` function called by the `main` function. We include the `<stdlib.h>` to access the definition of the `EXIT_SUCCESS` macro, which is used in the `return` statement.

This line defines the `main` function that's called at program startup:

```
int main() {
```

The `main` function defines the entry point for the program that's executed in a hosted environment when the program is invoked from the command line or from another program. C defines two possible execution environments: *freestanding* and *hosted*. A *freestanding* environment may not provide an operating system and is typically used in embedded programming. These implementations provide a minimal set of library functions, and the name and type of the function called at program startup are implementation defined. Most of the examples in this book work on the assumption that the `main` function is the one and only entry point.

Like other procedural languages, C programs contain *functions* that can accept arguments and return values. Each function is a reusable unit of work that you can invoke as often as required in your program. The `puts` function is invoked from the `main` function to print out the line `Hello, world!`:

```
    puts("Hello, world!");
```

The `puts` function is a C standard library function that writes a string argument to the `stdout` stream and appends a newline character to the output. The `stdout` stream typically represents the console or terminal window. `"Hello, world!"` is a string literal that behaves like a read-only string. This function

invocation outputs `Hello, world!` to the terminal.

Once your program has completed, you'll want it to exit. The `return` statement will exit `main` and return an integer value to the host environment or invocation script:

```
return EXIT_SUCCESS;
```

`EXIT_SUCCESS` is an object-like macro that may be defined as follows:

```
#define EXIT_SUCCESS 0
```

Each occurrence of `EXIT_SUCCESS` is replaced by a `0`, which is then returned to the host environment from the call to `main`. The script that invokes the program can then check its status to determine whether the invocation was successful. A return from the initial call to the `main` function is equivalent to calling the C standard library `exit` function with the value returned by the `main` function as its argument.

The final line of this program consists of a closing brace `}`, which closes the code block we opened with the declaration of the `main` function:

```
int main() {  
    // --snip--  
}
```

You can place the opening brace on the same line as the declaration or on its own line, as follows:

```
int main()  
{  
    // --snip--  
}
```

This decision is strictly a stylistic one, because whitespace characters (including newlines) are generally not syntactically meaningful. In this book, I usually place the opening brace on the line with the function declaration because it's stylistically more compact.

For now, save this file as *hello.c*. The file extension `.c` indicates that the file contains C language source code.

NOTE

If you've purchased an ebook, cut and paste the program into the editor. Using cut and paste can reduce transcription errors.

Compiling and Running a Program

Next, you need to compile and run the program, which involves two steps. The command to compile the program depends on which compiler you're using. On Linux and other Unix-like operating systems, enter **cc** on the command line followed by the name of the file you want to compile:

```
$ cc hello.c
```

If you enter the program correctly, the compile command will create a new file called *a.out* in the same directory as your source code.

NOTE

Compilers are invoked differently on other operating systems such as Windows or macOS. Refer to the documentation for your specific compiler.

Inspect your directory with the following command:

```
$ ls
a.out  hello.c
```

The *a.out* file in the output is the executable program, which you can now run on the command line:

```
$ ./a.out
Hello, world!
```

If everything goes right, the program should print `Hello, world!` to the terminal window. If it doesn't, compare the program text from [Listing 1-1](#) to your program and make sure they are the same.

The `cc` command accepts numerous compiler options. The `-o file` compiler option, for example, lets you give the executable file a memorable name instead of *a.out*. The following compiler invocation names the executable *hello*:

```
$ cc -o hello hello.c
$ ./hello
Hello, world!
```

We'll introduce other compiler and linker options (aka flags) throughout the book and dedicate a section to them in [Chapter 11](#).

Function Return Values

Functions will often return a value that's the result of a computation or that signifies whether the function successfully completed its task. For example, the `puts` function we used in our “Hello, world!” program takes a string to print and returns a value of type `int`. The `puts` function returns the value of the macro `EOF` (a negative integer) if a write error occurs; otherwise, it returns a nonnegative integer value.

Although it's unlikely that the `puts` function will fail and return `EOF` for our simple program, it's possible. Because the call to `puts` can fail and return `EOF`, it means that your first C program has a bug or, at least, can be improved as follows:

```
#include <stdio.h>
#include <stdlib.h>
int main() {
    if (puts("Hello, world!") == EOF) {
        return EXIT_FAILURE;
        // code here never executes
    }
    return EXIT_SUCCESS;
    // code here never executes
}
```

This revised version of the “Hello, world!” program checks whether the `puts` call returns the value `EOF`, indicating a write error. If the function returns `EOF`, the program returns the value of the `EXIT_FAILURE` macro (which evaluates to a nonzero value). Otherwise, the function succeeds, and the program returns `EXIT_SUCCESS`. The script that invokes the program can check its status to determine whether it was successful. Code following a return statement is *dead code* that never executes. This is indicated by a single-line comment in the revised program. Everything following `//` is disregarded by the compiler.

Formatted Output

The `puts` function is a simple way to write a string to `stdout`, but to print arguments other than strings, you'll need the `printf` function. The `printf` function takes a format string that defines how the output is formatted, followed by a variable number of arguments that are the actual values you want to print. For example, if you want to use the `printf` function to print out `Hello, world!`, you could write it like this:

```
printf("%s\n", "Hello, world!");
```

The first argument is the format string `"%s\n"`. The `%s` is a conversion specification that instructs the `printf` function to read the second argument (a string literal) and print it to `stdout`. The `\n` is an alphabetic escape sequence used to represent nongraphic characters and tells the function to include a new line after the string. Without the newline sequence, the next characters printed (likely the command prompt) would appear on the same line. This function call outputs the following:

```
Hello, world!
```

Take care not to pass user-supplied data as part of the first argument to the `printf` function, because doing so can result in a formatted output security vulnerability (Seacord 2013).

The simplest way to output a string is to use the `puts` function, as previously shown. If you use `printf` instead of `puts` in the revised version of the “Hello, world!” program, however, you’ll find it no longer works because the `printf` function returns the status differently than the `puts` function. The `printf` function returns the number of characters printed when successful, or a negative value if an output or encoding error occurred. Try modifying the “Hello, world!” program to use the `printf` function as an exercise.

Editors and Integrated Development Environments

You can use a variety of editors and integrated development environments (IDEs) to develop your C programs. [Figure 1-1](#) shows the most used editors, according to a 2023 JetBrains survey (<https://www.jetbrains.com/lp/devecosystem-2023/c/>).




Figure 1-1: Popular IDEs and editors

The exact tools available depend on which system you're using.

For Microsoft Windows, Microsoft's Visual Studio IDE (<https://visualstudio.microsoft.com>) is an obvious choice. Visual Studio comes in three editions: Community, Professional, and Enterprise. The Community edition has the advantage of being free, while the other editions add features at a cost. For this book, you'll need only the Community edition.

For Linux, the choice is less obvious as there are a variety of options. A popular choice is Visual Studio Code (VS Code). VS Code is a streamlined code editor with support for development operations such as debugging, task running, and version control (covered in [Chapter 11](#)). It provides just the tools a developer needs for a quick code-build-debug cycle. VS Code runs on macOS, Linux, and Windows and is free for private or commercial use. Installation instructions are available for Linux and other platforms (<https://code.visualstudio.com>).

[Figure 1-2](#) shows VS Code being used to develop the “Hello, world!” program on Ubuntu. The debug console shows that the program exited with status code 0 as expected.




Figure 1-2: Visual Studio Code running on Ubuntu

Vim is the editor of choice for many developers and power users. It is a text editor based on the vi editor written by Bill Joy in the 1970s for a version of Unix. It inherits the key bindings of vi but also adds functionality and extensibility that are missing from the original vi. You can optionally install Vim plug-ins such as YouCompleteMe (<https://github.com/ycm-core/YouCompleteMe>) or deoplete (<https://github.com/Shougo/deoplete.nvim>) that provide native semantic code completion for C programming.

GNU Emacs is an extensible, customizable, and free text editor. At its core, it's an interpreter for Emacs Lisp, a dialect of the Lisp programming language with extensions to support text editing—although I've never found this to be a problem. Full disclosure: Almost all the production C code I've developed was edited in Emacs.

Compilers

Many C compilers are available, so I won't discuss them all here. Different compilers implement different versions of the C standard. Many compilers for embedded systems support only C89/C90. Popular compilers for Linux and Windows work harder to support modern versions of the C standard, up to and including support for C23.

GNU Compiler Collection

The GNU Compiler Collection (GCC) includes frontends for C, C++, and Objective-C, as well as other languages (<https://gcc.gnu.org>). GCC follows a well-defined development plan under the guidance of the GCC steering committee.

GCC has been adopted as the standard compiler for Linux systems, although versions are also available for Microsoft Windows, macOS, and other platforms. Installing GCC on Linux is easy. The following command, for example, should install GCC on Ubuntu:

```
$ sudo apt-get install gcc
```

You can test the version of GCC you're using with the following command:

```
$ gcc --version
```

The output will display the version and copyright information for the installed GCC version.

Clang

Another popular compiler is Clang (<https://clang.llvm.org>). Installing Clang on Linux is also easy. You can use the following command to install Clang on Ubuntu:

```
$ sudo apt-get install clang
```

You can test the version of Clang you're using with the following command:

```
$ clang --version
```

This displays the installed Clang version.

Microsoft Visual Studio

As previously mentioned, the most popular development environment for Windows is Microsoft Visual Studio, which includes both the IDE and the compiler. Visual Studio (<https://visualstudio.microsoft.com/downloads/>) is bundled with Visual C++ 2022, which includes both the C and C++ compilers.

You can set options for Visual Studio on the Project Property pages. On the Advanced tab under C/C++, make sure you compile as C code by using the Compile as C Code (/TC) option and not the Compile as C++ Code (/TP) option. By default, when you name a file with a .c extension, it's compiled with /TC. If the file is named with .cpp, .cxx, or a few other extensions, it's compiled with /TP.

Portability

Every C compiler implementation is a little different. Compilers continually evolve—so, for example, a compiler like GCC might provide full support for C17 but be working toward support for C23, in which case it might have some C23 features implemented but not others. Consequently, compilers support a full spectrum of C standard versions (including in-between versions). The overall evolution of C implementations is slow, with many compilers significantly lagging behind the C standard.

Programs written for C can be considered *strictly conforming* if they use only those features of the language and library specified in the standard. These programs are intended to be maximally portable. However, because of the range of implementation behaviors, no real-world C program is strictly conforming, nor will it ever be (and probably shouldn't be). Instead, the C standard allows you to write *conforming* programs that may depend on nonportable language and library features.

It's common practice to write code for a single reference implementation, or sometimes several implementations, depending on the platforms to which you plan to deploy your code. The C standard ensures that these implementations don't differ too much, and it allows you to target several at once without having to learn a new language each time.

Five kinds of portability issues are enumerated in Annex J of the C standard documents:

- Implementation-defined behavior
- Unspecified behavior
- Undefined behavior
- Locale-specific behavior
- Common extensions

As you learn about the C language, you'll encounter examples of all five kinds of behaviors, so it's important to understand precisely what these are.

Implementation-Defined Behavior

Implementation-defined behavior is program behavior that's not specified by the C standard and that may produce different results between implementations but has consistent, documented behavior within an implementation. An example of implementation-defined behavior is the number of bits in a byte.

Implementation-defined behaviors are mostly harmless but can cause defects when porting to different implementations. Whenever possible, avoid writing code that depends on implementation-defined behaviors that vary among the C implementations you plan to compile your code with. A complete list of implementation-defined behaviors is enumerated in Annex J.3 of the C standard. You can document your dependencies on these implementation-defined behaviors by using a `static_assert` declaration, as discussed in [Chapter 11](#).

Unspecified Behavior

Unspecified behavior is program behavior for which the standard provides two or more options but doesn't mandate which option is chosen in any instance. Each execution of a given expression may yield different results or produce a different value than a previous execution of the same expression. An example of unspecified behavior is function parameter storage layout, which can vary among function invocations within the same program. Unspecified behaviors are enumerated in Annex J.1 of the C standard.

Undefined Behavior

Undefined behavior is behavior that isn't defined by the C standard or, less circularly, "behavior, upon use of a nonportable or erroneous program construct or of erroneous data, for which the standard imposes no requirements" (ISO/IEC 9899:2024). Examples of undefined behavior include signed integer overflow and dereferencing an invalid pointer value. Code that has undefined behavior is often incorrect, but not always. Undefined behaviors are identified in the standard as follows:

- When a "shall" or "shall not" requirement is violated and that requirement appears outside a constraint, behavior is undefined.
- When behavior is explicitly specified by the words "undefined behavior."
- By the omission of any explicit definition of behavior.

The first two kinds of undefined behavior are frequently referred to as *explicit undefined behaviors*, while the third kind is referred to as *implicit undefined behavior*. There is no difference in emphasis among these three; they all describe behavior that is undefined. The C standard Annex J.2, "Undefined behavior," lists the explicit undefined behaviors in C.

Developers often misconstrue undefined behaviors as errors or omissions in the C standard, but the decision to classify a behavior as undefined is *intentional* and *considered*. Behaviors are classified as undefined by the C standards committee for one of the following reasons:

- Give the implementer license not to catch program errors that are difficult to diagnose
- Avoid defining obscure corner cases that would favor one implementation strategy over another
- Identify areas of possible conforming language extension in which the implementer may augment the language by providing a definition of the

officially undefined behavior

These three reasons are quite different but are all considered portability issues. We'll examine examples of all three as they come up over the course of this book. Upon encountering undefined behavior, compilers have the latitude to do the following:

- Ignore undefined behavior completely, giving unpredictable results
- Behave in a documented manner characteristic of the environment (with or without issuing a diagnostic)
- Terminate a translation or execution (and issue a diagnostic)

None of these options are great (particularly the first), so it's best to avoid undefined behaviors except when the compiler specifies that these behaviors are defined to allow you to invoke a language augmentation. Compilers sometimes have a *pedantic* mode that can help notify the programmer of these portability issues.

Locale-Specific Behavior and Common Extensions

Locale-specific behavior depends on local conventions of nationality, culture, and language that each implementation documents. *Common extensions* are widely used in many systems but are not portable to all implementations.

Summary

In this chapter, you learned how to write a simple C language program, compile it, and run it. We looked at several editors and integrated development environments as well as a few compilers that you can use to develop C programs on Windows, Linux, and macOS systems. You should use newer versions of the compilers and other tools, as they tend to support newer features of the C programming language and provide better diagnostics and optimizations. However, you may not want to use newer versions of compilers if they break your existing code or if you're getting ready to deploy your code to avoid introducing unnecessary changes into your already tested application. We concluded this chapter with a discussion of the portability of C language programs.

Subsequent chapters will examine specific features of the C language and library, starting with objects, functions, and types in the next chapter.

2

OBJECTS, FUNCTIONS, AND TYPES



In this chapter, you'll learn about objects, functions, and types. We'll examine how to declare variables (objects with named identifiers) and functions, take the addresses of objects, and dereference those object pointers. Each object or function instance has a type. You've already seen some types that are available to C programmers. The first thing you'll learn in this chapter is one of the last things that I learned: every type in C is either an object type or a function type.

Entities

An *object* is storage in which you can represent values. To be precise, an object is defined by the C standard (ISO/IEC 9899:2024) as a “region of data storage in the execution environment, the contents of which can represent values,” with the added note, “when referenced, an object can be interpreted as having a particular type.” A variable is an example of an object.

Variables have a declared *type* that tells you the kind of object its value represents. For example, an object with type `int` contains an integer value. Type is important because the collection of bits that represent one type of object can have a different value if interpreted as a different type of object. For example, the number 1 is represented in the IEEE Standard for Floating-Point Arithmetic by the bit pattern `0x3f800000` (IEEE 754-2019). But if you were to interpret this

same bit pattern as an integer, you'd get the value 1,065,353,216 instead of 1.

Functions are not objects but do have types. A function type is characterized by both its return type and the number and types of its parameters.

The C language also has *pointers*, which can be thought of as an *address*—a location in memory where an object or function is stored.

Just like objects and functions, object pointers and function pointers are different things and must not be interchanged. In the following section, you'll write a simple program that attempts to swap the values of two variables to help you better understand objects, functions, pointers, and types.

Declaring Variables

When you declare a variable, you assign it a type and provide it a name, or *identifier*, by which the variable is referenced. Optionally, you can also *initialize* the variable.

[Listing 2-1](#) declares two integer objects with initial values. This simple program also declares, but doesn't define, a swap function to swap those values.

```
#include <stdio.h>
#include <stdlib.h>
```

```
❶ void swap(int, int); // defined in Listing 2-2
```

```
int main() {
    int a = 21;
    int b = 17;
    ❷ swap(a, b);
    printf("main: a = %d, b = %d\n", a, b);
    return EXIT_SUCCESS;
}
```

Listing 2-1: A program meant to swap two integers

This example program shows a main function with a single *compound statement* that includes the {} characters and all the statements between them (also referred to as a *block*). We define two variables, a and b, within the main function. We declare the variables as having the type int and initialize them to 21 and 17, respectively. Each variable must have a declaration. The main function then calls the swap function ❷ to try to swap the values of the two integers. The swap function is declared in this program ❶ but not defined. We'll

look at some possible implementations of this function later in this section.

DECLARING MULTIPLE VARIABLES

You can declare multiple variables in any single declaration, but doing so can become confusing if the variables are pointers or arrays or if the variables are of different types. For example, the following declarations are all correct:

```
char *src, c;  
int x, y[5];  
int m[12], n[15][3], o[21];
```

The first line declares two variables, `src` and `c`, which have different types. The `src` variable has a type of `char *`, and `c` has a type of `char`. The second line again declares two variables, `x` and `y`, with different types. The variable `x` has a type `int`, and `y` is an array of five elements of type `int`. The third line declares three arrays (`m`, `n`, and `o`) with different dimensions and numbers of elements.

These declarations are easier to understand if each is on its own line:

```
char *src;    // src has a type of char *  
char c;       // c has a type of char  
int x;        // x has a type int  
int y[5];     // y is an array of 5 elements of type int  
int m[12];    // m is an array of 12 elements of type int  
int n[15][3]; // n is an array of 15 arrays of 3 elements of t  
ype int  
int o[21];    // o is an array of 21 elements of type int
```

Readable and understandable code is less likely to have defects.

Swapping Values, First Attempt

Each object has a storage duration that determines its *lifetime*, which is the time during program execution for which the object exists, has storage, has a constant address, and retains its last-stored value. Objects must not be referenced outside their lifetime.

Local variables such as `a` and `b` from [Listing 2-1](#) have *automatic storage duration*, meaning that they exist until execution leaves the block in which they're declared. We're going to try to swap the values stored in these two variables. [Listing 2-2](#) shows our first attempt to implement the swap function.

```
void swap(int a, int b) {  
    int t = a;  
    a = b;  
    b = t;  
    printf("swap: a = %d, b = %d\n", a, b);  
}
```

Listing 2-2: A first attempt at implementing the `swap` function

The `swap` function is declared with two parameters, `a` and `b`, that we use to pass arguments to this function. C distinguishes between *parameters*, which are objects declared as part of the function declaration that acquire a value on entry to the function, and *arguments*, which are comma-separated expressions we include in the function call expression. We also declare a temporary variable `t` of type `int` in the `swap` function and initialize it to the value of `a`. This variable is used to temporarily save the value stored in `a` so that it's not lost during the `swap`.

We can now run the generated executable to test the program:

```
% ./a.out  
swap: a = 17, b = 21  
main: a = 21, b = 17
```

This result may be surprising. The variables `a` and `b` were initialized to 21 and 17, respectively. The first call to `printf` within the `swap` function shows that these two values were swapped, but the second call to `printf` in `main` shows the original values unchanged. Let's examine what happened.

C is a *call-by-value* (also called a *pass-by-value*) language, which means that when you provide an argument to a function, the value of that argument is copied into a distinct variable for use within the function. The `swap` function assigns the values of the objects you pass as arguments to their respective parameters. When the parameter values in the function are changed, the argument values in the caller are unaffected because they are distinct objects. Consequently, the variables `a` and `b` retain their original values in `main` during the second call to `printf`. The goal of the program was to swap the values of these two objects. By testing the program, we've discovered it has a bug, or defect.

Swapping Values, Second Attempt

To repair this bug, we can use pointers to rewrite the `swap` function. We use the

indirection (*) operator to both declare pointers and dereference them, as shown in [Listing 2-3](#).

```
void swap(int *pa, int *pb) {  
    int t = *pa;  
    *pa = *pb;  
    *pb = t;  
}
```

Listing 2-3: The revised swap function using pointers

When used in a function declaration or definition, * acts as part of a pointer declarator indicating that the parameter is a pointer to an object or function of a specific type. In the rewritten swap function, we declare two parameters, pa and pb, both having the type pointer to int.

The unary * operator denotes indirection. If its operand has type pointer to τ , the result of the operation has type τ . For example, consider the following assignment:

```
pa = pb;
```

This replaces the value of the pointer pa with the value of the pointer pb. Now consider the assignment in the swap function:

```
*pa = *pb;
```

The *pb operation reads the value referenced by pb, while the *pa operation reads the location referenced by pa. The value referenced by pb is then written to the location referenced by pa.

When you call the swap function in main, you must also place an ampersand (&) character before each variable name:

```
swap(&a, &b);
```

The unary & (*address-of*) operator generates a pointer to its operand. This change is necessary because the swap function now accepts arguments of type pointer to int instead of type int.

[Listing 2-4](#) shows the entire swap program with comments describing the objects created during execution of this code and their values.

```
#include <stdio.h>
#include <stdlib.h>

void swap(int *pa, int *pb) {    // pa → a: 21    pb → b: 17
    int t = *pa;                // t: 21
    *pa = *pb;                  // pa → a: 17    pb → b: 17
    *pb = t;                    // pa → a: 17    pb → b: 21
}

int main() {
    int a = 21;                  // a: 21
    int b = 17;                  // b: 17
    swap(&a, &b);
    printf("a = %d, b = %d\n", a, b);    // a: 17    b: 21
    return EXIT_SUCCESS;
}
```

Listing 2-4: A simulated call by reference

Upon entering the main function block, the variables `a` and `b` are initialized to 21 and 17, respectively. The code then takes the addresses of these objects and passes them to the `swap` function as arguments.

Within the `swap` function, the parameters `pa` and `pb` are now both declared as type pointer to `int` and contain copies of the arguments passed to `swap` from the calling function (in this case, `main`). These address copies still refer to the exact same objects, so when the values of their referenced objects are swapped in the `swap` function, the contents of the original objects declared in `main` are also swapped. This approach simulates *call by reference* (also known as *pass by reference*) by generating object addresses, passing those by value, and then dereferencing the copied addresses to access the original objects.

Object Types

This section introduces object types in C. Specifically, we'll cover the Boolean type, character types, and arithmetic types (including both integer and floating types).

Boolean

A Boolean data type has one of two possible values (`true` or `false`) that represent the two truth values of logic and Boolean algebra. Objects declared as `bool` can store only the values `true` and `false`.

RESERVED IDENTIFIERS

A *Boolean type* was introduced in C99 starting with an underscore (`_Bool`) to differentiate it in existing programs that had already declared their own identifiers named `bool`. Identifiers that begin with an underscore and either an uppercase letter or another underscore are always reserved. The C standards committee often creates new keywords such as `_Bool` under the assumption that you have avoided the use of reserved identifiers. If you haven't, as far as the C standards committee is concerned, it's your fault for not having read the standard more carefully. C23 added the `bool` keyword but retained `_Bool` as an alternative spelling, and `bool` is now the preferred spelling. The keywords `false` and `true` are constants of type `bool` with a value of 0 for `false` and 1 for `true`. If you are using an older version of C, you can include the header `<stdbool.h>`, still spell this type as `bool`, and assign it the values `true` and `false`.

The following example declares a function called `arm_detonator` that takes a single `int` argument and returns a value of type `bool`:

```
bool arm_detonator(int);

void arm_missile(void) {
    bool armed = arm_detonator(3);
    if (armed) puts("missile armed");
    else puts("missile disarmed");
}
```

The `arm_missile` function calls the `arm_detonator` function and assigns the return value to the Boolean variable `armed`. This value can then be tested to determine whether the missile is armed.

Historically, Boolean values were represented by integers and still behave as integers. They can be stored in integer variables and used anywhere integers would be valid, including in indexing, arithmetic, parsing, and formatting. C guarantees that any two true values will compare equally (which was impossible to achieve before the introduction of the `bool` type). You should use the `bool` type to represent Boolean values.

Character

The C language defines the following character types: `char`, `signed char`, and `unsigned char`. Each compiler implementation defines `char` to have the same alignment, size, range, representation, and behavior as either `signed char` or

`unsigned char`. Regardless of the choice made, `char` is a separate type from the other two and is incompatible with both.

The `char` type is commonly used to represent character data in C language programs. Objects of type `char` can represent the *basic execution character set*—the minimum set of characters required in the execution environment—including upper- and lowercase letters, the 10 decimal digits, the space character, punctuation, and control characters. The `char` type is inappropriate for integer data; use `signed char` to represent small, signed integer values, and use `unsigned char` to represent small, unsigned integer values.

The size of objects of type `char` is always 1 byte, and its width is `CHAR_BIT` bits. The `CHAR_BIT` macro from `<limits.h>` defines the number of bits in a byte. The value of `CHAR_BIT` macro cannot be less than 8, and on most modern platforms, it is 8.

The basic execution character set suits the needs of many conventional data processing applications, but its lack of non-English letters is an obstacle to acceptance by international users. To address this need, the C standards committee specified a new wide type to allow large character sets. You can represent the characters of a large character set as *wide characters* by using the `wchar_t` type, which generally takes more space than a basic character. Typically, implementations choose 16 or 32 bits to represent a wide character. The C standard library provides functions that support both narrow and wide character types. The `wchar_t` type was not designed to support Unicode and has consequently fallen out of favor for most implementations with the notable exception of Microsoft Visual Studio.

Arithmetic

C provides several *arithmetic types* that can be used to represent integers, enumerators, and floating-point values. [Chapter 3](#) covers some of these in more detail, but here's a brief introduction.

Integer

Signed integer types can be used to represent negative numbers, positive numbers, and zero. The standard signed integer types include `signed char`, `short int`, `int`, `long int`, and `long long int`.

For each signed integer type, there is a corresponding *unsigned integer type* that uses the same amount of storage: `unsigned char`, `unsigned short int`, `unsigned int`, `unsigned long int`, and `unsigned long long int`. The

unsigned types can represent positive numbers and zero. These unsigned integer types along with type `bool` make up the standard unsigned integer types.

Except for `int` itself, the keyword `int` may be omitted in the declarations for these types, so you might, for example, declare a type by using `long long` instead of `long long int`.

The signed and unsigned integer types are used to represent integers of various widths. Each platform determines the width for each of these types, given some constraints. Each type has a minimum representable range. The types are ordered by width, guaranteeing that wider types are at least as large as narrower types. This means that an object of type `long long int` can represent all values that an object of type `long int` can represent, an object of type `long int` can represent all values that can be represented by an object of type `int`, and so forth. The implementation-defined minimum and maximum representable values for integer types are specified in the `<limits.h>` header file.

Extended integer types may be provided in addition to the standard integer types. They are implementation defined, meaning that their width, precision, and behavior are up to the compiler. Extended integer types are typically larger than the standard integer types (for example, `__int128`).

In addition to the standard and extended integer types, C23 adds *bit-precise integer types*. These types accept an operand specifying the width of the integer, so a `_BitInt(32)` is a signed 32-bit integer, and an unsigned `_BitInt(32)` is an unsigned 32-bit integer. Bit-precise integer types do not require their width to be a power of two; the maximum width supported is specified by `BITINT_MAXWIDTH` (which must be at least the same as the width of unsigned `long long`).

The `int` type is typically assigned the natural width suggested by the architecture of the execution environment (for example, 16 bits on a 16-bit architecture and 32 bits on a 32-bit or 64-bit architecture). You can specify actual-width integers by using type definitions from the `<stdint.h>` or `<inttypes.h>` header, like `uint32_t`. These headers also provide type definitions for the greatest-width integer types: `uintmax_t` and `intmax_t`. The `intmax_t` type, for example, can represent any value of any signed integer type with the possible exceptions of signed bit-precise integer types and of signed extended integer types.

[Chapter 3](#) covers integer types in excruciating detail.

enum

An *enumeration*, or `enum`, allows you to define a type that assigns names

(*enumerators*) to integer values in cases with an enumerable set of constant values. The following are examples of enumerations:

```
enum day {sun, mon, tue, wed, thu, fri, sat};
enum cardinal_points {north = 0, east = 90, south = 180, west = 270};
enum months {jan = 1, feb, mar, apr, may, jun, jul, aug, sep, oct, nov, dec};
```

If you don't specify a value to the first enumerator with the = operator, the value of its enumeration constant is 0, and each subsequent enumerator without an = adds 1 to the value of the previous enumeration constant. Consequently, the value of sun in the day enumeration is 0, mon is 1, and so forth.

You can also assign specific values to each enumerator, as shown by the cardinal_points enumeration. Using = with enumerators may produce enumeration constants with duplicate values, which can be a problem if you incorrectly assume that all the values are unique. The months enumeration sets the first enumerator at 1, and each subsequent enumerator that isn't specifically assigned a value will be incremented by 1.

Starting with C23, you can specify the underlying type of the enumeration. For portability and other reasons (Meneide and Pygott 2022), it is always better to specify the enumeration type. In the following example, the enumeration constant a0 can be assigned the value 0xFFFFFFFFFFFFFFFFULL because the type is specified as unsigned long long:

```
enum a : unsigned long long {
    a0 = 0xFFFFFFFFFFFFFFFFULL
};
```

An omitted type is implementation defined. Visual C++ uses a signed int for the type, and GCC uses an unsigned int.

Floating

Floating-point arithmetic is similar to, and often used as a model for, the arithmetic of real numbers. The C language supports a variety of floating-point representations including, on most systems, representations in the IEEE Standard for Floating-Point Arithmetic (IEEE 754-2019). ISO/IEC 60559:2011 has content identical to IEEE 754-2019 but is referenced by the C standard because

it is published by the same standards organization. The choice of floating-point representation is implementation defined. [Chapter 3](#) covers floating types in detail.

The C language supports three standard floating types: `float`, `double`, and `long double`. The set of values of the type `float` is a subset of the set of values of the type `double`; the set of values of the type `double` is a subset of the set of values of the type `long double`.

C23 adds three *decimal floating types* (ISO/IEC TS 18661-2:2015), designated as `_Decimal32`, `_Decimal64`, and `_Decimal128`. Respectively, these have the decimal32, decimal64, and decimal128 IEC 60559 formats.

The standard floating types and the decimal floating types are collectively called the *real floating types*.

There are also three *complex types*, designated as `float complex`, `double complex`, and `long double complex`.

The real floating and complex types are collectively called the *floating types*. [Figure 2-1](#) shows the hierarchy of floating types.




Figure 2-1: The hierarchy of floating types

Complex types and decimal floating types are not covered in detail in this book.

void

The `void` type is a rather strange type. The keyword `void` (by itself) means “cannot hold any value.” For example, you can use it to indicate that a function doesn’t return a value or as the sole parameter of a function to indicate that the function takes no arguments. On the other hand, the *derived type* `void *` means that the pointer can reference *any* object.

Derived Types

Derived types are constructed from other types. These include function types, pointer types, array types, type definitions, structure types, and union types—all of which are covered here.

Function

A *function type* is derived from the return type and the number and types of its parameters. A function can return any complete object type except for an array type.

When you declare a function, you use the *function declarator* to specify the name of the function and the return type. If the declarator includes a parameter type list and a definition, the declaration of each parameter must include an identifier, except parameter lists with only a single parameter of type `void`, which needs no identifier.

Here are a few function type declarations:

```
int f(void);
int fprime();
int *fip();
void g(int i, int j);
void h(int, int);
```

First, we declare two functions, `f` and `fprime`, with no parameter that returns an `int`. Next, we declare a function `fip` with no parameters that return a pointer to an `int`. Finally, we declare two functions, `g` and `h`, each returning `void` and taking two parameters of type `int`.

Specifying parameters with identifiers (as done here with `g`) can be problematic if an identifier is a macro. However, providing parameter names is good practice for self-documenting code, so omitting the identifiers (as done with `h`) is not typically recommended.

In a function declaration, specifying parameters is optional. However, failing to do so can be problematic. Prior to C23, `fip` declares a function accepting any number of arguments of any type and returning an `int *`. The same function declaration for `fip` in C++ declares a function accepting no arguments and returning an `int *`. Starting with C23, a function declarator with an empty parameter list declares a prototype for a function that takes no arguments (like it does in C++).

A function type is also known as a *function prototype*. A function prototype informs the compiler about the number and types of parameters a function

accepts. Compilers use this information to verify that the correct number and type of parameters are used in the function definition and any calls to the function.

The *function definition* provides the actual implementation of the function. Consider the following function definition:

```
int max(int a, int b)
{return a > b ? a : b;}
```

The return type specifier is `int`; the function declarator is `max(int a, int b)`; and the function body is `{return a > b ? a : b;}`. The specification of a function type must not include any type qualifiers (see “Type Qualifiers” on [page 31](#)). The function body itself uses the conditional operator (`? :`), which is explained in [Chapter 4](#). This expression states that if `a` is greater than `b`, return `a`; otherwise, return `b`.

Pointer

A *pointer type* is derived from a function or object type called the *referenced type*. A pointer type derived from the referenced type τ is called a *pointer to τ* . A pointer provides a reference to an entity of the referenced type.

The following three declarations declare a pointer to `int`, a pointer to `char`, and a pointer to `void`:

```
int *ip = 0; // compliant
char *cp = NULL; // good
void *vp = nullptr; // better
```

Each pointer is initialized to a null pointer constant. A null pointer constant can be specified as an integer constant expression with the value 0, `(void *)0`, or the predefined constant `nullptr`. The `NULL` macro is defined in `<stddef.h>`. If a null pointer constant is converted to a pointer type, the resulting null pointer is guaranteed to compare unequally to a pointer to any object or function.

The `nullptr` constant was introduced in C23 and has advantages to using `NULL` (Gustedt 2022). [Table 2-1](#) shows common values for `NULL` and their associated types.

Table 2-1: Common Values for `NULL` and Their Associated Types

Value	Type
-------	------

0	int
0L	long
(void *)0	void *

These different types can have surprising results when invoking a type-generic macro with a `NULL` argument. The conditional expression `(true ? 0 : NULL)` is always defined, regardless of the type of `NULL`. However, the conditional expression `(true ? 1 : NULL)` is a constraint violation if `NULL` has type `void *`.

A `NULL` argument passed as a sentinel value to a variadic function, such as the Portable Operating System Interface (POSIX) `exec1` function, which expects a pointer, can have unexpected results. On most modern architectures, the `int` and `void *` types have different sizes. If `NULL` is defined as `0` on such an architecture, an incorrectly sized argument is passed to the variadic function.

Earlier in the chapter, I introduced the address-of (`&`) and indirection (`*`) operators. You use the `&` operator to take the address of an object or function. For example, taking the address of an `int` object results in an address with the type pointer to `int`:

```
int i = 17;
int *ip = &i;
```

The second declaration declares the variable `ip` as a pointer to `int` and initializes it to the address of `i`. You can also use the `&` operator on the result of the `*` operator:

```
ip = &*ip;
```

Dereferencing `ip` using the `*` operator resolves to the actual object `i`. Taking the address of `*ip` using the `&` operator retrieves the pointer, so these two operations cancel each other out.

The unary `*` operator converts a pointer to a type `T` into a value of type `T`. It denotes *indirection* and operates only on pointers. If the operand points to a function, the result of using the `*` operator is the function designator, and if it points to an object, the result is a value of the designated object. For example, if the operand is a pointer to `int`, the result of the indirection operator has type `int`. If the pointer is not pointing to a valid object or function, the behavior is undefined.

Array

An *array* is a contiguously allocated sequence of objects that all have the same element type. Array types are characterized by their element types and the number of elements in the array. Here we declare an array of 11 elements of type `int` identified by `ia` and an array of 17 elements of type pointer to `float` identified by `afp`:

```
int ia[11];  
float *afp[17];
```

You can use square brackets (`[]`) to identify an element of an array. For example, the following contrived code snippet creates the string `"0123456789"` to demonstrate how to assign values to the elements of an array:

```
char str[11];  
for (unsigned int i = 0; i < 10; ++i) {  
    str[i] = '0' + i;  
}  
str[10] = '\0';
```

The first line declares an array of `char` with a bound of 11. This allocates sufficient storage to create a string with 10 characters plus a null character. The `for` loop iterates 10 times, with the values of `i` ranging from 0 to 9. Each iteration assigns the result of the expression `'0' + i` to `str[i]`. Following the end of the loop, the null character is copied to the final element of the array `str[10]`, and `str` now contains the string `"0123456789"`.

In the expression `str[i]`, `str` is automatically converted to a pointer to the first member of the array (pointer to `char`), and `i` has an unsigned integer type. The subscript (`[]`) and addition (`+`) operators are defined so that `str[i]` is identical to `*(str + i)`. When `str` is an array object (as it is here), the expression `str[i]` designates the *i*th element of the array (counting from 0). Because arrays are indexed starting at 0, the array `char str[11]` is indexed from 0 to 10, with 10 being the last element, as referenced on the last line of this example.

If the operand of the unary `&` operator is the result of a `[]` operator, the result is as if the `&` operator were removed and the `[]` operator were changed to a `+` operator. For example, `&str[10]` is the same as `str + 10`:

```
&str[10] → &*(str + 10) → str + 10
```

You can also declare multidimensional arrays. [Listing 2-5](#) declares `arr` in the function `main` as a two-dimensional 3×5 array of type `int`, also referred to as a *matrix*.

```
#include <stdlib.h>
void func(int arr[5]);
int main() {
    unsigned int i = 0;
    unsigned int j = 0;
    int arr[3][5];
    ❶ func(arr[i]);
    ❷ int x = arr[i][j];
    return EXIT_SUCCESS;
}
```

Listing 2-5: Matrix operations

More precisely, `arr` is an array of three elements, each of which is an array of five elements of type `int`. When you use the expression `arr[i]` ❶ (which is equivalent to `*(arr+i)`), the following occurs:

1. `arr` is converted to a pointer to the initial array of five elements of type `int` starting at `arr[i]`.
2. `i` is scaled to the type of `arr` by multiplying `i` by the size of one array of five `int` objects.
3. The results from steps 1 and 2 are added.
4. Indirection is applied to the sum to produce an array of five elements of type `int`.

When used in the expression `arr[i][j]` ❷, that array is converted to a pointer to the first element of type `int`, so `arr[i][j]` produces an object of type `int`.

TYPE DEFINITIONS

You use `typedef` to declare an alias for an existing type; it never creates a new type. For example, each of the following declarations creates at least one new type alias:

```
typedef unsigned int uint_type;
typedef signed char schar_type, *schar_p, (*fp)(void);
```

On the first line, we declare `uint_type` as an alias for the type `unsigned int`. On the second line, we declare `schar_type` as an alias for `signed char`, `schar_p` as an alias for `signed char *`, and `fp` as an alias for `signed char (*)(void)`. Identifiers that end in `_t` in the standard headers are type definitions (aliases for existing types). You should not follow this convention in your own code because the C standard reserves identifiers that match the patterns `int[0-9a-z]*_t` and `uint[0-9a-z]*_t`, and POSIX reserves all identifiers that end in `_t`. If you define identifiers that use these names, they may collide with names used by the implementation, which can cause problems that are difficult to debug.

Structure

A *structure type* (also known as a `struct`) contains sequentially allocated members. Each member has its own name and may have a distinct type—unlike array elements, which must all be of the same type. Structures are like record types found in other programming languages.

Structures are useful for declaring collections of related objects and may be used to represent things such as a date, customer, or personnel record. They are especially useful for grouping objects that are frequently passed together as arguments to a function, so you don't need to repeatedly pass individual objects separately.

[Listing 2-6](#) declares a `struct` named `sigline` with type `struct sigrecord` and a pointer to `struct sigrecord` named `sigline_p`.

```
struct sigrecord {  
    int signum;  
    char signame[20];  
    char sigdesc[100];  
} sigline, *sigline_p;
```

Listing 2-6: A struct sigrecord

The structure has three member objects: `signum` is an object of type `int`, `signame` is an array of type `char` consisting of 20 elements, and `sigdesc` is an array of type `char` consisting of 100 elements.

Once you have defined a structure, you'll likely want to reference its members. You reference members of an object of the structure type by using the structure member (`.`) operator. If you have a pointer to a structure, you can reference its members with the structure pointer (`->`) operator. [Listing 2-7](#) demonstrates the use of each operator.

```
sigline.signum = 5;
strcpy(sigline.signame, "SIGINT");
strcpy(sigline.sigdesc, "Interrupt from keyboard");
```

```
❶ sigline_p = &sigline;
```

```
sigline_p->signum = 5;
strcpy(sigline_p->signame, "SIGINT");
strcpy(sigline_p->sigdesc, "Interrupt from keyboard");
```

Listing 2-7: Referencing structure members

The first three lines of [Listing 2-7](#) directly access members of the `sigline` object by using the dot (`.`) operator. We assign the address of the `sigline` object to the `sigline_p` pointer ❶. In the final three lines of the program, we indirectly access the members of the `sigline` object by using the `->` operator through the `sigline_p` pointer.

Union

Union types are like structures, except that the memory used by the member objects overlaps. Unions provide multiple different ways to look at the same memory.

[Listing 2-8](#) shows a union that contains a single member `f` of type `float` and a struct that contains three bitfields of type `uint32_t`: `significand`, `exponent`, and `sign`.

```
static_assert(
    (__STDC_IEC_60559_BFP__ >= 202311L || __STDC_IEC_559__ == 1)
    && __STDC_ENDIAN_LITTLE__
);

union {
    float f;
    struct {
        uint32_t significand : 23;
        uint32_t exponent : 8;
        uint32_t sign : 1;
    };
} float_encoding;
```

Listing 2-8: Decomposing a float using a union

This allows a (low-level) C programmer to use the entire floating-point value and examine (and possibly modify) its constituent parts. This union is not portable because implementations may use a different floating-point representation or endianness. The `static_assert` tests to ensure this union matches the implementation.

[Listing 2-9](#) shows a struct `n` that contains a member `type` and a union `u` that itself contains four members: `inode`, `fnode`, `dnode`, and `ldnode`.

```
enum node_type {
    integer_type,
    float_type,
    double_float_type,
    long_double_type
};

struct node {
    enum node_type type;
    union {
        int inode;
        float fnode;
        double dnode;
        long double ldnode;
    } u;
} n;

n.type = double_type;
n.u.dnode = 3.14;
```

Listing 2-9: Saving memory with a union

This structure might be used in a tree, a graph, or some other data structure that contains differently typed nodes. The `type` member might contain a value between 0 and 3, which indicates the type of the value stored in the structure. It is declared directly in the struct `n` because it is common to all nodes.

As with structures, you can access union members via the `.` operator. Using a pointer to a union, you can reference its members with the `->` operator. In [Listing 2-9](#), the `dnode` member is referenced as `n.u.dnode`. Code that uses this union will typically check the type of the node by examining the value stored in `n.type` and then accessing the value using `n.u.inode`, `n.u.fnode`, `n.u.dnode`, or `n.u.ldnode`, depending on the value stored in `n.type`. Without the union, each node would contain separate storage for all four data types. The use of a union

allows the same storage to be used for all union members. On the x86-64 GCC version 13.2 compiler, using a union saved 16 bytes per node.

Unions are commonly used to describe network or device protocols in cases where you do not know in advance which protocol will be used.

Tags

Tags are a special naming mechanism for structures, unions, and enumerations. For example, the identifier `s` in the following structure is a tag:

```
struct s {  
    // --snip--  
};
```

By itself, a tag is not a type name and cannot be used to declare a variable (Saks 2002). Instead, you must declare variables of this type as follows:

```
struct s v;    // instance of struct s  
struct s *p;  // pointer to struct s
```

The names of unions and enumerations are also tags and not types, meaning that they cannot be used alone to declare a variable. For example:

```
enum day {sun, mon, tue, wed, thu, fri, sat};  
day today; // error  
enum day tomorrow; // OK
```

The tags of structures, unions, and enumerations are defined in a separate *namespace* from ordinary identifiers. This allows a C program to have both a tag and another identifier with the same spelling in the same scope:

```
enum status {ok, fail};    // enumeration  
enum status status(void);  // function
```

You can even declare an object `s` of type `struct s`:

```
struct s s;
```

This may not be good practice, but it is valid C. You can think of `struct`

tags as type names and define an alias for the tag by using a typedef. Here's an example:

```
typedef struct s {int x;} t;
```

This now allows you to declare variables of type `t` instead of `struct s`. The tag name in `struct`, `union`, and `enum` is optional, so you can just dispense with it entirely:

```
typedef struct {int x;} t;
```

This works fine except in the case of self-referential structures that contain pointers to themselves:

```
struct tnode {  
    int count;  
    struct tnode *left;  
    struct tnode *right;  
};
```

C requires the use of tag types (`struct`, `union`, or `enum`) to include the tag name. The compiler will emit a diagnostic if you do not use `struct tnode` in the declaration of the `left` and `right` pointers. Consequently, you must declare a tag for the structure.

You can create an alias for the structure using a typedef:

```
typedef struct tnode {  
    int count;  
    struct tnode *left;  
    struct tnode *right;  
} tnode;
```

The declaration of the `left` and `right` pointers must still use the tag name because the typedef name is not introduced until after the `struct` declaration is complete. You can use the same name for the tag and the typedef, but a common idiom is to name the tag something ugly such as `tnode_` to encourage programmers to use the type name:

```
typedef struct tnode_ {
```

```
int count;
struct tnode_ *left;
struct tnode_ *right;
} tnode;
```

You can also define this type before the structure so that you can use it to declare the `left` and `right` members that refer to other objects of type `tnode`:

```
typedef struct tnode tnode;
struct tnode {
    int count;
    tnode *left;
    tnode *right;
};
```

Type definitions can improve code readability beyond their use with structures. For example, given the following type definition

```
typedef void fv(int), (*pfv)(int);
```

these declarations of the `signal` function all specify the same type:

```
void (*signal(int, void (*)(int)))(int);
fv *signal(int, fv *);
pfv signal(int, pfv);
```

The last two declarations are clearly easier to read.

Type Qualifiers

All the types examined so far have been unqualified types. You can qualify types by using one or more of the following qualifiers: `const`, `volatile`, and `restrict`. Each of these qualifiers changes behaviors when accessing objects of the qualified type.

The qualified and unqualified versions of types can be used interchangeably as arguments to functions, return values from functions, and structure and union members.

NOTE

The `_Atomic` type qualifier, available since C11, supports concurrent programs.

const

Objects declared with the `const` qualifier (const-qualified types) are not assignable but can have constant initializers. This means the compiler can place objects with const-qualified types in read-only memory, and any attempt to write to them will result in a runtime error:

```
const int i = 1; // const-qualified int
i = 2; // error: i is const-qualified
```

It's possible to accidentally convince your compiler to change a const-qualified object for you. In the following example, we take the address of a const-qualified object `i` and tell the compiler that it's actually a pointer to an `int`:

```
const int i = 1; // object of const-qualified type
int *ip = (int *)&i;
*ip = 2; // undefined behavior
```

C does not allow you to cast away the `const` if the original was declared as a const-qualified object. This code might appear to work, but it's defective and may fail later. For example, the compiler might place the const-qualified object in read-only memory, causing a memory fault when trying to store a value in the object at runtime.

C allows you to modify an object that is referenced by a const-qualified pointer by casting the `const` away, provided that the original object was not declared `const`:

```
int i = 12;
const int j = 12;
const int *ip = &i;
const int *jp = &j;
*(int *)ip = 42; // OK
*(int *)jp = 42; // undefined behavior
```

Be careful not to pass a const-qualified pointer to a function that modifies the object.

volatile

Objects are given a `volatile`-qualified type to allow for processes that are *extrinsic* to the compiler. The values stored in these objects may change without the knowledge of the compiler, or a write may synchronize externally. For example, every time the value from a real-time clock is read, it may change, even if the value has not been written to by the C program. Using a `volatile`-qualified type lets the compiler know that the value may change without its knowledge and ensures that every access to the real-time clock occurs. Otherwise, an access to the real-time clock may be optimized away or replaced by a previously read and cached value.

A `volatile`-qualified type can be used to access memory-mapped registers, which are accessed through an address just like any other memory. Input/output (I/O) devices often have memory-mapped registers, where you can write to, or read from, a specific address to set or retrieve information or data. Each read and write operation must occur, even if the compiler can see no reason for it. Declaring an object as `volatile` ensures that each read or write of that object at runtime occurs the same number of times and in the same order as indicated by the source code. For example, if `port` is defined as a `volatile`-qualified `int`, the compiler must generate instructions to read the value from `port` and then write this value back to `port` in the assignment:

```
port = port;
```

Without the `volatile` qualification, the compiler would see this as a no-op (a programming statement that does nothing) and might eliminate both the read and the write. Reads and writes of `volatile` memory are touched exactly once. A `volatile` operation cannot be eliminated or fused with a subsequent one, even if the compiler believes it's useless. A `volatile` operation cannot be speculated, even if the compiler can undo or otherwise make that speculation benign.

Objects with `volatile`-qualified types are used when a compiler is not aware of external interactions. For example, `volatile`-qualified types can be used for memory shared with untrusted code to avoid time-of-check to time-of-use (ToCToU) vulnerabilities. Such types are used to access objects from a signal handler and with `setjmp/longjmp` (refer to the C standard for information on signal handlers and `setjmp/longjmp`). Unlike Java and other programming languages, `volatile`-qualified types should not be used in C for synchronization between threads.

Memory-mapped I/O ports are modeled by a `static volatile`-qualified

objects model. Memory-mapped input ports such as a real-time clock are modeled by `static const volatile`-qualified objects. A `const volatile`-qualified object models a variable that can be altered by a separate thread. The meaning of the `static` storage-class specifier is explained later in this chapter.

restrict

A `restrict`-qualified pointer is used to promote optimization. Objects indirectly accessed through a pointer frequently cannot be fully optimized because of potential aliasing, which occurs when more than one pointer refers to the same object. Aliasing can inhibit optimizations because the compiler can't tell whether an object can change values when another apparently unrelated object is modified, for example.

The following function copies `n` bytes from the storage referenced by `q` to the storage referenced by `p`. The function parameters `p` and `q` are both `restrict`-qualified pointers:

```
void f(unsigned int n, int * restrict p, int * restrict q) {
    while (n-- > 0) {
        *p++ = *q++;
    }
}
```

Because both `p` and `q` are `restrict`-qualified pointers, the compiler can assume that an object accessed through one of the pointer parameters is not also accessed through the other. The compiler can make this assessment based solely on the parameter declarations without analyzing the function body.

Although using `restrict`-qualified pointers can result in more efficient code, you must ensure that the pointers do not refer to overlapping memory to prevent undefined behavior.

Scope

Objects, functions, macros, and other C language identifiers have *scope* that delimits the contiguous region where they can be accessed. C has four types of scope: file, block, function prototype, and function.

The scope of an object or function identifier is determined by where it is declared. If the declaration is outside any block or parameter list, the identifier has *file scope*, meaning its scope is the entire text file in which it appears as well

as any included files.

If the declaration appears inside a block or within the list of parameters, it has *block scope*, meaning that the identifier is accessible only from within the block. The identifiers for `a` and `b` from [Listing 2-4](#) have block scope and can be referenced only from within the code block in the `main` function in which they're defined.

If the declaration appears within the list of parameter declarations in a function prototype (not part of a function definition), the identifier has *function prototype scope*, which terminates at the end of the function declarator. *Function scope* is the area between the opening `{` of a function definition and its closing `}`. A label name is the only kind of identifier that has function scope. *Labels* are identifiers followed by a colon, and they identify a statement in the same function to which control may be transferred. ([Chapter 5](#) covers labels and control transfer.)

Scopes also can be *nested*, with *inner* and *outer* scopes. For example, you can define a block scope inside another block scope, and every block scope is defined within a file scope. The inner scope has access to the outer scope, but not vice versa. As the name implies, any inner scope must be completely contained within any outer scope that encompasses it.

If you declare the same identifier in both an inner scope and an outer scope, the identifier declared in the outer scope is *hidden* (also known as *shadowed*) by the identifier declared in the inner scope. Referencing the identifier from the inner scope will refer to the object in the inner scope; the object in the outer scope is hidden and cannot be referenced by its name. The easiest way to prevent this from becoming a problem is to use different names. [Listing 2-10](#) demonstrates different scopes and how identifiers declared in inner scopes can hide identifiers declared in outer scopes.

```
int j; // file scope of j begins

void f(int i) { // block scope of i begins
    int j = 1; // block scope of j begins; hides file-sc
ope j
    i++; // i refers to the function parameter
    for (int i = 0; i < 2; i++) { // block scope of loop-local i b
egins
        int j = 2; // block scope of the inner j begins; hid
es outer j
        printf("%d\n", j); // inner j is in scope, prints 2
```

```
    }                // block scope of the inner i and j ends
    printf("%d\n", j); // the outer j is in scope, prints 1
} // the block scope of i and j ends

void g(int j);        // j has function prototype scope; hides
file-scope j
```

Listing 2-10: Identifiers declared in inner scopes hiding identifiers declared in outer scopes

There is nothing wrong with this code, provided the comments accurately describe your intent. However, it's better to use different names for different identifiers to avoid confusion, which leads to bugs. Using short names such as `i` and `j` is fine for identifiers with small scopes. Identifiers in large scopes should have longer, descriptive names that are unlikely to be hidden in nested scopes. Some compilers will warn about hidden identifiers.

Storage Duration

Objects have a storage duration that determines their lifetime. Four storage durations are available: automatic, static, thread, and allocated. You've already seen that objects declared within a block or as a function parameter have automatic storage duration. The lifetime of these objects starts when the block in which they're declared begins execution and ends when execution of this block completes. If the block is entered recursively, a new object is created each time the block is entered, and each object has its own storage.

NOTE

Scope and lifetime are entirely different concepts. Scope applies to identifiers, whereas lifetime applies to objects. The scope of an identifier is the code region where the object denoted by the identifier can be accessed by its name. The lifetime of an object is the period for which the object exists.

Objects declared in file scope have *static* storage duration. The lifetime of those objects is the entire execution of the program, and their stored value is initialized prior to program startup.

Thread storage duration is used in concurrent programming and is not covered in this book. *Allocated* storage duration involves dynamically allocated memory and is discussed in [Chapter 6](#). Finally, as described in the next section, a storage-class specifier can determine or influence storage duration.

Storage Class

You can specify the storage class of an object or functions using storage-class specifiers. For C23, these include `auto`, `constexpr`, `extern`, `register`, `static`, `thread_local`, and `typedef`. The `constexpr` storage-class specifier is new in C23, and the `auto` storage-class specifier is significantly changed.

Storage-class specifiers specify various properties of identifiers and declared features:

- Storage duration: `static` in block scope, `thread_local`, `auto`, and `register`
- Linkage: `extern`, `static` and `constexpr` in file scope, and `typedef`
- Value: `constexpr`
- Type: `typedef`

With a few exceptions, only one storage-class specifier is allowed for each declaration. For example, `auto` may appear with all the others except `typedef`.

static

The `static` storage-class specifier is used to specify both storage duration and linkage.

File scope identifiers specified as `static` or `constexpr`, or functions specified as `static`, have internal linkage.

You can also declare a variable with block scope to have static storage duration by using the storage-class specifier `static`, as shown in the counting example in [Listing 2-11](#). These objects persist after the function has exited.

```
#include <stdio.h>
#include <stdlib.h>

void increment(void) {
    static unsigned int counter = 0;
    counter++;
    printf("%d ", counter);
}

int main() {
    for (int i = 0; i < 5; i++) {
        increment();
    }
    return EXIT_SUCCESS;
```

```
}
```

Listing 2-11: A counting example

This program outputs 1 2 3 4 5. The static variable `counter` is initialized to 0 once at program startup and incremented each time the `increment` function is called. The lifetime of `counter` is the entire execution of the program, and it will retain its last-stored value throughout its lifetime. You could achieve the same behavior by declaring `counter` with file scope. However, it's good software engineering practice to limit the scope of an object whenever possible.

extern

The `extern` specifier specifies static storage duration and external linkage. It can be used with function and object declarations in both file and block scope (but not function parameter lists). If `extern` is specified for the redeclaration of an identifier that has already been declared with internal linkage, the linkage remains internal. Otherwise (if the prior declaration was external, has no linkage, or is not in scope), the linkage is external.

thread_local

An object whose identifier is declared with the `thread_local` storage-class specifier has *thread storage duration*. Its initializer is evaluated prior to program execution, its lifetime is the entire execution of the thread for which it is created, and its stored value is initialized with the previously determined value when the thread is started. There is a distinct object per thread, and use of the declared name in an expression refers to the object associated with the thread evaluating the expression. (The topic of threading is beyond the scope of this book.)

constexpr

A scalar object declared with the `constexpr` storage-class specifier is a constant and has its value permanently fixed at translation time. The `constexpr` storage-class specifier may appear with `auto`, `register`, or `static`. If not already present, a `const` qualification is implicitly added to the object's type. The resulting object cannot be modified at runtime in any way. The compiler can then use this value in any other constant expression.

Additionally, the constant expression used for the initializer of such a constant is checked at compile time. Before the introduction of `constexpr` in

C23, a very large object constant might be declared as follows:

```
static size_t const BF0 = 0x100000000;
```

The initializer may or may not fit into `size_t`; a diagnostic is not required. In C23, this same object can be declared using `constexpr` as follows:

```
constexpr size_t BF0 = 0x100000000;
```

Now, a diagnostic is required on implementations where `size_t` has a width of 32 or less.

Static objects must be initialized with a constant value and not a variable:

```
int *func(int i) {  
    const int j = i; // ok  
    static int k = j; // error  
    return &k;  
}
```

Arithmetic constant expressions are allowed in initializers. Constant values are literal constants (for example, 1, 'a', or 0xFF), enum members, a scalar object declared with the `constexpr` storage-class specifier, and the result of operators such as `alignof` or `sizeof` (provided the operand does not have a variable-length array type). Unfortunately, `const`-qualified objects are not constant values. Starting with C23, an implementation may accept other forms of constant expressions; it is implementation defined whether they are integer constant expressions.

register

The `register` storage-class specifier suggests that access to an object be as fast as possible. The extent to which such suggestions are effective is implementation defined. Frequently, compilers can make better decisions about register allocation and ignore these programmer suggestions. The `register` storage class can be used only for an object that never has its address taken. A compiler can treat any register declaration simply as an `auto` declaration. However, whether addressable storage is used, the address of any part of an object declared with a storage-class specifier `register` cannot be computed, either explicitly by use of the unary `&` operator or implicitly by converting an array name to a pointer.

typedef

The typedef storage-class specifier defines an identifier to be a typedef name that denotes the type specified for the identifier. The typedef storage-class specifier was discussed earlier in the “Type Definitions” box.

auto

Prior to C23, the auto specifier was allowed only for objects declared at block scope (except function parameter lists). It indicates automatic storage duration and no linkage, which are the defaults for these kinds of declarations.

C23 introduced type inference into the C language by expanding the definition of the existing auto storage-class specifier. Prior to C23, declaring a variable in C requires the user to name a type. However, when the declaration includes an initializer, the type can be derived directly from the type of the expression used to initialize the variable. This has been a C++ feature since 2011.

The auto storage duration class specifier has similar behavior to C++ in that it allows the type to be inferred from the type of the assignment value. Take the following file scope definitions, for example:

```
static auto a = 3;
auto p = &a;
```

Because the integer literal 3 has an implicit type of int, these declarations are interpreted as if they had been written as:

```
static int a = 3;
int * p = &a;
```

Effectively, a is an int, and p is an int *. Type inference is extremely useful when implementing or invoking type-generic macros, as we’ll see in [Chapter 9](#).

typeof Operators

C23 introduced the typeof operators typeof and typeof_unqual. The typeof operators can operate on an expression or a type name and yield the type of their operand. If the type of the operand is a variably modified type, the operand is evaluated; otherwise, the operand is not evaluated.

The `typeof` operators and the auto storage duration class specifier both perform automatic type inference. They can both be used to determine the type of expression.

The auto storage duration class specifier is commonly used to declare initialized variables where the type can be inferred from the initial value. However, to form a derived type, you must use the `typeof` operator:

```
_Atomic(typeof(x)*) apx = &x;
```

The auto storage duration class specifier cannot be used with `_Generic` (described in [Chapter 9](#)) and `typedef` (described later in this chapter).

The result of the `typeof_unqual` operator is the nonatomic, unqualified version of the type that would result from the `typeof` operator. The `typeof` operator preserves all qualifiers.

The `typeof` operator is like the `sizeof` operator, which executes the expression in an unevaluated context to understand the final type. You can use the `typeof` operator anywhere you can use a type name. The following example illustrates the use of both `typeof` operators:

```
#include <stdlib.h>
const _Atomic int asi = 0;
const int si = 1;
const char* const beatles[] = {
    "John",
    "Paul",
    "George",
    "Ringo"
};

❶ typeof_unqual(si) main() {
    ❷ typeof_unqual(asi) plain_si;
    ❸ typeof(_Atomic ❹ typeof(si)) atomic_si;
    ❺ typeof(beatles) beatles_array;
    ❻ typeof_unqual(beatles) beatles2_array;
    return EXIT_SUCCESS;
}
```

At the first use of the `typeof_unqual` operator ❶, the operand is `si`, which has the type `const int`. The `typeof_unqual` operator strips the `const` qualifier,

resulting in just plain `int`. This use of the `typeof_unqual` operator is illustrative and not meant for production code. The `typeof_unqual` operator is used again on operand `asi` ❷, which has the type `const _Atomic int`. All qualifiers are once again stripped, resulting in a plain `int`. The operand to the `typeof` specifier at ❸ includes another `typeof` specifier. If the `typeof` operand is itself a `typeof` specifier, the operand is evaluated before evaluating the current `typeof` operator. This evaluation happens recursively until a `typeof` specifier is no longer the operand. In this case, the `typeof` specifier at ❸ does nothing and can be omitted. The `typeof` operator at ❹ is evaluated before the `typeof` operator at ❸ and returns `const int`. The `typeof` operator at ❸ is now evaluated and returns `const _Atomic int`. The `typeof` operator at ❺ returns a `const` array of four `const char` pointers. The `typeof_unqual` operator at ❻ strips the qualifier and returns an array of four `const char` pointers. The qualifiers, in this case, are stripped only from the array and not the element types the array contains.

The following main function is equivalent but doesn't use `typeof` operators:

```
int main() {
    int plain_si;
    const _Atomic int atomic_si;
    const char* const beatles_array[4];
    const char* beatles2_array[4];
    return EXIT_SUCCESS;
}
```

You can use the `typeof` operator to refer to a macro parameter to construct objects with the required types without specifying the type names explicitly as macro arguments.

Alignment

Object types have alignment requirements that place restrictions on the addresses at which objects of that type may be allocated. An *alignment* represents the number of bytes between successive addresses at which a given object can be allocated. Central processing units (CPUs) may have different behavior when accessing aligned data (for example, where the data address is a multiple of the data size) versus unaligned data.

Some machine instructions can perform multibyte accesses on nonword boundaries, but with a performance penalty. A *word* is a natural, fixed-sized unit

of data handled by the instruction set or the hardware of the processor. Some platforms cannot access unaligned memory. Alignment requirements may depend on the CPU word size (typically, 16, 32, or 64 bits).

Generally, C programmers need not concern themselves with alignment requirements, because the compiler chooses suitable alignments for its various types. However, on rare occasions, you might need to override the compiler's default choices—for example, to align data on the boundaries of the memory cache lines that must start at power-of-two address boundaries or to meet other system-specific requirements. Traditionally, these requirements were met by linker commands or similar operations involving other nonstandard facilities.

C11 introduced a simple, forward-compatible mechanism for specifying alignments. Alignments are represented as values of the type `size_t`. Every valid alignment value is a nonnegative integral power of two. An object type imposes a default alignment requirement on every object of that type: a stricter alignment (a larger power of two) can be requested using the alignment specifier (`alignas`). You can include an alignment specifier in a declaration. [Listing 2-12](#) uses the alignment specifier to ensure that `good_buff` is properly aligned (`bad_buff` may have incorrect alignment for member-access expressions).

```
struct S {
    double d; int i; char c;
};

int main() {
    unsigned char bad_buff[sizeof(struct S)];
    alignas(struct S) unsigned char good_buff[sizeof(struct S)];
    struct S *bad_s_ptr = (struct S *)bad_buff;
    struct S *good_s_ptr = (struct S *)good_buff; // correct alignment
    good_s_ptr->i = 12;
    return good_s_ptr->i;
}
```

Listing 2-12: Use of the `alignas` keyword

Although `good_buff` has proper alignment to be accessed through an lvalue of type `struct S`, this program still has undefined behavior. This undefined behavior stems from the underlying object `good_buff` being declared as an array of objects of type `unsigned char` and being accessed through an lvalue of a different type. The cast to `(struct S *)`, like any pointer cast, doesn't change

the effective type of the storage allocated to each array. Because it is an established practice to use areas of character type for low-level storage management, I co-authored a paper to make such code conforming in a future revision of the C standard (Seacord et al. 2024).

Alignments are ordered from weaker to stronger (also called *stricter*) alignments. Stricter alignments have larger alignment values. An address that satisfies an alignment requirement also satisfies any valid, weaker alignment requirement.

Alignment of dynamically allocated memory is covered in [Chapter 6](#).

Variably Modified Types

Variably modified types (VMTs) define a base type and an extent (number of elements), which is determined at runtime. VMTs are a mandatory feature of C23.

VMTs can be used as function parameters. Remember from earlier in this chapter that, when used in an expression, an array is converted to a pointer to the first element of the array. This means that we must add an explicit parameter to specify the size of the array—for example, the *n* parameter in the signature for `memset`:

```
void *memset(void *s, int c, size_t n);
```

When you call such a function, *n* should accurately represent the size of the array referenced by *s*. Undefined behavior results if this size is larger than the array.

When declaring a function to take an array as an argument that specifies a size, we must declare the size of the array before referencing the size in the array declaration. We could, for example, modify the signature for the `memset` function as follows to take the number of elements *n* and an array of at least *n* elements:

```
void *memset_vmt(size_t n, char s[n], int c);
```

For arrays of character type, the number of elements is equal to the size. In this function signature, `s[n]` is a variably modified type because `s[n]` depends on the runtime value of *n*.

We’ve changed the order of the parameters so that the size parameter *n* is declared before we use it in the array declaration. The array argument *s* is still

adjusted to a pointer, and no storage is allocated because of this declaration (except for the pointer itself). When calling this function, you must declare the actual storage for the array referenced by `s` and ensure that `n` is a valid size for it. Just like a non-VMT parameter, the actual array storage may be a fixed-size array, variable-length array (covered in [Chapter 6](#)), or dynamically allocated storage.

VMTs can generalize your functions, making them more useful. For example, the `matrix_sum` function sums all the values in a two-dimensional array. The following version of this function accepts a matrix with a fixed column size:

```
int matrix_sum(size_t rows, int m[][4]);
```

When passing a multidimensional array to a function, the number of elements in the initial dimension of the array (the rows) is lost and needs to be passed in as an argument. The `rows` parameter provides this information in this example. You can call this function to sum the values of any matrix with exactly four columns, as shown in [Listing 2-13](#).

```
int main(void) {
    int m1[5][4];
    int m2[100][4];
    int m3[2][4];
    printf("%d.\n", matrix_sum(5, m1));
    printf("%d.\n", matrix_sum(100, m2));
    printf("%d.\n", matrix_sum(2, m3));
}
```

Listing 2-13: Summing matrices with four columns

This is fine until you need to sum the values of a matrix that does not have four columns. For example, changing `m3` to have five columns would result in a warning such as this:

```
warning: incompatible pointer types passing 'int [2][5]' to parameter of type 'int (*)(*)[4]'
```

To accept this argument, you would have to write a new function with a signature that matches the new dimensions of the multidimensional array. The

problem with this approach, then, is that it fails to generalize sufficiently.

Instead of doing that, we can rewrite the `matrix_sum` function to use a VMT, as shown in [Listing 2-14](#). This change allows us to call `matrix_sum` with matrices of any dimension.

```
int matrix_sum(size_t rows, size_t cols, int m[rows][cols]) {
    int total = 0;

    for (size_t r = 0; r < rows; r++)
        for (size_t c = 0; c < cols; c++)
            total += ❶ m[r][c];
    return total;
}
```

Listing 2-14: Using a VMT as a function parameter

The compiler performs the matrix indexing ❶. Without VMTs, this would require either manual indexing or double indirection, which are both error prone.

Again, no storage is allocated by either the function declaration or the function definition. As with a non-VMT parameter, you need to allocate the storage for the matrix separately, and its dimensions must match those passed to the function as the `rows` and `cols` arguments. Failing to do so can result in undefined behavior.

Attributes

Starting with C23, you can use *attributes* to associate additional information with a declaration, statement, or type. This information can be used by the implementation to improve diagnostics, improve performance, or modify the behavior of the program in other ways. A comma-delimited list of zero or more attributes is specified within a pair of double square brackets, for example, `[[foo]]` or `[[foo, bar]]`.

Declarations attributes are specified in two ways. If the attribute specifier is at the start of a declaration, the attributes are applied to all declarations in the declaration group. Otherwise, the attributes are applied to the declaration to the immediate left of the attribute specifier. For example, in the following declaration group, the `foo` attribute is applied to `x`, `y`, and `z`:

```
[[foo]] int x, y, *z;
```

While in the second declaration group, the `foo` and `bar` attributes are applied only to `b`:

```
int a, b [[foo, bar]], *c;
```

C23 defines several attributes that apply to declarations, such as `nodiscard` and `deprecated`. The `nodiscard` attribute is used with function declarations to denote that the value returned by the function is expected to be used within an expression or initializer. The `deprecated` attribute is used with the declaration of a function or a type to denote that use of the function or type should be diagnosed as discouraged.

In addition to standard attributes, the implementation may provide nonportable attributes. Such attributes are also specified within double square brackets, but they include a vendor prefix to distinguish between attributes from different vendors. For example, the `[[clang::overloadable]]` attribute is used on a function declaration to specify that it can use C++-style function overloading in C, and the `[[gnu::packed]]` attribute is used on a structure declaration to specify that the member declarations of the structure should avoid using padding between member declarations whenever possible for a more space-efficient layout. Vendors typically use their own prefixes, and they may use whatever prefixes they choose. For example, Clang implements many attributes with the `gnu` prefix for improved compatibility with GCC. Your compiler should ignore unknown attributes, although they may still be diagnosed so you know that the attribute has no effect. Refer to your compiler's documentation for the full list of supported attributes.

EXERCISES

1. Add a `retrieve` function to the counting example from [Listing 2-6](#) to retrieve the current value of `counter`.
2. Declare an array of three pointers to functions and invoke the appropriate function based on an index value passed in as an argument.
3. Repair the following program with the appropriate use of the `volatile` type qualifier:

```
#include <stdlib.h>
#include <stdio.h>
#include <setjmp.h>

static jmp_buf buf;
```

```
int main() {
    int foo = 5;
    if (setjmp(buf) != 2) {
        if (foo != 5) {puts("hi"); longjmp(buf, 2);}
        foo = 6;
        longjmp(buf, 1);
    }
    return EXIT_SUCCESS;
}
```

Hint: The problem may only manifest for optimized builds.

Summary

In this chapter, you learned about objects and functions and how they differ. You learned how to declare objects and functions, take the addresses of objects, and dereference those object pointers. You also learned about most of the object types available to C programmers as well as derived types.

We'll return to these types in later chapters to further explore how they can be best used to implement your designs. In the next chapter, I provide detailed information about the two kinds of arithmetic types: integers and floating-point.

3

ARITHMETIC TYPES



In this chapter, you'll learn about the two kinds of *arithmetic types*: integers and floating types. Most operators in C operate on arithmetic types. Because C is a system-level language, performing arithmetic operations correctly can be difficult, resulting in frequent defects. This is partially because arithmetic operations in digital systems with limited range and precision do not always produce the same result as they would in ordinary mathematics. Performing basic arithmetic correctly is an essential foundation to becoming a professional C programmer.

We'll dive deep into how arithmetic works in the C language so that you have a firm grasp of these fundamental concepts. We'll also look at how to convert one arithmetic type to another, which is necessary for performing operations on mixed types.

Integers

As mentioned in [Chapter 2](#), each integer type represents a finite range of integers. Signed integer types represent values that can be negative, zero, or positive; unsigned integers represent values that can be only zero or positive. The range of values that each integer type can represent depends on your

implementation.

The *value* of an integer object is the ordinary mathematical value stored in the object. The *representation* of a value for an integer object is the particular encoding of the value in the bits of the object's allocated storage. We'll look at the representation in more detail later.

Padding, Width, and Precision

All integer types except `char`, `signed char`, and `unsigned char` may contain unused bits, called *padding*, that allow implementations to accommodate hardware quirks (such as skipping over a sign bit in the middle of a multiple-word representation) or to optimally align with a target architecture. The number of bits used to represent a value of a given type, excluding padding but including the sign, is called the *width* and is often denoted by *N*. The *precision* is the number of bits used to represent values, excluding sign and padding bits.

Integer Ranges

A *representable value* is a value that can be represented in the number of bits available to an object of a particular type. Values that cannot be represented will be diagnosed by the compiler or converted to a representable but different (incorrect) value. The `<limits.h>` header file defines object-like macros that expand to various limits and parameters of the standard integer types. To write portable code, you should use these macros rather than integer literals such as `+2147483647` (the maximum value representable as a 32-bit integer) that represent a specific limit and may change when porting to a different implementation.

The C standard imposes only three constraints on integer sizes. First, storage for every data type occupies an integral number of adjacent `unsigned char` objects (which may include padding). Second, each integer type must support a minimum range of values, allowing you to depend on a portable range of values across any implementation. Third, smaller types cannot be wider than larger types. So, for example, `short` cannot be wider than `int`, but both types may have the same width.

Integer Declarations

Unless declared as `unsigned`, integer types are assumed to be signed (except for `char`, which the implementation can define as either a signed or unsigned integer type). The following are valid declarations of unsigned integers:

```
unsigned int ui; // unsigned is required
unsigned u; // int can be omitted
unsigned long long ull2; // int can be omitted
unsigned char uc; // unsigned is required
```

When declaring signed integer types, you can omit the signed keyword—except for signed char, which requires the keyword to distinguish signed char from plain char.

You can also omit int when declaring variables of type short, long, or long long. For example:

```
int i; // signed can be omitted
long long int sll; // signed can be omitted
long long sll2; // signed and int can be omitted
signed char sc; // signed is required
```

These are all valid signed integer declarations.

Unsigned Integers

Unsigned integers have ranges that start at 0, and their upper bound is greater than that of the corresponding signed integer type. Unsigned integers are frequently used for counting items that may have large, nonnegative quantities.

Representation

Unsigned integer types are easier to understand and to use than signed integer types. They represent values using a pure binary system with no offset: the least significant bit has the weight 2^0 , the next least significant has the weight 2^1 , and so forth. The value of the binary number is the sum of all the weights of the set bits. [Table 3-1](#) shows some examples of unsigned values using an unpadded 8-bit representation.

Table 3-1: 8-Bit Unsigned Values

Decimal	Binary	Hexadecimal
0	0b00000000	0x00
1	0b00000001	0x01
42	0b00101010	0x2A
255	0b11111111	0xFF

Unsigned integer types do not have a sign bit, allowing for 1-bit greater

precision than the corresponding signed integer types. Unsigned integer values range from 0 to a maximum value that depends on the width of the type. This maximum value is $2^N - 1$, where N is the width. For example, most x86 architectures use 32-bit integers with no padding bits, so an object of type `unsigned int` has a range of 0 to $2^{32} - 1$ (4,294,967,295). The constant expression `UINT_MAX` from `<limits.h>` specifies the implementation-defined upper range for this type. [Table 3-2](#) shows the constant expressions from `<limits.h>` for each unsigned type and the minimum magnitude required by the standard.

Table 3-2: Unsigned Integer Minimum Magnitudes

Constant expression	Minimum magnitude	Maximum value for an object of type
<code>UCHAR_MAX</code>	$255 // 2^8 - 1$	<code>unsigned char</code>
<code>USHRT_MAX</code>	$65,535 // 2^{16} - 1$	<code>unsigned short int</code>
<code>UINT_MAX</code>	$65,535 // 2^{16} - 1$	<code>unsigned int</code>
<code>ULONG_MAX</code>	$4,294,967,295 // 2^{32} - 1$	<code>unsigned long int</code>
<code>ULLONG_MAX</code>	$18,446,744,073,709,551,615 // 2^{64} - 1$	<code>unsigned long long int</code>

Your compiler will replace these values with implementation-defined magnitudes.

Wraparound

C23 defines *wraparound* as “the process by which a value is reduced modulo 2^N , where N is the width of the resulting type.” Wraparound occurs when you perform arithmetic operations that result in values too small (less than 0) or too large (greater than $2^N - 1$) to be represented as a particular unsigned integer type. In this case, the value is reduced modulo the number that is one greater than the largest value that can be represented in the resulting type. Wraparound is well-defined behavior in the C language. Whether it is a defect in your code depends on the context. If you are counting something and the value wraps, it is likely to be an error. However, the use of wraparound in certain algorithms is intentional.

The code in [Listing 3-1](#) illustrates wraparound by initializing an unsigned integer value `ui` to its maximum value and incrementing it.

```
unsigned int ui = UINT_MAX; // 4,294,967,295 on x86
ui++;
printf("ui = %u\n", ui); // ui is 0
ui--;
```

```
printf("ui = %u\n", ui); // ui is 4,294,967,295
```

Listing 3-1: Unsigned integer wraparound

The resulting value cannot be represented as an unsigned int, so it wraps around to 0. If the resulting value is decremented, it falls outside the range again and will wrap around back to `UINT_MAX`.

Because of wraparound, an unsigned integer expression can never evaluate to less than 0. It's easy to lose track of this and implement comparisons that are always true or always false. For example, the `i` in the following for loop can never take on a negative value, so this loop will never terminate:

```
for (unsigned int i = n; i >= 0; --i)
```

This behavior has caused some notable real-world bugs. For example, all six power-generating systems on a Boeing 787 are managed by a corresponding generator control unit. Boeing's laboratory testing discovered that an internal software counter in the generator control unit wraps around after running continuously for 248 days (see the Federal Aviation Administration Rule at <https://www.federalregister.gov/documents/2015/05/01/2015-10066/airworthiness-directives-the-boeing-company-airplanes>). This defect causes all six generator control units on the engine-mounted generators to enter fail-safe mode at the same time.

To avoid unplanned behavior (such as having your airplane fall from the sky), it's important to check for wraparound by using the limits from `<limits.h>`. Be careful when implementing these checks because it's easy to make mistakes. For example, the following code contains a defect, as `sum + ui` can never be larger than `UINT_MAX`:

```
extern unsigned int ui, sum;
// assign values to ui and sum
if (sum + ui > UINT_MAX)
    too_big();
else
    sum = sum + ui;
```

If the result of adding `sum` and `ui` is larger than `UINT_MAX`, it is reduced modulo `UINT_MAX + 1`. Therefore, the test is useless, and the generated code will unconditionally perform the summation. Quality compilers might issue a

warning pointing this out, but not all do so. To remedy this, you can subtract `sum` from both sides of the inequality to form the following effective test:

```
extern unsigned int ui, sum;
// assign values to ui and sum
if (ui > UINT_MAX - sum)
    too_big();
else
    sum = sum + ui;
```

The `UINT_MAX` macro is the largest representable unsigned `int` value, and `sum` is a value between 0 and `UINT_MAX`. If `sum` is equal to `UINT_MAX`, the result of the subtraction is 0, and if `sum` is equal to 0, the result of the subtraction is `UINT_MAX`. Because the result of this operation will always fall in the allowable range of 0 to `UINT_MAX`, it can never wrap around.

The same problem occurs when checking the result of an arithmetic operation against 0, the minimum unsigned value:

```
extern unsigned int i, j;
// assign values to i and j
if (i - j < 0) // cannot happen
    negative();
else
    i = i - j;
```

Because unsigned integer values can never be negative, the subtraction will be performed unconditionally. Quality compilers may warn about this mistake as well. Instead of this useless test, you can check for wraparound by testing whether `j` is greater than `i`:

```
if (j > i) // correct
    negative();
else
    i = i - j;
```

If `j > i`, the difference would wrap around, so the possibility of wraparound is prevented. By eliminating the subtraction from the test, the possibility of wraparound is also eliminated.

WARNING

Keep in mind that the width used during wraparound depends on the implementation, which means you can obtain different results on different platforms. Your code won't be portable if you fail to account for this.

Signed Integers

Each unsigned integer type (excluding `bool`) has a corresponding signed integer type that occupies the same amount of storage. Use signed integers to represent negative, zero, and positive values, the range of which depends on the number of bits allocated to the type and the representation.

Representation

Representing signed integer types is more complicated than representing unsigned integer types. Historically, the C language has supported three different schemes for representing negative numbers: sign and magnitude, one's complement, and two's complement.

Starting with C23, only two's-complement representation is supported. In two's complement representation, the sign bit is given the weight $-(2^{N-1})$, and the other value bits have the same weights as for unsigned. The remainder of this book assumes the two's complement representation.

Signed two's-complement integer types with a width of N can represent any integer value in the range of -2^{N-1} to $2^{N-1} - 1$. This means, for example, that an 8-bit value of type `signed char` has a range of -128 to 127 . Compared to other signed integer representations, two's complement can represent an additional *most negative* value. The most negative value for an 8-bit `signed char` is -128 , and its absolute value $|-128|$ cannot be represented as this type. This leads to some interesting edge cases, which we'll soon examine in more detail.

[Table 3-3](#) shows the constant expressions from `<limits.h>` for each signed type and the minimum magnitudes required by the standard. Your compiler will replace these values with implementation-defined magnitudes.

Table 3-3: Signed Integer Minimum Magnitudes

Constant expression	Minimum magnitude	Type
<code>SCHAR_MIN</code>	$-128 // -2^7$	signed char
<code>SCHAR_MAX</code>	$+127 // 2^7 - 1$	signed char
<code>SHRT_MIN</code>	$-32,768 // -2^{15}$	short int
<code>SHRT_MAX</code>	$+32,767 // 2^{15} - 1$	short int
<code>INT_MIN</code>	$-32,768 // -2^{15}$	int
<code>INT_MAX</code>	$+32,767 // 2^{15} - 1$	int

LONG_MIN	$-2,147,483,648 // -2^{31}$	long int
LONG_MAX	$+2,147,483,647 // 2^{31} - 1$	long int
LLONG_MIN	$-9,223,372,036,854,775,808 // -2^{63}$	long long int
LLONG_MAX	$+9,223,372,036,854,775,807 // 2^{63} - 1$	long long int

To negate a value in two's-complement representation, simply toggle each nonpadding bit and then add 1 (with carries as necessary), as shown in [Figure 3-1](#).




Figure 3-1: Negating an 8-bit value in two's-complement representation

[Table 3-4](#) shows the binary and decimal representations for an 8-bit two's-complement signed integer type with no padding (that is, $N = 8$).

Table 3-4: 8-Bit Two's-Complement Values

Binary	Decimal	Weighting	Constant
00000000	0	0	
00000001	1	2^0	
01111110	126	$2^6 + 2^5 + 2^4 + 2^3 + 2^2 + 2^1$	
01111111	127	$2^8 - 1 - 1$	SCHAR_MAX
10000000	-128	$-(2^8 - 1) + 0$	SCHAR_MIN
10000001	-127	$-(2^8 - 1) + 1$	
11111110	-2	$-(2^8 - 1) + 126$	
11111111	-1	$-(2^8 - 1) + 127$	

It is not necessary to know the binary representations of numbers, but as a C programmer, you will likely find it useful.

Integer Overflow

Integer overflow occurs when a signed integer operation results in a value that cannot be represented in the resulting type. Signed integer overflow and unsigned integer wraparound are often confused. The primary difference is that

signed integer overflow is undefined behavior, while unsigned integer wraparound is well-defined behavior. Unsigned integers cannot overflow.

Consider the following function-like macro that returns the absolute value of an arithmetic operand:

```
// undefined or wrong for the most-negative value
#define ABS(i) ((i) < 0 ? -(i) : (i))
```

We'll examine macros in detail in [Chapter 9](#). For now, think of function-like macros as functions that operate on generic types. On the surface, this macro appears to correctly implement the absolute value function by returning the nonnegative value of `i` without regard to its sign. We use the conditional (`? :`) operator (which I'll cover in more detail in the next chapter) to test whether the value of `i` is negative. If so, `i` is negated to `-(i)`; otherwise, it evaluates to the unmodified value `(i)`.

Because we've implemented `ABS` as a function-like macro, it can take an argument of any type. This macro can overflow when passed a signed integer argument of type `int` or a larger signed integer type. Of course, invoking this macro with an unsigned integer is pointless because unsigned integers can never be negative, so the macro's output would just reproduce the argument. Let's explore the behavior of the `ABS` macro when passed a signed integer argument:

```
signed int si = -25;
signed int abs_si = ABS(si);
printf("%d\n", abs_si); // prints 25
```

In this example, we pass an object of type `signed int` with the value `-25` as an argument to the `ABS` macro. This invocation expands to the following:

```
signed int si = -25;
signed int abs_si = ((si) < 0 ? -(si) : (si));
printf("%d\n", abs_si); // prints 25
```

The macro correctly returned the absolute value of 25. So far, so good. The problem is that the negative of the two's-complement most negative value for a given type cannot be represented in that type, so this use of the `ABS` macro results in signed integer overflow. Consequently, this implementation of `ABS` is defective and can do anything, including unexpectedly returning a negative value:

```
signed int si = INT_MIN;
signed int abs_si = ABS(si); // undefined behavior
printf("%d\n", abs_si);
```

What should `ABS(INT_MIN)` return to fix this behavior? Signed integer overflow is undefined behavior in C, allowing implementations to silently wrap around (the most common behavior), trap, or both (for example, some operations wrap around while other operations trap). *Traps* interrupt execution of the program so that no further operations are performed. Common architectures like x86 do a combination of both. Because the behavior is undefined, no universally correct solution to this problem exists, but we can at least test for the possibility of undefined behavior before it occurs and take appropriate action.

To make the absolute-value macro useful for a variety of types, we'll add a type-dependent `flag` argument to it. The flag represents the `*_MIN` macro, which matches the type of the first argument. This value is returned in the following problematic case:

```
#define ABSM(i, flag) ((i) >= 0 ? (i) : ((i)==(flag) ? (flag) : -(i)))
signed int si = -25; // try INT_MIN to trigger the undefined behavior
signed int abs_si = ABSM(si, INT_MIN);
if (abs_si == INT_MIN)
    overflow(); // handle special case
else
    printf("%d\n", abs_si); // prints 25
```

The `ABSM` macro tests for the most negative value and simply returns it if found instead of triggering the undefined behavior by negating it.

On some systems, the C standard library implements the following `int`-only absolute-value function to avoid overflow when the function is passed `INT_MIN` as an argument:

```
int abs(int i) {
    return (i >= 0) ? i : -(unsigned)i; // avoids overflow
}
```

In this case, `i` is converted to an unsigned `int` and negated. (I'll discuss conversions in more detail later in this chapter.)

Perhaps surprisingly, the unary minus (-) operator is defined for unsigned integer types. The resulting unsigned integer value is reduced modulo the number that is one greater than the largest value that the resulting type can represent. Finally, `i` is implicitly converted back to `signed int` as required by the return statement. Because `-INT_MIN` can't be represented as a `signed int`, the result is implementation defined, which is why this implementation is used only on *some* systems, and even on these systems, the `abs` function returns an incorrect value.

The `ABS` and `ABSM` function-like macros evaluate their parameters more than once, which can cause surprises when their arguments change the program state. These are called *side effects* (covered in detail in [Chapter 4](#)). Function calls, on the other hand, evaluate each argument only once.

Unsigned integers have well-defined wraparound behavior. Signed integer overflow, or the possibility of it, should always be considered a defect.

Bit-Precise Integer Types

As noted in [Chapter 2](#), bit-precise integer types accept an operand specifying the width of the integer, so a `_BitInt(32)` is a signed 32-bit integer and unsigned `_BitInt(32)` is an unsigned 32-bit integer. Bit-precise integer types can have *any* width up to `BITINT_MAXWIDTH`. Bit-precise integer types are useful in application domains, such as using 256-bit integer values in cryptographic symmetric ciphers like Advanced Encryption Standard (AES), calculating Secure Hash Algorithm (SHA)-256 hashes, representing a 24-bit color space, or describing the layout of network or serial protocols.

Bit-precise integer types are also very useful when programming field-programmable gate arrays (FPGAs). *FPGAs* are integrated circuits often sold off-the-shelf that provide customers with the ability to reconfigure the hardware to meet specific use-case requirements after the manufacturing process. In the case of FPGA hardware, using normal integer types for small value ranges where the full bit-width isn't used is extremely wasteful and creates severe performance and space concerns. At the other extreme, FPGAs can support wide integers, essentially providing arbitrary precision, and existing FPGA applications make use of large integers—for example, up to 2,031 bits. Prior to C23, programmers must pick an integer data type of the next larger size and manually perform mask and shifting operations. However, this is error prone because integer widths are implementation defined.

A bit-precise signed integer type is designated as `_BitInt(N)`, where *N* is

an integer constant expression that specifies the width of the type. Because bit-precise integer types are specified including the sign bit, a signed `_BitInt(1)` is invalid because it has one sign bit and no value bits. Unsigned bit-precise integer types do not include a sign bit, so the correct way to specify a 1-bit integer is `unsigned _BitInt(1)`.

The `_BitInt` types follow the usual C standard integer conversion ranks, as detailed in the “Integer Conversion Rank” section on [page 65](#). The usual arithmetic conversions also work the same, where the smaller ranked integer is converted to the larger. However, `_BitInt` types are excepted from integer promotions.

Overflow occurs when a value exceeds the allowable range of a given data type. For example, `(_BitInt(3))7 + (_BitInt(3))2` overflows, and the result is undefined as with other signed integer types. To avoid the overflow, the operation type can be widened to 4 bits by casting one of the operands to `_BitInt(4)`. Unsigned `_BitInt` wraparound is well-defined, and the value wraps around with two’s complement semantics.

To avoid overflow, you can cast one of the operands to a sufficient width to represent all possible values. For example, the following function casts one of the operands to 32 bits:

```
_BitInt(32) multiply(_BitInt(8) a8, _BitInt(24) a24) {  
    _BitInt(32) a32 = a8 * (_BitInt(32))a24;  
    return a32;  
}
```

This guarantees that the product can be represented.

Integer Constants

Integer constants (or *integer literals*) introduce integer values into a program. For example, you might use them in a declaration to initialize a counter to 0. C has four kinds of integer constants that use different number systems: decimal constants, binary constants, octal constants, and hexadecimal constants.

Decimal constants always begin with a nonzero digit. For example, the following code uses two decimal constants:

```
unsigned int ui = 71;  
int si;  
si = -12;
```

In this example code, we initialize `ui` to the decimal constant 71 and assign `si` the decimal constant value -12. (Formally, -12 is the negation operator [-] followed by an integer constant [12]. However, the expression -12 is usable as an integer constant expression and therefore effectively indistinguishable from an integer constant whose value is -12.) Use decimal constants when introducing regular integer values into your code.

If a constant starts with a 0, optionally followed by digits 0 through 7, it is an *octal constant*. Here's an example:

```
int agent = 007;  
int permissions = 0777;
```

In this example, 007 octal equals 7 decimal, and the octal constant 0777 equals the decimal value 511. Octal constants are convenient when dealing with 3-bit fields such as POSIX file permissions.

You can also create a *hexadecimal constant* by prefixing a sequence of decimal digits and the letters a (or A) through f (or F) with 0x or 0X. For example:

```
int burger = 0xDEADBEEF;
```

Use hexadecimal constants when the constant you are introducing is meant to represent a bit pattern more than a particular value—for example, when representing an address. Idiomatically, most hexadecimal constants are written like 0xDEADBEEF because it resembles a typical hex dump. It's probably a good idea for you to write all your hexadecimal constants like this.

Starting with C23, you can also specify a binary constant by appending a sequence of 1 and 0 decimal digits to 0b. For example:

```
int mask = 0b110011;
```

Binary constants can be more readable than octal or hexadecimal constants, especially when the value is used as a bitmask.

You can also append a suffix to your constant to specify its type. Without a suffix, a decimal constant is given the `int` type if it can be represented as a value in that type. If it can't be represented as an `int`, it will be represented as a `long int` or `long long int`. The `L` suffix specifies the `long` type, and `LL` specifies the

long long type. You can combine these suffixes with U for unsigned. For example, the ULL suffix specifies the unsigned long long type. Here are some examples:

```
unsigned int ui = 71U;  
signed long int sli = 9223372036854775807L;  
unsigned long long int ui = 18446744073709551615ULL;
```

These suffixes can be either uppercase or lowercase. Uppercase is generally preferred for readability, as a lowercase letter l might be mistaken for the number 1.

Bit-precise constants were added in C23 to specify `_BitInt` literals. The suffixes `wb` and `uwb` designate a constant of type `_BitInt(N)` and unsigned `_BitInt(N)`, respectively. The width N is the smallest N greater than 1 that can accommodate the value and the sign bit (when present).

The `wb` suffix results in a `_BitInt` that includes space for the sign bit even if the value of the constant is positive or was specified in binary, octal, or hexadecimal notation:

- 3wb** Yields a `_BitInt(3)` that is then negated; two value bits, one sign bit
- 0x3wb** Yields a `_BitInt(3)` that is then negated; two value bits, one sign bit
- 3wb** Yields a `_BitInt(3)`; two value bits, one sign bit
- 3ub** Yields an unsigned `_BitInt(2)`
- 3uwb** Yields an unsigned `_BitInt(2)` that is then negated, resulting in wraparound

If we don't use a suffix and the integer constant isn't of the required type, it may be implicitly converted. (We'll discuss implicit conversion in the "Arithmetic Conversion" section on [page 64](#).) This may result in a surprising conversion or a compiler diagnostic, so it's best to specify an integer constant of an appropriate type. Section 6.4.4.1 of the C standard contains more information on integer constants (ISO/IEC 2024).

Floating-Point Representation

Floating-point representation is the most common digital representation of real numbers. Floating-point representation is a technique that uses scientific notation

to represent numbers with a mantissa and an exponent for a given base. For example, the decimal number 123.456 can be represented as 1.23456×10^2 , while the binary number 0b10100.11 can be represented as 1.010011×2^4 .

The C standard defines a general floating-point model for floating-point numbers. However, it does not require all implementations to use the same representation schemes or formats, and it allows implementations to provide values that are not in the C model. To keep things simple, we'll assume conformance to Annex F. Annex F includes the most common floating-point formats specified in IEC 60559. You can test the value of the `__STDC_IEC_559__` macro, or of the `__STDC_IEC_60559_BFP__` macro in newer compilers, to determine whether the implementation conforms to Annex F.

This section explains floating types, arithmetic, values, and constants, so you will know how and when to use them to emulate math on real numbers and when to avoid them.

Floating Types and Encodings

C has three standard floating types: `float`, `double`, and `long double`.

The `float` type can be used for floating-point data and results that can be adequately represented with the precision and exponent range of the type. Using `float` arithmetic to compute `float` results from `float` data is particularly vulnerable to roundoff error. The common IEC 60559 `float` type encodes values using 1 sign bit, 8 exponent bits, and 23 significand bits. The value has a 24-bit significand, which is encoded in 23 bits (with help from the exponent field to determine the implicit leading bit).

The `double` type provides greater precision and exponent range but requires additional storage. Arithmetic with the `double` type greatly increases the reliability of computation of `float` results from `float` data. The IEC 60559 `double` type encodes values using 1 sign bit, 11 exponent bits, and 52 significand bits. The value has a 53-bit significand, which is encoded in 52 bits (with help from the exponent field to determine the implicit leading bit).

These encodings for `float` and `double` are illustrated in [Figure 3-2](#).




Figure 3-2: The `float` and `double` types

Let's illustrate with an example encoding in type `float`:

1	1000 0001	011 0000 0000 0000 0000 0000
---	-----------	------------------------------

The sign bit is 1, the exponent field is 1000 0001, and the significand field is 011 0000 0000 0000 0000 0000. The sign bit encodes the sign of the number, where 0 is used for a positive sign and 1 is used for a negative sign. Consequently, the number represented in this example has a negative sign.

Because the exponent field is neither all 0s nor all 1s, the bits in the significand field are interpreted as bits to the right of a binary point where an implicit 1 bit is to the left of the binary point. In this example, the significand of the encoded number is $1.011\ 0000\ 0000\ 0000\ 0000\ 0000 = 1 + 2^{-2} + 2^{-3} = 1.375$.

Putting this all together produces the following real number: $-2^2(1 + 2^{-2} + 2^{-3}) = -5$.

C Floating-Point Model

The following formula represents a number in the `float` type using the C model:

$$x_f = s2^e \sum_{k=1}^{24} f_k 2^{-k}$$

$$-125 \leq e \leq +128$$

The s is the sign, which may be 1 or -1 . The e is the exponent, and f_1 through f_{24} are the significand bits. Note that the exponent in the C model representation is 1 greater than the exponent we determined from the encoding because the C model places the (explicit) leading bit to the right of the binary point, while the encoding has the (implicit) leading bit to the left of the binary point.

The following formula represents a number in the double type:

$$x_d = s2^e \sum_{k=1}^{53} f_k 2^{-k}$$

$$-1021 \leq e \leq +1024$$

In general, the C model defines floating-point numbers in each floating type by the parameters b , p , $e_{\min}$, and $e_{\max}$. The parameter b is the radix (the base for the exponent and the significand digits). The radix b for all the standard floating types is represented by the FLT_RADIX macro defined in <float.h>. For Annex F, the value of FLT_RADIX is 2. The parameter p is the number of base- b digits in the floating-point significand. The $e_{\min}$ parameter is the minimum negative integer such that b raised to one less than that power is a normalized floating-point number. Finally, the $e_{\max}$ parameter is the maximum integer such that b raised to one less than that power is a representable finite floating-point number, provided that representable finite floating-point number is normalized (as it will be for all IEC 60559 types). [Table 3-5](#) shows the actual macro names.

Table 3-5: Standard Type Characterization Macros in <float.h>

Parameter	float	double	long double
p	FLT_MANT_DIG	DBL_MANT_DIG	LDBL_MANT_DIG
$e_{\min}$	FLT_MIN_EXP	DBL_MIN_EXP	LDBL_MIN_EXP
$e_{\max}$	FLT_MAX_EXP	DBL_MAX_EXP	LDBL_MAX_EXP

Each implementation assigns the long double type one of the following formats:

- IEC 60559 quadruple (or binary128) format (IEC 60559 added binary128 to its basic formats in the 2011 revision)
- IEC 60559 binary64-extended format
- A non-IEC 60559 extended format
- IEC 60559 double (or binary64) format

Recommended practice for compiler implementers is to match the `long double` type with the IEC 60559 binary128 format or an IEC 60559 binary64-extended format. IEC 60559 binary64-extended formats include the common 80-bit IEC 60559 format.

Arithmetic with the `long double` type should be considered for computations whose reliability might benefit from the maximum range and precision that the implementation provides for a standard floating type. However, the extra range and precision in the `long double` type (compared with `double`) varies considerably among implementations, as does the performance (speed) of `long double` arithmetic. Consequently, the `long double` type is unsuitable for data interchange or reproducible results (across implementations) or for portable high performance.

Larger types have greater precision but require more storage. Any value that can be represented as a `float` can also be represented as a `double`, and any value that can be represented as a `double` can be represented as a `long double`. The header `<float.h>` defines several macros that define the characteristics of floating types.

Annex H of C23 specifies additional floating types that have the arithmetic interchange and extended floating-point formats specified in IEC 60559. These include a sequence of types with unbounded precision and range and a 16-bit type. Future versions of C may include other floating types.

Floating-Point Arithmetic

Floating-point arithmetic is similar to, and used to model, the arithmetic of real numbers. However, there are differences to consider. Unlike in real number arithmetic, floating-point numbers are bounded in magnitude and have finite precision. Addition and multiplication operations are *not* associative; the distributive property *doesn't* hold, nor do many other properties of real numbers.

Floating types cannot represent all real numbers exactly, even when they can be represented in a small number of decimal digits. For example, common decimal constants such as 0.1 can't be represented exactly as binary floating-

point numbers. Floating types may lack the necessary precision for various applications such as loop counters or performing financial calculations. See CERT C rule FLP30-C (do not use floating-point variables as loop counters) for more information.

Floating-Point Values

A floating-point representation whose significand is 0 (all $f_k = 0$) represents a floating-point zero. Zeros are signed according to the sign (s), and there are two floating-point zero values: $+0$ and -0 . They are equal but behave differently in a few operations. A notable example is $1.0/0.0$ yields positive infinity and $1.0/(-0.0)$ yields negative infinity.

There are no leading zeros in the significand of a *normalized* floating-point number ($f_1 = 1$); leading zeros are removed by adjusting the exponent. These are *normal* numbers, and they use the full precision of the significand. Therefore, `float` has 24 significant bits of precision, `double` has 53 significant bits of precision, and `long double` has 113 significant bits of precision (assuming the IEC 60559 binary128 format).

Subnormal numbers are positive and negative numbers (but not 0) of very small magnitude whose normalized representation would result in an exponent that is less than the smallest exponent for the type. Their representations have exponent $e = e_{\min}$ and leading significand bit $f_1 = 0$. [Figure 3-3](#) is a number line showing the range of subnormal values around 0. The precision of subnormal numbers is less than that of normalized numbers.




Figure 3-3: The domain of subnormal numbers

Floating types can also represent values that are not floating-point numbers, such as negative and positive infinity and not-a-number (NaN) values. NaNs are values that do not represent a number.

Having infinity available as a specific value allows operations to continue past overflow and divide-by-zero situations and produce a useful result without requiring special treatment. Dividing any nonzero number by (positive or

negative) zero yields an infinity. Operations with infinite values are well-defined in the IEEE floating-point standard.

A *quiet NaN* propagates through almost every arithmetic operation without raising a floating-point exception and is typically tested after a selected sequence of operations. An arithmetic operation with a *signaling NaN* operand generally raises a floating-point exception immediately. Floating-point exceptions are an advanced topic not covered here. For more information, refer to Annex F of the C standard.

In C23, the `NAN` and `INFINITY` macros in `<float.h>` and the `nan` functions in `<math.h>` provide designations for IEC 60559 quiet NaNs and infinities. The `FLT_SNAN`, `DBL_SNAN`, and `LDBL_SNAN` macros in `<float.h>` provide designations for IEC 60559 signaling NaNs. C Annex F doesn't require full support for signaling NaNs.

You can identify the class of a floating-point value using the `fpclassify` function-like macro, which classifies its argument value as NaN, infinite, normal, subnormal, or zero:

```
#include <math.h>
int fpclassify(real-floating x);
```

In [Listing 3-2](#), we use the `fpclassify` macro in the `show_classification` function to determine whether a floating-point value of type `double` is a normal value, subnormal value, zero, infinity, or NaN.

```
const char *show_classification(double x) {
    switch(fpclassify(x)) {
        case FP_INFINITE: return "Inf";
        case FP_NAN:      return "NaN";
        case FP_NORMAL:   return "normal";
        case FP_SUBNORMAL: return "subnormal";
        case FP_ZERO:     return "zero";
        default:          return "unknown";
    }
}
```

Listing 3-2: The `fpclassify` macro

The function argument `x` (a `double` in this example) is passed to the `fpclassify` macro, which switches on the return value. The `show_classification` function returns a string corresponding to the class of value stored in `x`.

There are also a variety of other classification macros including `isinf`, `isnan`, `isnormal`, `issubnormal`, `iszero`, and so forth that may be more useful than the `fpclassify` macro in many applications.

Floating Constants

A *floating constant* is a decimal or hexadecimal number that represents a real number. You should use floating-point constants to represent floating-point values that cannot be changed. The following are some examples of floating-point constants:

```
15.75
1.575E1    /* 15.75 */
1575e-2    /* 15.75 */
25E-4      /* 0.0025 */
```

The following illustrates constants defined two ways: with a decimal floating constant and with a hexadecimal floating constant. The hexadecimal constants have values that can be represented exactly in their (binary) type. The decimal constants require conversion to binary and might be slightly affected by rounding direction modes and evaluation methods. (Rounding modes and evaluation methods are not covered in this book.) Hexadecimal constants should be used in such cases if you want a specific (to the last bit) value.

```
DBL_EPSILON 2.2204460492503131E-16 // decimal constant
DBL_EPSILON 0X1P-52                // hex constant
DBL_MIN 2.2250738585072014E-308    // decimal constant
DBL_MIN 0X1P-1022                  // hex constant
DBL_MAX 1.7976931348623157E+308    // decimal constant
DBL_MAX 0X1.fffffffffffffP1023     // hex constant
```

All floating-point constants have a type. The type is `double` if unsuffixed, `float` if suffixed by the letter `f` or `F`, or `long double` if suffixed by the letter `l` or `L`, as shown here:

```
10.0F /* type float */
10.0  /* type double */
10.0L /* type long double */
```

The decimal point is mandatory in these examples, but the trailing zero is

not.

Arithmetic Conversion

Frequently, a value represented in one type (for example, `float`) must be represented in a different type (for example, `int`). This might occur when you have an object of type `float` and need to pass it as an argument to a function that accepts an object of type `int`. When such conversions are necessary, you should always ensure that the value is adequately representable in the new type. I'll discuss this further in "Safe Conversions" on [page 70](#).

Values can be implicitly or explicitly converted from one arithmetic type to another. You can use the `cast` operator to perform *explicit* conversions. [Listing 3-3](#) shows two examples of casts.

```
int si = 5;
short ss = 8;
long sl = (long)si;
unsigned short us = (unsigned short)(ss + sl);
```

Listing 3-3: Cast operators

To perform a cast, place a type name in parentheses just before the expression. The cast converts the expression to the unqualified version of the type name in parentheses. Here, we cast the value of `si` to the type `long`. Because `si` is of type `int`, this cast is guaranteed to be safe because the value can always be represented in a larger integer type of the same signedness.

The second cast in this example casts the result of the expression `(ss + sl)` to type `unsigned short`. Because the value is converted to an unsigned type (`unsigned short`) with less precision, the result of the conversion might not be equal to the original value. (Some compilers might warn about this; others won't.) In this example, the result of the expression (13) can be correctly represented in the resulting type.

Implicit conversion, also known as *coercion*, occurs automatically in expressions as required. Values are coerced, for example, when operations are performed on mixed types. In [Listing 3-3](#), implicit conversions are used to convert `ss` to the type of `sl` so that the addition `ss + sl` can be performed on a common type. The rules concerning which values are implicitly converted to which types are somewhat complicated and involve three concepts: integer conversion rank, integer promotions, and the usual arithmetic conversions.

Integer Conversion Rank

Integer conversion rank is a standard rank ordering of integer types used to determine a common type for computations. Every integer type has an integer conversion rank that determines when and how conversions are implicitly performed.

The C standard, section 6.3.1.1, paragraph 1 (ISO/IEC 9899:2024), states that every integer type has an integer conversion rank where the following applies:

- No two signed integer types have the same rank, even if they have the same representation.
- The rank of a signed integer type is greater than the rank of any signed integer type with less precision.
- The rank of `long long int` is greater than the rank of `long int`, which is greater than the rank of `int`, which is greater than the rank of `short int`, which is greater than the rank of `signed char`.
- The rank of a bit-precise signed integer type is greater than the rank of any standard integer type with less width or any bit-precise integer type with less width.
- The rank of any unsigned integer type equals the rank of the corresponding signed integer type, if any.
- The rank of any standard integer type is greater than the rank of any extended integer type with the same width or bit-precise integer type with the same width.
- The rank of any bit-precise integer type relative to an extended integer type of the same width is implementation defined.
- The rank of `char` equals the rank of `signed char` and `unsigned char`.
- The rank of `bool` is less than the rank of all other standard integer types.
- The rank of any enumerated type equals the rank of the compatible integer type. Each enumerated type is compatible with `char`, a signed integer type, or an unsigned integer type.

The rank of any extended signed integer type relative to another extended signed integer type with the same precision is implementation defined but still subject to the other rules for determining the integer conversion rank.

Integer Promotions

A *small type* is an integer with a lower conversion rank than `int` or `unsigned int`. *Integer promotion* is the process of converting values of small types to an `int` or `unsigned int`. Integer promotions allow you to use an expression of a small type in any expression where an `int` or `unsigned int` may be used. For example, you could use a lower-ranked integer type—typically, `char` or `short`—on the right-hand side of an assignment or as an argument to a function.

Integer promotions serve two primary purposes. First, they encourage operations to be performed in a natural size (`int`) for the architecture, which improves performance. Second, they help avoid arithmetic errors from the overflow of intermediate values, for example:

```
signed char cresult, c1, c2, c3;  
c1 = 100; c2 = 3; c3 = 4;  
cresult = c1 * c2 / c3;
```

Without integer promotion, `c1 * c2` would result in an overflow of the `signed char` type on platforms where `signed char` is represented by an 8-bit two's-complement value. This is because 300 is outside the range of values (–128 to 127) that can be represented by an object of this type. However, because of integer promotion, `c1`, `c2`, and `c3` are implicitly converted to objects of type `signed int`, and the multiplication and division operations take place in this size. There is no possibility of overflow while performing these operations because the resulting values can be represented by this wider type. In this specific example, the result of the entire expression is 75, which is within range of the `signed char` type, so the value is preserved when stored in `cresult`.

Prior to the first C standard, compilers used one of two approaches for integer promotions: the unsigned-preserving approach or the value-preserving approach. In the *unsigned-preserving approach*, the compiler promotes small, unsigned types to `unsigned int`. In the *value-preserving approach*, if all values of the original type can be represented as an `int`, the value of the original small type is converted to `int`. Otherwise, it's converted to `unsigned int`. When developing the original version of the standard (C89), the C standards committee decided on value-preserving rules, because they produce incorrect results less often than the unsigned-preserving approach. If necessary, you can override this behavior by using explicit type casts, as in [Listing 3-3](#).

The result of promoting small unsigned types depends on the precision of

the integer types, which is implementation defined. For example, the x86-32 and x86-64 architectures have an 8-bit `char` type, a 16-bit `short` type, and a 32-bit `int` type. For implementations that target one of these architectures, values of both `unsigned char` and `unsigned short` are promoted to `signed int`, because all the values that can be represented in these smaller types can be represented as a `signed int`. However, 16-bit architectures, such as Intel 8086/8088 and the IBM Series/1, have an 8-bit `char` type, a 16-bit `short` type, and a 16-bit `int` type. For implementations that target these architectures, values of type `unsigned char` are promoted to `signed int`, while values of type `unsigned short` are promoted to `unsigned int`. This is because all the values that can be represented as an 8-bit `unsigned char` type can be represented as a 16-bit `signed int`, but some values that can be represented as a 16-bit `unsigned short` cannot be represented as a 16-bit `signed int`.

The `_BitInt` types are exempt from integer promotions. Integer promotions might inflate the size of required hardware on some platforms, so `_BitInt` types aren't subject to the integer promotion rules. For example, in a binary expression involving a `_BitInt(12)` and an `unsigned _BitInt(3)`, the usual arithmetic conversions would not promote either operand to an `int` before determining the common type. Because one type is signed and one is unsigned and because the signed type has greater rank than the unsigned type (due to the bit-widths of the types), the `unsigned _BitInt(3)` will be converted to `_BitInt(12)` as the common type.

Usual Arithmetic Conversions

The *usual arithmetic conversions* are rules for yielding a *common real type* for the operands and result of an arithmetic operation. Ignoring complex or imaginary types, each operand is converted to the common real type. Many operators that accept integer operands (including `*`, `/`, `%`, `+`, `-`, `<`, `>`, `<=`, `>=`, `==`, `!=`, `&`, `^`, `|`, and `?:`) perform conversions using the usual arithmetic conversions. The usual arithmetic conversions are applied to the promoted operands.

The usual arithmetic conversions first check whether one of the operands in the balancing conversion is a floating type. If so, it applies the following rules:

1. If one type of either operand is `long double`, the other operand is converted to `long double`.
2. Otherwise, if one type of either operand is `double`, the other operand is converted to `double`.

3. Otherwise, if the type of either operand is `float`, the other operand is converted to `float`.
4. Otherwise, the integer promotions are performed on both operands.

If one operand has the type `double` and the other operand has the type `int`, for example, the operand of type `int` is converted to an object of type `double`. If one operand has the type `float` and the other operand has the type `double`, the operand of type `float` is converted to an object of type `double`. Particularly notable is the case of `int` and `float`, which converts the `int` operand to `float`, although `int` typically has greater precision than `float`.

If neither operand is a floating type, the following usual arithmetic conversion rules are applied to the promoted integer operands:

1. If both operands have the same type, no further conversion is needed.
2. Otherwise, if both operands have signed integer types or both have unsigned integer types, the operand with the type that has the lesser integer conversion rank is converted to the type of the operand with greater rank. If one operand has the type `int` and the other operand has the type `long`, for example, the operand of type `int` is converted to an object of type `long`.
3. Otherwise, if the operand that has the unsigned integer type has a rank greater than or equal to the rank of the other operand's type, then the operand with the signed integer type is converted to the type of the operand with the unsigned integer type. For example, if one operand has the type `signed int` and the other operand has the type `unsigned int`, the operand of type `signed int` is converted to an object of type `unsigned int`.
4. Otherwise, if the type of the operand with the signed integer type can represent all the values of the type of the operand with unsigned integer type, then the operand with unsigned integer type is converted to the type of the operand with signed integer type. For example, if one operand has the type `unsigned int` and the other operand has the type `signed long long` and the `signed long long` type can represent all the values of the `unsigned int` type, then the operand of type `unsigned int` is converted to an object of type `signed long long`. This is the case for implementations with a 32-bit `int` type and a 64-bit `long long` type, such as x86-32 and x86-64.
5. Otherwise, both operands are converted to the unsigned integer type corresponding to the type of the operand with signed integer type.

One consequence of `_BitInt` being exempt from the integer promotion rules

is that a `_BitInt` operand of a binary operator is not always promoted to an `int` or unsigned `int` as part of the usual arithmetic conversions. Instead, a lower-ranked operand is converted to the higher-rank operand type, and the result of the operation is the higher-ranked type. For example, given the following declarations

```
_BitInt(2) a2 = 1;  
_BitInt(3) a3 = 2;  
_BitInt(33) a33 = 1;  
signed char c = 3;
```

the `a2` operand in the following expression is converted to `_BitInt(3)` as part of the multiplication, and the resulting type is `_BitInt(3)`:

```
a2 * a3;
```

As part of the following multiplication, `c` is promoted to `int`, `a2` is converted to `int`, and the resulting type is `int`:

```
a2 * c;
```

Finally, as part of the following multiplication, `c` is promoted to `int`. Then, provided the width of `int` is not greater than 32, it is converted to `_BitInt(33)` and the resulting type is `_BitInt(33)`:

```
a33 * c;
```

These conversion rules, which evolved as new types were added, take some getting used to. The irregularities in these patterns resulted from varying architectural properties (notably, the PDP-11's automatic promotion of `char` to `int`) coupled with a desire to avoid changing the behavior of existing programs and (subject to those constraints) a desire for uniformity. When in doubt, use type casts to explicitly force the conversions that you intend. That said, try not to overuse explicit conversions because casts can disable important diagnostics.

An Example of Implicit Conversion

The following example illustrates the use of integer conversion rank, integer promotions, and the usual arithmetic conversions. This code compares the

signed char value `c` for equality with the unsigned int value `ui`. We'll assume this code is being compiled for the x86 architecture:

```
unsigned int ui = UINT_MAX;
signed char c = -1;
if (c == ui) {
    printf("%d equals %u\n", c, ui);
}
```

The variable `c` is of type `signed char`. Because `signed char` has a lower integer conversion rank than `int` or `unsigned int`, the value stored in `c` is promoted to an object of type `signed int` when used in the comparison. This is accomplished by sign-extending the original value of `0xFF` to `0xFFFFFFFF`. *Sign extension* is used to convert a signed value to a larger-width object. The sign bit is copied into each bit position of the expanded object. This operation preserves the sign and magnitude when converting a value from a smaller to a larger signed integer type.

Next, the usual arithmetic conversions are applied. Because the operands to the equal (`==`) operator have different signedness and equal rank, the operand with the signed integer type is converted to the type of the operand with the unsigned integer type. The comparison is then performed as a 32-bit unsigned operation. Because `UINT_MAX` has the same values as the promoted and converted value of `c`, the comparison yields 1, and the code snippet prints the following:

```
-1 equals 4294967295
```

This result should no longer be surprising.

Safe Conversions

Both implicit and explicit conversions (the result of a cast operation) can produce values that can't be represented in the resulting type. It's preferable to perform operations on objects of the same type to avoid conversions. However, conversions are unavoidable when a function returns or accepts an object of a different type. In those cases, we must ensure that the conversion is performed correctly.

Integer Conversions

Integer conversions occur when a value of an integer type is converted to a

different integer type. Conversions to larger types of the same signedness are always safe and don't need to be checked. Most other conversions can produce unexpected results if the resulting value cannot be represented in the resulting type. To perform these conversions correctly, you must test that the value stored in the original integer type is within the range of values that can be represented in the resulting integer type. As an example, the `do_stuff` function shown in [Listing 3-4](#) accepts a signed `long` argument value that needs to be used in a context in which only a signed `char` is appropriate.

```
#include <errno.h>
#include <limits.h>

errno_t do_stuff(signed long value) {
    if ((value < SCHAR_MIN) || (value > SCHAR_MAX)) {
        return ERANGE;
    }
    signed char sc = (signed char)value; // cast quiets warning
    // --snip--
}
```

Listing 3-4: Safe conversion

To perform this conversion safely, the function checks that `value` can be represented as a signed `char` in the range `[SCHAR_MIN, SCHAR_MAX]` and returns an error if it cannot.

The specific range tests vary based on the conversion. See CERT C rule INT31-C (“Ensure that integer conversions do not result in lost or misinterpreted data”) for more information.

Integer-Type to Floating-Type Conversions

Floating types that conform to Annex F support positive and negative infinities, so all integer values are in range. The usual IEC 60559 conversion rules apply. See CERT C rule FLP36-C (“Preserve precision when converting integral values to floating type”) for more information.

Floating-Type to Integer-Type Conversions

When a finite value of a floating type is converted to an integer type (other than `bool`), the fractional part is discarded. If the value of the integral part cannot be represented by the integer type, Annex F specifies that the “invalid” floating-point exception is raised, and the result is unspecified.

Floating-Type Demotions

Converting a floating-point value to a larger floating type is always safe. Demoting a floating-point value (that is, converting to a smaller floating type) is like converting an integer value to a floating type. Floating types that conform to Annex F support positive and negative infinities. Demoting values of floating types for these implementations will always succeed because any out-of-range values are converted to infinities. See CERT C rule FLP34-C (“Ensure that floating-point conversions are within range of the new type”) for more information on floating-point conversions.

EXERCISES

1. Write a `main` function that calls the `show_classification` function from [Listing 3-2](#) with different values of type `double`. Try to exercise all the `switch` cases.
2. Find an example where a conversion from `int` to `float` results in a loss of precision.

Summary

In this chapter, you learned about integers and floating types. You also learned about implicit and explicit conversions, integer conversion rank, integer promotions, and the usual arithmetic conversions.

The use of these basic types, particularly integers, is unavoidable and ubiquitous in C programming. Even the “Hello, world!” program returns an `int` and prints a string—an array of type `char`—which, of course, is an integer type. Because integer types are used so often, you can’t simply reread this chapter each time you need to use them. You must understand their behavior so you can program effectively.

In the next chapter, you’ll learn about operators and how to write simple expressions to perform operations on these arithmetic types as well as other object types.

4

EXPRESSIONS AND OPERATORS



In this chapter, you'll learn about operators and how to write simple expressions to perform operations on various object types. An *operator* is a keyword or one or more punctuation characters used to perform an operation. When an operator is applied to one or more operands, it becomes an expression that computes a value and that might have side effects. *Expressions* are sequences of operators and operands that compute a value or accomplish another purpose. The operands can be identifiers, constants, string literals, and other expressions.

In this chapter, we discuss simple assignment before stepping back to examine the mechanics of expressions (operators and operands, value computations, side effects, precedence, and order of evaluation). We then consider specific operators including `sizeof`, arithmetic, bitwise, cast, conditional, alignment, relational, compound assignment, and the comma operator. We've introduced many of these operators and expressions in previous chapters; here, we detail their behavior and how best to use them. Finally, we end the chapter with a discussion of pointer arithmetic.

Simple Assignment

A *simple assignment* replaces the value stored in the object designated by the left operand with the right operand. The value of the right operand is converted to the type of the assignment expression. Simple assignment has three components: the left operand, the assignment (=) operator, and the right operand, as shown in the following example:

```
int i = 21; // declaration with initializer
int j = 7;  // declaration with initializer
i = j;      // simple assignment
```

The first two lines are *declarations* that define and initialize *i* with the value 21 and *j* with the value 7. Initialization is different from simple assignment despite having similar syntax. An *initializer* is an optional part of a declaration; when present, it provides the initial value for the object. If the initializer is not present, objects (with automatic storage duration) are uninitialized.

The third line contains a simple assignment. You must define or declare all identifiers that appear in an expression (such as a simple assignment) for your code to compile.

The left operand in a simple assignment is always an expression (with an object type other than `void`), referred to as an *lvalue*. The *l* in lvalue originally comes from it being the *left* operand, but it may be more correct to think of it as standing for *locator value*, because it must designate an object. The right operand is also an expression, but it can simply be a value and doesn't need to identify an object. We refer to this value as an *rvalue* (*right* operand) or *expression value*. In this example, the identifiers for both objects *i* and *j* are lvalues. An lvalue can also be an expression, such as `*(p + 4)`, provided it references an object in memory.

In a simple assignment, the rvalue is converted to the type of the lvalue and then stored in the object designated by the lvalue. In the assignment `i = j`, the value is read from *j* and written to *i*. Because both *i* and *j* are the same type (`int`), no conversion is necessary. The assignment expression has the value of the result of the assignment and the type of the lvalue.

The rvalue does not need to refer to an object, as you can see in the following statement, which uses the types and values from the preceding example:

```
j = i + 12; // j now has the value 19
```

The expression `i + 12` is not an lvalue, because there is no underlying object storing the result. Instead, `i` by itself is an lvalue that is automatically converted into an rvalue to be used as an operand to the addition operation. The resulting value from the addition operation (which has no memory location associated with it) is also an rvalue. C constrains where lvalues and rvalues may appear. The following statements illustrate the correct and incorrect use of lvalues and rvalues:

```
int i;  
i = 5;      // i is an lvalue, 5 is an rvalue  
int j = i;  // lvalues can appear on the right side of an assignment  
7 = i;      // error: rvalues can't appear on the left side of an assignment
```

The assignment `7 = i` doesn't compile because the rvalue must always appear on the right side of the operator.

In the following example, the right operand has a different type from the assignment expression, so the value of `i` is first converted to a signed `char` type. The value of the expression enclosed in parentheses is then converted to the `long int` type and assigned to `k`:

```
signed char c;  
int i = INT_MAX;  
long k;  
k = (c = i);
```

Assignment must deal with real-world constraints. Specifically, simple assignment can result in truncation if a value is converted to a narrower type. As mentioned in [Chapter 3](#), each object requires a fixed number of bytes of storage. The value of `i` can always be represented by `k` (a larger type of the same signedness). However, in this example, the value of `i` is converted to signed `char` (the type of the assignment expression `c = i`). The value of the expression enclosed in parentheses is then converted to the type of the outer assignment expression—that is, `long int` type. If your implementation's signed `char` type has insufficient width to fully represent the value stored in `i`, values greater than `SCHAR_MAX` are truncated, and the value stored in `k` (`-1`) is truncated. To prevent values from being truncated, make sure that you choose sufficiently wide types that can represent any value that might arise.

Evaluations

Now that we've looked at simple assignment, let's step back for a moment and look at how expressions are evaluated. *Evaluation* mostly means simplifying an expression down to a single value. The evaluation of an expression can include both value computations and the initiation of side effects.

A *value computation* is the calculation of the value that results from the evaluation of the expression. Computing the final value may involve determining the identity of the object or reading the value previously assigned to an object. For example, the following expression contains several value computations to determine the identity of `i`, `a`, and `a[i]`:

```
a[i] + f() + 9
```

Because `f` is a function and not an object, the expression `f()` doesn't involve determining the identity of `f`. The value computations of operands must occur before the value computation of the result of the operator. In this example, separate value computations read the value of `a[i]` and determine the value returned by the call to the `f` function. A third computation then sums these values to obtain the value returned by the overall expression. If `a[i]` is an array of `int` and `f()` returns an `int`, the result of the expression will have the `int` type.

Side effects are changes to the state of the execution environment. Side effects include writing to an object, accessing (reading or writing) a *volatile*-qualified object, input/output (I/O), assignment, or calling a function that does any of these things. We can slightly modify the previous example to add an assignment. Updating the stored value of `j` is a side effect of the assignment:

```
int j;  
j = a[i] + f() + 9;
```

The assignment to `j` is a side effect that changes the state of the execution environment. Depending on the definition of the `f` function, the call to `f` may also have side effects.

Function Invocation

A *function designator* is an expression that has a function type and is used to invoke a function. In the following function invocation, `max` is the function

designator:

```
int x = 11;
int y = 21;
int max_of_x_and_y = max(x, y);
```

The `max` function returns the larger of its two arguments. In an expression, a function designator is converted to a *pointer-to-function returning type* at compile time. The value of each argument must be of a type that can be assigned to an object with (the unqualified version of) the type of its corresponding parameter. The number and type of each argument must agree with the number and type of each parameter accepted by the function. Here, that means two integer arguments. C also supports *variadic functions*, such as `printf`, which accept a variable number of arguments.

We can also pass one function to another, as shown by [Listing 4-1](#).

```
int f() {
    // --snip--
    return 0;
}
void g(int (*func)()) {
    // --snip--
    if (func() != 0)
        printf("g failed\n");
    // --snip--
}
// --snip--
g(f); // call g with function-pointer argument
// --snip--
```

Listing 4-1: Passing one function to another function

This code passes the address of a function designated by `f` to another function, `g`. The function `g` accepts a pointer to a function that accepts no arguments and returns `int`. A function passed as an argument is implicitly converted to a function pointer. The definition of `g` makes this explicit; an equivalent declaration is `void g(int func(void))`.

Increment and Decrement Operators

The *increment* (`++`) and *decrement* (`--`) operators increment and decrement a

modifiable lvalue, respectively. Both are *unary operators* because they take a single operand.

These operators can be used as either *prefix operators*, which appear before the operand, or *postfix operators*, which come after the operand. The prefix and postfix operators have different behaviors, which means they are commonly used as trick questions in quizzes and interviews. A prefix increment performs the increment before returning the value, whereas a postfix increment returns the value and then performs the increment. [Listing 4-2](#) illustrates these behaviors by performing a prefix or postfix increment or decrement operation and then assigning the result to `e`.

```
int i = 5;
int e;    // expression result
e = i++;  // postfix increment: e ← 5, i ← 6
e = i--;  // postfix decrement: e ← 6, i ← 5
e = ++i;  // prefix increment: e ← 6, i ← 6
e = --i;  // prefix decrement: e ← 5, i ← 5
```

Listing 4-2: The prefix and postfix increment and decrement operators

The `i++` operation in this example returns the unchanged value 5, which is then assigned to `e`. The value of `i` is then incremented as a side effect of the operation.

The *prefix* increment operator increments the value of the operand, and the expression returns the new value of the operand after it has been incremented. Consequently, the expression `++i` is equivalent to `i = i + 1`, except that `i` is evaluated only once. The `++i` operation in this example returns the incremented value 6, which is then assigned to `e`.

Operator Precedence and Associativity

In mathematics and computer programming, the *order of operations* (or *operator precedence*) is a collection of rules that dictates the order in which operations are performed during the evaluation of an expression. For example, multiplication is granted a higher precedence than addition. Therefore, the expression $2 + 3 \times 4$ is interpreted to have the value $2 + (3 \times 4) = 14$, not $(2 + 3) \times 4 = 20$.

Associativity determines how operators of the same precedence are grouped when no parentheses are used. C associativity differs from mathematics associativity. For example, while floating-point addition and multiplication are

both commutative ($a + b = b + a$ and $a \times b = b \times a$), they are not necessarily associative. If adjacent operators have equal precedence, the choice of which operation to apply first is determined by the associativity. *Left-associative* operators cause the operations to be grouped from the left, while *right-associative* operators cause the operations to be grouped from the right. You can think of grouping as the implicit introduction of parentheses. For example, the addition (+) operator has left associativity, so the expression $a + b + c$ is interpreted as $((a + b) + c)$. The assignment operator is right-associative, so the expression $a = b = c$ is interpreted as $(a = (b = c))$.

[Table 4-1](#), derived from the C Operator Precedence table at the C++ References website (https://en.cppreference.com/w/c/language/operator_precedence), lists the precedence and associativity of C operators, as specified by the language syntax. Operators are listed in order of descending precedence (that is, higher rows have higher precedence).

Table 4-1: Operator Precedence and Associativity

Precedence	Operator	Description	Associativity
0	(. . .)	Forced grouping	Left
1	++ -- () [] . -> (type){list}	Postfix increment and decrement Function call Array subscripting Structure and union member access Structure and union member access through pointer Compound literal	Left
2	++ -- + - ! ~ (type) * & sizeof _Alignof	Prefix increment and decrement Unary plus and minus Logical NOT and bitwise NOT Type cast Indirection (dereference) Address-of Size of Alignment requirement	Right
3	* / %	Multiplication, division, and remainder	Left
4	+ -	Addition and subtraction	
5	<< >>	Bitwise left shift and right shift	
6	< <= > >=	Relational operators < and ≤ Relational operators > and ≥	
7	== !=	Equal to and not equal to	
8	&	Bitwise AND	
9	^	Bitwise XOR (exclusive or)	

10		Bitwise OR (inclusive or)	
11	&&	Logical AND	
12		Logical OR	
13	? :	Conditional operator	Right
14	=	Simple assignment	
	+= -=	Assignment by sum and difference	
	*= /= %=	Assignment by product, quotient, and remainder	
	<<= >>=	Assignment by bitwise left shift and right shift	
	&= ^= =	Assignment by bitwise AND, XOR, and OR	
15	,	Expression sequencing	Left

Sometimes operator precedence can be intuitive, and sometimes it can be misleading. For example, the postfix ++ and -- operators have higher precedence than both the prefix ++ and -- operators, which in turn have the same precedence as the unary * operator. For example, if p is a pointer, then *p++ is equivalent to *(p++) and ++*p is equivalent to ++(*p), because both the prefix ++ operator and the unary * operator are right-associative. If two operators have the same precedence and associativity, they are evaluated from left to right. [Listing 4-3](#) illustrates the precedence rules among these operators.

```
char cba[] = "cba";
char *p = cba;
printf("%c", ++*p);

char xyz[] = "xyz";
char *q = xyz;
printf("%c", *q++);
```

Listing 4-3: The operator precedence rules

The pointer in the expression ++*p is first dereferenced, producing the c character. This value is then incremented, resulting in the character d. In this case, the prefix ++ operator operates on the object of type char pointed to by p and not the pointer. On the other hand, the pointer in the expression *q++ is incremented first, so it refers to the y character. However, the result of *postfix* increment operators is the value of the operand so that the original pointer value is dereferenced, producing the x character. Consequently, this code prints out the characters dx. You can use parentheses to change or clarify the order of operations.

Order of Evaluation

The *order of evaluation* of the operands of any C operator, including the order of evaluation of any subexpressions, is generally unspecified. The compiler will evaluate them in any order and may choose a different order when the same expression is evaluated again. This latitude allows the compiler to produce faster code by choosing the most efficient order. The order of evaluation is constrained by operator precedence and associativity.

[Listing 4-4](#) demonstrates the order of evaluation for function arguments.

```
int glob; // static storage initialized to 0

int f(void) {
    return glob + 10;
}
int g(void) {
    glob = 42;
    return glob;
}
int main(void) {
    int max_value = max(f(), g());
    // --snip--
}
```

Listing 4-4: The order of evaluation for function arguments

Both functions `f` and `g` access the global variable `glob`, meaning they rely on a shared state. Because the order of evaluation of functions `f` and `g` is unspecified, the arguments passed to `max` may differ between compilations. If `f` is called first, it will return 10, but if it's called last, it will return 52. Function `g` always returns 42 regardless of the order of evaluation. Consequently, the `max` function (which returns the greater of the two values) may return either 42 or 52, depending on the order of evaluation of its arguments. The only *sequencing guarantees* provided by this code are that both `f` and `g` are called before `max` and that the executions of `f` and `g` do not interleave.

We can rewrite this code as follows to ensure it always behaves in a predictable, portable manner:

```
int f_val = f();
int g_val = g();
int max_value = max(f_val, g_val);
```

In this revised program, `f` is called to initialize the `f_val` variable. This is guaranteed to be sequenced before the execution of `g`, which is called in the subsequent declaration to initialize the variable `g_val`. If one evaluation is *sequenced before* another evaluation, the first evaluation must complete before the second evaluation can begin. You can use sequence points (discussed in the following subsections) to guarantee that an object will be written before it is read. The execution of `f` is guaranteed to be sequenced before the execution of `g` because a sequence point exists between the evaluation of one full expression and the next full expression.

Unsequenced and Indeterminately Sequenced Evaluations

The executions of unsequenced evaluations can *interleave*, meaning that the instructions can be executed in any order, provided that reads and writes are performed in the order specified by the program. A program that performs reads and writes in the order specified by the program is *sequentially consistent* (Lamport 1979).

Some evaluations are *indeterminately sequenced*, which means they cannot interleave but can still be executed in any order. For example, the following statement contains several value computations and side effects:

```
printf("%d\n", ++i + ++j * --k);
```

The values of `i`, `j`, and `k` must be read before their values can be incremented or decremented. This means that the reading of `i` must be sequenced before the increment side effect, for example. Similarly, all side effects for the operands of the multiplication operation need to complete before the multiplication can occur. The multiplication must complete before the addition because of operator precedence rules. Finally, all side effects on the operands of the addition operation must complete before the addition can occur. These constraints produce a partial ordering among these operations because they don't require that `j` be incremented before `k` is decremented, for example. Unsequenced evaluations in this expression can be performed in any order, which allows the compiler to both reorder operations and to cache values in registers, allowing for faster overall execution. Function executions, on the other hand, are indeterminately sequenced and do not interleave with each other.

Sequence Points

A *sequence point* is the juncture at which all side effects will have completed. These are implicitly defined by the language, but you can control when they occur by how you code.

The sequence points are enumerated in Annex C of the C Standard. A sequence point occurs between the evaluation of one *full expression* (an expression that is not part of another expression or declarator) and the next full expression to be evaluated. A sequence point also occurs upon entering or exiting a called function.

If a side effect is unsequenced relative to either a different side effect on the same scalar or a value computation that uses the value of the same scalar object, the code has undefined behavior. A *scalar type* is either an arithmetic type or pointer type. The expression `i++ * i++` performs two unsequenced operations on `i` as the following code snippet shows:

```
int i = 5;
printf("Result = %d\n", i++ * i++);
```

You might think this code will produce the value 30, but because it has undefined behavior, that outcome isn't guaranteed.

Conservatively, we can ensure that side effects have completed before the value is read by placing every side-effecting operation in its own full expression. We can rewrite that code as follows to eliminate the undefined behavior:

```
int i = 5;
int j = i++;
int k = i++;
printf("Result = %d\n", j * k);
```

This example now contains a sequence point between every side-effecting operation. However, it's impossible to tell whether this rewritten code represents the programmer's original intent because the original code had no defined meaning. If you choose to omit sequence points, you must be sure you completely understand the sequencing of side effects. We also can write this same code as follows without changing the behavior:

```
int i = 5;
int j = i++;
printf("Result = %d\n", j * i++);
```

Now that we have described the mechanics of expressions, we'll return to discussing specific operators.

sizeof Operator

We can use the `sizeof` operator to find the size in bytes of its operand; specifically, it returns an unsigned integer of `size_t` type that represents the size. Knowing the correct size of an operand is necessary for most memory operations, including allocating and copying storage. The `size_t` type is defined in `<stddef.h>` as well as in other header files. We need to include one of these header files to compile any code that references the `size_t` type.

We can pass the `sizeof` operator an unevaluated expression of a complete object type or a parenthesized name of such a type:

```
int i;
size_t i_size = sizeof i;      // the size of the object i
size_t int_size = sizeof(int); // the size of the type int
```

It's always safe to parenthesize the operand to `sizeof`, because parenthesizing an expression doesn't change the way the size of the operand is calculated. The result of invoking the `sizeof` operator is a constant expression unless the operand is a variable-length array. The operand to `sizeof` is not evaluated.

If you need to determine the number of bits of storage available, you can multiply the size of an object by `CHAR_BIT`, which gives the number of bits contained in a byte. For example, the expression `CHAR_BIT * sizeof(int)` will produce the number of bits in an object of type `int`.

Arithmetic Operators

Operators that perform arithmetic operations on arithmetic types are detailed in the following sections. We can also use some of these operators with nonarithmetic operands.

Unary + and –

The *unary +* and *– operators* operate on a single operand of arithmetic type. The *– operator* returns the negative of its operand (that is, it behaves as though the operand were multiplied by -1). The *unary + operator* just returns the value.

These operators exist primarily to express positive and negative numbers.

If the operand has a small integer type, it's promoted (see [Chapter 3](#)), and the result of the operation has the result of the promoted type. As a point of trivia, C has no negative integer literals. A value such as `-25` is actually an rvalue of type `int` with the value 25 preceded by the unary `-` operator. However, the expression `-25` is guaranteed to be a constant integer expression.

Logical Negation

The result of the unary logical negation operator (`!`) is as follows:

- 0 if the evaluated value of its operand is not 0
- 1 if the evaluated value of its operand is 0

The operand is a scalar type. The result has type `int` for historical reasons. The expression `!E` is equivalent to `(0 == E)`. The logical negation operator is frequently used to check for null pointers; for example, `!p` is equivalent to `(nullptr == p)`. Null pointers may not hold the value zero but are guaranteed to evaluate to false.

Additive

The binary additive operators include addition (`+`) and subtraction (`-`). We can apply addition and subtraction to two operands of arithmetic types, but we can also use them to perform scaled pointer arithmetic. I'll discuss pointer arithmetic near the end of this chapter.

The binary `+` operator sums its two operands. The binary `-` operator subtracts the right operand from the left operand. The usual arithmetic conversions are performed on operands of arithmetic type for both operations.

Multiplicative

The binary multiplicative operators include multiplication (`*`), division (`/`), and remainder (`%`). The usual arithmetic conversions are implicitly performed on multiplicative operands to find a common type. You can multiply and divide both floating-point and integer operands, but remainder operates only on integer operands.

Various programming languages implement different kinds of integer division operations, including Euclidean, flooring, and truncating. In *Euclidean division*, the remainder is always nonnegative (Boute 1992). In *flooring division*, the quotient is rounded toward negative infinity (Knuth 1997). In *truncating*

division, the fractional part of the quotient is discarded, which is often referred to as *truncation toward zero*.

The C programming language implements truncating division, meaning that the remainder always has the same sign as the dividend, as shown in [Table 4-2](#).

Table 4-2: Truncating Division

/	Quotient	%	Remainder
10 / 3	3	10 % 3	1
10 / -3	-3	10 % -3	1
-10 / 3	-3	-10 % 3	-1
-10 / -3	3	-10 % -3	1

To generalize, if the quotient a / b is representable, then the expression $(a / b) * b + a \% b$ equals a . Otherwise, if the value of the divisor is equal to 0 or a / b overflows, both a / b and $a \% b$ will result in undefined behavior.

It's worth taking the time to understand the behavior of the `%` operator to avoid surprises. For example, the following code defines a faulty function called `is_odd` that attempts to test whether an integer is odd:

```
bool is_odd(int n) {  
    return n % 2 == 1;  
}
```

Because the result of the remainder operation always has the sign of the dividend n , when n is negative and odd, $n \% 2$ returns -1 , and the function returns `false`.

A correct, alternative solution is to test that the remainder is not 0 (because a remainder of 0 is the same regardless of the sign of the dividend):

```
bool is_odd(int n) {  
    return n % 2 != 0;  
}
```

Many central processing units (CPUs) implement remainder as part of the division operator, which can overflow if the dividend is equal to the minimum negative value for the signed integer type and the divisor is equal to -1 . This occurs even though the mathematical result of such a remainder operation is 0.

The C standard library provides floating-point remainder, truncation, and

rounding functions, including `fmod`, among others.

Bitwise Operators

We use *bitwise operators* to manipulate the bits of an object or any integer expression. Bitwise operators (`|`, `&`, `^`, `~`) treat the bits as a pure binary model without concern for the values represented by these bits. Typically, they're used on objects that represent *masks* or *bitmaps* where each bit indicates that something is “on” or “off,” “enabled” or “disabled,” or some other binary pairing. Using a mask, multiple bits can be set, unset, or inverted in a single bitwise operation. Masks and bitmaps are best represented as unsigned integer types, as the sign bit can be better used as a value and unsigned operations are less prone to undefined behavior.

Complement

The *unary complement operator* (`~`) works on a single operand of integer type and returns the *bitwise complement* of its operand—that is, a value in which each bit of the original value is flipped. The complement operator is used in applying the POSIX `umask`, for example. The `umask` masks or subtracts permissions. For example, a `umask` of `077` turns off read, write, and execute permissions for the group and others. A file's permission mode is the result of a logical AND operation between the complement of the mask and the process's requested permission mode setting.

Integer promotions are performed on the operand of the complement operator, and the result has the promoted type. For example, the following code snippet applies the `~` operator to a value of unsigned `char` type:

```
unsigned char uc = UCHAR_MAX; // 0xFF
int i = ~uc;
```

On an architecture with an 8-bit `char` type and 32-bit `int` type, `uc` is assigned the value `0xFF`. When `uc` is used as the operand to the `~` operator, `uc` is promoted to signed `int` by zero-extending it to 32 bits, `0x000000FF`. The complement of this value is `0xFFFFF000`. Therefore, on this platform, complementing an unsigned `char` type always results in a negative value of type signed `int`. As a general policy and to avoid surprises such as this, bitwise operations should operate only on values of sufficiently wide unsigned integer types.

Shift

Shift operations shift the value of each bit of an operand of integer type by a specified number of positions. Shifting is commonly performed in system programming, where bitmasks are common. Shift operations may also be used in code that manages network protocols or file formats to pack or unpack data. They include left-shift operations of the form

shift expression << *additive expression*

and right-shift operations of the form:

shift expression >> *additive expression*

The *shift expression* is the value to be shifted, and the *additive expression* is the number of bits by which to shift the value. [Figure 4-1](#) illustrates a logical left shift of 1 bit.




Figure 4-1: A logical left shift of 1 bit

The additive expression determines the number of bits by which to shift the value. For example, the result of $E1 \ll E2$ is the value of $E1$ left-shifted $E2$ bit positions; vacated bits are filled with zeros. If $E1$ has an unsigned type, the resulting value is $E1 \times 2^{E2}$. Values that cannot be represented in the resulting type will wrap around. If $E1$ has a signed type and nonnegative value and if $E1 \times 2^{E2}$ is representable in the result type, then that is the resulting value; otherwise, it is undefined behavior. Similarly, the result of $E1 \gg E2$ is the value of $E1$ right-shifted $E2$ bit positions. If $E1$ has an unsigned type or if $E1$ has a signed type and a nonnegative value, the value of the result is the integral part of the quotient of $E1/2^{E2}$. If $E1$ has a signed type and a negative value, the resulting value is implementation defined and may be either an arithmetic (sign-extended) shift or a logical (unsigned) shift, as shown in [Figure 4-2](#).




Figure 4-2: An arithmetic (signed) right shift and a logical (unsigned) right shift of 1 bit

In both shift operations, the integer promotions are performed on the operands, each of which has an integer type. The type of the result is that of the promoted left operand. The usual arithmetic conversions are *not* performed.

[Listing 4-5](#) shows how to perform right-shift operations on signed and unsigned integers that are free from undefined behavior.

```

extern int si1, si2, sresult;
extern unsigned int ui1, ui2, uresult;
// --snip--
❶ if ((si2 < 0) || (si2 >= sizeof(si1)*CHAR_BIT)) {
    /* error */
}
else {
    sresult = si1 >> si2;
}
❷ if (ui2 >= sizeof(unsigned int)*CHAR_BIT) {
    /* error */
}
else {
    uresult = ui1 >> ui2;
}

```

Listing 4-5: Correct right-shift operations

For signed integers ❶, you must ensure that the number of bits shifted is not negative, greater than, or equal to the width of the promoted left operand. For unsigned integers ❷, you omit the test for negative values, as unsigned integers can never be negative. You can perform safe left-shift operations in a similar manner.

Bitwise AND

The binary *bitwise AND operator* (&) returns the bitwise AND of two operands

of integer type. The usual arithmetic conversions are performed on both operands. Each bit in the result is set if and only if each of the corresponding bits in the converted operands is set, as shown in [Table 4-3](#).

Table 4-3: Bitwise AND Truth Table

x	y	x & y
0	0	0
0	1	0
1	0	0
1	1	1

Bitwise Exclusive OR

The *bitwise exclusive OR operator* (^) returns the bitwise exclusive OR of the operands of integer type. The operands must be integers, and the usual arithmetic conversions are performed on both. Each bit in the result is set if and only if exactly one of the corresponding bits in the converted operands is set, as shown in [Table 4-4](#). You can also think of this operation as “one or the other, but not both.”

Table 4-4: Bitwise Exclusive OR Truth Table

x	y	x ^ y
0	0	0
0	1	1
1	0	1
1	1	0

Exclusive OR is equivalent to the addition operation on the integers modulo 2—that is, because of wraparound $1 + 1 \bmod 2 = 0$ (Lewin 2012).

Beginners commonly mistake the exclusive OR operator for an exponent operator, erroneously believing that the expression $2 \wedge 7$ will compute 2 raised to the power of 7. The correct way to raise a number to a certain power in C is to use the pow functions defined in `<math.h>`, as shown in [Listing 4-6](#). The pow functions operate on floating-point arguments and return a floating-point result, so be aware that these functions might fail to produce the expected results because of truncation or other errors.

```
#include <math.h>
#include <stdio.h>
```

```
int main(void) {
    int i = 128;
    if (i == pow(2, 7)) {
        puts("equal");
    }
}
```

Listing 4-6: Using the pow functions

This code calls the `pow` function to compute 2 raised to the power of 7. Because 2^7 equals 128, this program will print `equal`.

Bitwise Inclusive OR

The *bitwise inclusive OR* (`|`) operator returns the bitwise inclusive OR of two operands. Each bit in the result is set if and only if at least one of the corresponding bits in the converted operands is set, as shown in [Table 4-5](#).

Table 4-5: Bitwise Inclusive OR Truth Table

x	y	x y
0	0	0
0	1	1
1	0	1
1	1	1

The operands must be integers, and the usual arithmetic conversions are performed on both.

Logical Operators

The *logical AND* (`&&`) and *OR* (`||`) operators are used primarily for logically joining two or more expressions of scalar type. They're commonly used in condition tests to combine multiple comparisons together, such as in the first operand of the conditional operator, the controlling expression of an `if` statement, or the controlling expression of a `for` loop. You shouldn't use logical operators with bitmap operands, as they are intended primarily for Boolean logic.

The `&&` operator returns 1 if neither of its operands is equal to 0 and returns 0 otherwise. Logically, this means that a `&& b` is true only if both `a` is true and `b` is true.

The `||` operator returns 1 if either of its operands is not equal to 0 and returns 0 otherwise. Logically, this means that `a || b` is true if `a` is true, `b` is true, or both `a` and `b` are true.

The C standard defines both operations in terms of “not equal to zero” because the operands can have values other than 0 and 1. Both operators accept operands of scalar type (integers, floats, and pointers), and the result of the operation has type `int`.

Unlike the corresponding bitwise binary operators, the logical AND operator and logical OR operator guarantee left-to-right evaluation. Both operators *short-circuit*: The second operand is not evaluated if the result can be deduced solely by evaluating the first operand. If the second operand is evaluated, there is a sequence point between the evaluations of the first and second operands. For example, the expression `0 && unevaluated` returns 0 regardless of the value of `unevaluated` because there is no possible value for `unevaluated` that produces a different result. Because of this, `unevaluated` is not evaluated to determine its value. The same is true for `1 || unevaluated` because this expression always returns 1.

Short-circuiting is commonly used in operations with pointers:

```
bool isN(int* ptr, int n) {  
    return ptr && *ptr == n; // avoid a null pointer dereference  
}
```

This code tests the value of `ptr`. If `ptr` is null, the second `&&` operand is not evaluated, preventing a null pointer dereference.

Short-circuiting can also be used to avoid unnecessary computing. In the following expression, the `is_file_ready` predicate function returns true if the file is ready:

```
is_file_ready() || prepare_file()
```

When the `is_file_ready` function returns true, the second `||` operand is not evaluated, as there is no need to prepare the file. This avoids potential errors or, when `prepare_file` is idempotent, unnecessary computing, assuming the cost of determining whether the file is ready is less than the cost of preparing the file.

Programmers should exercise caution if the second operand contains side

effects, because it may not be apparent whether these side effects occur. For example, in the following code, the value of `i` is incremented only when `i >= 0`:

```
enum {max = 15};
int i = 17;

if ((i >= 0) && ((i++) <= max)) {
    // --snip--
}
```

This code may be correct, but it's likely a programmer error.

Cast Operators

Casts (also known as *type casts*) explicitly convert a value of one type to a value of another type. To perform a cast, we precede an expression with a parenthesized type name, which converts the value of the expression to the unqualified version of the named type. The following code illustrates an explicit conversion, or cast, of `x` from type `double` to type `int`:

```
double x = 1.2;
int sum = (int)x + 1; // explicit conversion from double to int
```

Unless the type name specifies a void type, the type name must be a qualified or unqualified scalar type. The operand must also have scalar type; a pointer type cannot be converted to any floating-point type, and vice versa.

Casts are extremely powerful and must be used carefully. For one thing, casts may reinterpret the existing bits as a value of the specified type without changing the bits:

```
intptr_t i = (intptr_t)a_pointer; // reinterpret bits as an integer
```

Casts may also change these bits into whatever bits are needed to represent the original value in the resulting type:

```
int i = (int)a_float; // change bits to an integer representation
```

Casts can also disable diagnostics. For example:

```
char c;
// --snip--
while ((c = fgetc(in)) != EOF) {
    // --snip--
}
```

This generates the following diagnostic when compiled with Visual C++ 2022 with warning level /w4:

Severity	Code	Description
Warning	C4244	'=': conversion from 'int' to 'char', possible loss of data

Adding a cast to char disables the diagnostic without fixing the problem:

```
char c;
while ((c = (char)fgetc(in)) != EOF) {
    // --snip--
}
```

To mitigate these risks, C++ defines its own casts, which are less powerful.

Conditional Operator

The *conditional operator* (`? :`) is the only C operator that takes three operands. It returns a result based on the condition. You can use the conditional operator like this:

```
result = condition ? valueReturnedIfTrue : valueReturnedIfFalse;
```

The conditional operator evaluates the first operand, called the *condition*. The second operand (`valueReturnedIfTrue`) is evaluated if the condition is true, or the third operand (`valueReturnedIfFalse`) is evaluated if the condition is false. The result is the value of either the second or third operand (depending on which operand was evaluated).

This result is converted to a common type based on the second and third operands. There is a sequence point between the evaluation of the first operand and the evaluation of the second or third operand (whichever is evaluated) so that the compiler will ensure that all side effects resulting from evaluating the

condition have completed before the second or third operand is evaluated.

The conditional operator is similar to an `if...else` control flow block but returns a value as a function does. Unlike with an `if...else` control flow block, you can use the conditional operator to initialize a `const`-qualified object:

```
const int x = (a < b) ? b : a;
```

The first operand to the conditional operator must have scalar type. The second and third operands must have compatible types (roughly speaking). For more details on the constraints for this operator and the specifics of determining the return type, refer to Section 6.5.15 of the C standard (ISO/IEC 9899:2024).

alignof Operator

The `alignof` operator yields an integer constant representing the alignment requirement of its operand's declared complete object type. It does not evaluate the operand. When applied to an array type, it returns the alignment requirement of the element type. An alternative spelling of `_Alignof` is available for this operator. Prior to C23, the `alignof` spelling was available through a convenience macro provided in the header `<stdalign.h>`. The `alignof` operator is useful in static assertions that are used to verify assumptions about your program (discussed further in [Chapter 11](#)). The purpose of these assertions is to diagnose situations in which your assumptions are invalid. [Listing 4-7](#) demonstrates the use of the `alignof` operator.

```
#include <stdio.h>
#include <stddef.h>
#include <stdalign.h>
#include <assert.h>

int main(void) {
    int arr[4];
    static_assert(alignof(arr) == 4, "unexpected alignment");
    static_assert(alignof(max_align_t) == 16, "unexpected alignment");
    printf("Alignment of arr = %zu\n", alignof(arr));
    printf("Alignment of max_align_t = %zu\n", alignof(max_align_t));
}
```

Listing 4-7: The `alignof` operator

This simple program doesn't accomplish anything particularly useful. It declares an array `arr` of four integers followed by a static assertion concerning the alignment of the array and a runtime assertion concerning the alignment of `max_align_t` (an object type whose alignment is the greatest fundamental alignment). It then prints out these values. This program will not compile if either static assertion is false, or it will output the following:

```
Alignment of arr = 4
Alignment of max_align_t = 16
```

These alignments are characteristic of the x86-64 architecture.

Relational Operators

The *relational operators* include equal to (`==`), not equal to (`!=`), less than (`<`), greater than (`>`), less than or equal to (`<=`), and greater than or equal to (`>=`). Each returns 1 if the specified relationship is true and 0 if it is false. The result has type `int`, again, for historical reasons.

Note that C does not interpret the expression `a < b < c` to mean that `b` is greater than `a` but less than `c`, as ordinary mathematics does. Instead, the expression is interpreted to mean `(a < b) < c`. In English, if `a` is less than `b`, the compiler should compare 1 to `c`; otherwise, it compares 0 to `c`. If this is your intent, include the parentheses to make that clear to any potential code reviewer. Some compilers such as GCC and Clang provide the `-Wparentheses` flag to diagnose those problems. To determine whether `b` is greater than `a` but less than `c`, you can write this test: `(a < b) && (b < c)`.

The equality and inequality operators have lower precedence than the relational operators—and assuming otherwise is a common mistake. This means that the expression `a < b == c < d` is evaluated the same as `(a < b) == (c < d)`. In both cases, the comparisons `a < b` and `c < d` are evaluated first, and the resulting values (either 0 or 1) are compared for equality.

We can use these operators to compare arithmetic types or pointers. When we compare two pointers, the result depends on the relative locations in the address space of the objects pointed to. If both pointers point to the same object, they are equal.

Equality and inequality operators differ from the other relational operators. For example, you cannot use the other relational operators on two pointers to unrelated objects, because doing so makes no sense and is consequently

undefined behavior:

```
int i, j;
bool b1 = &i < &j; // undefined behavior
bool b2 = &i == &j; // OK, but tautologically false
```

You might compare pointers, for example, to determine whether you have reached the too-far element of an array.

Compound Assignment Operators

Compound assignment operators, shown in [Table 4-6](#), modify the current value of an object by performing an operation on it.

Table 4-6: Compound Assignment Operators

Operator	Assignment by
<code>+= -=</code>	Sum and difference
<code>*= /= %=</code>	Product, quotient, and remainder
<code><<= >>=</code>	Bitwise left shift and right shift
<code>&= ^= =</code>	Bitwise AND, XOR, and OR

A compound assignment of the form `E1 op = E2` is equivalent to the simple assignment expression `E1 = E1 op (E2)`, except that `E1` is evaluated only once. Compound assignments are primarily used as shorthand notation. There are no compound assignment operators for logical operators.

Comma Operator

In C, we use commas in two distinct ways: as operators and to separate items in a list (such as arguments to functions or lists of declarations). The *comma* (`,`) *operator* is a way to evaluate one expression before another. First, the left operand of a comma operator is evaluated as a void expression. There is a sequence point between the evaluation of the left operand and the evaluation of the right operand. Then, the right operand is evaluated after the left. The comma operation has the type and value of the right operand—mostly because it is the last expression evaluated.

You can't use the comma operator in contexts in which a comma might separate items in a list. Instead, you would include a comma within a parenthesized expression or within the second expression of a conditional

operator. For example, assume that `a`, `t`, and `c` each have type `int` in the following call to `f`:

```
f(a, (t=3, t+2), c)
```

The first comma separates the first and second arguments to the function. The second comma is a comma operator. The assignment is evaluated first, followed by the addition. Because of the sequence point, the assignment is guaranteed to complete before the addition takes place. The result of the comma operation has the type `int` and value 5. The third comma separates the second and third arguments to the function.

Pointer Arithmetic

Earlier in this chapter, we mentioned that the additive operators (addition and subtraction) can be used with either arithmetic operands or object pointers. In this section, we discuss adding a pointer and an integer, subtracting two pointers, and subtracting an integer from a pointer.

Adding or subtracting an expression that has an integer type to or from a pointer returns a value with the type of the pointer operand. If the pointer operand points to an element of an array, then the result points to an element offset from the original element. If the resulting pointer is beyond the bounds of the array, undefined behavior occurs. The difference of the array subscripts of the resulting and original array elements equals the integer expression:

```
int arr[100];
int *arrp1 = &arr[40];
int *arrp2 = arrp1 + 20; // arrp2 points to arr[60]
printf("%td\n", arrp2 - arrp1); // prints 20
```

Pointer arithmetic is automatically *scaled* to the size of the array element, rather than individual bytes. C allows a pointer to be formed to each element of an array, including one past the last element of the array object (also referred to as the *too-far* pointer). While this might seem unusual or unnecessary, many early C programs incremented a pointer until it was equal to the too-far pointer, and the C standards committee didn't want to break all this code, which is also idiomatic in C++ iterators. [Figure 4-3](#) illustrates forming the *too-far* pointer.




Figure 4-3: One past the last element of an array object

If both the pointer operand and the result point to elements of the same array object or the too-far pointer, the evaluation did not overflow; otherwise, the behavior is undefined. To satisfy the too-far requirement, an implementation need only provide one extra byte (which can overlap another object in the program) just after the end of the object.

C also allows objects to be treated as an array containing only a single element, allowing you to obtain a too-far pointer from a scalar.

The too-far special case allows us to advance a pointer until it is equal to the too-far pointer, as in the following function:

```
int m[2] = {1, 2};

int sum_m_elems(void) {
    int *pi; int j = 0;
    for (pi = &m[0]; pi < &m[2]; ++pi) j += *pi;
    return j;
}
```

Here, the for statement (explained in detail in the next chapter) in the `sum_m_elems` function loops while `pi` is less than the address of the too-far pointer for the array `m`. The pointer `pi` is incremented at the end of each iteration of the loop until the too-far pointer is formed, causing the loop condition to evaluate to 0 when tested.

When we subtract one pointer from another, both must point to elements of the same array object or the too-far element. This operation returns the difference of the subscripts of the two array elements. The type of the result is `ptrdiff_t` (a signed integer type). You should take care when subtracting pointers, because the range of `ptrdiff_t` may not be sufficient to represent the difference of pointers to elements of very large character arrays.

Summary

In this chapter, you learned how to use operators to write simple expressions that

perform operations on various object types. Along the way, you learned about some core C concepts, such as lvalues, rvalues, value computations, and side effects, which determine how expressions are evaluated. You also learned how operator precedence, associativity, order of evaluation, sequencing, and interleaving can affect the total order in which a program is executed. In the next chapter, you'll learn more about how to control the execution of your program by using selection, iteration, and jump statements.

5

CONTROL FLOW



In this chapter, you'll learn how to control the order in which individual statements are evaluated. We'll start by going over expression statements and compound statements that define the work to be performed. We'll then cover three kinds of statements that determine which code blocks are executed and in what order: selection, iteration, and jump statements.

Expression Statements

An *expression statement* is an optional expression terminated by a semicolon. It's one of the most common statements and a basic unit of work. The following examples show different expression statements.

Assigns a value to a:

```
a = 6;
```

Assigns the sum of a and b to c:

```
c = a + b;
```

A null statement:

```
; // null statement, does nothing
```

You can use a null statement when the syntax of the language requires a statement but no expression needs to be evaluated. Null statements are commonly used as placeholders in iteration statements.

The following expression statement increments the value of count:

```
++count;
```

After each full expression has been evaluated, its value (if any) is discarded (including assignment expressions in which the assignment itself is a side effect of the operation) so that any useful results occur as the consequence of side effects (as discussed in [Chapter 4](#)). Three of the four expression statements in this example have side effects (the null statement does nothing). Once all side effects have completed, execution proceeds to the statement following the semicolon.

Compound Statements

A *compound statement*, or *block*, is a list of zero or more statements, surrounded by braces. The statements in the block may be any kind of statement described throughout this chapter. Some of these statements may be declarations. (In early versions of C, declarations within the block had to precede all nondeclarations, but that restriction no longer applies.) Each statement in the block is executed in sequence unless modified by a control statement. After the final statement has been evaluated, execution proceeds to after the closing brace:

```
{  
    static int count = 0;  
    c += a;  
    ++count;  
}
```

This example declares a static variable of type `int` called `count`. The second line increases a variable `c` declared in an outer scope by the value stored in `a`. Finally, `count` is incremented to track how many times this block has been executed.

Compound statements can be nested so that one compound statement fully

encloses another. You may also have blocks with no statements at all (just the empty braces).

CODE STYLE

Competing coding styles disagree on when and where to place braces. If you're modifying existing code, it would be wise to follow the style already in use for the project. Otherwise, look at styles you see in code written by experienced C programmers and choose one that seems clear. For example, some programmers line up the opening and closing braces to make it easy to find the mate for a given brace. Others follow the style used in *The C Programming Language*, 2nd edition, by Brian Kernighan and Dennis Ritchie (Pearson, 1988), wherein the opening brace is placed at the end of the preceding line and the closing brace gets a line to itself. Once you have chosen a style, use it consistently.

Selection Statements

Selection statements allow you to conditionally execute a substatement depending on the value of a controlling expression. The *controlling expression* determines which statements are executed based on a condition. Selection statements include the `if` statement and the `switch` statement.

if

The `if` statement allows a programmer to execute a substatement based on the value of a controlling expression of scalar type.

There are two kinds of `if` statements. The first conditionally determines whether the substatement is executed:

```
if (expression)  
    substatement
```

In this case, the *substatement* is executed if the *expression* is not equal to 0. Only the single *substatement* of the `if` statement is conditionally executed, although it can be a compound statement.

[Listing 5-1](#) shows a division function that uses `if` statements. It divides a specified dividend by a specified divisor and returns the result in the object referenced by `quotient`. The function tests for both division by zero and signed integer overflow and returns `false` in either case.

```
bool safediv(int dividend, int divisor, int *quotient) {  
  ❶ if (!quotient) return false;  
  ❷ if ((divisor == 0) || ((dividend == INT_MIN) && (divisor == -1  
  )))  
    ❸ return false;  
  ❹ *quotient = dividend / divisor;  
    return true;  
}
```

Listing 5-1: A safe division function

The first line of this function ❶ tests `quotient` to ensure that it's not null. If it is null, the function returns `false` to indicate that it is unable to return a value. (We cover return statements later in this chapter.)

The second line of the function ❷ contains a more complex `if` statement. Its controlling expression tests whether the divisor is 0 or whether the division would result in signed integer overflow if unchecked. If the result of this expression is not equal to 0, the function returns `false` ❸ to indicate that it is unable to produce a quotient. If the controlling expression of the `if` statement evaluates to 0, the function does not return, and the remaining statements ❹ are executed to calculate the quotient and return `true`.

The second kind of `if` statement includes an `else` clause, which selects an alternative substatement to execute when the initial substatement is not selected:

```
if (expression)  
    substatement1  
else  
    substatement2
```

In this form, `substatement1` is executed if `expression` is not equal to 0, and `substatement2` is executed if `expression` is equal to 0. One of these substatements is executed, but never both.

For either form of the `if` statement, the conditionally executed substatement may also be an `if` statement. A common use of this is the `if...else` ladder, shown in [Listing 5-2](#).

```
if (expr1)  
    substatement1  
else if (expr2)
```

```
    substatement2
else if (expr3)
    substatement3
else
    substatement4
```

Listing 5-2: The if...else ladder syntax

One (and only one) of the four statements in an if...else ladder will execute:

- *substatement1* executes if *expr1* does not equal 0.
- *substatement2* executes if *expr1* equals 0 and if *expr2* does not equal 0.
- *substatement3* executes if both *expr1* and *expr2* equal 0 and *expr3* does not equal 0.
- *substatement4* executes only if the preceding conditions are all equal to 0.

The example shown in [Listing 5-3](#) uses an if...else ladder to print grades.

```
void printgrade(unsigned int marks) {
    if (marks >= 90) {
        puts("YOUR GRADE : A");
    } else if (marks >= 80) {
        puts("YOUR GRADE : B");
    } else if (marks >= 70) {
        puts("YOUR GRADE : C");
    } else {
        puts("YOUR GRADE : Failed");
    }
}
```

Listing 5-3: Using an if...else ladder to print grades

In this if...else ladder, the printgrade function tests the value of the unsigned int parameter marks to determine whether it is greater than or equal to 90. If so, the function prints YOUR GRADE : A. Otherwise, it tests whether marks is greater than or equal to 80, and so forth down the if...else ladder. If marks is not greater than or equal to 70, the function prints YOUR GRADE : Failed. This example uses a coding style in which the closing brace is followed by the else clause on the same line.

Only a single statement following the if statement is executed. For example, in the following code snippet, conditionally_executed is executed

only if condition is not equal to 0, but `unconditionally_executed` is always executed:

```
if (condition)
    conditionally_executed();
unconditionally_executed(); // always executed
```

Attempting to add another conditionally executed function is a common source of errors:

```
if (condition)
    conditionally_executed();
    also_conditionally_executed(); // ???
unconditionally_executed(); // always executed
```

In this code snippet, `also_conditionally_executed` is *unconditionally* executed. The name and indented formatting are deceptive because whitespace (in general) and indentation (in particular) are meaningless to the syntax. This code can be fixed by adding braces to delimit a single compound statement or block. This block is then executed as the single conditionally executed statement:

```
if (condition) {
    conditionally_executed();
    also_conditionally_executed(); // fixed it
}
unconditionally_executed(); // always executed
```

While the original code snippet was not incorrect, many coding guidelines recommend always including braces to avoid this kind of error:

```
if (condition) {
    conditionally_executed();
}
unconditionally_executed(); // always executed
```

My personal style is to omit the braces only when I can include the conditionally executed statement on the same line as the `if` statement:

```
if (!quotient) return false;
```

This issue is less problematic when you let your integrated development environment (IDE) format your code for you, as it won't be fooled by code indentation when formatting your code. The GCC and Clang compilers provide a `-Wmisleading-indentation` compiler flag that checks code indentation and warns when it doesn't correspond to the control flow.

switch

The `switch` statement works just like the `if...else` ladder, except that the controlling expression must have an integer type. For example, the `switch` statement in [Listing 5-4](#) performs the same function as the `if...else` ladder from [Listing 5-3](#), provided that `marks` is an integer in the range of 0 to 109. If `marks` is greater than 109, it will result in a failed grade because the resulting quotient will be greater than 10 and will consequently be caught by the default case.

```
switch (marks/10) {
    case 10:
    case 9:
        puts("YOUR GRADE : A");
        break;
    case 8:
        puts("YOUR GRADE : B");
        break;
    case 7:
        puts("YOUR GRADE : C");
        break;
    default:
        puts("YOUR GRADE : Failed");
}
```

Listing 5-4: Using a `switch` statement to print out grades

The `switch` statement causes control to jump to one of the three substatements, depending on the value of the controlling expression and the constant expressions in each case label. Following the jump, code is executed sequentially until the next control flow statement is reached. In our example, a jump to case `10` (which is empty) flows through and executes the subsequent statements in case `9`. This is necessary to the logic so that a perfect grade of 100 results in an A and not an F.

You can terminate the execution of the switch by inserting a break statement, causing control to jump to the execution of the statement directly following the overall switch statement. (We discuss break statements in more detail later in this chapter.) Make sure you remember to include a break statement before the next case label. If omitted, the control flow falls through to the next case in the switch statement—a common source of errors. Because the break statement isn't required, omitting it doesn't typically produce compiler diagnostics. GCC issues a warning for fall-through if you use the `-Wimplicit-fallthrough` flag. The C23 standard introduces the `[[fallthrough]]` attribute as a way for a programmer to specify that fall-through behavior is desirable, under the assumption that silent fall-through is an accidental omission of a break statement.

Integer promotions are performed on the controlling expression. The constant expression in each case label is converted to the promoted type of the controlling expression. If a converted value matches that of the promoted controlling expression, control jumps to the statement following the matched case label. Otherwise, if there is no match but there is a default label, control jumps to the labeled statement. If no converted case constant expression matches and there is no default label, no part of the switch body is executed. When switch statements are nested, a case or default label is accessible only within the closest enclosing switch statement.

There are best practices regarding the use of switch statements. [Listing 5-5](#) shows a less-than-ideal implementation of a switch statement that assigns interest rates to an account based on the account type. The `AccountType` enumeration represents the fixed number of account types offered by the bank.

```
typedef enum {Savings, Checking, MoneyMarket} AccountType;
void assignInterestRate(AccountType account) {
    double interest_rate;
    switch (account) {
        case Savings:
            interest_rate = 3.0;
            break;
        case Checking:
            interest_rate = 1.0;
            break;
        case MoneyMarket:
            interest_rate = 4.5;
            break;
    }
```

```
    printf("Interest rate = %g.\n", interest_rate);  
}
```

Listing 5-5: A switch statement without a default label

The `assignInterestRate` function defines a single parameter of the enumeration type `AccountType` and switches on it to assign the appropriate interest rate associated with each account type.

Nothing is wrong with the code as written, but it requires programmers to update the code in at least two separate places if they want to make any changes. Let's say the bank introduces a new type of account: a certificate of deposit. A programmer updates the `AccountType` enumeration as follows:

```
typedef enum {Savings, Checking, MoneyMarket, CD} AccountType;
```

However, if the programmer neglects to modify the switch statement in the `assignInterestRate` function, the `interest_rate` isn't assigned, resulting in an uninitialized read when the function attempts to print that value. This problem is common because the enumeration may be declared far from the switch statement, and the program may contain many similar switch statements that all reference an object of type `AccountType` in their controlling expression.

Both Clang and GCC help diagnose these problems at compilation time when you use the `-Wswitch-enum` flag. Alternatively, you can protect against such errors and improve the testability of this code by adding this default case to the switch statement, as shown in [Listing 5-6](#).

```
typedef enum {Savings, Checking, MoneyMarket, CD} AccountType;  
void assignInterestRate(AccountType account) {  
    double interest_rate;  
    switch (account) {  
        case Savings:  
            interest_rate = 3.0;  
            break;  
        case Checking:  
            interest_rate = 1.0;  
            break;  
        case MoneyMarket:  
            interest_rate = 4.5;  
            break;  
        case CD:  
            interest_rate = 7.5;
```

```
        break;
    default: abort();
}
printf("Interest rate = %g.\n", interest_rate);
return;
}
```

Listing 5-6: A switch statement with a default label

The switch statement now includes a case for CD, and the default clause is unused. However, retaining the default clause is good practice, in case another account type is added in the future.

Including a default clause does have the drawback of suppressing compiler warnings and not diagnosing the problem until runtime. Compiler warnings (if supported by your compiler) are therefore a better approach.

Iteration Statements

Iteration statements cause substatements (or compound statements) to be executed zero or more times, subject to termination criteria. The English word *iteration* means “the repetition of a process.” Iteration statements are more informally and commonly referred to as loops. A *loop* is “a process, the end of which is connected to the beginning.”

while

The while statement causes the loop body to execute repeatedly until the controlling expression is equal to 0. The evaluation of the controlling expression occurs before each execution of the loop body. Consider the following example:

```
void f(unsigned int x) {
    while (x > 0) {
        printf("%d\n", x);
        --x;
    }
    return;
}
```

If *x* is not initially greater than 0, the while loop exits without executing the loop body. If *x* is greater than 0, its value is output and then decremented. Once the end of the loop is reached, the controlling expression is tested again. This

pattern repeats until the expression evaluates to 0. Overall, this loop will count down from `x` to 1.

A `while` loop is an entry-controlled loop that executes until its controlling expression evaluates to 0. [Listing 5-7](#) shows an implementation of the C standard library `memset` function. This function copies the value of `val` (converted to an `unsigned char`) into each of the first `n` characters of the object pointed to by `dest`.

```
void *memset(void *dest, int val, size_t n) {
    unsigned char *ptr = (unsigned char*)dest;
    while (n-- > 0)
        *ptr++ = (unsigned char)val;
    return dest;
}
```

Listing 5-7: The C standard library `memset` function

The first line of the `memset` function converts `dest` to a pointer to an `unsigned char` and assigns the resulting value to the `unsigned char` pointer `ptr`. This lets us preserve the value of `dest` to return on the last line of the function. The remaining two lines of the function form a `while` loop that copies the value of `val` (converted to an `unsigned char`) into each of the first `n` characters of the object `dest` points to. The controlling expression of the `while` loop tests that `n-- > 0`.

The `n` argument is a loop counter that's decremented on each iteration of the loop as a side effect of the evaluation of the controlling expression. The loop counter in this case monotonically decreases until it reaches the minimum value (0). The loop performs `n` repetitions, where `n` is less than or equal to the bound of the memory that `ptr` references.

The `ptr` pointer designates a sequence of objects of type `unsigned char`, from `ptr` through `ptr + n - 1`. The value of `val` is converted to an `unsigned char` and written to each object in turn. If `n` is greater than the bound of the object that `ptr` references, the `while` loop writes to memory outside the bounds of this object. This is undefined behavior and a common security flaw, referred to as a *buffer overflow*, or *overflow*. Provided the object referenced by `ptr` has at least `n` bytes, the `while` loop terminates without undefined behavior. In the final iteration of the loop, the controlling expression `n-- > 0` evaluates to 0, causing the loop to terminate.

It's possible to write an *infinite loop*—a loop that never terminates. To avoid

writing a `while` loop that inadvertently runs forever, be sure you initialize any objects referenced by the controlling expression before the start of the `while` loop. Also make sure that the controlling expression changes during the `while` loop's execution in a manner that causes the loop to terminate after iterating an appropriate number of times.

do...while

The `do...while` statement is similar to the `while` statement, except that the evaluation of the controlling expression occurs after each execution of the loop body rather than before. As a result, the loop body is guaranteed to execute once before the condition is tested. The `do...while` iteration statement has the following syntax:

```
do
    statement
while (expression);
```

In a `do...while` iteration, *statement* is unconditionally executed once, after which *expression* is evaluated. If *expression* is not equal to 0, control returns to the top of the loop and *statement* is executed again. Otherwise, execution passes to the statement following the loop.

The `do...while` iteration statement is commonly used in input/output (I/O), where it makes sense to read from a stream before testing the state of the stream, as shown in [Listing 5-8](#).

```
#include <stdio.h>
// --snip--
int count;
float quant;
char units[21], item[21];
do {
    count = fscanf(stdin, "%f%20s of %20s", &quant, units, item);
    fscanf(stdin, "%*[^\\n]");
} while (!feof(stdin) && !ferror(stdin));
// --snip--
```

Listing 5-8: An input loop

This code inputs a floating-point quantity, a unit of measure (as a string), and an item name (also as a string) from the standard input stream `stdin` until

the end-of-file indicator has been set or a read error has occurred. We'll discuss I/O in detail in [Chapter 8](#).

for

The `for` statement might be the most C-like thing about C. The `for` statement repeatedly executes a statement and is typically used when the number of iterations is known before entering the loop. It has the following syntax:

```
for (clause1; expression2; expression3)  
    statement
```

The controlling expression (*expression2*) is evaluated before each execution of the loop body, and *expression3* is evaluated after each execution of the loop body. If *clause1* is a declaration, the scope of any identifiers it declares is the remainder of the declaration and the entire loop, including the other two expressions.

The purpose of *clause1*, *expression2*, and *expression3* is apparent when we translate the `for` statement into an equivalent `while` loop, as shown in [Figure 5-1](#).




Figure 5-1: Translating a `for` loop into a `while` loop

[Listing 5-9](#) shows a modified version of the `memset` implementation from [Listing 5-7](#); we have replaced the `while` loop with a `for` loop.

```
void *memset(void *dest, int val, size_t n) {
```

```
unsigned char *ptr = (unsigned char *)dest;
for (❶ size_t i = 0; ❷ i < n; ❸ ++i) {
    *(ptr + i) = (unsigned char)val;
}
return dest;
}
```

Listing 5-9: Filling a character array by using a for loop

The for loop is popular among C programmers because it provides a convenient location for declaring and/or initializing the loop counter ❶, specifying the controlling expression for the loop ❷, and incrementing the loop counter ❸, all on the same line. The `*(ptr + i)` lvalue expression could be written equivalently using the index operator as `ptr[i]`.

The for loop can also be somewhat misleading. Let's take the example of a singly linked list in C that declares a node structure consisting of a data element and a pointer to the next node in the list. We also define a pointer `p` to the node structure:

```
struct node {
    int data;
    struct node *next;
};
struct node *p;
```

Using the definition of `p`, the following example (used to deallocate the storage for a linked list) erroneously reads the value `p` after it has been freed:

```
for (p = head; p != nullptr; p = p->next) {
    free(p);
}
```

Reading `p` after it has been freed is undefined behavior.

If this loop were rewritten as a while loop, it would become apparent that the code reads `p` after it was freed:

```
p = head;
while (p != nullptr) {
    free(p);
    p = p->next;
}
```

The for loop can be confusing because it evaluates *expression3* after the main body of the loop, even though lexically it appears before the loop's body.

The correct way to perform this operation is to save the required pointer before freeing it, like this:

```
for (p = head; p != nullptr; p = q) {  
    q = p->next;  
    free(p);  
}
```

You can read more about dynamic memory management in [Chapter 6](#).

Jump Statements

A *jump statement* unconditionally transfers control to another section of the same function when encountered. These are the lowest-level control flow statements and generally correspond closely to the underlying assembly language code.

goto

Any statement may be preceded by a *label*, which is an identifier followed by a colon. C23 also allows you to place labels in front of declarations and at the end of a compound statement, which was not allowed in previous versions of C. A goto statement causes a jump to the statement prefixed by the named label in the enclosing function. The jump is unconditional, meaning it happens every time the goto statement is executed. Here's an example of a goto statement:

```
/* executed statements */  
goto location;  
/* skipped statements */  
location:  
/* executed statements */
```

Execution continues until the goto statement is reached, at which point control jumps to the statement following the location label, where execution continues. Statements between the goto statement and the label are passed over.

The goto statement has had a bad reputation since Edsger Dijkstra wrote a paper titled “Go To Statement Considered Harmful” (1968). His criticism was

that goto statements can result in *spaghetti code* (code with a complex and tangled control structure, resulting in a program flow that's conceptually twisted and tangled like a bowl of spaghetti) if used haphazardly. However, goto statements can also make code easier to read if used in a clear, consistent manner.

One helpful way to use goto statements is to chain them together to release allocated resources (such as dynamic allocated memory or an open file) when an error occurs and you must leave a function. This scenario happens when a program allocates multiple resources; each allocation can fail, and resources must be released to prevent leaking. If the first resource allocation fails, no cleanup is needed, because no resources have been allocated. However, if the second resource cannot be allocated, the first resource needs to be released. Similarly, if the third resource cannot be allocated, the second and first resources allocated need to be released, and so forth. This pattern results in duplicated cleanup code, and it can be error-prone because of the duplication and additional complexity.

One solution is to use nested if statements, which can also become difficult to read if nested too deeply. Instead, we can use a goto chain as shown in [Listing 5-10](#) to release resources.

```
int do_something(void) {
    FILE *file1, *file2;
    object_t *obj;
    int ret_val = 0; // initially assume a successful return value

    file1 = fopen("a_file", "w");
    if (file1 == nullptr) {
        return -1;
    }

    file2 = fopen("another_file", "w");
    if (file2 == nullptr) {
        ret_val = -1;
        goto FAIL_FILE2;
    }

    obj = malloc(sizeof(*obj));
    if (obj == nullptr) {
        ret_val = -1;
        goto FAIL_OBJ;
    }
```

```
// operate on allocated resources

// clean up everything
free(obj);
FAIL_OBJ: // otherwise, close only the resources we opened
    fclose(file2);
FAIL_FILE2:
    fclose(file1);
    return ret_val;
}
```

Listing 5-10: Using a goto chain to release resources

The code follows a simple pattern: resources are allocated in a certain order, operated upon, and then released in reverse (last in, first out) order. If an error occurs while allocating a resource, the code uses a goto to jump to the appropriate location in cleanup code and releases only those resources that have been allocated.

Used like this, goto statements can make code easier to read. A real-world example is the `copy_process` function from `kernel/fork.c` from v6.7 of the Linux kernel (<https://elixir.bootlin.com/linux/v6.7/source/kernel/fork.c#L2245>), which uses 20 goto labels to perform cleanup code when an internal function fails.

continue

You can use a `continue` statement inside a loop to jump to the end of the loop body, skipping the execution of the remaining statements inside the loop body for the current iteration. For example, the `continue` statement is equivalent to `goto END_LOOP_BODY`; in each of the loops shown in [Listing 5-11](#).

```
while (/* _ */) {
    // --snip--
    continue;
    // --snip--
END_LOOP_BODY: ;
}
```

```
do {
    // --snip--
    continue;
```

```
    // --snip--  
END_LOOP_BODY: ;  
} while (/* _ */);
```

```
for (/* _ */) {  
    // --snip--  
    continue;  
    // --snip--  
END_LOOP_BODY: ;  
}
```

Listing 5-11: Using the `continue` statement

The `continue` statement is used in conjunction with a conditional statement so that processing may continue with the subsequent loop iteration after the objective of the current loop iteration has been achieved.

break

A `break` statement terminates execution of a switch or iteration statement. We used `break` within a switch statement in [Listing 5-4](#). Within a loop, a `break` statement causes the loop to terminate and the program execution to resume at the statement following the loop. For instance, the `for` loop in the following example exits only when the uppercase or lowercase Q key is pressed on the keyboard:

```
#include <stdio.h>  
int main(void) {  
    char c;  
    for(;;) {  
        puts("Press any key, Q to quit: ");  
        c = toupper(getchar());  
        if (c == 'Q') break;  
        // --snip--  
    }  
} // loop exits when either q or Q is pressed
```

We typically use `break` statements to discontinue the execution of the loop when the work it was performing has been completed. For example, the `break` statement in [Listing 5-12](#) exits the loop after it finds the specified key in an array. Assuming that key is unique in `arr`, the `find_element` function would behave the same without the `break` statement but, depending on the length of the

array and the point at which key is discovered, could run much slower.

```
size_t find_element(size_t len, int arr[len], int key) {
    size_t pos = (size_t)-1;
    // traverse arr and search for key
    for (size_t i = 0; i < len; ++i) {
        if (arr[i] == key) {
            pos = i;
            break; // terminate loop
        }
    }
    return pos;
}
```

Listing 5-12: Breaking out of a loop

Because `continue` and `break` bypass part of a loop body, use these statements carefully: the code following these statements is not executed.

return

A `return` statement terminates execution of the current function and returns control to its caller. You've already seen many examples of `return` statements in this book. A function may have zero or more `return` statements.

A `return` statement can simply `return`, or it can `return` an expression. Within a void function (a function that doesn't return a value), the `return` statement should simply `return`. When a function returns a value, the `return` statement should `return` an expression that produces a value of the return type. If a `return` statement with an expression is executed, the value of the expression is returned to the caller as the value of the function call expression:

```
int sum(int x, int y, int z) {
    return x + y + z;
}
```

This simple function sums its parameters and returns the sum. The `return` expression `x + y + z` produces a value of type `int`, which matches the return type of the function. If this expression produced a different type, it would be implicitly converted to an object having the return type of the function. The `return` expression can also be as simple as returning 0 or 1. The function result may then be used in an expression or assigned to a variable.

Be aware that if control reaches the closing brace of a non-void function (a function declared to return a value) without evaluating a return statement with an expression, using the return value of the function call is undefined behavior. For example, the following function fails to return a value when `a` is nonnegative because the condition `a < 0` is false:

```
int absolute_value(int a) {
    if (a < 0) {
        return -a;
    }
}
```

We can easily repair this defect by providing a return value when `a` is nonnegative, as shown in [Listing 5-13](#).

```
int absolute_value(int a) {
    if (a < 0) {
        return -a;
    }
    return a;
}
```

Listing 5-13: The `absolute_value` function returns a value along all paths.

However, this code still has a bug (see [Chapter 3](#)). Identifying this bug is left as an exercise for you.

EXERCISES

1. Modify the function from [Listing 5-10](#) to make it clear to the caller which file could not be opened.
2. Fix the remaining bug in the `absolute_value` function in [Listing 5-13](#).

Summary

In this chapter, you learned about control flow statements. Control flow statements allow you to create flexible programs that can repeat tasks and alter their execution based on program inputs:

- Selection statements, such as `if` and `switch`, allow you to select from a set

of statements depending on the value of a controlling expression.

- Iteration statements repeatedly execute a loop body until a controlling expression equals 0.
- Jump statements unconditionally transfer control to a new location.

In the next chapter, you'll learn about dynamically allocated memory. Similar to control flow statements, you can use dynamic memory to create flexible programs that allocate memory based on program inputs.

6

DYNAMICALLY ALLOCATED MEMORY



In [Chapter 2](#), you learned that every object has a storage duration that determines its lifetime and that C defines four storage durations: static, thread, automatic, and allocated. In this chapter, you'll learn about *dynamically allocated memory*, which is allocated from the heap at runtime. Dynamically allocated memory is useful when the exact storage requirements for a program are unknown before runtime.

We'll first describe the differences between allocated, static, and automatic storage duration. We'll skip thread storage allocation as this involves parallel execution, which we don't cover here. We'll then explore the functions you can use to allocate and deallocate dynamic memory, common memory allocation errors, and strategies for avoiding them. The terms *memory* and *storage* are used interchangeably in this chapter, similar to the way they're used in practice.

Storage Duration

Objects occupy *storage*, which might be read-write memory, read-only memory, or central processing unit (CPU) registers. Storage of allocated duration has significantly different properties from storage of either automatic or static storage duration. First, we'll review automatic and static storage duration.

Objects of *automatic storage duration* are declared within a block or as a function parameter. The lifetime of these objects begins when the block in which

they are declared begins execution and ends when execution of the block ends. If the block is entered recursively, a new object is created each time, each with its own storage.

Objects declared at file scope have *static storage duration*. The lifetime of these objects is the entire execution of the program, and their stored value is initialized prior to program startup. You can also declare a variable within a block to have static storage duration by using the `static` storage-class specifier.

The Heap and Memory Managers

Dynamically allocated memory has *allocated storage duration*. The lifetime of an allocated object extends from the allocation until the deallocation.

Dynamically allocated memory is allocated from the *heap*, which is simply one or more large, subdividable blocks of memory managed by the memory manager.

Memory managers are libraries that manage the heap for you by providing implementations of the standard memory management functions described in this chapter. A memory manager runs as part of the client process. The memory manager requests one or more blocks of memory from the operating system (OS) and then allocates this memory to the client process when it invokes a memory allocation function. Allocation requests don't go directly to the OS because it's slower and works only in big chunks of memory, whereas allocators split up those big chunks into little chunks and are faster.

Memory managers manage unallocated and deallocated memory only. Once memory has been allocated, the caller manages the memory until it's returned. It's the caller's responsibility to ensure that the memory is deallocated, although most implementations will reclaim dynamically allocated memory when the program terminates.

MEMORY MANAGER IMPLEMENTATIONS

Memory managers frequently implement a variant of a dynamic storage allocation algorithm described by Donald Knuth (1997). This algorithm uses *boundary tags*, which are size fields that appear before and after the block of memory returned to the programmer. This size information allows all memory blocks to be traversed from any known block in either direction so that the memory manager can coalesce two bordering unused blocks into a larger block to minimize fragmentation.

When a program starts, the free memory areas are long and contiguous. During the program's lifetime, memory is allocated and deallocated. Eventually, the long contiguous regions fragment into smaller and smaller contiguous areas. As a result,

larger allocations can fail even though the total amount of free memory is sufficient for the allocation. Memory allocated for the client process and memory allocated for internal use within the memory manager are all within the addressable memory space of the client process.

When to Use Dynamically Allocated Memory

As previously mentioned, dynamically allocated memory is used when the exact storage requirements for a program are unknown before runtime. Dynamically allocated memory is less efficient than statically allocated memory because the memory manager needs to find appropriately sized blocks of memory in the runtime heap, and the caller must explicitly free those blocks when no longer needed, all of which requires additional processing. Dynamically allocated memory also requires additional processing for housekeeping operations such as *defragmentation* (the consolidation of adjacent free blocks), and the memory manager often uses extra storage for control structures to facilitate these processes.

Memory leaks occur when dynamically allocated memory that's no longer needed isn't returned to the memory manager. If these memory leaks are severe, the memory manager eventually won't be able to satisfy new requests for storage.

By default, you should declare objects with either automatic or static storage duration for objects whose sizes are known at compilation time. Dynamically allocate memory when the size of the storage or the number of objects is unknown before runtime. For example, you might use dynamically allocated memory to read a table from a file at runtime, especially if you do not know the number of rows in the table at compile time. Similarly, you might use dynamically allocated memory to create linked lists, hash tables, binary trees, or other data structures for which the number of data elements held in each container is unknown at compile time.

Memory Management

The C standard library defines memory management functions for allocating (for example, `malloc`, `calloc`, and `realloc`) and deallocating (`free`) dynamic memory. The OpenBSD `reallocarray` function is not defined by the C standard library but can be useful for memory allocation. C23 added two additional deallocation functions: `free_sized` and `free_aligned_sized`.

Dynamically allocated memory is required to be suitably aligned for objects

up to the requested size, including arrays and structures. C11 introduced the `aligned_alloc` function for hardware with stricter-than-normal memory alignment requirements.

malloc

The `malloc` function allocates space for an object of a specified size. The representation of the returned storage is indeterminate. In [Listing 6-1](#), we call the `malloc` function to dynamically allocate storage for an object the size of `struct widget`.

```
#include <stdlib.h>

typedef struct {
    double d;
    int i;
    char c[10];
} widget;

❶ widget *p = malloc(sizeof *p);
❷ if (p == nullptr) {
    // handle allocation error
}
// continue processing
```

Listing 6-1: Allocating storage for a widget with the `malloc` function

All memory allocation functions accept an argument of type `size_t` that specifies the number of bytes of memory to be allocated ❶. For portability, we use the `sizeof` operator when calculating the size of objects, because the size of objects of distinct types, such as `int` and `long`, may differ among implementations.

The `malloc` function returns either a null pointer to indicate an error or a pointer to the allocated space. Therefore, we check whether `malloc` returns a null pointer ❷ and appropriately handles the error.

After the function successfully returns the allocated storage, we can reference members of the `widget` structure through the `p` pointer. For example, `p->i` accesses the `int` member of `widget`, while `p->d` accesses the `double` member.

Allocating Memory Without Declaring a Type

You can store the return value from `malloc` as a `void` pointer to avoid declaring a type for the referenced object:

```
void *p = malloc(size);
```

Alternatively, you can use a `char` pointer, which was the convention before the `void` type was introduced to C:

```
char *p = malloc(size);
```

In either case, the object that `p` points to has no type until an object is copied into this storage. Once that occurs, the object has the *effective type* of the last object copied into this storage, which imprints the type onto the allocated object.

In the following example, the storage that `p` references has an effective type of `widget`:

```
widget w = {3.2, 9, "abc",};  
memcpy(p, &w, sizeof(w));
```

Following the call to `memcpy`, the change of effective type influences optimizations and nothing else.

Because allocated memory can store any sufficiently small object type, pointers returned by allocation functions, including `malloc`, must be sufficiently aligned. For example, if an implementation has objects with 1-, 2-, 4-, 8-, and 16-byte alignments and 16 or more bytes of storage are allocated, the alignment of the returned pointer is a multiple of 16.

Reading Uninitialized Memory

The contents of memory returned from `malloc` are *uninitialized*, which means it has an indeterminate representation. Reading uninitialized memory is never a good idea; think of it as undefined behavior. If you'd like to know more, I wrote an in-depth article on uninitialized reads (Seacord 2017). The `malloc` function doesn't initialize the returned memory because you are expected to overwrite this memory anyway.

Even so, beginners commonly make the mistake of assuming that the memory `malloc` returns contains zeros. The program shown in [Listing 6-2](#) makes this exact error.

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

int main() {
    char *str = (char *)malloc(16);
    if (str) {
        strncpy(str, "123456789abcdef", 15);
        printf("str = %s.\n", str);
        free(str);
        return EXIT_SUCCESS;
    }
    return EXIT_FAILURE;
}
```

Listing 6-2: An initialization error

This program calls `malloc` to allocate 16 bytes of memory and then uses `strncpy` to copy the first 15 bytes of a string into the allocated memory. The programmer attempts to create a properly null-terminated string by copying one fewer byte than the size of the allocated memory. In doing so, the programmer assumes that the allocated storage already contains a zero value to serve as the null byte. However, the storage could easily contain nonzero values, in which case the string wouldn't be properly null terminated, and the call to `printf` would result in undefined behavior.

A common solution is to write a null character into the last byte of the allocated storage, as follows:

```
strncpy(str, "123456789abcdef", 15);
❶ str[15] = '\0';
```

If the source string (the string literal "123456789abcdef" in this example) has fewer than 15 bytes, the null termination character will be copied, and the assignment ❶ is unnecessary. If the source string has 15 bytes or longer, adding this assignment ensures that the string is properly null terminated.

aligned_alloc

The `aligned_alloc` function is similar to the `malloc` function, except that it requires you to provide an alignment as well as a size for the allocated object. The function has the following signature, where `size` specifies the object's size

and alignment specifies its alignment:

```
void *aligned_alloc(size_t alignment, size_t size);
```

Although C requires the dynamically allocated memory from `malloc` to be sufficiently aligned for all standard types, including arrays and structures, you might occasionally need to override the compiler's default choices. You can use the `aligned_alloc` function to request stricter alignment than the default (in other words, a larger power of two). If the value of `alignment` is not a valid alignment supported by the implementation, the function fails by returning a null pointer. See [Chapter 2](#) for more information on alignment.

calloc

The `calloc` function allocates storage for an array of `nmemb` objects, each of whose size is `size` bytes. It has the following signature:

```
void *calloc(size_t nmemb, size_t size);
```

This function initializes the storage to all zero-valued bytes. These zero values might not be the same one used to represent floating-point zero or null-pointer constants. You can also use the `calloc` function to allocate storage for a single object, which can be thought of as an array of one element.

Internally, the `calloc` function works by multiplying `nmemb` by `size` to determine the required number of bytes to allocate. Historically, some `calloc` implementations failed to validate that these values wouldn't wrap around when multiplied. C23 requires this test, and modern implementations of `calloc` return a null pointer if the space cannot be allocated or if the product `nmemb * size` would wrap around.

realloc

The `realloc` function increases or decreases the size of previously allocated storage. It takes a pointer to memory allocated by an earlier call to `aligned_alloc`, `malloc`, `calloc`, or `realloc` (or a null pointer) and a size and has the following signature:

```
void *realloc(void *ptr, size_t size);
```

You can use the `realloc` function to grow or (less commonly) shrink the size of an array.

Avoiding Memory Leaks

To avoid introducing bugs when you use `realloc`, you should understand how this function is specified. If the newly allocated storage is larger than the old contents, `realloc` leaves the additional storage uninitialized. If `realloc` succeeds in allocating the new object, it calls `free` to deallocate the old object. The pointer to the new object may have the same value as a pointer to the old object. If the allocation fails, the `realloc` function retains the old object data at the same address and returns a null pointer. A call to `realloc` can fail, for example, when insufficient memory is available to allocate the requested number of bytes. The following use of `realloc` might be erroneous:

```
size += 50;
if ((p = realloc(p, size)) == nullptr) return nullptr;
```

In this example, `size` is incremented by 50 before calling `realloc` to increase the size of the storage that `p` references. If the call to `realloc` fails, `p` is assigned a null pointer value, but `realloc` doesn't deallocate the storage that `p` references, resulting in this memory being leaked.

[Listing 6-3](#) demonstrates the correct use of the `realloc` function.

```
void *p = malloc(100);
void *p2;

// --snip--
if ((nsize == 0) || (p2 = realloc(p, nsize)) == nullptr) {
    free(p);
    return nullptr;
}
p = p2;
```

Listing 6-3: An example of the correct use of the `realloc` function

[Listing 6-3](#) declares two variables, `p` and `p2`. The variable `p` refers to the dynamically allocated memory `malloc` returns, and `p2` starts out uninitialized. Eventually, this memory is resized, which we accomplish by calling the `realloc` function with the `p` pointer and the new `nsize` size. The return value from `realloc` is assigned to `p2` to avoid overwriting the pointer stored in `p`. If `realloc`

returns a null pointer, the memory `p` references is freed, and the function returns a null pointer. If `realloc` succeeds and returns a pointer to an allocation of size `newsize`, `p` is assigned the pointer to the newly reallocated storage, and execution continues.

This code also includes a test for a zero-byte allocation. Avoid passing the `realloc` function a value of 0 as the size argument, as that is undefined behavior (as clarified in C23).

If the following call to the `realloc` function doesn't return a null pointer, the address stored in `p` is invalid and can no longer be read:

```
newp = realloc(p, ...);
```

In particular, the following test is not allowed:

```
if (newp != p) {  
    // update pointers to reallocated memory  
}
```

Any pointers that reference the memory `p` previously pointed to must be updated to reference the memory `newp` pointed to after the call to `realloc` regardless of whether `realloc` kept the same address for the storage.

One solution to this problem is to go through an extra indirection, sometimes called a *handle*. If all uses of the reallocated pointer are indirect, they'll all be updated when that pointer is reassigned.

Calling `realloc` with a Null Pointer

Calling `realloc` with a null pointer is equivalent to calling `malloc`. Provided `newsize` isn't equal to 0, we can replace the following code

```
if (p == nullptr)  
    newp = malloc(newsize);  
else  
    newp = realloc(p, newsize);
```

with this:

```
newp = realloc(p, newsize);
```

The first, longer version of this code calls `malloc` for the initial allocation and `realloc` to adjust the size later as required. Because calling `realloc` with a null pointer is equivalent to calling `malloc`, the second version concisely accomplishes the same thing.

reallocarray

As we have seen in previous chapters, both signed integer overflow and unsigned integer wraparound are serious problems that can result in buffer overflows and other security vulnerabilities. In the following code snippet, for example, the expression `num * size` might wrap around before being passed as the size argument in the call to `realloc`:

```
if ((newp = realloc(p, num * size)) == nullptr) {  
    // --snip--
```

The OpenBSD `reallocarray` function can reallocate storage for an array, but like `calloc`, it checks for wraparound during array size calculations, which saves you from having to perform these checks. The `reallocarray` function has the following signature:

```
void *reallocarray(void *ptr, size_t nmemb, size_t size);
```

The `reallocarray` function allocates storage for `nmemb` members of size `size` and checks for wraparound in the `nmemb * size` calculation. Other platforms, including the GNU C Library (libc), have adopted this function, and it has been proposed for inclusion in the next revision of the POSIX standard. The `reallocarray` function does not zero out the allocated storage.

The `reallocarray` function is useful when two values are multiplied to determine the size of the allocation:

```
if ((newp = reallocarray(p, num, size)) == nullptr) {  
    // --snip--
```

This call to the `reallocarray` function will fail and return a null pointer if `num * size` would wrap around.

free

When it's no longer needed, memory can be deallocated using the `free` function. Deallocating memory allows that memory to be reused, reducing the chances that you'll exhaust the available memory and often providing more efficient use of the heap.

We can deallocate memory by passing a pointer to that memory to the `free` function, which has the following signature:

```
void free(void *ptr);
```

The `ptr` value must have been returned by a previous call to `malloc`, `aligned_alloc`, `calloc`, or `realloc`. CERT C rule MEM34-C, “Only free memory allocated dynamically,” discusses what happens when the value is not returned. Memory is a limited resource and so must be reclaimed.

If we call `free` with a null-pointer argument, nothing happens, and the `free` function simply returns:

```
char *ptr = nullptr;  
free(ptr);
```

Freeing the same pointer twice, on the other hand, is a serious error.

free_sized

C23 introduced two new memory deallocation functions. The `free_sized` function has the following signature:

```
void free_sized(void *ptr, size_t size);
```

If `ptr` is a null pointer or the result obtained from a call to `malloc`, `realloc`, or `calloc`, where `size` is equal to the requested allocation size, this function behaves the same as `free(ptr)`. You cannot pass the result of `aligned_alloc` to this function; you must use the `free_aligned_sized` function (described in the next section). By *reminding* the allocator of the size of that allocation, you can reduce deallocation cost and allow extra security-hardening functionality. However, if you specify the size incorrectly, the behavior is undefined.

Using the `free_sized` function, we could improve the performance and safety of the following code

```
void *buf = malloc(size);
use(buf, size);
free(buf);
```

by rewriting it as:

```
void *buf = malloc(size);
use(buf, size);
free_sized(buf, size);
```

This is feasible and practical when the size of the allocation is retained or can be inexpensively re-created.

free_aligned_sized

The second of the two new memory deallocation functions that C23 introduced is the `free_aligned_sized` function. The `free_aligned_sized` function has the following signature:

```
void free_aligned_sized(void *ptr, size_t alignment, size_t size)
;
```

If `ptr` is a null pointer or the result obtained from a call to `aligned_alloc`, where `alignment` is equal to the requested allocation alignment and `size` is equal to the requested allocation size, this function is equivalent to `free(ptr)`. Otherwise, the behavior is undefined. In other words, this function may be used only for deallocating explicitly aligned memory.

Using the `free_aligned_sized` function, we could improve the performance and safety of the following code

```
void *aligned_buf = aligned_alloc(alignment, size);
use_aligned(buf, size, alignment);
free(buf);
```

by rewriting it as:

```
void *aligned_buf = aligned_alloc(size, alignment);
use_aligned(buf, size, alignment);
free_aligned_sized(buf, alignment, size);
```

This is feasible and practical when the alignment and size of the allocation is retained or can be inexpensively re-created.

Dealing with Dangling Pointers

If you call one of the free functions on the same pointer more than once, undefined behavior occurs. These defects can result in a security flaw known as a *double-free vulnerability*. One consequence is that they might be exploited to execute arbitrary code with the permissions of the vulnerable process. The full effects of double-free vulnerabilities are beyond the scope of this book, but I discuss them in detail in *Secure Coding in C and C++* (Seacord 2013). Double-free vulnerabilities are especially common in error-handling code, as programmers attempt to free allocated resources.

Another common error is to access memory that has already been freed. This type of error frequently goes undetected because the code might appear to work but then fails in an unexpected manner away from the actual error. In [Listing 6-4](#), taken from an actual application, the argument to `close` is invalid because the second call to `free` has reclaimed the storage `dirp` formerly pointed to.

```
#include <dirent.h>
#include <stdlib.h>
#include <unistd.h>

int closedir(DIR *dirp) {
    free(dirp->d_buf);
    free(dirp);
    return close(dirp->d_fd); // dirp has already been freed
}
```

Listing 6-4: Accessing already freed memory

We refer to pointers to already freed memory as *dangling pointers*. Dangling pointers are a potential source of errors (like a banana peel on the floor). Every use of a dangling pointer (not just dereferencing) is undefined behavior. When used to access memory that has already been freed, dangling pointers can result in use-after-free vulnerabilities (CWE 416). When passed to the `free` function, dangling pointers can result in double-free vulnerabilities (CWE 415). See CERT C rule MEM30-C, “Do not access freed memory,” for more information on these topics.

Setting the Pointer to Null

To limit the opportunity for defects involving dangling pointers, set the pointer to `nullptr` after completing a call to `free`:

```
char *ptr = malloc(16);  
// --snip--  
free(ptr);  
ptr = nullptr;
```

Any future attempt to dereference the pointer will usually result in a crash, increasing the likelihood that the error is detected during implementation and testing. If the pointer is set to `nullptr`, the memory can be freed multiple times without consequence. Unfortunately, the `free` function cannot set the pointer to `nullptr` itself because it's passed a copy of the pointer and not the actual pointer.

Memory States

Dynamically allocated memory can exist in one of three states shown in [Figure 6-1](#): unallocated and uninitialized within the memory manager, allocated but uninitialized, and allocated and initialized. Calls to the `malloc` and `free` functions, as well as writing the memory, cause the memory to transition between states.




Figure 6-1: Memory states

Different operations are valid depending on the state of the memory. Avoid any operation on memory that's not shown as valid or explicitly listed as invalid. Following execution of the `memset` function in this code snippet

```
char *p = malloc(100);  
memset(p, 0, 50);
```

the first 50 bytes are allocated and initialized, while the last 50 bytes are allocated but uninitialized. Initialized bytes can be read, but uninitialized bytes must not be read.

Flexible Array Members

Allocating storage for a structure that contains an array has always been a little tricky in C. There's no problem if the array has a fixed number of elements, as the size of the structure can easily be determined. Developers, however, frequently need to declare an array whose size is unknown until runtime, and originally, C offered no straightforward way to do so.

Flexible array members let you declare and allocate storage for a structure with any number of fixed members, where the last member is an array of unknown size. Starting with C99, the last member of a struct with more than

one member can have an *incomplete array type*, which means that the array has an unknown size that you can specify at runtime. A flexible array member allows you to access a variable-length object.

For example, [Listing 6-5](#) shows the use of a flexible array member data in widget. We dynamically allocate storage for the object by calling the malloc function.

```
#include <stdlib.h>

constexpr size_t max_elem = 100;

typedef struct {
    size_t num;
    ❶ int data[];
} widget;

widget *alloc_widget(size_t num_elem) {
    if (num_elem > max_elem) return nullptr;
    ❷ widget *p = (widget *)malloc(sizeof(widget) + sizeof(int) * num_elem);
    if (p == nullptr) return nullptr;

    p->num = num_elem;
    for (size_t i = 0; i < p->num; ++i) {
        ❸ p->data[i] = 17;
    }
    return p;
}
```

Listing 6-5: Flexible array members

We first declare a struct whose last member, the data array ❶, is an incomplete type (with no specified size). We then allocate storage for the entire struct ❷. When computing the size of a struct containing a flexible array member using the sizeof operator, the flexible array member is ignored. Therefore, we must explicitly include an appropriate size for the flexible array member when allocating storage. To accomplish that, we allocate additional bytes for the array by multiplying the number of elements in the array (num_elem) by the size of each element (sizeof(int)). This program assumes that the value of num_elem is such that when multiplied by sizeof(int), wraparound won't occur.

We can access this storage by using a `.` or `->` operator ❸, as if the storage had been allocated as `data[num_elem]`. See CERT C rule MEM33-C, “Allocate and copy structures containing a flexible array member dynamically,” for more information on allocating and copying structures containing flexible array members.

Prior to C99, multiple compilers supported a similar “struct hack” using a variety of syntaxes. CERT C rule DCL38-C, “Use the correct syntax when declaring a flexible array member, ” is a reminder to use the syntax specified in C99 and later versions of the C standard.

Other Dynamically Allocated Storage

C has language and library features beyond the memory management functions that support dynamically allocated storage. This storage is typically allocated in the stack frame of the caller (the C standard does not define a stack, but it’s a common implementation feature). A *stack* is a last-in-first-out (LIFO) data structure that supports nested invocation of functions at runtime. Each function invocation creates a *stack frame* in which local variables (of automatic storage duration) and other data specific to that invocation of the function can be stored.

alloca

For performance reasons, `alloca` (a nonstandard function supported by some implementations) allows dynamic allocation at runtime from the stack rather than the heap. This memory is automatically released when the function that called `alloca` returns. The `alloca` function is an *intrinsic* (or *built-in*) function, which is specially handled by the compiler. This allows the compiler to substitute a sequence of automatically generated instructions for the original function call. For example, on the x86 architecture, the compiler substitutes a call to `alloca` with a single instruction to adjust the stack pointer to accommodate the additional storage.

The `alloca` function originated in an early version of the Unix operating system from Bell Laboratories but is not defined by the C standard library or POSIX. [Listing 6-6](#) shows an example function called `prnterr` that uses the `alloca` function to allocate storage for an error string before printing it out to `stderr`.

```
void prnterr(errno_t errnum) {
```

```

❶ rsize_t size = strerrorlen_s(errno) + 1;
❷ char *msg = (char *)alloca(size);
  if (❸ strerror_s(msg, size, errno) != 0) {
    ❹ fputs(msg, stderr);
  }
  else {
    ❺ fputs("unknown error", stderr);
  }
}

```

Listing 6-6: The `printerr` function

The `printerr` function takes a single argument, `errno`, of `errno_t` type. We call the `strerrorlen_s` function ❶ to determine the length of the error string associated with this error number. Once we know the size of the array that we need to allocate to hold the error string, we can call the `alloca` function ❷ to efficiently allocate storage for the array. We then retrieve the error string by calling the `strerror_s` function ❸ and store the result in the newly allocated storage `msg` references. Assuming the `strerror_s` function succeeds, we output the error message ❹; otherwise, we output `unknown error` ❺. This `printerr` function is written to demonstrate the use of `alloca` and is more complicated than it needs to be.

The `alloca` function can be tricky to use. First, the call to `alloca` can make allocations that exceed the bounds of the stack. However, the `alloca` function doesn't return a null pointer value, so there's no way to check for the error. For this reason, it's critically important to avoid using `alloca` with large or unbounded allocations. The call to `strerrorlen_s` in this example should return a reasonable allocation size.

A further problem with the `alloca` function is that programmers may become confused by having to free memory allocated by `malloc` but not `alloca`. Calling `free` on a pointer not obtained by calling `aligned_alloc`, `calloc`, `realloc`, or `malloc` is a serious error. Due to those issues, the use of `alloca` is discouraged.

Both GCC and Clang provide a `-walloca` compiler flag that diagnoses all calls to the `alloca` function. GCC also provides a `-walloca-larger-than=size` compiler flag that diagnoses any call to the `alloca` function when the requested memory is more than `size`.

Variable-Length Arrays

Variable-length arrays (VLAs) were introduced in C99. A VLA is an object of a variably modified type (covered in [Chapter 2](#)). Storage for the VLA is allocated at runtime and is equal to the size of the base type of the variably modified type multiplied by the runtime extent.

The size of the array cannot be modified after you create it. All VLA declarations must be at *block scope*.

The following example declares the VLA `vla` of size `size` as an automatic variable in function `func`:

```
void func(size_t size) {  
    int vla[size];  
    // --snip--  
}
```

VLAs are useful when you don't know the number of elements in the array until runtime. Unlike the `alloca` function, VLAs are freed when the corresponding block ends, just like any other automatic variable. [Listing 6-7](#) replaces the call to `alloca` in the `printerr` function from [Listing 6-6](#) with a VLA. The change modifies just a single line of code (shown in bold).

```
void print_error(int errnum) {  
    size_t size = strerrorlen_s(errnum) + 1;  
    char msg[size];  
    if (strerror_s(msg, size, errnum) != 0) {  
        fputs(msg, stderr);  
    }  
    else {  
        fputs("unknown error", stderr);  
    }  
}
```

Listing 6-7: The `print_error` function rewritten to use a VLA

The main advantage of using VLAs instead of the `alloca` function is that the syntax matches the programmer's model of how arrays with automatic storage duration work. VLAs work just like automatic variables (because they are). Another advantage is that memory does not accumulate while iterating (which can accidentally happen with `alloca` because memory is released at the end of the function).

VLAAs share some of the problems of the `alloca` function, in that they can attempt to make allocations that exceed the bounds of the stack. Unfortunately, there's no portable way to determine the remaining stack space to detect such an error. Also, the calculation of the array's size could wrap around when the size you provide is multiplied by the size of each element. For those reasons, it's important to validate the size of the array before declaring it to avoid overly large or incorrectly sized allocations. This can be especially important in functions that are called with untrusted inputs or are called recursively, because a complete new set of automatic variables for the functions (including these arrays) will be created for each recursion. Untrusted inputs must be validated before being used for any allocations, including from the heap.

You should determine whether you have sufficient stack space in the worst-case scenario (maximum-sized allocations with deep recursions). On some implementations, it's also possible to pass a negative size to the VLA, so make sure your size is represented as a `size_t` or other unsigned type. See CERT C rule ARR32-C, "Ensure size arguments for variable-length arrays are in a valid range," for more information. VLAAs reduce stack usage when compared to using worst-case fixed-sized arrays.

The following file-scope declarations demonstrate another confusing aspect of VLAAs:

```
static const unsigned int num_elem = 12;
double array[num_elem];
```

Is this code valid? If `array` is a VLA, then the code is invalid, because the declaration is at file scope. If `array` is a constant-sized array, then the code is valid. GCC currently rejects this example because `array` is a VLA. However, C23 allows implementations to extend the definition of an integer constant expression, which Clang does by making `array` a constant-sized array.

We can rewrite these declarations to be portable using `constexpr` on all C23-conforming implementations:

```
constexpr unsigned int num_elem = 12;
double array[num_elem];
```

Finally, another interesting and unexpected behavior occurs when calling `sizeof` on a VLA. The compiler usually performs the `sizeof` operation at

compile time. However, if the expression changes the size of the array, it will be evaluated at runtime, including any side effects. The same is true of typedef, as the program in [Listing 6-8](#) shows.

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    size_t size = 12;
    (void)(sizeof(size++));
    printf("%zu\n", size); // prints 12
    (void)sizeof(int[size++]);
    printf("%zu\n", size); // prints 13
    typedef int foo[size++];
    printf("%zu\n", size); // prints 14
    typeof(int[size++]) f;
    printf("%zu\n", size); // prints 15
    return EXIT_SUCCESS;
}
```

Listing 6-8: Unexpected side effects

In this simple test program, we declare a variable `size` of type `size_t` and initialize it to 12. The operand in `sizeof(size++)` isn't evaluated because the type of the operand is not a VLA. Consequently, the value of `size` doesn't change. We then call the `sizeof` operator with `int[size++]` as the argument. Because this expression changes the size of the array, `size` is incremented and is now equal to 13. The `typedef` similarly increments the value of `size` to 14. Finally, we declare `f` to be of `typeof(int[size++]`), which further increments `size`. Because these behaviors aren't well understood, avoid using expressions with side effects with the `typeof` operator, `sizeof` operators, and `typedefs` to improve understandability.

Debugging Allocated Storage Problems

As noted earlier in this chapter, improper memory management can lead to errors like leaking memory, reading from or writing to freed memory, and freeing memory more than once. One way to avoid some of these problems is to set pointers to a null pointer value after calling `free`, as we've already discussed. Another strategy is to keep your dynamic memory management as simple as possible. For example, you should allocate and free memory in the same module,

at the same level of abstraction, rather than freeing memory in subroutines, which leads to confusion about if, when, and where memory is freed.

A third option is to use *dynamic analysis tools*, such as AddressSanitizer, Valgrind, or dmalloc to detect and report memory errors. AddressSanitizer, as well as general approaches to debugging, testing, and analysis, are discussed in [Chapter 11](#), while dmalloc is covered in this section. AddressSanitizer or Valgrind are effective tools and better choices if they are available for your environment.

dmalloc

The *debug memory allocation* (dmalloc) library Gray Watson created replaces malloc, realloc, calloc, free, and other memory management features with routines that provide debugging facilities that you can configure at runtime. The library has been tested on a variety of platforms.

Follow the installation directions provided at <https://dmalloc.com> to configure, build, and install the library. [Listing 6-9](#) contains a short program that prints out usage information and exits (it would typically be part of a longer program). This program has several intentional errors and vulnerabilities.

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

#ifdef DMALLOC
#include "dmalloc.h"
#endif

void usage(char *msg) {
    fprintf(stderr, "%s", msg);
    free(msg);
    return;
}

int main(int argc, char *argv[]) {
    if (argc != 3 && argc != 4) {
        // the error message is less than 80 chars
        char *errmsg = (char *)malloc(80);
        sprintf(
            errmsg,
            "Sorry %s,\nUsage: caesar secret_file keys_file [output_fil
e]\n",
            getenv("USER")
```

```
    );  
    usage(errmsg);  
    free(errmsg);  
    return EXIT_FAILURE;  
}  
// --snip--  
  
return EXIT_SUCCESS;  
}
```

Listing 6-9: Catching a memory bug with dmalloc

Recent versions of glibc will detect at least one of the vulnerabilities from this program:

```
Sorry (null),  
Usage: caesar secret_file keys_file [output_file]  
free(): double free detected in tcache 2  
Program terminated with signal: SIGSEGV
```

After repairing this error, include the lines shown in bold font to allow dmalloc to report the file and line numbers of calls that cause problems.

I'll show the output later, but we need to discuss a few things first. The dmalloc distribution also comes with a command line utility. You can run the following command to get further information on how to use this utility:

```
% dmalloc --usage
```

Before debugging your program with dmalloc, enter the following at the command line:

```
% dmalloc -l logfile -i 100 low
```

This command sets the logfile name to *logfile* and instructs the library to perform a check after 100 invocations, as the *-i* argument specifies. If you specify a larger number as the *-i* argument, dmalloc will check the heap less often, and your code will run faster; lower numbers are more likely to catch memory problems. The third argument enables a low number of debug features. Other options include runtime for minimal checking or medium or high for more extensive heap verification.

After executing this command, we can compile the program by using GCC as follows:

```
% gcc -DDMALLOC caesar.c -ocaesar -ldmalloc
```

When you run the program, you should see the following error:

```
% ./caesar
Sorry student,
Usage: caesar secret_file keys_file [output_file]
debug-malloc library: dumping program, fatal error
  Error: tried to free previously freed pointer (err 61)
Aborted (core dumped)
```

And if you examine the logfile, you'll find the following information:

```
% more logfile
1571549757: 3: Dmalloc version '5.5.2' from 'https://dmalloc.com/'
1571549757: 3: flags = 0x4e48503, logfile 'logfile'
1571549757: 3: interval = 100, addr = 0, seen # = 0, limit = 0
1571549757: 3: starting time = 1571549757
1571549757: 3: process pid = 29531
1571549757: 3:   error details: finding address in heap
1571549757: 3:   pointer '0x7ff010812f88' from 'caesar.c:29' prev
  access 'unknown'
1571549757: 3: ERROR: free: tried to free previously freed pointer (err 61)
```

These messages indicate that we've attempted to free the storage referenced by `errmsg` twice, first in the `usage` function and then in `main`, which constitutes a double-free vulnerability. Of course, this is just a single example of the types of bugs that `dmalloc` can detect, and other defects exist in the simple program we are testing.

Safety-Critical Systems

Systems with high safety requirements frequently ban the use of dynamic memory because memory managers can have unpredictable behavior that significantly impacts performance. Forcing all applications to live within a fixed, preallocated area of memory can eliminate these problems and make it easier to

verify memory use. In the absence of recursion, `alloca`, and VLAs (also prohibited in safety-critical systems), an upper bound on the use of stack memory can be derived statically, making it possible to prove that sufficient storage exists to execute the functionality of the application for all possible inputs.

Both GCC and Clang have a `-wvla` flag that warns if a VLA is used. GCC also has a `-wvla-larger-than=byte-size` flag that warns for declarations of VLAs whose size is either unbounded or bounded by an argument that allows the array size to exceed *byte-size* bytes.

EXERCISES

1. Repair the use-after-free defect from [Listing 6-4](#).
2. Use `dmalloc` to perform additional testing of the program from [Listing 6-9](#). Try varying inputs to the program to identify other memory management defects.

Summary

In this chapter, you learned about working with memory that has allocated storage duration and how this differs from objects of either automatic or static storage duration. We described the heap and memory managers and each of the standard memory management functions. We identified some common causes of errors when using dynamic memory, such as leaks and double-free vulnerabilities, and introduced some mitigations to help avoid these problems.

We also covered some more specialized memory allocation topics such as flexible array members, the `alloca` function, and VLAs. We concluded the chapter with a discussion of debugging allocated storage problems by using the `dmalloc` library.

In the next chapter, you'll learn about characters and strings.

7

CHARACTERS AND STRINGS



Strings are such an important and useful data type that nearly every programming language implements them in some form. Often used to represent text, strings constitute most of the data exchanged between an end user and a program, including text input fields, command line arguments, environment variables, and console input.

In C, the string data type is modeled on the idea of a formal string (Hopcroft and Ullman 1979):

Let Σ be a non-empty finite set of characters, called the alphabet. A string over Σ is any finite sequence of characters from Σ . For example, if $\Sigma = \{0, 1\}$, then 01011 is a string over Σ .

In this chapter, we'll talk about the various character sets, including ASCII and Unicode, that can be used to compose strings (the *alphabet* from the formal definition). We'll cover how strings are represented and manipulated using the legacy functions from the C standard library, the bounds-checked interfaces, and POSIX and Windows application programming interfaces (APIs).

Characters

The characters that people use to communicate aren't naturally understood by digital systems, which operate on bits. To process characters, digital systems use *character encodings* that assign unique integer values, referred to as *code points*, to designate specific characters. As you'll see, there are multiple ways to encode

the same notional character in your program. Common standards used by C implementations for encoding characters include Unicode, ASCII, Extended ASCII, ISO 8859-1 (Latin-1), Shift-JIS, and EBCDIC.

NOTE

The C standard explicitly references Unicode and ASCII.

ASCII

The *7-bit American Standard Code for Information Interchange*, better known as *7-bit ASCII*, specifies a set of 128 characters and their coded representation (ANSI X3.4-1986). Characters from `0x00` to `0x1f` and character `0x7f` are control characters, such as null, backspace, horizontal tab, and DEL. Characters from `0x20` through `0x7e` are all printable characters such as letters, digits, and symbols.

We often refer to this standard with the updated name *US-ASCII* to clarify that this system was developed in the United States and focuses on the typographical symbols predominantly used in this country. Most modern character encoding schemes are based on US-ASCII, although they support many additional characters.

Characters in the `0x80` to `0xFF` range are not defined by US-ASCII but are part of the 8-bit character encoding known as *Extended ASCII*. Numerous encodings for these ranges exist, and the actual mapping depends on the code page. A *code page* is a character encoding that maps a set of printable characters and control characters to unique numbers.

Unicode

Unicode has become the universal character encoding standard for representing text in computer processing. It supports a much wider range of characters than ASCII does; the current Unicode Standard (Unicode 2023) encodes characters in the range `U+0000` to `U+10FFFF`, which amounts to a 21-bit code space. An individual Unicode value is expressed as `U+` followed by four or more hexadecimal digits in printed text. The Unicode characters `U+0000` to `U+007F` are identical to those in US-ASCII, and the range `U+0000` to `U+00FF` is identical to ISO 8859-1, consisting of characters from the Latin script used throughout the Americas, Western Europe, Oceania, and much of Africa.

Unicode organizes code points into *planes*, which are continuous groups of 65,536 code points. There are 17 planes, identified by the numbers 0 to 16. The

most used characters, including those found in major, older encoding standards, have been placed into the first plane (0x0000 to 0xFFFF), which is called the *basic multilingual plane (BMP)*, or Plane 0.

Unicode also specifies several *Unicode transformation formats (UTFs)*, which are character encoding formats that assign each Unicode scalar value to a unique code unit sequence. A *Unicode scalar value* is any Unicode code point except high-surrogate and low-surrogate code points. (Surrogate pairs are explained later in this section.) A *code unit* is the minimal bit combination that can represent encoded text for processing or interchange. The Unicode standard defines three UTFs to allow for code units of various sizes:

UTF-8 Represents each character as a sequence of one to four 8-bit code units

UTF-16 Represents each character as a sequence of one or two 16-bit code units

UTF-32 Represents each character as a single 32-bit code unit

The UTF-8 encoding is the dominant encoding for POSIX operating systems. It has the following desirable properties:

- It encodes US-ASCII characters (U+0000 to U+007F) as single bytes in the range 0x00 to 0x7F. This means that files and strings that contain only 7-bit ASCII characters have the same encoding under both ASCII and UTF-8.
- Using a null byte to terminate a string (a topic we'll discuss later) works the same as for an ASCII string.
- All currently defined Unicode code points can be encoded using between 1 and 4 bytes.
- Unicode allows character boundaries to be easily identified by scanning for well-defined bit patterns in either direction.

On Windows, you can compile and link your programs with the Visual C++ `/utf8` flag to set the source and execution character sets as UTF-8. You'll also need to configure Windows to use Unicode UTF-8 for worldwide language support.

UTF-16 is currently the dominant encoding for Windows operating systems. Like UTF-8, UTF-16 is a variable-width encoding. As just mentioned, the BMP consists of characters from U+0000 to U+FFFF. Characters whose code points are greater than U+FFFF are called *supplementary characters*. Supplementary

characters are defined by a pair of code units called *surrogates*. The first code unit is from the high-surrogates range (U+D800 to U+DBFF), and the second code unit is from the low-surrogates range (U+DC00 to U+DFFF).

Unlike other variable-length UTFs, UTF-32 is a fixed-length encoding. The main advantage of UTF-32 is that the Unicode code points can be directly indexed, meaning that you can find the *n*th code point in a sequence of code points in constant time. In contrast, a variable-length encoding requires accessing each code point to find the *n*th code point in a sequence.

Source and Execution Character Sets

No universally accepted character encoding existed when C was originally standardized, so it was designed to work with a wide variety of character representations. Instead of specifying a character encoding like Java, each C implementation defines both a *source character set* in which source files are written and an *execution character set* used for character and string literals at compile time.

Both the source and execution character sets must contain encodings for the uppercase and lowercase letters of the Latin alphabet; the 10 decimal digits; 29 graphic characters; and the space, horizontal tab, vertical tab, form feed, and newline characters. The execution character set also includes alert, backspace, carriage return, and null characters.

Character conversion and classification functions (such as `isdigit`) are evaluated at runtime, based on the locale-determined encoding in effect at the time of the call. A program's *locale* defines its code sets, date and time formatting conventions, monetary conventions, decimal formatting conventions, and sort order.

Data Types

C defines several data types to represent character data, some of which we have already seen. C offers the unadorned `char` type to represent *narrow characters* (those that can be represented in as few as 8 bits) and the `wchar_t` type to represent *wide characters* (those that may require more than 8 bits).

`char`

As I have already mentioned, `char` is an integer type, but each implementation defines whether it's signed or unsigned. In portable code, you can assume neither.

Use the `char` type for character data (where signedness has no meaning) and not for integer data (where signedness is important). You can safely use the `char` type to represent 7-bit character encodings, such as US-ASCII. For these encodings, the high-order bits are always 0, so you don't have to be concerned about sign extension when a value of type `char` is converted to `int` and implementation defined as a signed type.

You can also use the `char` type to represent 8-bit character encodings, such as Extended ASCII, ISO/IEC 8859, EBCDIC, and UTF-8. These 8-bit character encodings can be problematic on implementations that define `char` as an 8-bit signed type. For example, the following code prints the string `end of file` when an EOF is detected:

```
char c = 'ÿ'; // extended character
if (c == EOF) puts("end of file");
```

Assuming the implementation-defined execution character set is ISO/IEC 8859-1, the Latin small letter `y` with diaeresis (`ÿ`) is defined to have the representation 255 (`0xFF`). For implementations in which `char` is defined as a signed type, `c` will be sign-extended to the width of signed `int`, making the `ÿ` character indistinguishable from EOF because they have the same representation.

A similar problem occurs when using the character classification functions defined in `<ctype.h>`. These library functions accept a character argument as an `int` or the value of the macro `EOF`. They return a nonzero value if the character belongs to the set of characters that the function's description defines and zero if the value doesn't belong to it. For example, the `isdigit` function tests whether the character is a decimal-digit character in the current locale. Any argument value that isn't a valid character or EOF results in undefined behavior.

To avoid undefined behavior when invoking these functions, cast `c` to `unsigned char` before the integer promotions, as shown here:

```
char c = 'ÿ';
if (isdigit((unsigned char)c)) {
    puts("c is a digit");
}
```

The value stored in `c` is zero-extended to the width of signed `int`, eliminating the undefined behavior by ensuring that the resulting value is representable as an `unsigned char`. Note that the initialization of `c` to `'ÿ'` may

result in a warning or error.

int

Use the `int` type for data that could be either EOF (a negative value) or character data interpreted as unsigned `char` and then converted to `int`. Functions that read character data from a stream, such as `fgetc`, `getc`, and `getchar`, return the `int` type. As we've seen, character-handling functions from `<ctype.h>` also accept this type because they might be passed the result of `fgetc` or related functions.

wchar_t

The `wchar_t` type is an integer type added to C to process the characters of a large character set. It can be a signed or unsigned integer type, depending on the implementation, and it has an implementation-defined inclusive range of `WCHAR_MIN` to `WCHAR_MAX`. Most implementations define `wchar_t` to be either a 16- or 32-bit unsigned integer type, but implementations that don't support localization may define `wchar_t` to have the same width as `char`. C does not permit a variable-length encoding for wide strings (despite UTF-16 being used this way in practice on Windows). Implementations can conditionally define the macro `__STDC_ISO_10646__` as an integer constant of the form `yyymmL` (for example, `199712L`) to mean that the `wchar_t` type is used to represent Unicode characters corresponding to the specified version of the standard.

Implementations that chose a 16-bit type for `wchar_t` cannot meet the requirements for defining `__STDC_ISO_10646__` for ISO/IEC 10646 editions more recent than Unicode 3.1. Consequently, the requirement for defining `__STDC_ISO_10646__` is either a `wchar_t` type larger than 20 bits or a 16-bit `wchar_t` and a value for `__STDC_ISO_10646__` earlier than `200103L`. The `wchar_t` type can be used for encodings other than Unicode, such as wide EBCDIC.

Writing portable code using `wchar_t` can be difficult because of the range of implementation-defined behavior. For example, Windows uses a 16-bit unsigned integer type, while Linux typically uses a 32-bit unsigned integer type. Code that calculates the lengths and sizes of wide-character strings is error prone and must be performed with care.

char16_t and char32_t

Other languages (including Ada95, Java, TCL, Perl, Python, and C#) have data

types for Unicode characters. C11 introduced the 16- and 32-bit character data types `char16_t` and `char32_t`, declared in `<uchar.h>`, to provide data types for UTF-16 and UTF-32 encodings, respectively. C doesn't provide standard library functions for the new data types, except for one set of character conversion functions. Without library functions, these types have limited usefulness.

C defines two environment macros that indicate how characters represented in these types are encoded. If the environment macro `__STDC_UTF_16__` has the value 1, values of type `char16_t` are UTF-16 encoded. If the environment macro `__STDC_UTF_32__` has the value 1, values of type `char32_t` are UTF-32 encoded. If the macro isn't defined, another implementation-defined encoding is used. Visual C++ does not define these macros, suggesting that you can't use `char16_t` on Windows for UTF-16.

Character Constants

C allows you to specify *character constants*, also known as *character literals*, which are sequences of one or more characters enclosed in single quotes, such as `'ÿ'`. Character constants allow you to specify character values in the source code of your program. [Table 7-1](#) shows the types of character constants that can be specified in C.

Table 7-1: Types of Character Constants

Prefix	Type
None	unsigned char
u8'a'	char8_t
L'a'	The unsigned type corresponding to wchar_t
u'a'	char16_t
U'a'	char32_t

The value of a character constant containing more than one character (for example, `'ab'`) is implementation defined. So is the value of a source character that cannot be represented as a single code unit in the execution character set. The earlier example of `'ÿ'` is one such case. If the execution character set is UTF-8, the value might be `0xC3BF` to reflect the UTF-8 encoding of the two code units needed to represent the `U+00FF` code-point value. C23 adds the `u8` prefix for character literals to represent a UTF-8 encoding. A UTF-8, UTF-16, or UTF-32 character constant cannot contain more than one character. The value must be representable with a single UTF-8, UTF-16, or UTF-32 code unit, respectively. Because UTF-8 encodes US-ASCII characters (`U+0000` to `U+007F`) as single

bytes in the range 0x00 to 0x7F, the u8 prefix can be used to create ASCII characters, even on implementations where the locale-dependent character encoding is some other encoding, such as EBCDIC.

Escape Sequences

The single quote (') and backslash (\) have special meanings, so they cannot be directly represented as characters. Instead, to represent the single quote, we use the escape sequence \', and to represent the backslash, we use \\. We can represent other characters, such as the question mark (?), and arbitrary integer values by using the escape sequences shown in [Table 7-2](#).

Table 7-2: Escape Sequences

Character	Escape sequence
Single quote	\'
Double quote	\"
Question mark	\?
Backslash	\\
Alert	\a
Backspace	\b
Form feed	\f
Newline	\n
Carriage return	\r
Horizontal tab	\t
Vertical tab	\v
Octal character	\<up to three octal digits>
Hexadecimal character	\<x hexadecimal digits>

The following nongraphical characters are represented by escape sequences consisting of the backslash followed by a lowercase letter: \a (alert), \b (backspace), \f (form feed), \n (newline), \r (carriage return), \t (horizontal tab), and \v (vertical tab).

Octal digits can be incorporated into an octal escape sequence to construct a single character for a character constant or a single wide character for a wide-character constant. The numerical value of the octal integer specifies the value of the desired character or wide character. *A backslash followed by numbers is always interpreted as an octal value.* For example, you can represent the backspace character (8 decimal) as the octal value \10 or, equivalently, \010.

You can also incorporate the hexadecimal digits that follow the \x to construct a single character or wide character for a character constant. The

numerical value of the hexadecimal integer forms the value of the desired character or wide character. For example, you can represent the backspace character as the hexadecimal value `\x8` or, equivalently, `\x08`.

Linux

Character encodings have evolved differently on various operating systems. Before UTF-8 emerged, Linux typically relied on various language-specific extensions of ASCII. The most popular of these were ISO 8859-1 and ISO 8859-2 in Europe, ISO 8859-7 in Greece, KOI-8/ISO 8859-5/CP1251 in Russia, EUC and Shift-JIS in Japan, and BIG5 in Taiwan. Linux distributors and application developers are phasing out these older legacy encodings in favor of UTF-8 to represent localized text strings (Kuhn 1999).

GCC has several flags that allow you to configure character sets. Here are a couple of flags you may find useful:

```
-fexec-charset=charset
```

The `-fexec-charset` flag sets the execution character set that's used to interpret string and character constants. The default is UTF-8. The *charset* can be any encoding supported by the system's `iconv` library routine described later in this chapter. For example, setting `-fexec-charset=IBM1047` instructs GCC to interpret string constants hardcoded in source code, such as `printf` format strings, according to EBCDIC code page 1,047.

To select the wide-execution character set, used for wide-string and character constants, use the `-fwide-exec-charset` flag:

```
-fwide-exec-charset=charset
```

The default is UTF-32 or UTF-16, corresponding to the width of `wchar_t`.

The input character set defaults to your system locale (or UTF-8 if the system locale is not configured). To overwrite the input character set used for translating the character set of the input file to the source character set used by GCC, use the `-finput-charset` flag:

```
-finput-charset=charset
```

Clang has `-fexec-charset` and `-finput-charset` but not `-fwide-exec-`

charset. Clang allows you to set *charset* to UTF-8 only and rejects any attempt to set it to something else.

Windows

Support for character encodings in Windows has irregularly evolved. Programs developed for Windows can handle character encodings using either Unicode interfaces or interfaces that rely implicitly on locale-dependent character encodings. For most modern applications, you should choose the Unicode interfaces by default to ensure that the application behaves as you expect when processing text. Generally, this code will have better performance, as narrow strings passed to Windows library functions are frequently converted to Unicode strings.

The main and wmain Entry Points

Visual C++ supports two entry points to your program: `main`, which allows you to pass narrow-character arguments, and `wmain`, which allows you to pass wide-character arguments. You declare formal parameters to `wmain` by using a similar format to `main`, as shown in [Table 7-3](#).

Table 7-3: Windows Program Entry Point Declarations

Narrow-character arguments	Wide-character arguments
<code>int main();</code>	<code>int wmain();</code>
<code>int main(int argc, char *argv[]);</code>	<code>int wmain(int argc, wchar_t *argv[]);</code>
<code>int main(int argc, char *argv[], char *envp[]);</code>	<code>int wmain(int argc, wchar_t *argv[], wchar_t *envp[]);</code>

For either entry point, the character encoding ultimately depends on the calling process. However, by convention, the `main` function receives its optional arguments and environment as pointers to text encoded with the current Windows (also called ANSI) code page, while the `wmain` function receives UTF-16-encoded text.

When you run a program from a shell such as the command prompt, the shell's command interpreter converts the arguments into the proper encoding for that entry point. A Windows process starts with a UTF-16-encoded command line. The startup code emitted by the compiler/linker calls the `CommandLineToArgvW` function to convert the command line to the `argv` form required to call `main` or passes the command line arguments directly to the `argv` form required to call `wmain`. In a call to `main`, the results are then transcoded to

the current Windows code page, which can vary from system to system. The ASCII character ? is substituted for characters that lack representation in the current Windows code page.

The Windows console uses an original equipment manufacturer (OEM) code page when writing data to the console. The actual encoding used varies from system to system but is often different from the Windows code page. For example, on a US English version of Windows, the Windows code page may be Windows Latin 1, while the OEM code page may be DOS Latin US. In general, writing textual data to stdout or stderr requires the text to be converted to the OEM code page first, or requires setting the console's output code page to match the encoding of the text being written out. Failure to do so may cause unexpected output to be printed to the console. However, even if you carefully match the character encodings between your program and the console, the console might still fail to display the characters as expected because of other factors, such as the current font selected for the console not having the appropriate glyphs required to represent the characters. Additionally, the Windows console has historically been unable to display characters outside of the Unicode BMP because it stores only a 16-bit value for character data for each cell.

Narrow vs. Wide Characters

There are two versions of all system APIs in the Win32 software development kit (SDK): a narrow Windows (ANSI) version with an A suffix and a wide-character version with a W suffix:

```
int SomeFuncA(LPSTR SomeString);  
int SomeFuncW(LPWSTR SomeString);
```

Determine whether your application is going to use wide (UTF-16) or narrow characters and then code accordingly. The best practice is to explicitly call the narrow- or wide-string version of each function and pass a string of the appropriate type:

```
SomeFuncW(L"String");  
SomeFuncA("String");
```

Examples of actual functions from the Win32 SDK include the MessageBoxA/MessageBoxW and CreateWindowExA/CreateWindowExW functions.

Character Conversion

Although international text is increasingly encoded in Unicode, it's still encoded in language- or country-dependent character encodings, making it necessary to convert between these encodings. Windows still operates in locales with traditional, limited character encodings, such as IBM EBCDIC and ISO 8859-1. Programs frequently need to convert between the Unicode and traditional encoding schemes when performing input/output (I/O).

It's not possible to convert all strings to each language- or country-dependent character encoding. This is obvious when the encoding is US-ASCII, which can't represent characters requiring more than 7 bits of storage. Latin-1 will never encode the character properly, and many kinds of non-Japanese letters and words cannot be converted to Shift-JIS without losing information.

C Standard Library

The C standard library provides a handful of functions to convert between narrow-code units (`char`) and wide-code units (`wchar_t`). The `mbtowc` (multibyte to wide character), `wctomb` (wide character to multibyte), `mbrtowc` (multibyte restartable to wide character), and `wcrtomb` (wide-character restartable to multibyte) functions convert one code unit at a time, writing the result to an output object or buffer. The `mbstowcs` (multibyte string to wide-character string), `wcstombs` (wide-character string to multibyte string), `mbsrtowcs` (multibyte-string restartable to multibyte string), and `wcsrtombs` (wide-character-string restartable to wide-character string) functions convert a sequence of code units, writing the result to an output buffer.

Conversion functions need to store data to properly process a sequence of conversions between function calls. The *nonrestartable* forms store the state internally and are consequently unsuitable for multithreaded processing. The *restartable* versions have an additional parameter that's a pointer to an object of type `mbstate_t` that describes the current conversion state of the associated multibyte character sequence. This object holds the state data that makes it possible to restart the conversion where it left off after another call to the function to perform an unrelated conversion. The *string* versions are for performing bulk conversions of multiple code units at once.

These functions have a few limitations. As discussed earlier, Windows uses 16-bit code units for `wchar_t`. This can be a problem, because the C standard requires an object of type `wchar_t` to be capable of representing any character in the current locale, and a 16-bit code unit can be too small to do so. C technically

doesn't allow you to use multiple objects of type `wchar_t` to represent a single character. Consequently, the standard conversion functions may result in a loss of data. On the other hand, most POSIX implementations use 32-bit code units for `wchar_t`, allowing the use of UTF-32. Because a single UTF-32 code unit can represent a whole code point, conversions using standard functions cannot lose or truncate data.

The C standards committee added the following functions to C11 to address the potential loss of data using standard conversion functions:

`mbrtoc16`, `c16rtomb` Converts between a sequence of narrow-code units and one or more `char16_t` code units

`mbrtoc32`, `c32rtomb` Converts a sequence of narrow-code units to one or more `char32_t` code units

The first two functions convert between locale-dependent character encodings, represented as an array of `char`, and UTF-16 data stored in an array of `char16_t` (assuming `__STDC_UTF_16__` has the value 1). The second two functions convert between the locale-dependent encodings and UTF-32 data stored in an array of `char32_t` encoded data (assuming `__STDC_UTF_32__` has the value 1). The program shown in [Listing 7-1](#) uses the `mbrtoc16` function to convert a UTF-8 input string to a UTF-16-encoded string.

```
#include <locale.h>
#include <uchar.h>
#include <stdio.h>
#include <wchar.h>

static_assert(__STDC_UTF_16__ == 1, "UTF-16 is not supported");
❶

size_t utf8_to_utf16(size_t utf8_size, const char utf8[utf8_size]
, char16_t *utf16) {
    size_t code, utf8_idx = 0, utf16_idx = 0;
    mbstate_t state = {0};
    while ((code = ❷ mbrtoc16(&utf16[utf16_idx],
        &utf8[utf8_idx], utf8_size - utf8_idx, &state))) {
        switch(code) {
            case (size_t)-1: // invalid code unit sequence detected
            case (size_t)-2: // code unit sequence missing elements
                return 0;
            case (size_t)-3: // high surrogate from a surrogate pair
                utf16_idx++;
```

```

        break;
    default:          // one value written
        utf16_idx++;
        utf8_idx += code;
    }
}
return utf16_idx + 1;
}

int main() {
    setlocale(LC_ALL, "es_MX.utf8"); ❸
    char utf8[] = u8"I ♥ s!";
    char16_t utf16[sizeof(utf8)]; // UTF-16 requires less code units
    size_t output_size = utf8_to_utf16(sizeof(utf8), utf8, utf16);
    printf("%s\nConverted to %zu UTF-16 code units: [", utf8, output_size);
    for (size_t x = 0; x < output_size; ++x) {
        printf("%#x ", utf16[x]);
    }
    puts("]");
}

```

Listing 7-1: Converting a UTF-8 string to a `char16_t` string with the `mbrtoc16` function

We call the `setlocale` function ❸ to set the multibyte character encoding to UTF-8 by passing an implementation-defined string. The static assertion ❶ ensures that the macro `__STDC_UTF_16__` has the value 1. (Refer to [Chapter 11](#) for more information on static assertions.) As a result, each call to the `mbrtoc16` function converts a single code point from a UTF-8 representation to a UTF-16 representation. If the resulting UTF-16 code unit is a high surrogate (from a surrogate pair), the state object is updated to indicate that the next call to `mbrtoc16` will write out the low surrogate without considering the input string.

There is no string version of the `mbrtoc16` function, so we loop through a UTF-8 input string iteratively, calling the `mbrtoc16` function ❷ to convert it to a UTF-16 string. In the case of an encoding error, the `mbrtoc16` function returns `(size_t)-1`, and if the code unit sequence is missing elements, it returns `(size_t)-2`. If either situation occurs, the loop terminates and the conversion ends.

A return value of `(size_t)-3` means that the function has output the high surrogate from a surrogate pair and then stored an indicator in the state

parameter. The indicator is used the next time the `mbrtoc16` function is called so it can output the low surrogate from a surrogate pair to form a complete `char16_t` sequence that represents a single code point. All restartable encoding conversion functions in the C standard behave similarly with the state parameter.

If the function returns anything other than `(size_t)-1`, `(size_t)-2`, or `(size_t)-3`, the `utf16_idx` index is incremented and the `utf8_idx` index is increased by the number of code units read by the function, and the conversion of the string continues.

libiconv

GNU `libiconv` is a commonly used cross-platform, open source library for performing string-encoding conversions. It includes the `iconv_open` function that allocates a conversion descriptor you can use to convert byte sequences from one character encoding to another. The documentation for this function (<https://www.gnu.org/software/libiconv/>) defines strings you can use to identify a particular *charset* such as ASCII, ISO-8859-1, SHIFT_JIS, or UTF-8 to denote the locale-dependent character encoding.

Win32 Conversion APIs

The Win32 SDK provides two functions for converting between wide- and narrow-character strings: `MultiByteToWideChar` and `WideCharToMultiByte`.

The `MultiByteToWideChar` function maps string data that's encoded in an arbitrary character code page to a UTF-16 string. Similarly, the `WideCharToMultiByte` function maps string data encoded in UTF-16 to an arbitrary character code page. Because UTF-16 data cannot be represented by all code pages, this function can specify a default character to use in place of any UTF-16 character that cannot be converted.

Strings

C doesn't support a primitive string type and likely never will. Instead, it implements strings as arrays of characters. C has two types of strings: narrow and wide.

A *narrow string* has the type array of `char`. It consists of a contiguous sequence of characters that includes a terminating null character. A pointer to a string references its initial character. The size of a string is the number of bytes allocated to the backing array storage. The length of a string is the number of code units (bytes) preceding the first null character. In [Figure 7-1](#), the size of the

string is 7, and the length of the string is 5. Elements of the backing array beyond the last element must not be accessed. Elements of the array that haven't been initialized must not be read.




Figure 7-1: A sample narrow string

A *wide string* has the type array of `wchar_t`. It's a contiguous sequence of wide characters that includes a terminating null wide character. A *pointer* to a wide string references its initial wide character. The *length* of a wide string is the number of code units preceding the first null wide character. [Figure 7-2](#) illustrates both the UTF-16BE (big-endian) and UTF-16LE (little-endian) representations of *hello*. The size of the array is implementation defined. This array is 14 bytes and assumes an implementation that has an 8-bit byte and 16-bit `wchar_t` type. The length of this string is 5, as the number of characters has not changed.




Figure 7-2: Sample UTF-16BE and UTF-16LE wide strings

Elements of the backing array beyond the last element must not be accessed. Elements of the array that haven't been initialized must not be read.

String Literals

A *character string literal* is a string constant represented by a sequence of zero or more multibyte characters enclosed in double quotes—for example, "ABC". You can use various prefixes to declare string literals of different character types:

- char string literal type, such as "ABC"
- wchar_t string literal type with L prefix, such as L"ABC"
- UTF-8 string literal type with u8 prefix, such as u8"ABC"
- char16_t string literal type with u prefix, such as u"ABC"
- char32_t string literal type with U prefix, such as U"ABC"

The C standard doesn't mandate that an implementation use ASCII for string literals. However, you can use the u8 prefix to force a string literal to be UTF-8 encoded, and if all the characters in the literal are ASCII characters, the compiler will produce an ASCII string literal, even if the implementation would normally encode string literals in another encoding (for example, EBCDIC).

A string literal has a non-const array type. Modifying a string literal is undefined behavior and prohibited by the CERT C rule STR30-C, “Do not attempt to modify string literals.” This is because these string literals may be stored in read-only memory, or multiple string literals may share the same memory, resulting in multiple strings being altered if one string is modified.

String literals often initialize array variables, which you can declare with an explicit bound that matches the number of characters in the string literal. Consider the following declaration:

```
#define S_INIT "abc"  
// --snip--  
const char s[4] = S_INIT;
```

The size of the array s is four, the exact size required to initialize the array to the string literal, including the space for a trailing null byte.

If you add another character to the string literal used to initialize the array, however, the meaning of the code changes substantially:

```
#define S_INIT "abcd"  
// --snip--  
const char s[4] = S_INIT;
```

The size of the array `s` remains four, although the size of the string literal is now five. As a result, the array `s` is initialized to the character array `"abcd"` with the trailing null byte omitted. By design, this syntax allows you to initialize a character array and not a string. Therefore, it's unlikely that your compiler will diagnose this declaration as an error.

There is some risk that if the string literal changes during maintenance, a string could unintentionally be changed to a character array with no terminating null character, particularly when the string literal is defined apart from the declaration, as in this example. If your intent is to always initialize `s` to a string, you should omit the array bound. If you don't specify the bound of the array, the compiler will allocate sufficient space for the entire string literal, including the terminating null character:

```
const char s[] = S_INIT;
```

This approach simplifies maintenance because the size of the array can always be determined even if the size of the string literal changes.

The size of arrays declared using this syntax can be determined at compile time by using the `sizeof` operator:

```
size_t size = sizeof(s);
```

If, instead, we declared this string as follows

```
const char *foo = S_INIT;
```

we would need to invoke the `strlen` function to get the length

```
size_t length = strlen(foo) + 1U;
```

which may incur a runtime cost and is different from the size.

String-Handling Functions

Several approaches can be used to manage strings in C, the first of which are the C standard library functions. Narrow-string-handling functions are defined in the `<string.h>` header file and wide-string-handling functions in `<wchar.h>`. These legacy string-handling functions have been associated in recent years with

various security vulnerabilities. This is because they don't check the size of the array (frequently lacking the information needed to perform such checks) and trust you to provide adequately sized character arrays to hold the output. While it's possible to write safe, robust, and error-free code using these functions, they promote programming styles that can result in buffer overflows if a result is too large for the provided array. These functions aren't inherently insecure but are prone to misuse and need to be used carefully (or not at all).

As a result, C11 introduced the normative (but optional) Annex K bounds-checking interfaces. This annex provides alternative library functions intended to promote safer, more secure programming by requiring you to provide the length of output buffers, for example, and validating that these buffers are adequately sized to contain the output from these functions. For instance, Annex K defines the `strcpy_s`, `strcat_s`, `strncpy_s`, and `strncat_s` functions as close replacements for the C standard library's `strcpy`, `strcat`, `strncpy`, and `strncat` functions.

<string.h> and <wchar.h>

The C standard library includes well-known functions such as `strcpy`, `strncpy`, `strcat`, `strncat`, `strlen`, and so forth, as well as the `memcpy` and `memmove` functions that you can use to copy and move strings, respectively. The C standard also provides a wide-character interface that operates on objects of type `wchar_t` instead of `char`. (These function names are like the narrow-string function names, except that `str` is replaced with `wcs`, and a `w` is added in front of the memory function names.) [Table 7-4](#) gives some examples of narrow- and wide-character string functions. Refer to the C standard (ISO/IEC 9899:2024) or man pages for more information on how to use these functions.

Table 7-4: Narrow- and Wide-String Functions

Narrow (char)	Wide (wchar_t)	Description
<code>strcpy</code>	<code>wscpy</code>	String copy
<code>strncpy</code>	<code>wcsncpy</code>	Truncated, zero-filled copy
<code>memcpy</code>	<code>wmemcpy</code>	Copies a specified number of nonoverlapping code units
<code>memmove</code>	<code>wmemmove</code>	Copies a specified number of (possibly overlapping) code units
<code>strcat</code>	<code>wscat</code>	Concatenates strings
<code>strncat</code>	<code>wcsncat</code>	Concatenates strings with truncation
<code>strcmp</code>	<code>wscmp</code>	Compares strings
<code>strncmp</code>	<code>wcsncmp</code>	Compares truncated strings
<code>strchr</code>	<code>wcschr</code>	Locates a character in a string

strcspn	wscspn	Computes the length of a complementary string segment
strdup	wcsdup	Copies string into allocated storage
strndup	N/A	Truncated copy into allocated storage
strpbrk	wcspbrk	Finds the first occurrence of a set of characters in a string
strchr	wcsrchr	Finds the first occurrence of a character in a string
strspn	wcsspn	Computes the length of a string segment
strstr	wcsstr	Finds a substring
strtok	wcstok	String tokenizer (modifies the string being tokenized)
memchr	wmemchr	Finds a code unit in memory
strlen	wcslen	Computes string length
memset	wmemset	Fills memory with a specified code unit
memset_explicit	N/A	Like memset but always performed

These string-handling functions are considered efficient because they leave memory management to the caller and can be used with both statically and dynamically allocated storage. In the next couple of sections, I'll go into more detail on some of the more commonly used functions.

NOTE

The `wcsdup` function listed in [Table 7-4](#) is not a C standard library function but is defined by POSIX.

Size and Length

As mentioned earlier in this chapter, strings have both a size (which is the number of bytes allocated to the backing array storage) and a length. You can determine the size of a statically allocated backing array at compile time by using the `sizeof` operator:

```
char str[100] = "Here comes the sun";
size_t str_size = sizeof(str); // str_size is 100
```

You can compute the length of a string by using the `strlen` function:

```
char str[100] = "Here comes the sun";
size_t str_len = strlen(str); // str_len is 18
```

The `wcslen` function computes the length of a wide string measured by the number of code units preceding the terminating null wide character:

```
wchar_t str[100] = L"Here comes the sun";  
size_t str_len = wcslen(str); // str_len is 18
```

The length is a count of something, but what exactly is being counted can be unclear. Here are some of the things that *could* be counted when taking the length of a string:

Bytes Useful when allocating storage.

Code units Number of individual code units used to represent the string. This length depends on encoding and can also be used to allocate memory.

Code points Code points (characters) can take up multiple code units. This value is not useful when allocating storage.

Extended grapheme cluster A group of one or more Unicode scalar values that approximates a single user-perceived character. Many individual characters, such as é, , and 🇸🇪, may be constructed from multiple Unicode scalar values. Unicode's boundary algorithms combine these code points into extended grapheme clusters.

The `strlen` and `wcslen` functions count code units. For the `strlen` function, this corresponds to the number of bytes. Determining the amount of storage required by using the `wcslen` function is more complicated because the size of the `wchar_t` type is implementation defined. [Listing 7-2](#) contains examples of dynamically allocating storage for both narrow and wide strings.

```
// narrow strings  
char *str1 = "Here comes the sun";  
char *str2 = malloc(strlen(str1) + 1);  
  
// wide strings  
wchar_t *wstr1 = L"Here comes the sun";  
wchar_t *wstr2 = malloc((wcslen(wstr1) + 1) * sizeof(*wstr1));
```

Listing 7-2: Dynamically allocating storage for narrow- and wide-string functions

For narrow strings, we can determine the size of the string by adding 1 to the return value of the `strlen` function to account for the terminating null character. For wide strings, we can determine the size of the string by adding 1 to the return value of the `wcslen` function to account for the terminating null wide character and then multiply the sum by the size of the `wchar_t` type. Because

`str1` and `wstr1` are declared as pointers (and not arrays), it's not possible to use the `sizeof` operator to obtain their sizes.

Code point or extended grapheme cluster counts cannot be used for storage allocation because they consist of an unpredictable number of code units. (For an interesting exposition on string length, see “It’s Not Wrong that “”.length == 7” at <https://hsivonen.fi/string-length/>.) Extended grapheme clusters are used to determine where to truncate a string, for example, because of a lack of storage. Truncation at extended grapheme cluster boundaries avoids slicing user-perceived characters.

Calling the `strlen` function can be an expensive operation because it needs to traverse the length of the array looking for a null character. The following is a straightforward implementation of the `strlen` function:

```
size_t strlen(const char * str) {
    const char *s;
    for (s = str; *s; ++s) {}
    return s - str;
}
```

The `strlen` function has no way of knowing the size of the object referenced by `str`. If you call `strlen` with an invalid string that lacks a null character before the bound, the function will access the array beyond its end, resulting in undefined behavior. Passing a null pointer to `strlen` will also result in undefined behavior (a null-pointer dereference). This implementation of the `strlen` function also has undefined behavior for strings larger than `PTRDIFF_MAX`. You should refrain from creating such objects (in which case this implementation is fine).

strcpy

Calculating the size of dynamically allocated memory is not always easy. One approach is to store the size when allocating and reuse this value later. The code snippet in [Listing 7-3](#) uses the `strcpy` function to make a copy of `str` by determining the length and then adding 1 to accommodate the terminating null character.

```
char str[100] = "Here comes the sun";
size_t str_size = strlen(str) + 1;
char *dest = (char *)malloc(str_size);
if (dest) {
```

```
    strcpy(dest, str);  
}  
else {  
    /* handle error */  
}
```

Listing 7-3: Copying a string

We can then use the value stored in `str_size` to dynamically allocate the storage for the copy. The `strcpy` function copies the string from the source string (`str`) to the destination string (`dest`), including the terminating null character. The `strcpy` function returns the address of the beginning of the destination string, which is ignored in this example.

The following is a simple implementation of the `strcpy` function:

```
char *strcpy(char *dest, const char *src) {  
    char *save = dest;  
    while ((*dest++ = *src++));  
    return save;  
}
```

This code saves a pointer to the destination string in `save` (to use as the return value) before copying all the bytes from the source to the destination array. The `while` loop terminates when the first null byte is copied. Because `strcpy` doesn't know the length of the source string or the size of the destination array, it assumes that all the function's arguments have been validated by the caller, allowing the implementation to simply copy each byte from the source string to the destination array without performing any checks.

Argument Checking

Argument checking can be performed by either the calling function or the called function. Redundant argument testing by both the caller and the callee is a largely discredited style of defensive programming. The usual discipline is to require validation on only one side of each interface.

The most time-efficient approach is for the caller to perform the check, because the caller should have a better understanding of the program state. In [Listing 7-3](#), we can see that the arguments to `strcpy` are valid without introducing further redundant tests: the variable `str` references a statically allocated array that was properly initialized in the declaration, and the `dest`

parameter is a valid non-null pointer referencing dynamically allocated storage of sufficient size to hold a copy of `str`, including the null character. Therefore, the call to `strcpy` is safe, and the copy can be performed in a time-efficient manner. This approach to argument checking is commonly used by C standard library functions because it adheres to the “spirit of C,” in that it’s optimally efficient and trusts the programmer (to pass valid arguments).

The safer, more secure approach is for the callee to check the arguments. This approach is less error-prone because the library function implementer validates the arguments, so we no longer need to trust the programmer to pass valid ones. The function implementer is usually in a better position to understand which arguments need to be validated. If the input validation code is defective, the repair needs to be made in only one place. All the code to validate the arguments is in one place, so this approach can be more space efficient. However, because these tests run even when unnecessary, they can also be less time efficient. Frequently, the caller of these functions will place checks before suspect calls that may or may not already perform similar checks. This approach would also impose additional error handling on callees that don’t currently return error indications but would presumably need to if they validated arguments. For strings, the called function can’t always determine whether the argument is a valid null-terminated string or points to sufficient space to make a copy.

The lesson here is don’t assume that the C standard library functions validate arguments unless the standard explicitly requires them to.

memcpy

The `memcpy` function copies `size` characters from the object referenced by `src` into the object referenced by `dest`:

```
void *memcpy(void * restrict dest, const void * restrict src, size_t size);
```

You can use the `memcpy` function instead of `strcpy` to copy strings when the size of the destination array is larger than or equal to the `size` argument to `memcpy`, the source array contains a null character before the bound, and the string length is less than `size - 1` (so that the resulting string will be properly null terminated). The best advice is to use `strcpy` when copying a string and `memcpy` when copying only raw, untyped memory. Also remember that the assignment (`=`) operator can efficiently copy objects in many cases.

memccpy

Most of the C standard library's string-handling functions return a pointer to the beginning of the string passed as an argument so that you can nest calls to string functions. For example, the following sequence of nested function calls constructs a full name using a Western naming order by copying, then concatenating, the constituent parts:

```
strcat(strcat(strcat(strcat(strcpy(name, first), " "), middle), "
"), last);
```

However, piecing together the array name from its constituent substrings requires name to be scanned many more times than necessary; it would have been more useful for the string-handling functions to return pointers to the *end* of the modified string to eliminate this need for rescanning. C23 introduced the `memccpy` function with a better interface design. POSIX environments should already provide this, but you may need to enable its declaration as follows:

```
#define _XOPEN_SOURCE 700
#include <string.h>
```

The `memccpy` function has the following signature:

```
void *memccpy(void * restrict s1, const void * restrict s2, int c
, size_t n);
```

Like the `memchr` function, `memccpy` scans the source sequence for the first occurrence of a character specified by one of its arguments. The character can have any value, including zero. It copies (at most) the specified number of characters from the source to the destination, without writing beyond the end of the destination buffer. Finally, it returns a pointer just past the copy of the specified character if it exists.

[Listing 7-4](#) reimplements the preceding sequence of nested string-handling function calls using the `memccpy` function. This implementation is more performant and secure.

```
#include <stdarg.h>
#include <string.h>
#include <stdio.h>
```



```

#include <stdint.h>

constexpr size_t name_size = 18U;

char *vstrcat(char *buff, size_t buff_length, ...) {
    char *ret = buff;
    va_list list;
    va_start(list, buff_length);
    const char *part = nullptr;
    size_t offset = 0;
    while ((part = va_arg(list, const char *)) {
        ❶ buff = (char *)memcpy(buff, part, '\0', buff_length - offset);
    }
    if (buff == nullptr) {
        ret[0] = '\0';
        break;
    }
    ❷ offset = --buff - ret;
    }
    va_end(list);
    return ret;
}

int main() {
    char name[name_size] = "";
    char first[] = "Robert";
    char middle[] = "C.";
    char last[] = "Seacord";

    puts(vstrcat(
        name, sizeof(name), first, " ",
        middle, " ", last, nullptr
    ));
}

```

Listing 7-4: String concatenation with `memcpy`

[Listing 7-4](#) defines the variadic function `vstrcat` that accepts a buffer (`buff`) and buffer length (`buff_length`) as fixed arguments and a variable list of string arguments. A null pointer is used as the sentinel value to indicate the end of the variable-length argument list. The `memcpy` function is invoked ❶ to concatenate each string to the buffer. As previously noted, `memcpy` returns a pointer just past the copy of the specified character, which in this case is the null termination character `'\0'`.

Instead of nesting calls, we call `memcpy` for each string argument passed to `vstrcat` and store the return value in `buff`. This allows you to concatenate directly to the end of the string instead of having to find the null termination character each time, making this solution more performant.

If `buff` is a null pointer, we couldn't copy the entire string. Instead of returning a partial name in this case, we return an empty string. This empty string can be printed or treated as an error condition.

Because the `memcpy` function returns a pointer to the character *after* the copy of the null byte, we decrement `buff` using the prefix decrement operator and then subtract the value stored in `ret` to obtain a new offset ❷. The size argument to the `memcpy` function (which it uses to prevent buffer overflow) is calculated by subtracting `offset` from `buff_length`. This is a more secure approach than nested function calls, which are always suspect because there is no way to check for an error.

memset, memset_s, and memset_explicit

The `memset` function copies the value of `(unsigned char)c` into each of the first `n` characters of the object pointed to by `s`:

```
void *memset(void *s, int c, size_t n);
```

The `memset` function is frequently used to clear memory—for example, to initialize memory allocated by `malloc` to zero. However, in the following example, it's used incorrectly:

```
void check_password() {
    char pwd[64];
    if (get_password(pwd, sizeof(pwd))) {
        /* check password */
    }
    memset(pwd, 0, sizeof(pwd));
}
```

One problem with the `get_password` function is that it uses the `memset` function to clear an automatic variable after it has been read for the last time. This is being done for security reasons to make sure that the sensitive information stored here is inaccessible. However, the compiler doesn't know that and may perform a “dead store” optimization. This is when a compiler notices a

write is not followed by a read, and just like this book, there is no sense in writing it if no one is going to read it. Consequently, the call to `memset` in the `get_password` function is likely to be removed by the compiler.

This problem was meant to be addressed in C11 by the `memset_s` function from Annex K bounds-checking interfaces (discussed in the next section). Unfortunately, this function has not been implemented by any compiler mentioned in this book.

To solve this problem again, C23 introduced the `memset_explicit` function for making sensitive information inaccessible. In contrast to the `memset` function, the intention is that the memory store is always performed (that is, never elided), regardless of optimizations.

gets

The `gets` function is a flawed input function that accepts input without providing any way to specify the size of the destination array. For that reason, it cannot prevent buffer overflows. As a result, the `gets` function was deprecated in C99 and eliminated from C11. However, it has been around for many years, and most libraries still provide an implementation for backward compatibility, so you may see it in the wild. You should *never* use this function, and you should replace any use of the `gets` function you find in any code you're maintaining.

Because the `gets` function is so bad, we'll spend some time examining why it's so awful. The function shown in [Listing 7-5](#) prompts the user to enter either `y` or `n` to indicate whether they'd like to continue.

```
#include <stdio.h>
#include <stdlib.h>
#include <ctype.h>

void get_y_or_n(void) {
    char response[8];
    puts("Continue? [y] n: ");
    gets(response);
    if (tolower(response[0]) == 'n') exit(EXIT_SUCCESS);
    return;
}
```

Listing 7-5: Misuse of the obsolete `gets` function

This function has undefined behavior if more than eight characters are entered at the prompt. This undefined behavior occurs because the `gets` function

has no way of knowing how large the destination array is and will write beyond the end of the array object.

[Listing 7-6](#) shows a simplified implementation of the `gets` function. As you can see, the caller of this function has no way to limit the number of characters read.

```
char *gets(char *dest) {
    int c;
    char *p = dest;
    while ((c = getchar()) != EOF && c != '\n') {
        *p++ = c;
    }
    *p = '\0';
    return dest;
}
```

Listing 7-6: A `gets` function implementation

The `gets` function iterates reading a character at a time. The loop terminates if either an EOF or newline `'\n'` character is read. Otherwise, the function will continue to write to the `dest` array without concern for the boundaries of the object.

[Listing 7-7](#) shows the `get_y_or_n` function from [Listing 7-5](#) with the `gets` function inlined.

```
#include <stdio.h>
#include <stdlib.h>
void get_y_or_n(void) {
    char response[8];
    puts("Continue? [y] n: ");
    int c;
    char *p = response;
    ❶ while ((c = getchar()) != EOF && c != '\n') {
        *p++ = c;
    }
    *p = '\0';
    if (response[0] == 'n')
        exit(0);
}
```

Listing 7-7: A poorly written `while` loop

The size of the destination array is now available, but the while loop ❶ doesn't use this information. You should ensure that reaching the bound of the array is a loop termination condition when reading or writing to an array in a loop such as this one.

Annex K Bounds-Checking Interfaces

C11 introduced the Annex K bounds-checking interfaces with alternative functions that verify that output buffers are large enough for the intended result and return a failure indicator if they aren't. These functions are designed to prevent writing data past the end of an array and to null-terminate all string results. These string-handling functions leave memory management to the caller, and memory can be statically or dynamically allocated before the functions are invoked.

Microsoft created the C11 Annex K functions to help retrofit its legacy code base in response to numerous, well-publicized security incidents in the 1990s. These functions were then proposed to the C standards committee for standardization, published as ISO/IEC TR 24731-1 (ISO/IEC TR 24731-1:2007) and then later incorporated into C11 as an optional annex. Despite the improved usability and security provided by these functions, they aren't yet widely implemented at the time of writing.

`gets_s`

The Annex K bounds-checking interface has a `gets_s` function we can use to eliminate the undefined behavior caused by the call to `gets` in [Listing 7-5](#), as shown in [Listing 7-8](#).

```
#define __STDC_WANT_LIB_EXT1__ 1
#include <stdio.h>
#include <stdlib.h>
#include <ctype.h>

void get_y_or_n(void) {
    char response[8];
    puts("Continue? [y] n: ");
    gets_s(response, sizeof(response));
    if (tolower(response[0]) == 'n') {
        exit(EXIT_SUCCESS);
    }
}
```

Listing 7-8: Use of the `gets_s` function

The two functions are similar, except that the `gets_s` function checks the array bounds. The default behavior when the maximum number of characters input is exceeded is implementation defined, but typically the `abort` function is called. You can change this behavior via the `set_constraint_handler_s` function, which I'll explain further in “Runtime Constraints” on [page 163](#).

You need to define `__STDC_WANT_LIB_EXT1__` as a macro that expands to the integer constant 1 before including the header files that define the bounds-checking interfaces to allow them to be used in your program. Unlike the `gets` function, the `gets_s` function takes a size argument. Consequently, the revised function calculates the size of the destination array by using the `sizeof` operator and passes this value as an argument to the `gets_s` function. The implementation-defined behavior is the result of the runtime-constraint violation.

strcpy_s

The `strcpy_s` function is a close replacement for the `strcpy` function defined in `<string.h>`. The `strcpy_s` function copies characters from a source string to a destination character array up to and including the terminating null character. Here's the `strcpy_s` signature:

```
errno_t strcpy_s(  
    char * restrict s1, rsize_t s1max, const char * restrict s2  
);
```

The `strcpy_s` function has an extra argument of type `rsize_t` that specifies the maximum length of the destination buffer. The `rsize_t` type is like the `size_t` type except that functions that accept an argument of this type test to make sure that the value is not greater than `RSIZE_MAX`. The `strcpy_s` function succeeds only when it can fully copy the source string to the destination without overflowing the destination buffer. The `strcpy_s` function verifies that the following runtime constraints aren't violated:

- Neither `s1` nor `s2` are null pointers.
- `s1max` is not greater than `RSIZE_MAX`.
- `s1max` does not equal zero.
- `s1max` is greater than `strnlen_s(s2, s1max)`.
- Copying does not take place between overlapping objects.

To perform the string copy in a single pass, a `strcpy_s` function implementation retrieves a character (or characters) from the source string and copies it to the destination array until it has copied the entire string or the destination array is full. If it can't copy the entire string and `s1max` is not zero, the `strcpy_s` function sets the first byte of the destination array to the null character, creating an empty string.

Runtime Constraints

Runtime constraints are violations of a function's runtime requirements that the function will detect and diagnose by a call to a handler. If this handler returns, the functions will return a failure indicator to the caller.

The bounds-checking interfaces enforce runtime constraints by invoking a runtime-constraint handler, which may simply return. Alternatively, the runtime-constraint handler might print a message to `stderr` and/or abort the program. You can control which handler function is called via the `set_constraint_handler_s` function and make the handler simply return as follows:

```
int main(void) {
    constraint_handler_t oconstraint =
        set_constraint_handler_s(ignore_handler_s);
    get_y_or_n();
}
```

If the handler returns, the function that identified the runtime-constraint violation and invoked the handler indicates a failure to its caller by using its return value.

The bounds-checking interface functions typically check the conditions either immediately upon entry or as they perform their tasks and gather sufficient information to determine whether a runtime constraint has been violated. The runtime constraints of the bounds-checking interfaces are conditions that would otherwise be undefined behavior for C standard library functions.

Implementations have a default constraint handler that they invoke if no calls to the `set_constraint_handler_s` function have been made. The default handler's behavior may cause the program to exit or abort, but implementations are encouraged to provide reasonable behavior by default. This allows, for example, compilers customarily used to implement safety-critical systems to not abort by default. You must check the return value of calls to functions that can

return and not simply assume their results are valid. Implementation-defined behavior can be eliminated by invoking the `set_constraint_handler_s` function before invoking any bounds-checking interfaces or using any mechanism that invokes a runtime-constraint handler.

Annex K provides the `abort_handler_s` and `ignore_handler_s` functions, which represent two common strategies for handling errors. The C implementation's default handler need not be either of these handlers.

POSIX

POSIX also defines several string-handling functions, such as `strdup` and `strndup` (IEEE Std 1003.1:2018), that provide another set of string-related APIs for POSIX-compliant platforms such as GNU/Linux and Unix (IEEE Std 1003.1:2018). Both functions were adopted into the C standard library by C23.

These replacement functions use dynamically allocated memory to ensure that buffer overflows don't occur, and they implement a *callee allocates, caller frees* model. Each function ensures that enough memory is available (except when a call to `malloc` fails). The `strdup` function, for example, returns a pointer to a new string that contains a duplicate of the argument. The returned pointer should be passed to the C standard `free` function to reclaim the storage when it's no longer needed.

[Listing 7-9](#) contains a code snippet that uses the `strdup` function to make a copy of the string returned by the `getenv` function.

```
const char *temp = getenv("TMP");
if (temp != nullptr) {
    char *tmpvar = strdup(temp);
    if (tmpvar != nullptr) {
        printf("TMP = %s.\n", tmpvar);
        free(tmpvar);
    }
}
```

Listing 7-9: Copying a string using the `strdup` function

The C standard library `getenv` function searches an environment list, provided by the host environment, for a string that matches the string referenced by a specified name (`TMP` in this example). Strings in this environment list are referred to as *environment variables* and provide an additional mechanism for communicating strings to a process. These strings don't have a well-defined

encoding but typically match the system encoding used for command line arguments, stdin, and stdout.

The returned string (the value of the variable) may be overwritten by a subsequent call to the `getenv` function, so it's a good idea to retrieve any environmental variable you need before creating any threads to eliminate the possibility of a race condition. If later use is anticipated, you should copy the string so the copy can be safely referenced as needed, as illustrated by the idiomatic example shown in [Listing 7-9](#).

The `strndup` function is equivalent to `strdup`, except that `strndup` copies, at most, $n + 1$ bytes into the newly allocated memory (while `strdup` copies the entire string) and ensures that the newly created string is always properly terminated.

These POSIX functions can help prevent buffer overflows by automatically allocating storage for the resulting strings, but this requires introducing additional calls to `free` when this storage is no longer needed. This means matching a call to `free` to each call to `strdup` or `strndup`, for example, which can be confusing to programmers who are more familiar with the behavior of the string functions defined by `<string.h>`.

Microsoft

Visual C++ provides all the string-handling functions defined by the C standard library up to C99 but doesn't implement the full POSIX specification. However, sometimes the Microsoft implementation of these APIs differs from the requirements of a given standard or has a function name that conflicts with an identifier reservation in another standard. In these circumstances, Microsoft will often prefix the function name with an underscore. For instance, the POSIX function `strdup` isn't available on Windows, but the function `_strdup` is available and behaves the same way.

NOTE

For more information on Microsoft's POSIX support, see <https://docs.microsoft.com/en-us/cpp/c-runtime-library/compatibility>.

Visual C++ also supports many of the safe string-handling functions from Annex K and will diagnose the use of an unsafe variant unless you define `_CRT_SECURE_NO_WARNINGS` prior to including the header file that declares the function. Unfortunately, Visual C++ does not conform to Annex K of the C

standard, because Microsoft chose not to update its implementation based on changes to the APIs that occurred during the standardization process. For example, Visual C++ doesn't provide the `set_constraint_handler_s` function but instead retains an older function with similar behavior but an incompatible signature:

```
_invalid_parameter_handler _set_invalid_parameter_handler(_invalid_parameter_handler)
```

Microsoft also doesn't define the `abort_handler_s` and `ignore_handler_s` functions, the `memset_s` function (which was not defined by ISO/IEC TR 24731-1), or the `RSIZE_MAX` macro. Visual C++ also doesn't treat overlapping source and destination sequences as runtime-constraint violations and instead has undefined behavior in such cases. “Bounds-Checking Interfaces: Field Experience and Future Directions” (Seacord 2019) provides additional information on all aspects of the bounds-checked interfaces, including Microsoft's implementation.

Summary

In this chapter, you learned about character encodings, such as ASCII and Unicode. You also learned about the various data types used to represent characters in C language programs, such as `char`, `int`, `wchar_t`, and so forth. We then covered character conversion libraries, including C standard library functions, `libiconv`, and Windows APIs.

In addition to characters, you also learned about strings and the legacy functions and bounds-checked interfaces defined in the C standard library for handling strings, as well as some POSIX- and Microsoft-specific functions.

Manipulating character and string data is a common programming task in C as well as a frequent source of errors. We outlined various approaches to handling these data types; you should determine which approach is best suited to your application and apply that approach consistently.

In the next chapter, you'll learn about I/O, which, among other things, can be used to read and write characters and strings.

8

INPUT/OUTPUT



This chapter will teach you how to perform input/output (I/O) operations to read data from, or write data to, terminals and filesystems.

Information can enter a program via command line arguments or the environment and exit it via its return status. However, most information typically enters or exits a program through I/O operations. We'll cover techniques that use C standard streams and POSIX file descriptors. We'll start by discussing C standard text and binary streams. We'll then cover different ways of opening and closing files using C standard library and POSIX functions.

Next, we'll discuss reading and writing characters and lines, reading and writing formatted text, and reading from and writing to binary streams. We'll also cover stream buffering, stream orientation, and file positioning.

Many other devices and I/O interfaces (such as `ioctl`) are available but are beyond the scope of this book.

Standard I/O Streams

C provides streams to communicate with files stored on supported, structured storage devices and terminals. A *stream* is a uniform abstraction for

communicating with files and devices that consume or produce sequential data such as sockets, keyboards, Universal Serial Bus (USB) ports, and printers.

C uses the opaque `FILE` data type to represent streams. A `FILE` object holds the internal state information for the connection to the associated file, including the file position indicator, buffering information, an error indicator, and an end-of-file indicator. You should never allocate a `FILE` object yourself. C standard library functions operate on objects of type `FILE *` (that is, a pointer to the `FILE` type). As a result, streams are frequently referred to as *file pointers*.

C provides an extensive application programming interface (API), found in `<stdio.h>`, for operating on streams; we'll explore this API later in this chapter. However, because these I/O functions need to work with a wide variety of devices and filesystems across many platforms, they're highly abstracted, which makes them unsuitable for anything beyond the simplest applications.

For example, the C standard has no concept of directories, because it must be able to work with nonhierarchical filesystems. The C standard makes few references to filesystem-specific details, like file permissions or locking. However, function specifications frequently state that certain behaviors happen “to the extent that the underlying system supports it,” meaning that they will occur only if they're supported by your implementation.

As a result, you'll generally need to use the less portable APIs provided by POSIX, Windows, and other platforms to perform I/O in real-world applications. Frequently, applications will define their own APIs that, in turn, rely on platform-specific APIs to provide safe, secure, and portable I/O operations.

Error and End-of-File Indicators

As just mentioned, a `FILE` object holds the internal state information for the connection to the associated file, including an error indicator that records whether a read/write error has occurred and an end-of-file indicator that records whether the end of the file has been reached. When opened, the error and end-of-file indicators for the stream are cleared. The following C standard library functions all set the error indicator for the stream when an error occurs: the byte input functions (`getc`, `fgetc`, and `getchar`), byte output functions (`putc`, `fputc`, and `putchar`), `fflush`, `fseek`, and `fsetpos`. Input functions such as `fgetc` and `getchar` will also set the end-of-file indicator for the stream if the stream is at end-of-file. Certain functions such as `rewind` and `freopen` clear the error indicator for the stream, and functions such as `rewind`, `freopen`, `ungetc`, `fseek`, and `fsetpos` clear the end-of-file indicator for the stream. The wide character

I/O functions behave similarly.

These indicators may be tested and cleared explicitly:

- The `ferror` function tests the error indicator for the specified stream and returns nonzero if and only if the error indicator is set for the specified stream.
- The `feof` function tests the end-of-file indicator for the specified stream and returns nonzero if and only if the end-of-file indicator is set for the specified stream.
- The `clearerr` function clears the end-of-file and error indicators for the specified stream.

The following short program illustrates the interaction between these functions and the two indicators:

```
#include <stdio.h>
#include <assert.h>

int main() {
    FILE* tmp = tmpfile();
    fputs("Effective C\n", tmp);
    rewind(tmp);
    for (int c; (c = fgetc(tmp)) != EOF; putchar(c)) {}
    printf("%s", "End-of-file indicator ");
    puts(feof(tmp) ? "set" : "clear");
    printf("%s", "Error indicator ");
    puts(ferror(tmp) ? "set" : "clear");
    clearerr(tmp); // clear both indicators
    printf("%s", "End-of-file indicator ");
    puts(feof(tmp) ? "set" : "clear");
}
```

This program produces the following output on stdout:

```
Effective C
End-of-file indicator set
Error indicator clear
End-of-file indicator clear
```

The loop terminates by end-of-file, after which the end-of-file indicator is set. Both indicators are cleared by the call to the `clearerr` function.

Stream Buffering

Buffering is the process of temporarily storing data in memory that's passing between a process and a device or file. Buffering improves the throughput of I/O operations, which often have high latencies per individual I/O operation with the system. Similarly, when a program requests to write to block-oriented devices like disks, the driver can cache the data in memory until it has accumulated enough data for one or more device blocks, at which point it writes the data all at once to the disk, improving throughput. This strategy is called *flushing* the output buffer.

A stream can be in one of three states:

Unbuffered Characters are intended to appear from the source or at the destination as soon as possible. Streams where more than one program may be accessing the data concurrently are often best unbuffered. Streams used for error reporting or logging might be unbuffered.

Fully buffered Characters are intended to be transmitted to or from the host environment as a block when a buffer is filled. Streams used for file I/O are normally fully buffered to optimize throughput.

Line buffered Characters are intended to be transmitted to or from the host environment as a block when a newline character is encountered. Streams connected to interactive devices such as terminals are line-buffered when you open them.

In the next section, we'll introduce predefined streams and describe how they're buffered.

Predefined Streams

A C program has three *predefined text streams* open and available for use on startup. These predefined streams are declared in `<stdio.h>`:

```
extern FILE * stdin; // standard input stream
extern FILE * stdout; // standard output stream
extern FILE * stderr; // standard error stream
```

The *standard output stream* (`stdout`) is the conventional output destination from the program. This stream is usually associated with the terminal that initiated the program but can be redirected to a file or other stream. You can enter the following commands in a Linux or Unix shell:

```
$ echo fred
fred
$ echo fred > tempfile
$ cat tempfile
fred
```

Here, the output from the `echo` command is redirected to *tempfile*.

The *standard input stream* (`stdin`) is the conventional input source for the program. By default, `stdin` is associated with the keyboard but may be redirected to input from a file, for example, with the following commands:

```
$ echo "one two three four five six seven" > tempfile
$ wc < tempfile
1 7 34
```

The contents of the file *tempfile* are redirected to the `stdin` of the `wc` command, which outputs the newline (1), word (7), and byte (34) counts from *tempfile*. The `stdin` and `stdout` streams are fully buffered if and only if the stream doesn't refer to an interactive device.

The *standard error stream* (`stderr`) is for writing diagnostic output. The `stderr` stream isn't fully buffered so that error messages can be viewed as soon as possible.

[Figure 8-1](#) shows the predefined streams `stdin`, `stdout`, and `stderr` attached to the keyboard and display of the user's terminal.




Figure 8-1: Standard streams attached to I/O communication channels

The output stream of one program can be redirected to be another application's input stream by using POSIX pipes:

```
$ echo "Hello Robert" | sed "s/Hello/Hi/" | sed "s/Robert/robot/"
Hi robot
```

The stream editor `sed` is a Unix utility used for filtering and transforming text. The vertical bar character (`|`) is available on many platforms to chain commands.

Stream Orientation

Each stream has an *orientation* that indicates whether the stream contains narrow or wide characters. After a stream is associated with an external file, but before any operations are performed on it, the stream doesn't have an orientation. Once a wide-character I/O function has been applied to a stream without orientation, the stream becomes a *wide-oriented stream*. Similarly, once a byte I/O function has been applied to a stream without orientation, the stream becomes a *byte-oriented stream*. Multibyte character sequences or narrow characters that can be represented as an object of type `char` (which are required by the C standard to be 1 byte) can be written to a byte-oriented stream.

You can reset the orientation of a stream by using the `fwide` function or by closing and then reopening the file. Applying a byte I/O function to a wide-oriented stream or a wide-character I/O function to a byte-oriented stream results in undefined behavior. Never mix narrow-character data, wide-character data, and binary data in the same file.

All three predefined streams (`stderr`, `stdin`, and `stdout`) are unoriented at program startup.

Text and Binary Streams

The C standard supports both text streams and binary streams. A *text stream* is an ordered sequence of characters composed into lines, each of which consists of zero or more characters plus a terminating newline character sequence. You can denote a single line break on a Unix-like system by using a line feed (`\n`). Most Microsoft Windows programs use a carriage return (`\r`) followed by a line feed (`\n`).

The different newline conventions can cause text files that have been transferred between systems with different conventions to display or parse

incorrectly, though this is increasingly uncommon on recent systems that now understand foreign newline conventions.

A *binary stream* is an ordered sequence of arbitrary binary data. Data read in from a binary stream will be the same as data written out earlier to that same stream, under the same implementation. On non-POSIX systems, streams may have an implementation-defined number of null bytes appended to the end of the stream.

Binary streams are always more capable and more predictable than text streams. However, the easiest way to read or write an ordinary text file that can work with other text-oriented programs is through a text stream.

Opening and Creating Files

When you open or create a file, it's associated with a stream. The `fopen` and the POSIX `open` functions open or create a file.

fopen

The `fopen` function opens the file whose name is given as a string and pointed to by `filename` and then associates a stream with it:

```
FILE *fopen(  
    const char * restrict filename,  
    const char * restrict mode  
);
```

The mode argument points to one of the strings shown in [Table 8-1](#) to determine how to open the file.

Table 8-1: Valid File Mode Strings

Mode string	Description
r	Open existing text file for reading
w	Truncate to zero length or create text file for writing
a	Append, open, or create text file for writing at end-of-file
rb	Open existing binary file for reading
wb	Truncate file to zero length or create binary file for writing
ab	Append, open, or create binary file for writing at end-of-file
r+	Open existing text file for reading and writing
w+	Truncate to zero length or create text file for reading and writing
a+	Append, open, or create text file for update, writing at current end-of-file

r+b or rb+	Open existing binary file for reading and writing
w+b or wb+	Truncate to zero length or create binary file for reading and writing
a+b or ab+	Append, open, or create binary file for update, writing at current end-of-file

Opening a file with read mode (by passing r as the first character in the mode argument) fails if the file doesn't exist or cannot be read.

Opening a file with append mode (by passing a as the first character in the mode argument) causes all subsequent writes to the file to occur at the current end-of-file at the point of buffer flush or actual write, regardless of intervening calls to the `fseek`, `fsetpos`, or `rewind` functions. Incrementing the current end-of-file by the amount of data written is atomic with respect to other threads writing to the same file provided the file was also opened in append mode. If the implementation is incapable of incrementing the current end-of-file atomically, it will fail instead of performing nonatomic end-of-file writes. In some implementations, opening a binary file with append mode (by passing b as the second or third character in the mode argument) may initially set the file position indicator for the stream beyond the last data written because of null character padding.

You can open a file in update mode by passing + as the second or third character in the mode argument, allowing both read and write operations to be performed on the associated stream. Opening (or creating) a text file with update mode may instead open (or create) a binary stream in some implementations. On POSIX systems, text and binary streams have the exact same behavior.

The C11 standard added the *exclusive mode* for reading and writing binary and text files, as shown in [Table 8-2](#).

Table 8-2: Valid File Mode Strings Added by C11

Mode string	Description
wX	Create exclusive text file for writing
wbX	Create exclusive binary file for writing
w+X	Create exclusive text file for reading and writing
w+bx or wb+X	Create exclusive binary file for reading and writing

Opening a file with exclusive mode (by passing x as the last character in the mode argument) fails if the file already exists or cannot be created. The check for the existence of the file and the creation of the file if it doesn't exist are atomic with respect to other threads and concurrent program executions. If the

implementation is incapable of performing the check for the existence of the file and the creation of the file atomically, it fails rather than perform a nonatomic check and creation.

As a final note, make sure that you never copy a FILE object. The following program, for example, can fail because a by-value copy of stdout is being used in the call to fputs:

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    FILE my_stdout = *stdout;
    if (fputs("Hello, World!\n", &my_stdout) == EOF) {
        return EXIT_FAILURE;
    }
    return EXIT_SUCCESS;
}
```

This program has undefined behavior and typically crashes when run.

open

On POSIX systems, the open function (IEEE Std 1003.1:2018) establishes the connection between a file identified by path and a value called a *file descriptor*:

```
int open(const char *path, int oflag, ...);
```

The *file descriptor* is a nonnegative integer that refers to the structure representing the file (called the *open file description*). The file descriptor returned by the open function is the lowest numbered unused file descriptor and is unique to the calling process. The file descriptor is used by other I/O functions to refer to that file. The open function sets the file offset used to mark the current position within the file to the beginning of the file. For a file descriptor underlying a stream, this file offset is separate from the stream's file position indicator.

The value of the oflag parameter sets the open file description's *file access modes*, which specify whether the file is being opened for reading, writing, or both. Values for oflag are constructed by a bitwise-inclusive OR of a file access mode and any combination of access flags. Applications must specify exactly one of the following file access modes in the value of oflag:

- O_EXEC** Open for execute only (nondirectory files)
- O_RDONLY** Open for reading only
- O_RDWR** Open for reading and writing
- O_SEARCH** Open directory for search only
- O_WRONLY** Open for writing only

The value of the `oflag` parameter also sets the *file status flags*, which control the behavior of the `open` function and affect how file operations are performed. These flags include the following:

- O_APPEND** Sets the file offset to the end-of-file prior to each write
- O_TRUNC** Truncates the length to 0
- O_CREAT** Creates a file
- O_EXCL** Causes the open to fail if **O_CREAT** is also set and the file exists

The `open` function takes a variable number of arguments. The value of the argument following the `oflag` argument specifies the file-mode bits (the file permissions when you create a new file) and is of type `mode_t`.

[Listing 8-1](#) shows an example of using the `open` function to open a file for writing.

```
#include <err.h>
#include <fcntl.h>
#include <sys/stat.h>
#include <stdio.h>
#include <stdlib.h>

int main() {
    int fd;
    int flags = O_WRONLY | O_CREAT | O_TRUNC;
    ❶ mode_t mode = S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH;
    const char *pathname = "/tmp/file";
    if ((fd = open(pathname, flags, mode) ❷) == -1) {
        err(EXIT_FAILURE, "Can't open %s", pathname);
    }
    // --snip--
}
```

Listing 8-1: Opening a file as write-only

The call to `open` ❷ takes multiple arguments, including the pathname of the

file, the `oflag`, and the mode. We create a mode flag ❶ that is a bitwise-inclusive OR of the following mode bits for access permission:

S_IRUSR Read permission bit for the owner of the file

S_IWUSR Write permission bit for the owner of the file

S_IRGRP Read permission bit for the group owner of the file

S_IROTH Read permission bit for other users

The `open` function sets these permissions only if it creates the file. If the file already exists, its current permissions are kept. The file access mode is `O_WRONLY`, which means the file is opened for writing only. The `O_CREAT` file status flag informs `open` to create the file; the `O_TRUNC` file status flag informs `open` that if the file exists and is successfully opened, it should discard the previous contents of the file.

If the file was successfully opened, the `open` function returns a nonnegative integer representing the file descriptor. Otherwise, `open` returns `-1` and sets `errno` to indicate the error. [Listing 8-1](#) checks for a value of `-1`, writes a diagnostic message to the predefined `stderr` stream if an error occurred, and then exits.

In addition to `open`, POSIX has other useful functions for working with file descriptors, such as the `fileno` function to get the file descriptor associated with an existing file pointer and the `fdopen` function to create a new stream file pointer from an existing file descriptor. POSIX APIs available through the file descriptor allow access to features of POSIX filesystems that aren't normally exposed through the file pointer interfaces such as directories (`posix_getdents`, `fdopendir`, `readdir`), file permissions (`fchmod`), and file locks (`fcntl`).

Closing Files

Opening a file allocates resources. If you continually open files without closing them, you'll eventually run out of file descriptors or handles available for your process, and attempting to open more files will fail. Consequently, it's important to close files after you've finished using them.

fclose

The C standard library `fclose` function closes the file:

```
int fclose(FILE *stream);
```

Any unwritten buffered data for the stream is delivered to the host environment to be written to the file. Any unread buffered data is discarded.

It's possible for the `fclose` function to fail. When `fclose` writes the remaining buffered output, for example, it might return an error because the disk is full. Even if you know the buffer is empty, errors can still occur when closing a file if you're using the Network File System (NFS) protocol. Despite the possibility of failure, recovery is often impossible, so programmers commonly ignore errors returned by `fclose`. When closing the file fails, a common practice is to abort the process or to truncate the file so its contents are meaningful when next read.

To ensure your code is robust, make sure you check for errors. File I/O can fail for any number of reasons. The `fclose` function returns `EOF` if any errors were detected:

```
if (fclose(fp) == EOF) {  
    err(EXIT_FAILURE, "Failed to close file\n");  
}
```

You need to explicitly call `fflush` or `fclose` on any buffered stream the program has written to, instead of letting `exit` (or a return from `main`) flush it, to perform the error checking.

The value of a pointer to a `FILE` object is indeterminate after the associated file is closed. Whether a file of zero length (in which an output stream hasn't written any data) exists is implementation defined.

You can reopen a closed file in the same program or another one, and its contents can be reclaimed or modified. If the initial call to the `main` function returns or if the `exit` function is called, all open files close (and all output streams are flushed) before program termination.

Other paths to program termination, such as calling the `abort` function, may not close all files properly, which means that buffered data not yet written to a disk might be lost.

close

On POSIX systems, you can use the `close` function to deallocate the file descriptor specified by `fd`:

```
int close(int fd);
```

If an I/O error occurred while reading from or writing to the filesystem during `close`, it may return `-1` with `errno` set to the cause of failure. If an error is returned, the state of `fd` is unspecified, meaning you can no longer read or write data to the descriptor or attempt to close it again—effectively leaking the file descriptor. The `posix_close` function is being added to The Open Group Base Specifications Issue 8 to address this problem.

Once a file is successfully closed, the file descriptor no longer exists, because the integer corresponding to it no longer refers to a file. Files are also closed when the process owning that file descriptor terminates.

Except in rare circumstances, an application that uses `fopen` to open a file will use `fclose` to close it; an application that uses `open` to open a file will use `close` to close it (unless it passed the descriptor to `fdopen`, in which case it must close by calling `fclose`).

Reading and Writing Characters and Lines

The C standard defines functions for reading and writing specific characters or lines.

Most byte stream functions have counterparts that take a wide character (`wchar_t`) or wide-character string instead of a narrow character (`char`) or string, respectively (see [Table 8-3](#)). Byte-stream functions are declared in the header file `<stdio.h>`, and the wide-stream functions are declared in `<wchar.h>`. The wide-character functions operate on the same streams (such as `stdout`).

Table 8-3: Narrow- and Wide-String I/O Functions

char	wchar_t	Description
<code>fgetc</code>	<code>fgetwc</code>	Reads a character from a stream.
<code>getc</code>	<code>getwc</code>	Reads a character from a stream.
<code>getchar</code>	<code>getwchar</code>	Reads a character from <code>stdin</code> .
<code>fgets</code>	<code>fgetws</code>	Reads a line from a stream.
<code>fputc</code>	<code>fputwc</code>	Writes a character to a stream.
<code>putc</code>	<code>putwc</code>	Writes a character to a stream.
<code>fputs</code>	<code>fputws</code>	Writes a string to a stream.
<code>putchar</code>	<code>putwchar</code>	Writes a character to <code>stdout</code> .
<code>puts</code>	N/A	Writes a string to <code>stdout</code> .
<code>ungetc</code>	<code>ungetwc</code>	Returns a character to a stream.
<code>scanf</code>	<code>wscanf</code>	Reads formatted character input from <code>stdin</code> .
<code>fscanf</code>	<code>fwscanf</code>	Reads formatted character input from a stream.
<code>sscanf</code>	<code>swscanf</code>	Reads formatted character input from a buffer.

<code>printf</code>	<code>wprintf</code>	Prints formatted character output to <code>stdout</code> .
<code>fprintf</code>	<code>fwprintf</code>	Prints formatted character output to a stream.
<code>sprintf</code>	<code>swprintf</code>	Prints formatted character output to a buffer.
<code>snprintf</code>	N/A	This is the same as <code>sprintf</code> with truncation. The <code>swprintf</code> function also takes a length argument but behaves differently from <code>snprintf</code> in the way it interprets it.

In this chapter, we'll discuss the byte-stream functions only. You may want to avoid wide-character function variants altogether and work exclusively with UTF-8 character encodings, if possible, as these functions are less prone to programmer error and security vulnerabilities.

The `fputc` function converts the character `c` to the type `unsigned char` and writes it to stream:

```
int fputc(int c, FILE *stream);
```

It returns `EOF` if a write error occurs; otherwise, it returns the character it has written.

The `putc` function is just like `fputc`, except that most libraries implement it as a macro:

```
int putc(int c, FILE *stream);
```

If `putc` is implemented as a macro, it may evaluate its `stream` argument more than once. Using `fputc` is generally safer. See CERT C rule FIO41-C, “Do not call `getc()`, `putc()`, `getwc()`, or `putwc()` with a stream argument that has side effects,” for more information.

The `putchar` function is equivalent to the `putc` function, except that it uses `stdout` as the value of the stream argument.

The `fputs` function writes the string `s` to the stream `stream`:

```
int fputs(const char * restrict s, FILE * restrict stream);
```

This function doesn't write the null character from the string `s`—nor does it write a newline character—but outputs only the characters in the string. If a write error occurs, `fputs` returns `EOF`. Otherwise, it returns a nonnegative value. For example, the following statements output the text `I am Groot`, followed by a newline:

```
fputs("I ", stdout);  
fputs("am ", stdout);  
fputs("Groot\n", stdout);
```

The `puts` function writes the string `s` to the stream `stdout` followed by a newline:

```
int puts(const char *s);
```

The `puts` function is the most convenient function for printing simple messages because it takes only a single argument. Here's an example:

```
puts("This is a message.");
```

The `fgetc` function reads the next character as an unsigned `char` from a stream and returns its value, converted to an `int`:

```
int fgetc(FILE *stream);
```

If an end-of-file condition or read error occurs, the function returns `EOF`.

The `getc` function is equivalent to `fgetc`, except that if it's implemented as a macro, it may evaluate its stream argument more than once. Consequently, this argument should never be an expression with side effects. Analogous to the `fputc` function, using `fgetc` is generally safer and should be preferred to `getc`.

The `getchar` function is equivalent to the `getc` function, except that it uses `stdout` as the value of the stream argument.

You may recall that the `gets` function reads characters from `stdin` and writes them into a character array until a newline or `EOF` is reached. The `gets` function is inherently insecure. It was deprecated in C99 and removed from C11 and *should never be used*. If you need to read a string from `stdin`, consider using the `fgets` function instead. The `fgets` function reads at most one less than `n` characters from a stream into a character array pointed to by `s`:

```
char *fgets(char * restrict s, int n, FILE * restrict stream);
```

No additional characters are read after a (retained) newline character or `EOF`. A null character is written immediately following the last character read into the

array.

Stream Flushing

As described earlier in this chapter, streams can be fully or partially buffered, meaning that data you thought you wrote may not yet be delivered to the host environment. This can be a problem when the program terminates abruptly. The `fflush` function delivers any unwritten data for a specified stream to the host environment to be written to the file:

```
int fflush(FILE *stream);
```

The behavior is undefined if the last operation on the stream was input. If the stream is a null pointer, the `fflush` function performs this flushing action on all streams. Make sure that your file pointer isn't null before calling `fflush` if this isn't your intent.

Setting the Position in a File

Random-access files (which include a disk file, for example, but not a terminal) maintain a file position indicator associated with the stream. The *file position indicator* describes where in the file the stream is currently reading or writing.

When you open a file, the indicator is positioned at the file's start (unless you open it in append mode). You can position the indicator wherever you want to read or write any portion of the file. The `ftell` function obtains the current value of the file position indicator, while the `fseek` function sets the file position indicator. These functions use the `long int` type to represent offsets (positions) in a file and are therefore limited to offsets that can be represented as a `long int`. [Listing 8-2](#) demonstrates the use of the `ftell` and `fseek` functions.

```
#include <err.h>
#include <stdio.h>
#include <stdlib.h>

long int get_file_size(FILE *fp) {
    if (fseek(fp, 0, SEEK_END) != 0) {
        err(EXIT_FAILURE, "Seek to end-of-file failed");
    }
    long int fpi = ftell(fp);
    if (fpi == -1L) {
```

```

        err(EXIT_FAILURE, "ftell failed");
    }
    return fpi;
}

int main() {
    FILE *fp = fopen("fred.txt", "rb");
    if (fp == nullptr) {
        err(EXIT_FAILURE, "Cannot open fred.txt file");
    }
    printf("file size: %ld\n", get_file_size(fp));
    if (fclose(fp) != EOF) {
        err(EXIT_FAILURE, "Failed to close file");
    }
    return EXIT_SUCCESS;
}

```

Listing 8-2: Using the `ftell` and `fseek` functions

This program opens a file called *fred.txt* and calls the `get_file_size` function to find the file size. The `get_file_size` function calls `fseek` to set the file position indicator to the end of the file (indicated by `SEEK_END`) and the `ftell` function to retrieve the current value of the file position indicator for the stream as a long `int`. This value is returned by the `get_file_size` function and is printed in the `main` function. Finally, we close the file referenced by the `fp` file pointer.

The `fseek` function has different constraints for text and binary files. The offset must be either zero or a value previously returned by `ftell` for a text file, whereas you can use calculated offsets for a binary file.

To ensure your code is robust, make sure you check for errors. File I/O can fail for any number of reasons. The `fopen` function returns a null pointer when it fails. The `fseek` function returns nonzero only for a request that cannot be satisfied. On failure, the `ftell` function returns `-1L` and stores an implementation-defined value in `errno`. If the return value from `ftell` is equal to `-1L`, we use the `err` function to print the last component of the program name, a colon character, a space followed by an appropriate error message corresponding to the value stored in `errno`, and finally, a newline character. The `fclose` function returns `EOF` if any errors were detected. One of the unfortunate aspects of the C standard library demonstrated by this short program is that each function tends to report errors in its own unique way, so you normally need to refer to your documentation to see how to test for errors.

The `fgetpos` and `fsetpos` functions use the `fpos_t` type to represent offsets. This type can represent arbitrarily large offsets, meaning you can use `fgetpos` and `fsetpos` with arbitrarily large files. A wide-oriented stream has an associated `mbstate_t` object that stores the stream's current parse state. A successful call to `fgetpos` stores this multibyte state information as part of the value of the `fpos_t` object. A later successful call to `fsetpos` using the same stored `fpos_t` value restores the parse state as well as the position within the controlled stream. It's not possible to convert an `fpos_t` object to an integer byte or character offset within the stream except indirectly by calling `fsetpos` followed by `ftell`. The short program shown in [Listing 8-3](#) demonstrates the use of the `fgetpos` and `fsetpos` functions.

```
#include <err.h>
#include <stdio.h>
#include <stdlib.h>

int main() {
    FILE *fp = fopen("fred.txt", "w+");
    if (fp == nullptr) {
        err(EXIT_FAILURE, "Cannot open fred.txt file");
    }
    fpos_t pos;
    if (fgetpos(fp, &pos) != 0) {
        err(EXIT_FAILURE, "get position");
    }
    if (fputs("abcdefghijklmnopqrstuvwxyz", fp) == EOF) {
        fputs("Cannot write to fred.txt file\n", stderr);
    }
    if (fsetpos(fp, &pos) != 0) {
        err(EXIT_FAILURE, "set position");
    }
    long int fpi = ftell(fp);
    if (fpi == -1L) {
        err(EXIT_FAILURE, "ftell");
    }
    printf("file position = %ld\n", fpi);
    if (fputs("0123456789", fp) == EOF) {
        fputs("Cannot write to fred.txt file\n", stderr);
    }
    if (fclose(fp) == EOF) {
        err(EXIT_FAILURE, "Failed to close file\n");
    }
    return EXIT_SUCCESS;
}
```

Listing 8-3: Using the `fgetpos` and `fsetpos` functions

This program opens the *fred.txt* file for writing and then calls `fgetpos` to get the current file position within the file, which is stored in `pos`. We then write some text to the file before calling `fsetpos` to restore the file position indicator to the position stored in `pos`. At this point, we can use the `ftell` function to retrieve and print the file position, which should be 0. After running this program, *fred.txt* contains the following text:

```
0123456789klmnopqrstuvwxyz
```

You cannot write to a stream and then read from it again without an intervening call to the `fflush` function to write any unwritten data or to a file positioning function (`fseek`, `fsetpos`, or `rewind`). You also cannot read from a stream and then write to it without an intervening call to a file positioning function.

The `rewind` function sets the file position indicator to the beginning of the file:

```
void rewind(FILE *stream);
```

The `rewind` function is equivalent to invoking `fseek` followed by `clearerr` to clear the error indicator for the stream:

```
fseek(stream, 0L, SEEK_SET);  
clearerr(stream);
```

Because there is no way to determine if `rewind` failed, you should use `fseek` so that you can check for errors.

You shouldn't attempt to use file positions in files opened in append mode, because many systems don't modify the current file position indicator for append or will forcefully reset to the end of the file when writing. If using APIs that use file positions, the file position indicator is maintained by subsequent reads, writes, and positioning requests. Both POSIX and Windows have APIs that never use the file position indicator; for those, you always need to specify the offset into the file at which to perform the I/O. POSIX defines the `lseek` function, which behaves similarly to `fseek` but operates on an open file

description (IEEE Std 1003.1:2018).

Removing and Renaming Files

The C standard library provides a remove function to delete a file and a rename function to move or rename it:

```
int remove(const char *filename);  
int rename(const char *old, const char *new);
```

In POSIX, the file deletion function is `unlink`, and the directory removal function is `rmdir`:

```
int unlink(const char *path);  
int rmdir(const char *path);
```

POSIX also uses `rename` for renaming. One obvious difference between the C standard and POSIX is that C does not have a concept of directories, while POSIX does. Consequently, no specific semantics are defined in the C standard for dealing with directories.

The `unlink` function has better-defined semantics than the `remove` function because it's specific to POSIX filesystems. In POSIX and Windows, we can have any number of links to a file, including hard links and open file descriptors. The `unlink` function always removes the directory entry for the file but deletes the file only when there are no more links or open file descriptors referencing it. Even after deletion, the contents of the file may remain in permanent storage. The `rmdir` function removes a directory whose name is given by `path` only if it is an empty directory.

In POSIX, the `remove` function is required to behave the same as the `unlink` function when the argument is not a directory and to behave the same as the `rmdir` function when the argument is a directory. The `remove` function may behave differently on other operating systems.

The filesystem is shared with other programs running concurrently to yours. These other programs will modify the filesystem while your program runs. This means that a file entry can disappear or be replaced by a different file entry, which can be a source of security exploits and unexpected data loss. POSIX provides functions that let you `unlink` and `rename` files referred to by an open file descriptor or handle. These can be used to prevent security exploits and possible

unexpected data loss in a shared public filesystem.

Using Temporary Files

We frequently use *temporary files* as an interprocess communication mechanism or for temporarily storing information out to disk to free up random-access memory (RAM). For example, one process might write to a temporary file that another process reads from. These files are normally created in a temporary directory by using functions such as the C standard library's `tmpfile` and `tmpnam` or POSIX's `mkstemp`.

Temporary directories can be either global or user specific. In Unix and Linux, the `TMPDIR` environment variable is used to specify the location of the global temporary directories, which are typically `/tmp` and `/var/tmp`. Systems running Wayland or the X11 window system usually have user-specific temporary directories defined by the `$XDG_RUNTIME_DIR` environment variable, which is typically set to `/run/user/$uid`. In Windows, you can find user-specific temporary directories in the *AppData* section of the User Profile, typically `C:\Users\User Name\AppData\Local\Temp` (`%USERPROFILE%\AppData\Local\Temp`). On Windows, the global temporary directory is specified by either the `TMP` or `TEMP` environment variable. The `C:\Windows\Temp` directory is a system folder used by Windows to store temporary files.

For security reasons, it's best for each user to have their own temporary directory, because the use of global temporary directories frequently results in security vulnerabilities. The most secure function for creating temporary files is the POSIX `mkstemp` function. However, because accessing files in shared directories may be difficult or impossible to implement securely, we recommended that you not use any of the available functions and instead perform the interprocess communication by using sockets, shared memory, or other mechanisms designed for this purpose.

Reading Formatted Text Streams

In this section, we'll demonstrate the use of the `fscanf` function to read formatted input. The `fscanf` function is the corresponding input version of the `fprintf` function that we introduced all the way back in [Chapter 1](#) and has the following signature:

```
int fscanf(FILE * restrict stream, const char * restrict format,
...);
```

The `fscanf` function reads input from the stream pointed to by `stream`, under control of the format string that tells the function how many arguments to expect, their type, and how to convert them for assignment. Subsequent arguments are pointers to the objects receiving the converted input. The result is undefined if there are insufficient arguments for the format string. If you provide more arguments than conversion specifiers, the excess arguments are evaluated but otherwise ignored. The `fscanf` function has lots of functionality that we'll only touch upon here. For more information, refer to the C standard.

To demonstrate the use of `fscanf`, as well as some other I/O functions, we'll implement a program that reads in the *signals.txt* file shown in [Listing 8-4](#) and prints each line.

```
1 HUP Hangup
2 INT Interrupt
3 QUIT Quit
4 ILL Illegal instruction
5 TRAP Trace trap
6 ABRT Abort
7 EMT EMT trap
8 FPE Floating-point exception
```

Listing 8-4: The signals.txt file

Each line of this file contains the following: a signal number (a small positive-integer value), the signal ID (a small string of up to six alphanumeric characters), and a short string with a description of the signal. Fields are whitespace delimited except for the description field, which is delimited by one or more space or tab characters at the beginning and by a newline at the end.

[Listing 8-5](#) shows the signals program, which reads this file and prints out each line.

```
#include <err.h>
#include <stdio.h>
#include <string.h>
#include <stdlib.h>

#define TO_STR_HELPER(x) #x
```

```

#define TO_STR(x) TO_STR_HELPER(x)

#define DESC_MAX_LEN 99

int main() {
    int status = EXIT_SUCCESS;
    FILE *in;

    struct sigrecord {
        int signum;
        char signame[10];
        char sigdesc[DESC_MAX_LEN + 1];
    } rec;

    if ((in = fopen("signals.txt", "r")) == nullptr) {
        err(EXIT_FAILURE, "Cannot open signals.txt file");
    }

    while (true) {
        int n = fscanf(in, "%d%9s%*[\t]%" TO_STR(DESC_MAX_LEN) "[^\n]",
            &rec.signum, rec.signame, rec.sigdesc
        );
        if (n == 3) {
            printf(
                "Signal\n  number = %d\n  name = %s\n  description = %s\n",
                rec.signum, rec.signame, rec.sigdesc
            );
        }
        else if (ferror(in)) {
            perror("Error indicated");
            status = EXIT_FAILURE;
            break;
        }
        else if (n == EOF) {
            // normal end-of-file
            break;
        }
        else if (feof(in)) {
            fputs("Premature end-of-file detected\n", stderr);
            status = EXIT_FAILURE;
            break;
        }
        else {
            fputs("Failed to match signum, signame, or sigdesc\n\n", st

```

```

derr);
    int c;
    while ((c = getc(in)) != '\n' && c != EOF);
    status = EXIT_FAILURE;
}
}

❷ if (fclose(in) == EOF) {
    err(EXIT_FAILURE, "Failed to close file\n");
}

return status;
}

```

Listing 8-5: The signals program

We define several variables in the main function, including the `rec` structure ❶, which we use to store the signal information found on each line of the file.

The `rec` structure contains three members: a `signum` member of type `int` that will hold the signal number; a `signame` member that's an array of `char` and will hold the signal ID; and the `sigdesc` member, an array of `char` that will hold the description of the signal. Both arrays have fixed sizes that we determined were adequately sized for the strings being read from the file. If the strings read from the file are too long to fit in these arrays, the program will treat it as an error.

The call to `fscanf` ❸ reads each line of input from the file. It appears inside of an infinite `while (true)` loop ❷ that we must break out of for the program to terminate. We assign the return value from the `fscanf` function to a local variable `n`. The `fscanf` function returns `EOF` if an input failure occurs before the first conversion has completed. Otherwise, the function returns the number of input items assigned, which can be fewer than provided for, or even zero, in the event of an early matching failure. The call to `fscanf` assigns three input items, so we print the signal description only when `n` is equal to 3. Next, we call `ferror(in)` to determine if the call to `fscanf` set the error indicator. If it did, we print `errno` with a call to the `perror` function and then set the status to `EXIT_FAILURE`. Next, if `n` equals `EOF`, we break out of the loop because we have successfully processed all the input. The final possibility is that `fscanf` returned a value that is not the expected number of input items, nor is it `EOF` indicating an early matching failure. In this case, we treat the condition as a nonfatal error:

```

fputs("Failed to match signum, signame, or sigdesc\n\n", stderr);

```

```
int c;  
while ((c = getc(in)) != '\n' && c != EOF);  
status = EXIT_FAILURE;
```

We print a message to `stderr` to let the user know that there is a problem with one of the signal descriptions in the file, but we continue to process the remaining entries. The loop discards the defective line and `status` is assigned `EXIT_FAILURE` to indicate to the calling program that an error occurred. You'll notice that proper error handling in this program makes up the bulk of the code.

The `fscanf` function uses a *format string* that determines how the input text is assigned to each argument. In this case, the `"%d%9s%*[\t]%99[^\n]"` format string contains four *conversion specifiers*, which specify how the input read from the stream is converted into values stored in the objects referenced by the format string's arguments. We introduce each conversion specification with the percent character (%). After the %, the following may appear, in sequence:

- An optional character `*` that discards the input without assigning it to an argument
- An optional integer greater than zero that specifies the maximum field width (in characters)
- An optional length modifier that specifies the size of the object
- A conversion specifier character that specifies the type of conversion to be applied

The first conversion specifier in the format string is `%d`. This conversion specifier matches the first optionally signed decimal integer, which should correspond to the signal number in the file, and stores the value in the third argument referenced by `rec.signum`. Without an optional length modifier, the length of the input depends on the conversion specifier's default type. For the `d` conversion specifier, the argument must point to a signed `int`.

The second conversion specifier in this format string is `%9s`, which matches the next sequence of non-whitespace characters from the input stream—corresponding to the signal name—and stores these characters as a string in the fourth argument referenced by `rec.signame`. The length modifier prevents more than nine characters from being input and then writes a null character in `rec.signame` after the matched characters. A conversion specifier of `%10s` in this example would allow a buffer overflow to occur. A conversion specifier of `%9s` can still fail to read the entire string, resulting in a matching error. When reading

data into a fixed-size buffer as we are doing, you should test inputs that exactly match or slightly exceed the fixed buffer length to ensure buffer overflow does not occur and that the string is properly null terminated.

We're going to skip the third conversion specifier for a moment and talk about the fourth one: `%99[^\n]`. This fancy conversion specifier will match the signal description field in the file. The brackets (`[]`) contain a *scanset*, which is like a regular expression. This scanset uses the circumflex (`^`) to exclude `\n` characters. Put together, `%99[^\n]` reads all the characters until it reaches a `\n` (or EOF) and stores them in the fifth argument referenced by `rec.sigdesc`. C programmers commonly use this syntax to read an entire line. This conversion specifier also includes a maximum string length of 99 characters to avoid buffer overflows.

We can now revisit the third conversion specifier: `%*[\t]`. As we have just seen, the fourth conversion specifier reads all the characters, starting from the end of the signal ID. Unfortunately, this includes any whitespace characters between the signal ID and the start of the description. The purpose of the `%*[\t]` conversion specifier is to consume any space or horizontal tab characters between these two fields and suppress them by using the assignment-suppressing specifier `*`. It's also possible to include other whitespace characters in the scanset for this conversion specifier.

Finally, we call the `fclose` function ❹ to close the file.

Reading from and Writing to Binary Streams

The C standard library `fread` and `fwrite` functions can operate on both text and binary streams. The `fwrite` function has the following signature:

```
size_t fwrite(const void * restrict ptr, size_t size, size_t nmem  
b,  
FILE * restrict stream);
```

This function writes up to `nmem` elements of `size` bytes from the array pointed to by `ptr` to `stream`. The `fwrite` function behaves as if it converts each object to an array of `unsigned char` (every object can be converted to an array of this type) and then calls the `fputc` function to write the value of each character in the array in order. The file position indicator for the stream is advanced by the number of characters successfully written.

POSIX defines similar read and write functions that operate on file

descriptors instead of streams (IEEE Std 1003.1:2018).

[Listing 8-6](#) demonstrates the use of the `fwrite` function to write signal records to the *signals.bin* file.

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

typedef struct sigrecord {
    int signum;
    char signame[10];
    char sigdesc[100];
} rec;

int main() {
    int status = EXIT_SUCCESS;
    FILE *fp;

    ❶ if ((fp = fopen("signals.bin", "wb")) == nullptr) {
        fputs("Cannot open signals.bin file\n", stderr);
        return EXIT_FAILURE;
    }

    ❷ rec sigrec30 = {30, "USR1", "user-defined signal 1"};
    rec sigrec31 = {
        .signum = 31, .signame = "USR2", .sigdesc = "user-defined sig
        nal 2"
    };

    size_t size = sizeof(rec);

    ❸ if (fwrite(&sigrec30, size, 1, fp) != 1) {
        fputs("Cannot write sigrec30 to signals.bin file\n", stderr);
        status = EXIT_FAILURE;
        goto close_files;
    }

    if (fwrite(&sigrec31, size, 1, fp) != 1) {
        fputs("Cannot write sigrec31 to signals.bin file\n", stderr);
        status = EXIT_FAILURE;
    }

close_files:
    if (fclose(fp) == EOF) {
        fputs("Failed to close file\n", stderr);
```

```

    status = EXIT_FAILURE;
}

return status;
}

```

Listing 8-6: Writing to a binary file using direct I/O

We open the *signals.bin* file in `wb` mode ❶ to create a binary file for writing. We declare two `rec` structures ❷ and initialize them with the signal values we want to write to the file. For comparison, the `sigrec30` structure is initialized with positional initializers, and `sigrec31` is initialized using designated initializers. Both initialization styles have the same behavior; designated initializers make the declaration less terse but clearer. The actual writing begins at ❸. We check the return values from each call to the `fwrite` function to ensure that it wrote the correct number of elements.

[Listing 8-7](#) uses the `fread` function to read the data we just wrote from the *signals.bin* file.

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>

typedef struct rec {
    int signum;
    char signame[10];
    char sigdesc[100];
} rec;

int main() {
    int status = EXIT_SUCCESS;
    FILE *fp;
    rec sigrec;

    ❶ if ((fp = fopen("signals.bin", "rb")) == nullptr) {
        fputs("Cannot open signals.bin file\n", stderr);
        return EXIT_FAILURE;
    }

    // read the second signal
    ❷ if (fseek(fp, sizeof(rec), SEEK_SET) != 0) {
        fputs("fseek in signals.bin file failed\n", stderr);
        status = EXIT_FAILURE;
    }
}

```

```

        goto close_files;
    }

❸ if (fread(&sigrec, sizeof(rec), 1, fp) != 1) {
    fputs("Cannot read from signals.bin file\n", stderr);
    status = EXIT_FAILURE;
    goto close_files;
}

printf(
    "Signal\n  number = %d\n  name = %s\n  description = %s\n\n",
    sigrec.signum, sigrec.signame, sigrec.sigdesc
);

close_files:
    if (fclose(fp) == EOF) {
        fputs("Failed to close file\n", stderr);
        status = EXIT_FAILURE;
    }

    return status;
}

```

Listing 8-7: Reading from a binary file using direct I/O

We open the binary file by using the `rb` mode ❶ for reading. Next, to make this example a bit more interesting, the program reads and prints the information for a specific signal, rather than reading the entire file. We could indicate which signal to read by using an argument to the program, but for this example, we hardcoded it as the second signal. To accomplish this, the program invokes the `fseek` function ❷ to set the file position indicator for the stream referenced by `fp`. As mentioned earlier in this chapter, the file position indicator determines the file position for the subsequent I/O operation. For a binary stream, we set the new position by adding the offset (measured in bytes) to the position specified by the final argument (the beginning of the file, as indicated by `SEEK_SET`). The first signal is at position 0 in the file, and each subsequent signal is at an integer multiple of the size of the structure from the beginning of the file.

After the file position indicator is positioned at the start of the second signal, we call the `fread` function ❸ to read the data from the binary file into the structure referenced by `&sigrec`. The call to `fread` reads a single element whose size is specified by `sizeof(rec)` from the stream pointed to by `fp`. In most cases, this object has the size and type of the corresponding call to `fwrite`. The

file position indicator for the stream is advanced by the number of characters successfully read. We check the return value from the `fread` function to ensure the correct number of elements, here one, was read.

Endian

Object types other than character types can include padding as well as value representation bits. Different target platforms can pack bytes into multiple-byte words in different ways, called *endianness*.

NOTE

The term endianness is drawn from Jonathan Swift's 1726 satire, Gulliver's Travels, in which civil war erupts over whether the big end or the little end of a boiled egg is the proper end to crack open.

A *big-endian ordering* places the most significant byte first and the least significant byte last, while a *little-endian ordering* does the opposite. For example, consider the unsigned hexadecimal number `0x1234`, which requires at least two bytes to represent. In a big-endian ordering, these two bytes are `0x12`, `0x34`, while in a little-endian ordering, the bytes are arranged as `0x34`, `0x12`. Intel and AMD processors use the little-endian format, while the ARM and POWER series of processors can switch between the little- and big-endian formats. However, big-endian is the dominant ordering in network protocols such as Internet Protocol (IP), Transmission Control Protocol (TCP), and User Datagram Protocol (UDP). Endianness can cause problems when a binary file is created on one computer and is read on another computer with different endianness.

C23 has added a mechanism to determine your implementation's byte ordering at runtime using three macros that expand to integer constant expressions. The `__STDC_ENDIAN_LITTLE__` macro represents a byte order storage in which the least significant byte is placed first and the rest are in ascending order. The `__STDC_ENDIAN_BIG__` macro represents a byte order storage in which the most significant byte is placed first and the rest are in descending order.

The `__STDC_ENDIAN_NATIVE__` macro describes the endianness of the execution environment with respect to bit-precise integer types, standard integer types, and most extended integer types. The short program in [Listing 8-8](#) determines the byte ordering for the execution environment by testing the value of the `__STDC_ENDIAN_NATIVE__` macro. If the execution environment is neither

little-endian nor big-endian and has some other implementation-defined byte order, the macro `__STDC_ENDIAN_NATIVE__` will have a different value.

```
#include <stdbit.h>
#include <stdio.h>

int main (int argc, char* argv[]) {
    if (__STDC_ENDIAN_NATIVE__ == __STDC_ENDIAN_LITTLE__) {
        puts("little endian");
    }
    else if (__STDC_ENDIAN_NATIVE__ == __STDC_ENDIAN_BIG__) {
        puts("big endian");
    }
    else {
        puts("other byte ordering");
    }
    return 0;
}
```

Listing 8-8: Determining the byte ordering

All this variation between platforms implies that, for interhost communication, you should adopt a standard for the external format and use format conversion functions to *marshal* arrays of external data to and from multiple-byte native objects (using exact width types). POSIX has some suitable functions for this purpose, including `htonl`, `htons`, `ntohl`, and `ntohs`, that convert values between host and network byte order.

Endianness independence in binary data formats can be achieved by always storing the data in one fixed endianness or including a field in the binary file to indicate the endianness of the data.

Summary

In this chapter, you learned about streams, including stream buffering, the predefined streams, stream orientation, and the difference between text and binary streams.

You then learned how to create, open, and close files by using the C standard library and POSIX APIs. You also learned how to read and write characters and lines, read and write formatted text, and read and write from binary streams. You looked at how to flush a stream, set the position in a file, remove files, and rename files. Without I/O, communication with the user would

be limited to the program's return value. Finally, you learned about temporary files and how to avoid using them.

In the next chapter, you'll learn about the compilation process and the preprocessor, including file inclusion, conditional inclusion, and macros.

9

PREPROCESSOR

with Aaron Ballman



The preprocessor is the part of the C compiler that runs at an early phase of compilation and transforms the source code before it's translated, such as inserting code from one file (typically a header) into another (typically a source file). The preprocessor also allows you to specify that an identifier should be automatically substituted by a source code segment during macro expansion. In this chapter, you'll learn how to use the preprocessor to include files, define object- and function-like macros, conditionally include code based on implementation-specific features, and embed binary resources into your program.

The Compilation Process

Conceptually, the compilation process consists of a pipeline of eight phases, as shown in [Figure 9-1](#). We call these *translation phases* because each phase translates the code for processing by the next phase.




Figure 9-1: Translation phases

The preprocessor runs before the translator translates the source code into object code, which allows the preprocessor to modify the source code the user wrote *before* the translator operates on it. Consequently, the preprocessor has a limited amount of semantic information about the program being compiled. It doesn't understand functions, variables, or types. Only basic elements, such as header names, identifiers, literals, and punctuation characters like +, -, and ! are meaningful to the preprocessor. These basic elements, called *tokens*, are the smallest elements of a computer program that have meaning to a compiler.

The preprocessor operates on *preprocessing directives* that you include in the source code to program the behavior of the preprocessor. You spell preprocessing directives with a leading # token followed by a directive name, such as #include, #define, #embed, or #if. Each preprocessing directive is terminated by a newline character. You can indent directives by including whitespace between the beginning of the line and the #

```
#define THIS_IS_FINE 1
```

or between the # and the directive:

```
#  define SO_IS_THIS 1
```

Preprocessing directives instruct the preprocessor to alter the resulting translation unit. If your program contains preprocessing directives, the code the translator consumes is not the same code you wrote. Compilers usually provide a way to view the preprocessor output, called a *translation unit*, passed to the translator. Viewing the preprocessor output is unnecessary, but you may find it informative to see the actual code given to the translator. [Table 9-1](#) lists flags that common compilers use to output a translation unit.

Table 9-1: Outputting a Translation Unit

Compiler	Example command line
Clang	<code>clang other-options -E -o tu.i tu.c</code>
GCC	<code>gcc other-options -E -o tu.i tu.c</code>
Visual C++	<code>cl other-options /P /Fitu.i tu.c</code>

Preprocessed output files are commonly given a *.i* file extension.

File Inclusion

A powerful feature of the preprocessor is the ability to insert the contents of one source file into the contents of another source file by using the `#include` preprocessing directive. The included files are called *headers* to distinguish them from other source files. Headers typically contain declarations for use by other programs. This is the most common way to share external declarations of functions, objects, and data types with other parts of the program.

You've already seen many examples of including the headers for C standard library functions in the examples in this book. For instance, the program in [Table 9-2](#) is separated into a header named *bar.h* and a source file named *foo.c*. The source file *foo.c* does not directly contain a declaration of `func`, yet the function is successfully referenced by name within `main`. During preprocessing, the `#include` directive inserts the contents of *bar.h* into *foo.c* in the place of the `#include` directive itself.

Table 9-2: Header Inclusion

Original sources	Resulting translation unit
<i>bar.h</i> <code>int func(void);</code>	<code>int func(void);</code> <code>int main(void) {</code> <code> return func();</code> <code>}</code>
<i>foo.c</i> <code>#include "bar.h"</code> <code>int main(void) {</code> <code> return func();</code> <code>}</code>	

The preprocessor executes `#include` directives as it encounters them. Therefore, inclusion has transitive properties: if a source file includes a header that itself includes another header, the preprocessed output will contain the contents of both headers. For example, given the *baz.h* and *bar.h* headers and the *foo.c* source file, the output after running the preprocessor on the *foo.c* source code is shown in [Table 9-3](#).

Table 9-3: Transitive Header Inclusion

Original sources	Resulting translation unit
<i>baz.h</i> <code>int other_func(void);</code>	<code>int other_func(void);</code> <code>int func(void);</code> <code>int main(void) {</code> <code> return func();</code>
<i>bar.h</i> <code>#include "baz.h"</code>	

<pre>int func(void);</pre>	<pre>}</pre>
<pre>foo.c #include "bar.h" int main(void) { return func(); }</pre>	

Compiling the *foo.c* source file causes the preprocessor to include the "bar.h" header. The preprocessor then finds the include directive for the "baz.h" header and includes it as well, bringing the declaration for `other_func` into the resulting translation unit.

A best practice is to avoid relying on transitive includes because they make your code brittle. Consider using tools like include-what-you-use (<https://include-what-you-use.org>) to automatically remove reliance on transitive includes.

Starting with C23, you can test for the presence of an include file before the `#include` directive is executed with the `__has_include` preprocessor operator. It takes a header name as the only operand. The operator returns true if the specified file can be found and false otherwise. You can use it with conditional inclusion to provide an alternative implementation if a file cannot be included. For example, you can use the `__has_include` preprocessor operator to test for C standard library threading or POSIX threads support as follows:

```
#if __has_include(<threads.h>)
#   include <threads.h>
    typedef thrd_t thread_handle;
#elif __has_include(<pthread.h>)
    typedef pthread_t thread_handle;
#endif
```

You can use either a quoted include string (for example, `#include "foo.h"`) or an angle-bracketed include string (for example, `#include <foo.h>`) to specify the file to include. The difference between these syntaxes is implementation defined, but they typically influence the search path used to find the included files. For example, both Clang and GCC attempt to find files included with:

- Angle brackets on the *system include path*, specified using the `-isystem` flag
- Quoted strings on the *quoted include path*, specified using the `-iquote` and `-isystem` flags

Refer to your compiler's documentation for the specific differences between these two syntaxes. Normally, headers for standard or system libraries are found on the default system include path, and your own project headers are found on the quoted include path.

The header operand passed to the `__has_include` preprocessor operator is specified with either quotes or angle brackets. The operator uses the same search path heuristics as the `#include` directive. Consequently, you should use the same form for both the `#include` directive and corresponding `__has_include` operator to ensure a consistent result.

Conditional Inclusion

Frequently, you'll need to write different code to support different implementations. For example, you may want to provide alternative implementations of a function for different target architectures. One solution to this problem is to maintain two files with slight variations between them and compile the appropriate file for a particular implementation. A better solution is to either translate or refrain from translating the target-specific code based on a preprocessor definition.

You can conditionally include source code using a preprocessing directive such as `#if`, `#elif`, or `#else` with a predicate condition. A *predicate condition* is the controlling constant expression that's evaluated to determine which branch of the program the preprocessor should take. They're typically used along with the preprocessor defined operator, which determines whether a given identifier is the name of a defined macro.

The conditional inclusion directives are similar to the `if` and `else` statements. When the predicate condition is evaluated to a nonzero preprocessor value, the `#if` branch is processed, and all other branches are not. When the predicate condition evaluates to zero, the next `#elif` branch, if any, has its predicate tested for inclusion. If none of the predicate conditions evaluate to nonzero, then the `#else` branch, if there is one, is processed. The `#endif` preprocessing directive indicates the end of the conditionally included code.

The `defined` operator evaluates to 1 if the given identifier is defined as a macro or 0 otherwise. For example, the preprocessing directives shown in [Listing 9-1](#) conditionally determine which header contents to include in the resulting translation unit. The preprocessed output depends on whether `_WIN32` or `__ANDROID__` is a defined macro. If neither is a defined macro, the preprocessor output will be empty.

```
#if defined(_WIN32)
# include <Windows.h>
#elif defined(__ANDROID__)
# include <android/log.h>
#endif
```

Listing 9-1: An example of conditional inclusion

Unlike with the `if` and `else` keywords, preprocessor conditional inclusion cannot use braces to denote the block of statements controlled by the predicate. Instead, preprocessor conditional inclusion will include all the tokens from the `#if`, `#elif`, or `#else` directive (following the predicate) to the next balanced `#elif`, `#else`, or `#endif` token found, while skipping any tokens in a conditional inclusion branch not taken. Conditional inclusion directives can be nested. You can write

```
#ifdef identifier
```

as shorthand for:

```
#if defined identifier
```

Similarly, you can write

```
#ifndef identifier
```

as shorthand for:

```
#if !defined identifier
```

Starting in C23, you can write

```
#elifdef identifier
```

as shorthand for

```
#elif defined identifier
```

and you can write

```
#elifndef identifier
```

as shorthand for

```
#elif !defined identifier
```

or, equivalently:

```
#elif !defined(identifier)
```

The parentheses around the identifier are optional.

Generating Diagnostics

A conditional inclusion directive may need to generate an error if the preprocessor can't take any of the conditional branches because no reasonable fallback behavior exists. Consider the example in [Listing 9-2](#), which uses conditional inclusion to select between including the C standard library header `<threads.h>` or the POSIX threading library header `<pthread.h>`. If neither option is available, you should alert the programmer porting the system that the code must be repaired.

```
#if __STDC__ && __STDC_NO_THREADS__ != 1
#   include <threads.h>
#elif POSIX_THREADS == 200809L
#   include <pthread.h>
#else
    int compile_error[-1]; // induce a compilation error
#endif
```

Listing 9-2: Inducing a compilation error

Here, the code generates a diagnostic but doesn't describe the actual problem. For this reason, C has the `#error` preprocessing directive, which causes the implementation to produce a compile-time diagnostic message. You can optionally follow this directive with one or more preprocessor tokens to include in the resulting diagnostic message. Using these, we can replace the erroneous array declaration from [Listing 9-2](#) with an `#error` directive such as the one

shown in [Listing 9-3](#).

```
#if __STDC__ && __STDC_NO_THREADS__ != 1
#   include <threads.h>
#elif POSIX_THREADS == 200809L
#   include <pthread.h>
#else
#   error "Neither <threads.h> nor <pthread.h> is available"
#endif
```

Listing 9-3: An #error directive

This code generates the following error message if neither threading library header is available:

```
Neither <threads.h> nor <pthread.h> is available
```

In addition to the `#error` directive, C23 added the `#warning` directive. This directive is like the `#error` directive in that they both cause the implementation to generate a diagnostic. However, instead of terminating compilation, the diagnostic message is generated, and compilation continues as normal (unless other command line options disable the warnings or upgrade them into errors). The `#error` directive should be used for *fatal* problems such as a missing library with no fallback implementation, while the `#warning` directive should be used for *nonfatal* problems such as a missing library with a low-quality fallback implementation.

Using Header Guards

One problem you'll encounter when writing headers is preventing programmers from including the same file twice in a translation unit. Given that you can transitively include headers, you could easily include the same header multiple times by accident (possibly even leading to infinite recursion between headers).

Header guards ensure that a header is included only once per translation unit. A header guard is a design pattern that conditionally includes the contents of a header based on whether a header-specific macro is defined. If the macro is not already defined, you define it so that a subsequent test of the header guard won't conditionally include the code. In the program shown in [Table 9-4](#), *bar.h* uses a header guard (shown in bold) to prevent its (accidental) duplicate inclusion from *foo.c*.

Table 9-4: A Header Guard

Original sources	Resulting translation unit
<pre> bar.h #ifndef BAR_H #define BAR_H inline int func() { return 1; } #endif /* BAR_H */ </pre>	<pre> inline int func() { return 1; } extern inline int func(); int main() { return func(); } </pre>
<pre> foo.c #include "bar.h" #include "bar.h" // repeated inclusion is // usually not this obvious extern inline int func(); int main() { return func(); } </pre>	

The first time that "bar.h" is included, the `#ifndef` test to see that `BAR_H` is not defined will return `true`. We then define the macro `BAR_H` with an empty replacement list, which is sufficient to define `BAR_H`, and the function definition for `func` is included. The second time that "bar.h" is included, the preprocessor won't generate any tokens because the conditional inclusion test will return `false`. Consequently, `func` is defined only once in the resulting translation unit.

A common practice when picking the identifier to use as a header guard is to use the salient parts of the file path, filename, and extension, separated by an underscore and written in all capital letters. For example, if you had a header that would be included with `#include "foo/bar/baz.h"`, you might choose `FOO_BAR_BAZ_H` as the header guard identifier.

Some IDEs will automatically generate the header guard for you. Avoid using a reserved identifier as your header guard's identifier, which could introduce undefined behavior. Identifiers that begin with an underscore followed by a capital letter are reserved. For example, `_FOO_H` is a reserved identifier and a bad choice for a user-chosen header guard identifier, even if you're including a file named `_foo.h`. Using a reserved identifier can result in a collision with a macro defined by the implementation, leading to a compilation error or incorrect code.

Macro Definitions

The `#define` preprocessing directive defines a macro. You can use *macros* to define constant values or function-like constructs with generic parameters. The macro definition contains a (possibly empty) *replacement list*—a code pattern

that's injected into the translation unit when the preprocessor expands the macro:

```
#define identifier replacement-list
```

The `#define` preprocessing directive is terminated with a newline. In the following example, the replacement list for `ARRAY_SIZE` is `100`:

```
#define ARRAY_SIZE 100  
int array[ARRAY_SIZE];
```

Here the `ARRAY_SIZE` identifier is replaced by `100`. If no replacement list is specified, the preprocessor simply removes the macro name. You can typically specify a macro definition on your compiler's command line—for example, using the `-D` flag in Clang and GCC or the `/D` flag in Visual C++. For Clang and GCC, the command line option `-DARRAY_SIZE=100` specifies that the macro identifier `ARRAY_SIZE` is replaced by `100`, producing the same result as the `#define` preprocessing directive from the previous example. If you don't specify the macro replacement list on the command line, compilers will typically provide a replacement list. For example, `-DFOO` is typically identical to `#define FOO 1`.

The scope of a macro lasts until the preprocessor encounters either an `#undef` preprocessing directive specifying that macro or the end of the translation unit. Unlike variable or function declarations, a macro's scope is independent of any block structure.

You can use the `#define` directive to define either an object-like macro or a function-like macro. A *function-like* macro is parameterized and requires passing a (possibly empty) set of arguments when you invoke it, similar to how you would invoke a function. Unlike functions, macros let you perform operations using the program's symbols, which means you can create a new variable name or reference the source file and line number at which the macro is being invoked. An *object-like* macro is a simple identifier that will be replaced by a code fragment.

[Table 9-5](#) illustrates the difference between function-like and object-like macros. `FOO` is an object-like macro that is replaced by the tokens `(1 + 1)` during macro expansion, and `BAR` is a function-like macro that is replaced by the tokens `(1 + (x))`, where `x` is whatever parameter is specified when invoking `BAR`.

Table 9-5: Macro Definition

Original source	Resulting translation unit
<pre>#define F00 (1 + 1) #define BAR(x) (1 + (x)) int i = F00; int j = BAR(10); int k = BAR(2 + 2);</pre>	<pre>int i = (1 + 1); int j = (1 + (10)); int k = (1 + (2 + 2));</pre>

The opening parenthesis of a function-like macro definition must immediately follow the macro name, with no intervening whitespace. If a space appears between the macro name and the opening parenthesis, the parenthesis simply becomes part of the replacement list, as is the case with the object-like F00 macro. The macro replacement list terminates with the first newline character in the macro definition. However, you can join multiple source lines with the backslash (\) character followed by a newline to make your macro definitions easier to understand. For example, consider the following definition of the `cbrt` type-generic macro that computes the cube root of its floating-point argument:

```
#define cbrt(X) _Generic((X), \
    long double: cbrt1(X), \
    default: cbrt(X), \
    float: cbrtf(X) \
)
```

That definition is equivalent to, but easier to read, than the following:

```
#define cbrt(X) _Generic((X), long double: cbrt1(X), default: cbrt(X), float: cbrtf(X))
```

One danger when defining a macro is you can no longer use the macro's identifier in the rest of the program without inducing a macro replacement. For example, because of macro expansion, the following invalid program won't compile:

```
#define foo (1 + 1)
void foo(int i);
```

This is because the tokens the preprocessor receives from the translator result in the following invalid code:

```
void (1 + 1)(int i);
```

You can solve this problem by consistently adhering to an idiom throughout your program, such as defining macro names with all uppercase letters or prefixing all macro names with a mnemonic, as you might find in some styles of Hungarian notation.

NOTE

Hungarian notation is an identifier-naming convention in which the name of a variable or function indicates its intention or kind and, in some dialects, its type.

After you've defined a macro, the only way to redefine it is to invoke the `#undef` directive for the macro. Once you've undefined it, the named identifier no longer represents a macro. For example, the program shown in [Table 9-6](#) defines a function-like macro, includes a header that uses the macro, and then undefines the macro so that it can be redefined later.

Table 9-6: Undefining Macros

Original sources	Resulting translation unit
<pre>header.h NAME(first) NAME(second) NAME(third)</pre>	<pre>enum Names { first, second, third, };</pre>
<pre>file.c enum Names { #define NAME(X) X, #include "header.h" #undef NAME }; void func(enum Names Name) { switch (Name) { #define NAME(X) case X: #include "header.h" #undef NAME } }</pre>	<pre>void func(enum Names Name) { switch (Name){ case first: case second: case third: } }</pre>

The first use of the `NAME` macro declares the names of enumerators within the `Names` enumeration. The `NAME` macro is undefined and then redefined to generate the case labels in a `switch` statement.

Undefining a macro is safe even when the named identifier isn't the name of a macro. This macro definition works regardless of whether `NAME` is already defined. To keep examples short, we don't normally follow this practice in this

book.

Macro Replacement

Function-like macros look like functions but behave differently. When the preprocessor encounters a macro identifier, it invokes the macro, which expands the identifier to replace it with the tokens from the replacement list, if any, specified in the macro's definition.

For function-like macros, the preprocessor replaces all parameters in the replacement list with the corresponding arguments in the macro invocation after expanding them. Any parameter in the replacement list preceded by a `#` token is replaced with a string literal preprocessing token that contains the text of the argument preprocessing tokens (a process sometimes called *stringizing*). The `STRINGIZE` macro in [Table 9-7](#) stringizes the value of `x`.

Table 9-7: Stringizing

Original source	Resulting translation unit
<pre>#define STRINGIZE(x) #x const char *str = STRINGIZE(12);</pre>	<pre>const char *str = "12";</pre>

The preprocessor also deletes all instances of the `##` preprocessing token in the replacement list, concatenating the preceding preprocessing token with the following token, which is called *token pasting*. The `PASTE` macro in [Table 9-8](#) is used to create a new identifier by concatenating `foo`, the underscore character (`_`), and `bar`.

Table 9-8: Token Pasting

Original source	Resulting translation unit
<pre>#define PASTE(x, y) x ## _ ## y int PASTE(foo, bar) = 12;</pre>	<pre>int foo_bar = 12;</pre>

After expanding the macro, the preprocessor rescans the replacement list to expand additional macros within it. If the preprocessor finds the name of the macro being expanded while rescanning—including the rescanning of nested macro expansions within the replacement list—it won't expand the name again. Furthermore, if macro expansion results in a fragment of program text that's identical to a preprocessing directive, that fragment won't be treated as a preprocessing directive.

During macro expansion, a repeated parameter name in the replacement list

will be replaced multiple times by the argument given in the invocation. This can have surprising effects if the argument to the macro invocation involves side effects, as shown in [Table 9-9](#). This problem is explained in detail in CERT C rule PRE31-C, “Avoid side effects in arguments to unsafe macros.”

Table 9-9: Unsafe Macro Expansion

Original source	Resulting translation unit
<pre>#define bad_abs(x) (x >= 0 ? x : -x) int func(int i) { } return bad_abs(i++); }</pre>	<pre>int func(int i) { return (i++ >= 0 ? i++ : -i++); }</pre>

In the macro definition in [Table 9-9](#), each instance of the macro parameter *x* is replaced by the macro invocation argument *i++*, causing *i* to be incremented twice in a way that a programmer or reviewer reading the original source code can easily overlook. Parameters like *x* in the replacement list, as well as the replacement list itself, should usually be fully parenthesized as in `((x) >= 0 ? (x) : -(x))` to prevent portions of the argument *x* from associating with other elements of the replacement list in unexpected ways.

GNU *statement expressions* allow you to use loops, switches, and local variables within an expression. Statement expressions are a nonstandard compiler extension supported by GCC, Clang, and other compilers. Using statement expressions, you can rewrite the `bad_abs(x)` as follows:

```
#define abs(x) ({
    auto x_tmp = x;
    x_tmp >= 0 ? x_tmp : x_tmp;
})
```

You can safely invoke the `abs(x)` macro with side-effecting operands.

Another potential surprise is that a comma in a function-like macro invocation is always interpreted as a macro argument delimiter. The C standard `ATOMIC_VAR_INIT` macro (removed in C23) demonstrates the danger ([Table 9-10](#)).

Table 9-10: The `ATOMIC_VAR_INIT` Macro

Original sources	Resulting translation unit
<pre>stdatomic.h #define ATOMIC_VAR_INIT(value) (value)</pre>	<error>
<pre>foo.c</pre>	


```
#include <stdatomic.h>

struct S {
    int x, y;
};

Atomic struct S val = ATOMIC_VAR_INIT({1, 2});
```

This code fails to translate because the comma in `ATOMIC_VAR_INIT({1, 2})` is treated as a function-like macro argument delimiter, causing the preprocessor to interpret the macro as having two syntactically invalid arguments {1 and 2} instead of a single valid argument {1, 2}. This usability issue is one of the reasons the `ATOMIC_VAR_INIT` macro was deprecated in C17 and removed in C23.

Type-Generic Macros

The C programming language doesn't allow you to overload functions based on the types of the parameters passed to the function, as you can in other languages such as Java and C++. However, you might sometimes need to alter the behavior of an algorithm based on the argument types. For example, `<math.h>` has three `sin` functions (`sin`, `sinf`, and `sinl`) because each of the three floating-point types (double, float, and long double, respectively) has a different precision. Using generic selection expressions, you can define a single function-like identifier that delegates to the correct underlying implementation based on the argument type when called.

A *generic selection expression* maps the type of its unevaluated operand expression to an associated expression. If none of the associated types match, it can optionally map to a default expression. You can use *type-generic macros* (macros that include generic selection expressions) to make your code more readable. In [Table 9-11](#), we define a type-generic macro to select the correct variant of the `sin` function from `<math.h>`.

Table 9-11: A Generic Selection Expression as a Macro

Original source	Resulting _Generic resolution
<pre>#define singen(X) _Generic((X), \ float: sinf, \ double: sin \ long double: sinl \)(X) int main() { printf("%f, %Lf\n",</pre>	<pre>int main() { printf("%f, %Lf\n", sin(3.14159), sinl(1.5708L)); }</pre>

```
singen(3.14159),
singen(1.5708L)
);
}
```

The controlling expression (x) of the generic selection expression is unevaluated; the type of the expression selects a function from the list of type : expr mappings. The generic selection expression picks one of these function designators (either `sinf`, `sin`, or `sinl`) and then executes it. In this example, the argument type in the first call to `singen` is `double`, so the generic selection resolves to `sin`, and the argument type in the second call to `singen` is `long double`, so this resolves to `sinl`. Because this generic selection expression has no default association, an error occurs if the type of (x) doesn't match any of the associated types. If you include a default association for a generic selection expression, it will match every type not already used as an association, including types you might not expect, such as pointers or structure types.

Type-generic macro expansion can be difficult to use when the resulting value type depends on the type of an argument to the macro, as with the `singen` example in [Table 9-11](#). For instance, it can be a mistake to call the `singen` macro and assign the result to an object of a specific type or pass its result as an argument to `printf` because the necessary object type or format specifier will depend on whether `sin`, `sinf`, or `sinl` is called. You can find examples of type-generic macros for math functions in the C standard library `<tgmath.h>` header.

C23 partially addressed this problem with the introduction of automatic type inference using the `auto` type specifier, described in [Chapter 2](#). Consider using automatic type inference when initializing an object with a type-generic macro to avoid unintentional conversions on initialization. For example, the following file scope definitions

```
static auto a = sin(3.5f);
static auto p = &a;
```

are interpreted as if they had been written as:

```
static float a = sinf(3.5f);
static float *p = &a;
```

Effectively, `a` is a `float` and `p` is a `float *`.

In [Table 9-12](#), we replace the types of the two variables declared in `main` from [Table 9-11](#) with the `auto` type specifier. This makes it easier to invoke a type-generic macro, although it is not strictly necessary as the programmer can also deduce the types. The `auto` type specifier is useful when invoking a type-generic function-like macro where the type of the resulting value depends on a macro parameter to avoid accidental type conversion on initialization.

Table 9-12: Type-Generic Macros with Automatic Type Inference

Original source	Resulting <code>_Generic</code> resolution
<pre>#define singen(X) _Generic((X), \ float: sinf, \ double: sin, \ long double: sinl \)(X) int main(void) { auto f = singen(1.5708f); auto d = singen(3.14159); }</pre>	<pre>int main(void) { auto f = sinf(1.5708f); auto d = sin(3.14159); }</pre>

You can also use the `auto` type specifier for declaring variables in type-generic macros where you don't know the type of the arguments.

Embedded Binary Resources

You may find that you need to dynamically load digital resources, such as images, sounds, video, text files, or other binary data, at runtime. It may instead be beneficial to load those resources at compile time so they can be stored as part of the executable rather than dynamically loaded.

Prior to C23, there were two common approaches to embedding binary resources into your program. For limited amounts of binary data, the data could be specified as an initializer for a constant-size array. However, for larger binary resources, this approach could introduce significant compile-time overhead, so the use of a linker script or other postprocessing was necessary to keep compile times reasonable.

C23 added the `#embed` preprocessor directive to embed a digital resource directly into source code as if it were a comma-delimited list of integer constants. The new directive allows an implementation to optimize for better compile time efficiency when using the embedded constant data as an array initializer. Using `#embed`, the implementation does not need to parse each integer constant and comma token separately; it can inspect the bytes directly and use a

more efficient mapping of the resource.

[Table 9-13](#) shows an example of embedding the binary resource *file.txt* as the initializer for the buffer array declaration. For this example, *file.txt* contains the ASCII text *meow* to keep the code listing short. Significantly larger binary resources are typically embedded.

Table 9-13: Embedding Binary Resources

Original source	Resulting translation unit
<pre>unsigned char buffer[] = { #embed <file.txt> };</pre>	<pre>unsigned char buffer[] = { 109, 101, 111, 119 };</pre>

Like `#include`, the filename specified in the `#embed` directive can be listed within either angle brackets or double quotes. Unlike `#include`, there is no notion of *system* or *user* embedded resources, so the only difference between the two forms is that the double-quoted form will start searching for the resource from the same directory as the source file before trying other search paths. Compilers have a command line option to specify search paths for embedded resources; see your compiler documentation for more details.

The `#embed` directive supports several parameters to control what data is embedded into the source file: `limit`, `suffix`, `prefix`, and `if_empty`. The most useful of those parameters is the `limit` parameter, which specifies how much data to embed (in bytes). This can be helpful if the only data needed at compile time is in the header of the file or if the file is an *infinite* resource like */dev/urandom* in some operating systems. The `prefix` and `suffix` parameters insert tokens before or after the embedded resource, respectively, if the resource is found and is not empty. The `if_empty` parameter inserts tokens if the embedded resource is found but has no content (including when the `limit` parameter is explicitly set to 0).

Like `__has_include`, you can test whether an embedded resource can be found using the `__has_embed` preprocessor operator. This operator returns:

- `__STDC_EMBED_FOUND__` if the resource can be found and isn't empty
- `__STDC_EMBED_EMPTY__` if the resource can be found and is empty
- `__STDC_EMBED_NOT_FOUND__` if the resource cannot be found

Predefined Macros

The implementation defines some macros without requiring you to include a

header. These macros are called *predefined macros* because they're implicitly defined by the preprocessor rather than explicitly defined by the programmer. For example, the C standard defines various macros that you can use to interrogate the compilation environment or provide basic functionality. Some other aspects of the implementation (such as the compiler or the compilation target operating system) also automatically define macros. [Table 9-14](#) lists some of the common macros the C standard defines. You can obtain a full list of predefined macros from Clang or GCC by passing the `-E -dM` flags to those compilers. Check your compiler documentation for more information.

Table 9-14: Predefined Macros

Macro name	Replacement and purpose
<code>__DATE__</code>	A string literal of the date of translation of the preprocessing translation unit in the form <i>Mmm dd yyyy</i> .
<code>__TIME__</code>	A string literal of the time of translation for the preprocessing translation unit in the form <i>hh:mm:ss</i> .
<code>__FILE__</code>	A string literal representing the presumed filename of the current source file.
<code>__LINE__</code>	An integer constant representing the presumed line number of the current source line.
<code>__STDC__</code>	The integer constant 1 if the implementation conforms to the C standard.
<code>__STDC_HOSTED__</code>	The integer constant 1 if the implementation is a hosted implementation or the integer constant 0 if it is stand-alone. This macro is conditionally defined by the implementation.
<code>__STDC_VERSION__</code>	The integer constant representing the version of the C standard the compiler is targeting, such as 202311L for the C23 standard.
<code>__STDC_UTF_16__</code>	The integer constant 1 if values of type <code>char16_t</code> are UTF-16 encoded. This macro is conditionally defined by the implementation.
<code>__STDC_UTF_32__</code>	The integer constant 1 if values of type <code>char32_t</code> are UTF-32 encoded. This macro is conditionally defined by the implementation.
<code>__STDC_NO_ATOMICS__</code>	The integer constant 1 if the implementation doesn't support atomic types, including the <code>_Atomic</code> type qualifier, and the <code><stdatomic.h></code> header. This macro is conditionally defined by the implementation.
<code>__STDC_NO_COMPLEX__</code>	The integer constant 1 if the implementation doesn't support complex types or the <code><complex.h></code> header. This macro is conditionally defined by the implementation.
<code>__STDC_NO_THREADS__</code>	The integer constant 1 if the implementation doesn't support the <code><threads.h></code> header. This macro is conditionally defined by the implementation.
<code>__STDC_NO_VLA__</code>	The integer constant 1 if the implementation doesn't support variable-length arrays. This macro is conditionally defined by the implementation.

Summary

In this chapter, you learned about some of the features provided by the preprocessor. You learned how to include fragments of program text in a translation unit, conditionally compile code, embed binary resources into your program, and generate diagnostics on demand. You then learned how to define

and undefine macros, how macros are invoked, and about macros that are predefined by the implementation. The preprocessor is popular in C language programming but shunned in C++ programming. Use of the preprocessor can be error prone, so it is best to follow the recommendations and rules from *The CERT C Coding Standard*.

In the next chapter, you'll learn how to structure your program into more than one translation unit to create more maintainable programs.

10

PROGRAM STRUCTURE

with Aaron Ballman



Any real-world system is made up of multiple components, such as source files, headers, and libraries. Many contain resources including images, sounds, and configuration files. Composing a program from smaller logical components is good software engineering practice, because these components are easier to manage than a single large file.

In this chapter, you'll learn how to structure your program into multiple units consisting of both source and include files. You'll also learn how to link multiple object files together to create libraries and executable files.

Principles of Componentization

Nothing prevents you from writing your entire program within the `main` function of a single source file. However, as the function grows, that approach will quickly become unmanageable. For this reason, it makes sense to decompose your program into a collection of components that exchange information across a shared boundary, or *interface*. Organizing source code into components makes it easier to understand and allows you to reuse the code elsewhere in the program, or even with other programs.

Understanding how best to decompose a program typically requires experience. Many of the decisions programmers make are driven by

performance. For example, you may need to minimize communication over a high-latency interface. Bad hardware can only go so far; you need bad software to really screw up performance.

Performance is only one software quality attribute (ISO/IEC 25000:2014) and must be balanced with maintainability, code readability, understandability, safety, and security. For example, you may design a client application to handle input field validation from the user interface to avoid a round trip to the server. This helps performance but can hurt security if inputs to the server are not validated. A simple solution is to validate inputs in both locations.

Developers frequently do strange things for illusionary gains. The strangest of these is invoking the undefined behavior of signed integer overflow to improve performance. Frequently, these local code optimizations have no impact on overall system performance and are considered *premature optimizations*. Donald Knuth, author of *The Art of Computer Programming* (Addison-Wesley, 1997), described premature optimization as “the root of all evil.”

In this section, we’ll cover some principles of component-based software engineering.

Coupling and Cohesion

In addition to performance, the aim of a well-structured program is to achieve desirable properties like low coupling and high cohesion. *Cohesion* is a measure of the commonality between elements of a programming interface. Assume, for example, that a header exposes functions for calculating the length of a string, calculating the tangent of a given input value, and creating a thread. This header has low cohesion because the exposed functions are unrelated to each other. Conversely, a header that exposes functions to calculate the length of a string, concatenate two strings together, and search for a substring within a string has high cohesion because all the functionality is related. This way, if you need to work with strings, you need only to include the string header. Similarly, related functions and type definitions that form a public interface should be exposed by the same header to provide a highly cohesive interface of limited functionality. We’ll discuss public interfaces further in “Data Abstractions” on [page 215](#).

Coupling is a measure of the interdependency of programming interfaces. For example, a tightly coupled header can’t be included in a program by itself; instead, it must be included with other headers in a specific order. You may couple interfaces for a variety of reasons, such as a mutual reliance on data structures, interdependence between functions, or the use of a shared global

state. But when interfaces are tightly coupled, modifying program behavior becomes difficult because changes can have a ripple effect across the system. You should always strive to have loose coupling between interface components, regardless of whether they're members of a public interface or implementation details of the program.

By separating your program logic into distinct, highly cohesive components, you make it easier to reason about the components and test the program (because you can verify the correctness of each component independently). The result is a more maintainable, less buggy system.

Code Reuse

Code reuse is the practice of implementing functionality once and then reusing it in various parts of the program without duplicating the code. Code duplication can lead to subtly unexpected behavior, oversized and bloated executables, and increased maintenance costs. And anyway, why write the same code more than once?

Functions are the lowest-level reusable units of functionality. Any logic that you might repeat more than once is a candidate for encapsulating in a function. If the functionality has only minor differences, you might be able to create a parameterized function that serves multiple purposes. Each function should perform work that isn't duplicated by any other function. You can then compose individual functions to solve increasingly sophisticated problems.

Packaging reusable logic into functions can improve maintainability and eliminate defects. For example, though you could determine the length of a null-terminated string by writing a simple `for` loop, it's more maintainable to use the `strlen` function from the C standard library. Because other programmers are already familiar with the `strlen` function, they'll have an easier time understanding what that function is doing than what the `for` loop is doing. Furthermore, if you reuse existing functionality, you're less likely to introduce behavioral differences in ad hoc implementations, and you make it easier to globally replace the functionality with a better-performing algorithm or more secure implementation, for example.

When designing functional interfaces, a balance must be struck between *generality* and *specificity*. An interface that's specific to the current requirement may be lean and effective but hard to modify when requirements change. A general interface might allow for future requirements but be cumbersome for foreseeable needs.

Data Abstractions

A *data abstraction* is any reusable software component that enforces a clear separation between the abstraction's public interface and the implementation details. The *public interface* for each data abstraction includes the data type definitions, function declarations, and constant definitions required by the users of the data abstraction and is placed in headers. The implementation details of how the data abstraction is coded, as well as any private utility functions, are hidden within source files or in headers that are in a separate location from the public interface headers. This separation of the public interface from the private implementation allows you to change the implementation details without breaking code that depends on your component.

Header files typically contain function declarations and type definitions for the component. For example, the C standard library `<string.h>` header provides the public interface for string-related functionality, while `<threads.h>` provides utility functions for threading. This logical separation has low coupling and high cohesion, making it easier to access only the specific components you need and reduce compile time and the likelihood of name collisions. You don't need to know anything about threading application programming interfaces (APIs), for example, if all you need is the `strlen` function.

Another consideration is whether you should explicitly include the headers required by your header or require the users of the header to include them first. It's a good idea for data abstractions to be self-contained and include the headers they use. Not doing so is a burden on the users of the abstraction and leaks implementation details about the data abstraction. Examples in this book don't always follow this practice to keep these file listings concise.

Source files implement the functionality declared by a given header or the application-specific program logic used to perform whatever actions are needed for a given program. For example, if you have a *network.h* header that describes a public interface for network communications, you may have a *network.c* source file (or *network_win32.c* for Windows only and *network_linux.c* for Linux only) that implements the network communication logic.

It's possible to share implementation details between two source files by using a header, but the header file should be placed in a distinct location from the public interface to prevent accidentally exposing implementation details.

A *collection* is a good example of a data abstraction that separates the basic functionality from the implementation or underlying data structure. A collection groups data elements and supports operations such as adding elements to the

collection, removing data elements from the collection, and checking whether the collection contains a specific data element.

There are many ways to implement a collection. For example, a collection of data elements may be represented as a flat array, a binary tree, a directed (possibly acyclic) graph, or a different structure. The choice of data structure can impact an algorithm's performance, depending on what kind of data you're representing and how much data there is to represent. For example, a binary tree may be a better abstraction for a large amount of data that needs good lookup performance, whereas a flat array is likely a better abstraction for a small amount of data of fixed size. Separating the interface of the collection data abstraction from the implementation of the underlying data structure allows the implementation to change without requiring changes to code that relies on the collection interface.

Opaque Types

Data abstractions are most effective when used with opaque data types that hide information. In C, *opaque* (or *private*) data types are those expressed using an incomplete type, such as a forward-declared structure type. An *incomplete type* is a type that describes an identifier but lacks information needed to determine the size of objects of that type or their layout. Hiding internal-only data structures discourages programmers who use the data abstraction from writing code that depends on implementation details, which may change. The incomplete type is exposed to users of the data abstraction, while the fully defined type is accessible only to the implementation.

Say we want to implement a collection that supports a limited number of operations, such as adding an element, removing an element, and searching for an element. The following example implements `collection_type` as an opaque type, hiding the implementation details of the data type from the library's user. To accomplish this, we create two headers: an external *collection.h* header included by the user of the data type and an internal header included only in files that implement the functionality of the data type.

collection.h

```
typedef struct collection * collection_type;
// function declarations
extern errno_t create_collection(collection_type *result);
extern void destroy_collection(collection_type col);
extern errno_t add_to_collection(
```

```
    collection_type col, const void *data, size_t byteCount
);
extern errno_t remove_from_collection(
    collection_type col, const void *data, size_t byteCount
);
extern errno_t find_in_collection(
    const collection_type col, const void *data, size_t byteCount
);
// --snip--
```

The `collection_type` identifier is aliased to `struct collection_type` (an incomplete type). Consequently, functions in the public interface must accept a pointer to this type, instead of an actual value type, because of the constraints placed on the use of incomplete types in C.

In the internal header, `struct collection_type` is fully defined but not visible to a user of the data abstraction:

collection_priv.h

```
struct node_type {
    void *data;
    size_t size;
    struct node_type *next;
};

struct collection_type {
    size_t num_elements;
    struct node_type *head;
};
```

Users of the data abstraction include only the external *collection.h* file, whereas modules that implement the abstract data type also include the internal definitions *collection_priv.h* file. This allows the implementation of the `collection_type` data type to remain private.

Executables

In [Chapter 9](#), we learned that the compiler is a pipeline of translation phases and that the compiler's ultimate output is object code. The last phase of translation, called the *link phase*, takes the object code for all the translation units in the program and links them together to form a final executable. This can be an executable that a user can run, such as *a.out* or *foo.exe*, a library, or a more

specialized program such as a device driver or a firmware image (machine code to be burned onto read-only memory [ROM]). Linking allows you to separate your code into distinct source files that can be compiled independently, which helps to build reusable components.

Libraries are executable components that cannot be executed independently. Instead, you incorporate libraries into executable programs. You can invoke the functionality of the library by including the library's headers in your source code and calling the declared functions. The C standard library is an example of a library—you include the headers from the library, but you do not directly compile the source code that implements the library functionality. Instead, the implementation ships with a prebuilt version of the library code.

Libraries allow you to build on the work of others for the generic components of a program so you can focus on developing the logic that is unique to your program. For example, when writing a video game, reusing existing libraries should allow you to focus on developing the game logic, not worrying about the details of retrieving user input, network communications, or graphics rendering. Libraries compiled with one compiler can often be used by programs built with a different compiler.

Libraries are linked into your application and can be either static or dynamic. A *static library*, also known as an *archive*, incorporates its machine or object code directly into the resulting executable, which means that a static library is often tied to a specific release of the program. Because a static library is incorporated at link time, the contents of the static library can be further optimized for your program's use of the library. Library code used by the program can be made available for link-time optimizations (for example, using the `-flto` flag), while unused library code can be stripped from the final executable.

A *dynamic library*, also referred to as a *shared library* or a *dynamic shared object*, is an executable without the startup routines. It can be packaged with the executable or installed separately but must be available when the executable calls a function provided by the dynamic library. Many modern operating systems will load the dynamic library code into memory once and share it across all the applications that need it. You can replace a dynamic library with different versions as necessary after your application has been deployed.

Letting the library evolve separately from the program comes with its own set of benefits and risks. A developer can correct bugs in the library after an application has already shipped without requiring the application to be

recompiled, for instance. However, dynamic libraries provide the potential opportunity for a malicious attacker to replace a library with a nefarious one or an end user to accidentally use an incorrect version of the library. It's also possible to make a *breaking change* in a new library release that results in an incompatibility with existing applications that use the library. Static libraries might execute somewhat faster because the object code (binary) is included in the executable file, enabling further optimizations. The benefits of using dynamic libraries usually outweigh the disadvantages.

Each library has one or more headers that contain the public interface to the library and one or more source files that implement the logic for the library. You can benefit from structuring your code as a collection of libraries even if the components aren't turned into actual libraries. Using an actual library makes it harder to accidentally design a tightly coupled interface where one component has special knowledge of the internal details of another component.

Linkage

Linkage is a process that controls whether an interface is public or private and determines whether any two identifiers refer to the same entity. Ignoring macros and macro parameters that are replaced early in the translation phases, an *identifier* can denote a standard attribute, an attribute prefix, or an attribute name; an object; a function; a tag or a member of a structure, union, or enumeration; a typedef name; or a label name.

C provides three kinds of linkage: external, internal, or none. Each declaration of an identifier with *external linkage* refers to the same function or object everywhere in the program. Identifiers referring to declarations with internal linkage refer to the same entity only within the translation unit containing the declaration. If two translation units both refer to the same internal linkage identifier, they refer to different instances of the entity. If a declaration has *no linkage*, it's a unique entity in each translation unit.

The linkage of a declaration is either explicitly declared or implied. If you declare an entity at file scope without explicitly specifying `extern` or `static`, the entity is implicitly given external linkage. Identifiers that have no linkage include function parameters, block scope identifiers declared without an `extern` storage class specifier, or enumeration constants.

[Listing 10-1](#) shows examples of declarations of each kind of linkage.

```
static int i; // i has explicit internal linkage
```

```
extern void foo(int j) {  
    // foo has explicit external linkage  
    // j has no linkage because it is a parameter  
}
```

Listing 10-1: Examples of internal, external, and no linkage

If you explicitly declare an identifier with the static storage class specifier at file scope, it has internal linkage. The static keyword gives internal linkage only to file scope entities. Declaring a variable at block scope as static creates an identifier with no linkage, but it does give the variable static storage duration. As a reminder, static storage duration means its lifetime is the entire execution of the program, and its stored value is initialized only once, prior to program startup. The different meanings of *static* when used in different contexts are obviously confusing and consequently a common interview question.

You can create an identifier with external linkage by declaring it with the extern storage class specifier. This works only if you haven't previously declared the linkage for that identifier. The extern storage class specifier has no effect if a prior declaration gave the identifier linkage.

Declarations with conflicting linkage can lead to undefined behavior; see CERT C rule DCL36-C, “Do not declare an identifier with conflicting linkage classifications,” for more information.

[Listing 10-2](#) shows sample declarations with implicit linkage.

foo.c

```
void func(int i) { // implicit external linkage  
    // i has no linkage  
}  
static void bar(); // internal linkage, different bar from bar.c  
extern void bar() {  
    // bar still has internal linkage because the initial declarati  
on  
    // was declared as static; this extern specifier has no effect  
}
```

Listing 10-2: Examples of implicit linkage

[Listing 10-3](#) shows sample declarations with explicit linkage.

bar.c

```
extern void func(int i); // explicit external linkage
static void bar() { // internal linkage; different bar from foo.
c
    func(12); // calls func from foo.c
}
int i; // external linkage; doesn't conflict with i from foo.c or
bar.c
void baz(int k) { // implicit external linkage
    bar(); // calls bar from bar.c, not foo.c
}
```

Listing 10-3: Examples of explicit linkage

The identifiers in your public interface should have external linkage so that they can be called from outside their translation unit. Identifiers that are implementation details should be declared with internal or no linkage (provided they don't need to be referenced from another translation unit). A common approach to achieving this is to declare your public interface functions in a header with or without using the `extern` storage class specifier (the declarations implicitly have external linkage, but there is no harm in explicitly declaring them with `extern`) and define the public interface functions in a source file in a similar manner.

However, within the source file, all declarations that are implementation details should be explicitly declared `static` to keep them private—accessible to just that source file. You can include the public interface declared within the header by using the `#include` preprocessor directive to access its interface from another file. A good rule of thumb is that file-scope entities that don't need to be visible outside the file should be declared as `static`. This practice limits the global namespace pollution and decreases the chances of surprising interactions between translation units.

Structuring a Simple Program

To learn how to structure a complex, real-world program, let's develop a simple program to determine whether a number is prime. A *prime number* (or a *prime*) is a natural number that cannot be formed by multiplying two smaller natural numbers. We'll write two separate components: a static library containing the testing functionality and a command line application that provides a user interface for the library.

The *primetest* program accepts a whitespace-delimited list of integer values

as input and then outputs whether each value is a prime number. If any of the inputs are invalid, the program will output a helpful message explaining how to use the interface.

Before exploring how to structure the program, let's examine the user interface. First, we print the help text for the command line program, as shown in [Listing 10-4](#).

```
// print command line help text
static void print_help() {
    puts("primetest num1 [num2 num3 ... numN]\n");
    puts("Tests positive integers for primality.");
    printf("Tests numbers in the range [2-%llu].\n", ULLONG_MAX);
}
```

Listing 10-4: Printing help text

The `print_help` function prints usage information about how to use the command to the standard output.

Next, because the command line arguments are passed to the program as textual input, we define a utility function to convert them to integer values, as shown in [Listing 10-5](#).

```
// converts a string argument arg to an unsigned long long value
// referenced by val
// returns true if the argument conversion succeeds and false if
// it fails
static bool convert_arg(const char *arg, unsigned long long *val)
{
    char *end;

    // strtoull returns an in-band error indicator; clear errno before the call
    errno = 0;
    *val = strtoull(arg, &end, 10);

    // check for failures where the call returns a sentinel value and sets errno
    if ((*val == ULLONG_MAX) && errno) return false;
    if (*val == 0 && errno) return false;
    if (end == arg) return false;

    // If we got here, the argument conversion was successful.
    // However, we want to allow only values greater than one,
```

```
// so we reject values <= 1.
if (*val <= 1) return false;
return true;
}
```

Listing 10-5: Converting a single command line argument

The `convert_arg` function accepts a string argument as input and uses an output parameter to report the converted argument. An *output parameter* returns a function result to the caller via a pointer, allowing multiple values to be returned in addition to the function return value. The function returns `true` if the argument conversion succeeds and `false` if it fails. The `convert_arg` function uses the `strtoull` function to convert the string to an unsigned long long integer value and takes care to properly handle conversion errors. Additionally, because the definition of a prime number excludes 0, 1, and negative values, the `convert_arg` function treats those as invalid inputs.

We use the `convert_arg` utility function in the `convert_cmd_line_args` function, shown in [Listing 10-6](#), which loops over all the command line arguments provided and attempts to convert each argument from a string to an integer.

```
static unsigned long long *convert_cmd_line_args(int argc,
                                                  const char *argv
[],
                                                  size_t *num_args
) {
    *num_args = 0;

    if (argc <= 1) {
        // no command line arguments given (the first argument is the
        // name of the program being executed)
        print_help();
        return nullptr;
    }

    // We know the maximum number of arguments the user could have
    // passed,
    // so allocate an array large enough to hold all the elements.
    Subtract
    // one for the program name itself. If the allocation fails, tr
    eat it as
    // a failed conversion (it is OK to call free(nullptr)).
    unsigned long long *args =
```

```

        (unsigned long long *)malloc(sizeof(unsigned long long) * (
argc - 1));
    bool failed_conversion = (args == nullptr);
    for (int i = 1; i < argc && !failed_conversion; ++i) {
        // Attempt to convert the argument to an integer. If we
        // couldn't convert it, set failed_conversion to true.
        unsigned long long one_arg;
        failed_conversion |= !convert_arg(argv[i], &one_arg);
        args[i - 1] = one_arg;
    }

    if (failed_conversion) {
        // free the array, print the help, and bail out
        free(args);
        print_help();
        return nullptr;
    }

    *num_args = argc - 1;
    return args;
}

```

Listing 10-6: Processing all the command line arguments

If any argument fails to convert, it calls the `print_help` function to report the proper command line usage to the user and then returns a null pointer. This function is responsible for allocating a sufficiently large buffer to hold the array of integers. It also handles all error conditions, such as running out of memory or failing to convert an argument. If the function succeeds, it returns an array of integers to the caller and writes the converted number of arguments into the `num_args` parameter. The returned array is allocated storage and must be deallocated when no longer needed.

There are several ways to determine whether a number is prime. The naive approach is to test a value N by determining whether it is evenly divisible by $[2..N - 1]$. This approach has poor performance characteristics as the value of N gets larger. Instead, we'll use one of the many algorithms designed for testing primality. [Listing 10-7](#) shows a nondeterministic implementation of the Miller-Rabin primality test that's suitable for quickly testing whether a value is probably prime (Schoof 2008). Please see the Schoof paper for an explanation of the mathematics behind the Miller-Rabin primality test algorithm.

```

static unsigned long long power(unsigned long long x, unsigned lo

```

```

ng long y,
                                unsigned long long p) {
    unsigned long long result = 1;
    x %= p;

    while (y) {
        if (y & 1) result = (result * x) % p;
        y >>= 1;
        x = (x * x) % p;
    }
    return result;
}

static bool miller_rabin_test(unsigned long long d, unsigned long
long n) {
    unsigned long long a = 2 + rand() % (n - 4);
    unsigned long long x = power(a, d, n);

    if (x == 1 || x == n - 1) return true;

    while (d != n - 1) {
        x = (x * x) % n;
        d *= 2;

        if (x == 1) return false;
        if (x == n - 1) return true;
    }
    return false;
}

```

Listing 10-7: The Miller-Rabin primality test algorithm

The interface to the Miller-Rabin primality test is the `is_prime` function shown in [Listing 10-8](#). This function accepts two arguments: the number to test (`n`) and the number of times to perform the test (`k`). Larger values of `k` provide a more accurate result but worsen performance. We'll place the algorithm from [Listing 10-6](#) in a static library, along with the `is_prime` function, which will provide the library's public interface.

```

bool is_prime(unsigned long long n, unsigned int k) {
    if (n <= 1 || n == 4) return false;
    if (n <= 3) return true;

    unsigned long long d = n - 1;
    while (d % 2 == 0) d /= 2;

```

```
    for (; k != 0; --k) {
        if (!miller_rabin_test(d, n)) return false;
    }
    return true;
}
```

Listing 10-8: The interface to the Miller-Rabin primality test algorithm

Finally, we need to compose these utility functions into a program. [Listing 10-9](#) shows the implementation of the main function. It uses a fixed number of iterations of the Miller-Rabin test and reports whether the input values are probably prime or definitely not prime. It also handles deallocating the memory allocated by `convert_cmd_line_args`.

```
int main(int argc, char *argv[]) {
    size_t num_args;
    unsigned long long *vals = convert_cmd_line_args(argc, argv, &num_args);

    if (!vals) return EXIT_FAILURE;

    for (size_t i = 0; i < num_args; ++i) {
        printf("%llu is %s.\n", vals[i],
            is_prime(vals[i], 100) ? "probably prime" : "not prime");
    }

    free(vals);
    return EXIT_SUCCESS;
}
```

Listing 10-9: The main function

The main function calls the `convert_cmd_line_args` function to convert the command line arguments into an array of unsigned long long integers. For each argument in this array, the program loops, calling `is_prime` to determine whether each value is probably prime or not prime using the Miller-Rabin primality test implemented by the `is_prime` function.

Now that we've implemented the program logic, we'll produce the required build artifacts. Our goal is to produce a static library containing the Miller-Rabin implementation and a command line application driver.

Building the Code

Create a new file named *isprime.c* with the code from Listings 10-8 and 10-9 (in that order), adding the `#include` directives for `"isprime.h"` and `<stdlib.h>` at the top of the file. The quotes and angle brackets surrounding the headers are important for telling the preprocessor where to search for those files, as discussed in [Chapter 9](#). Next, create a header named *isprime.h* with the code from [Listing 10-10](#) to provide the public interface for the static library, with a header guard.

```
#ifndef PRIMETEST_IS_PRIME_H
#define PRIMETEST_IS_PRIME_H

bool is_prime(unsigned long long n, unsigned k);

#endif // PRIMETEST_IS_PRIME_H
```

Listing 10-10: The public interface for the static library

Create a new file named *driver.c* with the code from Listings 10-5, 10-6, 10-7, and 10-10 (in that order), adding the `#include` directives for the following: `"isprime.h"`, `<assert.h>`, `<errno.h>`, `<limits.h>`, `<stdio.h>`, and `<stdlib.h>` at the top of the file. All three files are in the same directory in our example, but in a real-world project, you would likely put the files in different directories, depending on the conventions of your build system. Create a local directory named *bin*, which is where the build artifacts from this example will be created.

We use Clang to create the static library and executable program, but both GCC and Clang support the command line arguments in the example, so either compiler will work. First, compile both C source files into object files placed in the *bin* directory:

```
% cc -c -std=c23 -Wall -Wextra -pedantic isprime.c -o bin/isprime.o
% cc -c -std=c23 -Wall -Wextra -pedantic driver.c -o bin/driver.o
```

For older compilers, it may be necessary to replace `-std=c23` with `-std=c2x`.

If you execute the command and get an error such as

```
unable to open output file 'bin/isprime.o': 'No such file or dire
```

ctory'

then create the local *bin* directory and try the command again. The `-c` flag instructs the compiler to compile the source into an object file without invoking the linker to produce executable output. We'll need the object files to create a library. The `-o` flag specifies the pathname of the output file.

After executing the commands, the *bin* directory should contain two object files: *isprime.o* and *driver.o*. These files contain the object code for each translation unit. You could link them together directly to create the executable program. However, in this case, we'll make a static library. To do this, execute the `ar` command to generate the static library named *libPrimalityUtilities.a* in the *bin* directory:

```
% ar rcs bin/libPrimalityUtilities.a bin/isprime.o
```

The `r` option instructs the `ar` command to replace any existing files in the archive with the new files, the `c` option creates the archive, and the `s` option writes an object-file index into the archive (which is equivalent to running the `ranlib` command). This creates a single archive file that's structured to allow retrieval of the original object files used to create the archive, like a compressed tarball or ZIP file. By convention, static libraries on Unix systems are prefixed with *lib* and have a *.a* file extension.

You can now link the driver object file to the *libPrimalityUtilities.a* static library to produce an executable named *primetest*. This can be accomplished either by invoking the compiler without the `-c` flag, which invokes the default system linker with the appropriate arguments, or by invoking the linker directly. Invoke the compiler to use the default system linker as follows:

```
% cc bin/driver.o -Lbin -lPrimalityUtilities -o bin/primetest
```

The `-L` flag instructs the linker to look in the local *bin* directory for libraries to link, and the `-l` flag instructs the linker to link the *libPrimalityUtilities.a* library to the output. Omit the *lib* prefix and the *.a* suffix from the command line argument because the linker adds them implicitly. For example, to link against the *libm* math library, specify `-lm` as the link target. As with compiling source files, the output of the linked files is specified by the `-o` flag.

You can now test the program to see whether values are probably prime or

definitely not prime. Be sure to try out cases like negative numbers, known prime and nonprime numbers, and incorrect input, as shown in [Listing 10-11](#).

```
% ./bin/primetest 899180
899180 is not prime
% ./bin/primetest 8675309
8675309 is probably prime
% ./bin/primetest 0
primetest num1 [num2 num3 ... numN

Tests positive integers for primality.
Tests numbers in the range [2-18446744073709551615].
```

Listing 10-11: Running the primetest program with sample input

The number 8,675,309 is prime.

Summary

In this chapter, you learned about the benefits of loose coupling, high cohesion, data abstractions, and code reuse. Additionally, you learned about related language constructs such as opaque data types and linkage. You were introduced to some best practices on how to structure the code in your projects and saw an example of building a simple program with different types of executable components. These skills are important as you transition from writing practice programs to developing and deploying real-world systems.

In the next chapter, we'll learn how to use various tools and techniques to create high-quality systems, including assertions, debugging, testing, and static and dynamic analysis. These skills are all necessary to develop safe, secure, and performant modern systems.

11

DEBUGGING, TESTING, AND ANALYSIS



This final chapter describes tools and techniques for producing correct, effective, safe, secure, and robust programs, including static (compile-time) and runtime assertions, debugging, testing, static analysis, and dynamic analysis. The chapter also discusses which compiler flags are recommended for use in different phases of the software development process.

This chapter marks a transition point from learning to program in C to professional C programming. Programming in C is relatively easy, but mastering C programming is a lifetime endeavor. Modern C programming requires a disciplined approach to develop and deploy safe, secure, and performant systems.

Assertions

An *assertion* is a function with a Boolean value, known as a *predicate*, which expresses a logical proposition about a program. You use an assertion to verify that a specific assumption you made during the implementation of your program remains valid. C supports static assertions that can be checked at compile time using `static_assert` and runtime assertions that are checked during program execution using `assert`. The `assert` macro is defined in the `<assert.h>` header. In C23, `static_assert` is a keyword. In C11, `static_assert` was provided as a

macro in `<assert.h>`. Prior to that, C did not have static assertions.

Static Assertions

Static assertions can be expressed using the `static_assert` keyword as follows:

```
static_assert(integer-constant-expression, string-literal);
```

Since C23, `static_assert` also accepts a single-argument form:

```
static_assert(integer-constant-expression);
```

If the value of the integer constant expression is not equal to 0, the `static_assert` declaration has no effect. If the integer constant expression is equal to 0, the compiler will produce a diagnostic message with the text of the string literal, if present.

You can use static assertions to validate assumptions at compile time, such as specific implementation-defined behaviors. Any change in implementation-defined behavior will then be diagnosed at compilation.

Let's look at three examples of using static assertions. First, in [Listing 11-1](#), we use `static_assert` to verify that `struct packed` has no padding bytes.

```
struct packed {  
    int i;  
    char *p;  
};  
  
static_assert(  
    sizeof(struct packed) == sizeof(int) + sizeof(char *),  
    "struct packed must not have any padding"  
);
```

Listing 11-1: Asserting the absence of padding bytes

The predicate for the static assertion in this example tests that the size of the packed structure is the same as the combined size of its `int` and `char *` members. For example, on x86-32, both `int` and `char *` are 4 bytes, and the structure is not padded, but on x86-64, `int` is 4 bytes, `char *` is 8 bytes, and the compiler adds 4 padding bytes between the two fields.

A good use of static assertions is to document all your assumptions

concerning implementation-defined behavior. This will prevent the code from compiling when porting to another implementation where those assumptions are invalid.

Because a static assertion is a declaration, it can appear at file scope, immediately following the definition of the struct whose property it asserts.

For the second example, the `clear_stdin` function, shown in [Listing 11-2](#), calls the `getchar` function to read characters from `stdin` until the end of the file is reached.

```
#include <stdio.h>
#include <limits.h>

void clear_stdin() {
    int c;

    do {
        c = getchar();
        static_assert(
            sizeof(unsigned char) < sizeof(int),
            "FI034-C violation"
        );
    } while (c != EOF);
}
```

Listing 11-2: Using `static_assert` to verify integer sizes

Each character is obtained as an `unsigned char` converted to an `int`. It's common practice to compare the character returned by the `getchar` function with `EOF`, often in a `do...while` loop, to determine when all the available characters have been read. For this function loop to work correctly, the terminating condition must be able to differentiate between a character and `EOF`. However, the C standard allows for `unsigned char` and `int` to have the same range, meaning that on some implementations, this test for `EOF` could return false positives, in which case the `do...while` loop may terminate early. Because this is an unusual condition, you can use `static_assert` to validate that the `do...while` loop can properly distinguish between valid characters and `EOF`.

In this example, the static assertion verifies that `sizeof(unsigned char) < sizeof(int)`. The static assertion is placed near the code that depends on this assumption being true so that you can easily locate the code that will need to be repaired if the assumption is violated. Because static assertions are evaluated at

compile time, placing them within executable code has no impact on the runtime efficiency of the program. See the CERT C rule FIO34-C, “Distinguish between characters read from a file and EOF or WEOF,” for more information on this topic.

Finally, in [Listing 11-3](#), we use `static_assert` to perform compile-time bounds checking.

```
static const char prefix[] = "Error No: ";
constexpr int size = 14;
char str[size];

// ensure that str has sufficient space to store at
// least one additional character for an error code
static_assert(
    sizeof(str) > sizeof(prefix),
    "str must be larger than prefix"
);
strcpy(str, prefix);
```

Listing 11-3: Using `static_assert` to perform bounds checking

This code snippet uses `strcpy` to copy a constant string prefix to a statically allocated array `str`. The static assertion ensures that `str` has sufficient space to store at least one additional character for an error code following the call to `strcpy`.

This assumption may become invalid if a developer, for example, reduced size or changed the prefix string to "Error Number: " during maintenance. Having added the static assertion, the maintainer would now be warned about the problem.

Remember that the string literal is a message for the developer or maintainer and not an end user of the system. It's intended to provide information useful for debugging.

Runtime Assertions

The `assert` macro injects runtime diagnostic tests into programs. It's defined in the `<assert.h>` header file and takes a scalar expression as a single argument:

```
#define assert(scalar-expression) /* implementation-defined */
```

The `assert` macro is implementation defined. If the scalar expression is equal to 0, the macro expansion typically writes information about the failing

call (including the argument text, the name of the source file `__FILE__`, the source line number `__LINE__`, and the name of the enclosing function `__func__`) to the standard error stream `stderr`. After writing this information to `stderr`, the `assert` macro calls the `abort` function.

The `dup_string` function shown in [Listing 11-4](#) uses runtime assertions to check that the `size` argument is less than or equal to `LIMIT` and that `str` is not a null pointer.

```
void *dup_string(size_t size, char *str) {
    assert(size <= LIMIT);
    assert(str != nullptr);
    // --snip--
}
```

Listing 11-4: Using `assert` to verify program conditions

The messages from these assertions might take the following form:

```
Assertion failed: size <= LIMIT, function dup_string, file foo.c,
line 122.
Assertion failed: str != nullptr, function dup_string, file foo.c
, line 123.
```

The implicit assumption is that the caller validates arguments before calling `dup_string` so that the function is never called with invalid arguments. The runtime assertions are then used to validate this assumption during the development and test phases.

The assertion's predicate expression is often reported in a failed assertion message, which allows you to use `&&` on a string literal with the assertion predicate to generate additional debugging information when an assertion fails. Doing so is always safe because string literals in C can never have a null pointer value. For example, we can rewrite the assertions in [Listing 11-4](#) to have the same functionality but provide additional context when the assertion fails, as shown in [Listing 11-5](#).

```
void *dup_string(size_t size, char *str) {
    assert(size <= LIMIT && "size is larger than the expected limit
");
    assert(str != nullptr && "the caller must ensure str is not nul
l");
}
```

```
// --snip--  
}
```

Listing 11-5: Using `assert` with additional contextual information

You should disable assertions before code is deployed by defining the `NDEBUG` macro (typically as a flag passed to the compiler). If `NDEBUG` is defined as a macro name at the point in the source file where `<assert.h>` is included, the `assert` macro is defined as follows:

```
#define assert(ignore) ((void)0)
```

The reason the macro does not expand empty is because if it did, then code such as

```
assert(thing1) // missing semicolon  
assert(thing2);
```

would compile in release mode but not in debug mode. The reason it expands to `((void) 0)` rather than just `0` is to prevent warnings about statements with no effect. The `assert` macro is redefined according to the current state of `NDEBUG` each time that `<assert.h>` is included.

Use static assertions to check assumptions that can be checked at compile time, and use runtime assertions to detect invalid assumptions during testing. Because runtime assertions are typically disabled before deployment, avoid using them to check for conditions that can come up during normal operations, such as the following:

- Invalid input
- Errors opening, reading, or writing streams
- Out-of-memory conditions from dynamic allocation functions
- System call errors
- Invalid permissions

You should instead implement these checks as normal error-checking code that's always included in the executable. Assertions should be used only to validate preconditions, postconditions, and invariants designed into the code (programming errors).

Compiler Settings and Flags

Compilers typically don't enable optimization or security hardening by default. Instead, you can enable optimization, error detection, and security hardening using build flags (Weimer 2018). I recommend specific flags for GCC, Clang, and Visual C++ in the next section, after first describing how and why you might want to use them.

Select your build flags based on what you're trying to accomplish. Distinct phases of software development call for different compiler and linker configurations. Some flags, such as the warnings, will be common to all phases. Other flags, such as the debug or the optimization level, are specific to each phase.

Build The goal of the build phase is to take full advantage of compiler analysis to eliminate defects before debugging. Dealing with numerous diagnostics at this stage can seem bothersome but is much better than having to find these problems through debugging and testing, or not finding them until after the code has shipped. During the build phase, you should use compiler options that maximize diagnostics to help you eliminate as many defects as possible.

Debug During debugging, you're typically trying to determine why your code isn't working. To best accomplish this, use a set of compiler flags that includes debug information, allows assertions to be useful, and enables a quick turnaround time for the inevitable edit-compile-debug cycle.

Test You may want to retain debug information and leave assertions enabled during testing to assist in identifying the root cause of any problems that are discovered. Runtime instrumentation can be injected to help detect errors.

Profile-Guided Optimization This configuration defines compiler and linker flags that control how the compiler adds runtime instrumentation to the code it normally generates. One purpose of instrumentation is to collect profiling statistics, which can be used to find program hot spots for profile-guided optimizations.

Release The final phase is to build the code for deployment to its operational environment. Before deploying the system, make sure you thoroughly test your release configuration, because using a different set of compilation flags can trigger new defects, for example, from latent

undefined behaviors or timing effects caused by optimization.

I'll now cover some specific compiler and linker flags you might want to use for your compiler and software development phase.

GCC and Clang Flags

[Table 11-1](#) lists recommended compiler and linker options (aka *flags*) for both GCC and Clang. You can find documentation for compiler and linker options in the GCC manual (<https://gcc.gnu.org/onlinedocs/gcc/Invoking-GCC.html>) and the Clang Compiler User's Manual (<https://clang.llvm.org/docs/UsersManual.html#command-line-options>).

Table 11-1: Recommended Compiler and Linker Flags for GCC and Clang

Flag	Purpose
-D_FORTIFY_SOURCE=2	Detect buffer overflows
-fpie -Wl,-pie	Required for address space layout randomization
-fpic -shared	Disable text relocations for shared libraries
-g3	Generate abundant debugging information
-O2	Optimize your code for speed/space efficiency
-Wall	Turn on recommended compiler warnings
-Wextra	Turn on even more recommended compiler warnings
-Werror	Turn warnings into errors
-std=c23	Specify the language standard
-pedantic	Issue warnings demanded by strict conformance to the standard
-Wconversion	Warn for implicit conversions that may alter a value
-Wl,-z,noexecstack	Mark the stack segments as nonexecutable
-fstack-protector-strong	Add stack protection to functions

-O

The uppercase letter `-O` flag controls *compiler optimization*. Most optimizations are completely disabled at optimization level 0 (`-O0`). This is the default when no optimization level has been set by a command line option. Similarly, the `-Og` flag suppresses optimization passes that may hinder the debugging experience.

Many diagnostics are issued by GCC only at higher optimization levels, such as `-O2` or `-O3`. To ensure that issues are identified during development, use the same (higher) optimization level you plan to adopt in production during the compilation and analysis phase. Clang, on the other hand, does not require the optimizer to issue diagnostics. As a result, Clang can be run with optimizations disabled during the compilation/analysis and debug phases.

The `-Os` compiler option optimizes for size, enabling all `-O2` optimizations except those that often increase code size. The `-Oz` compiler option optimizes aggressively for size rather than speed, which may increase the number of instructions executed if those instructions require fewer bytes to encode. The `-Oz` option behaves similarly to `-Os`, and it may be used in Clang but only in conjunction with `-mno-outline`. The `-Oz` compiler option may be used in GCC versions 12.1 or greater.

-glevel

The `-glevel` flag produces debugging information in the operating system's native format. You can specify how much information to produce by setting the debug *level*. The default level is `-g2`. Level 3 (`-g3`) includes extra information, such as all the macro definitions present in the program. Level 3 also allows you to expand macros in debuggers that support the capability.

Different settings are appropriate for debugging. Optimization levels should be low or disabled so that the machine instructions correspond closely to the source code. Symbols should also be included to assist in debugging. The `-O0 -g3` compiler flags are a good default, although other options are acceptable.

Consider the following program:

```
#include <stdio.h>
#include <stdlib.h>

#define HELLO "hello world!"

int main()
{
    puts(HELLO);

    return EXIT_SUCCESS;
}
```

The `-Og` compiler option affects only the optimization level without enabling debug symbols:

```
$ gcc -Og hello.c -o hello
$ gdb hello
(...)
(No debugging symbols found in hello)
(gdb)
```

Compiling with `-Og -g` provides some symbols:

```
$ gcc -Og -g hello.c -o hello
$ gdb hello
(...)
Reading symbols from hello...
(gdb) break main
Breakpoint 1 at 0x1149: file hello.c, line 6.
(gdb) start
Temporary breakpoint 2 at 0x1149: file hello.c, line 6.
Starting program: /home/test/Documents/test/hello

Breakpoint 1, main () at hello.c:6
6      int main()
(gdb) print HELLO
No symbol "HELLO" in current context.
(gdb)
```

Compiling with `-Og -g3` adds more symbols:

```
$ gcc -Og -g3 hello.c -o hello
$ gdb hello
(...)
Reading symbols from hello...
(gdb) break main
Breakpoint 1 at 0x1149: file hello.c, line 6.
(gdb) start
Temporary breakpoint 2 at 0x1149: file hello.c, line 6.
Starting program: /home/test/Documents/test/hello

Breakpoint 1, main () at hello.c:6
6      int main()
(gdb) print HELLO
$1 = "hello world!"
(gdb)
```

The `-g3` option causes debugging information to be generated in the operating system's native format, but the `-ggdb3` option tells GCC to use the most expressive format available for use by the GNU Project debugger (GDB). As a result, if you are only debugging with GDB, `-Og -ggdb3` is also a good choice of options.

The `-O0 -g3` options are recommended for the standard edit-compile-debug cycle.

-Wall and -Wextra

Compilers typically enable by default only the most conservatively correct diagnostic messages. Additional diagnostics can be enabled to check source code more aggressively for issues. Use the following flags to enable additional diagnostic messages when compiling code with GCC and Clang: `-Wall` and `-Wextra`.

The `-Wall` and `-Wextra` compiler flags enable predefined sets of compile-time warnings. The warnings in the `-Wall` set are generally easy to avoid or eliminate by modifying the diagnosed code. The warnings in the `-Wextra` set either are situational or indicate problematic constructs that are harder to avoid and, in some cases, may be necessary.

Despite their names, the `-Wall` and `-Wextra` options do not enable all possible warning diagnostics; they enable only a predefined subset. For a complete list of specific warnings enabled by the `-Wall` and `-Wextra` compiler flags, on GCC run:

```
$ gcc -Wall -Wextra -Q --help=warning
```

Alternatively, you can consult the documentation for GCC warning options and Clang diagnostic flags.

-Wconversion

Data type conversions can alter data values in unexpected ways. Memory safety violations may result from adding or subtracting these values from a pointer. The `-Wconversion` compiler option warns about:

- Implicit conversions that may alter a value, including conversions between floating-point and integer values
- Conversions between signed and unsigned integers, for example:

```
unsigned ui = -1;
```

- Conversions to smaller types

Warnings about conversions between signed and unsigned integers can be disabled by using `-Wno-sign-conversion`, but they're often useful in finding certain classes of defects and security vulnerabilities. The `-Wconversion` command line option should remain enabled.

-Werror

The `-Werror` flag turns all warnings into errors, requiring you to address them before you can begin debugging. This flag simply encourages good programming discipline.

-std=

The `-std=` flag can be used to specify the language standard as `c89`, `c90`, `c99`, `c11`, `c17`, or `c23` (you may need to specify `-std=c2x` when using an older compiler). If no C language dialect options are given, the default for GCC 13 is `-std=gnu17`, which provides extensions to the C language that, on rare occasions, conflict with the C standard. For portability, specify the standard you're using. For access to new language features, specify a recent standard. A good choice if you are reading this second edition of *Effective C* is `-std=c23`.

-pedantic

The `-pedantic` flag issues warnings when code deviates from strict conformance to the standard. This flag is typically used in conjunction with the `-std=` flag to improve the code's portability.

-D_FORTIFY_SOURCE=2

The `_FORTIFY_SOURCE` macro provides lightweight support for detecting buffer overflows in functions that perform operations on memory and strings. This macro can't detect all types of buffer overflows, but compiling your source with `-D_FORTIFY_SOURCE=2` provides an extra level of validation for functions that copy memory and are a potential source of buffer overflows such as `memcpy`, `memset`, `strcpy`, `strcat`, and `sprintf`. Some of the checks can be performed at compile time and result in diagnostics; others occur at runtime and can result in a runtime error.

The `_FORTIFY_SOURCE` macro requires optimizations to be enabled. Consequently, it must be disabled for unoptimized debug builds.

To overwrite a predefined `_FORTIFY_SOURCE` value, turn it off with `-U_FORTIFY_SOURCE` and on again with `-D_FORTIFY_SOURCE=2`. This will eliminate the warning that macros are being redefined.

The `_FORTIFY_SOURCE=3` macro has improved compiler checks for buffer overflows since version 12 of GCC and version 2.34 of the GNU C Library (glibc). The `-D_FORTIFY_SOURCE={1, 2, 3}` macro for glibc relies heavily on GCC-specific implementation details. Clang implements its own style of

fortified function calls.

Specify either `-D_FORTIFY_SOURCE=2` (recommended) or `-D_FORTIFY_SOURCE=1` for analysis, testing, and production builds using Clang and GCC versions prior to 12.0 and `_FORTIFY_SOURCE=3` for GCC version 12.0 and later.

-fpie -Wl, -pie, and -fpic -shared

Address space layout randomization (ASLR) is a security mechanism that randomizes the process's memory space to prevent attackers from predicting the location of the code they're trying to execute. You can learn more about ASLR and other security mitigations in *Secure Coding in C and C++* (Seacord 2013).

You must specify the `-fpie`, `-Wl, -pie`, and `-fpic -shared` flags to create position-independent executable programs and make it possible to enable ASLR for your main program (executable). However, while code emitted for your main program with these options is position independent, it does use some relocations that cannot be used in shared libraries (dynamic shared objects). For those, use `-fpic` and link with `-shared` to avoid text relocations on architectures that support position-dependent shared libraries. Dynamic shared objects are always position independent and therefore support ASLR.

-Wl,-z,noexecstack

Several operating systems, including OpenBSD, Windows, Linux, and macOS, enforce reduced privileges in the kernel to prevent any part of the process address space from being both writable and executable. This policy is called W^X .

The `-Wl, -z, noexecstack` linker option tells the linker to mark the stack segments as nonexecutable, which enables the operating system (OS) to configure memory access rights when the program executable is loaded into memory.

-fstack-protector-strong

The `-fstack-protector-strong` option protects applications from the most common forms of stack buffer overflow exploits by adding a stack canary. The `-fstack-protector` option is often viewed as insufficient and the `-fstack-protector-all` option as excessive. The `-fstack-protector-strong` option was introduced as a compromise between these two extremes.

Visual C++ Options

Visual C++ provides a wide assortment of compiler options, many of which are similar to the options available for GCC and Clang. One obvious difference is that Visual C++ generally uses the forward slash (/) character instead of a hyphen (-) to indicate a flag. [Table 11-2](#) lists recommended compiler and linker flags for Visual C++. (For more information on Visual C++ options, see <https://docs.microsoft.com/en-us/cpp/build/reference/compiler-options-listed-by-category>.)

Table 11-2: Recommended Compiler Flags for Visual C++

Flag	Purpose
/guard:cf	Add control flow guard security checks
/analyze	Enable static analysis
/sdl	Enable security features
/permissive-	Specify standards conformance mode to the compiler
/O2	Set optimization to level 2
/W4	Set compiler warnings to level 4
/WX	Turn warnings into errors
/std:clatest	Select the latest/greatest language version

Several of these options are similar to options provided by the GCC and Clang compilers. The /O2 optimization level is appropriate for deployed code, while /Od disables optimization to speed compilation and simplify debugging. The /W4 warning level is appropriate for new code, as it's roughly equivalent to -Wall in GCC and Clang. The /Wall option in Visual C++ isn't recommended because it produces a high number of false positives. The /WX option turns warnings into errors and is equivalent to the -Werror flag in GCC and Clang. I cover the remaining flags in further detail in the following sections.

/guard:cf

When you specify the *control flow guard (CFG)* option, the compiler and linker insert extra runtime security checks to detect attempts to compromise your code. The /guard:cf option must be passed to both the compiler and the linker.

/analyze

The /analyze flag enables static analysis, which provides information about possible defects in your code. I discuss static analysis in more detail in “Static Analysis” on [page 251](#).

/sdl

The `/sdl` flag enables additional security features, including treating extra security-relevant warnings as errors and additional secure code-generation features. It also enables other security features from the Microsoft *Security Development Lifecycle (SDL)*. The `/sdl` flag should be used in all production builds where security is a concern.

/permissive-

You can use `/permissive-` to help identify and fix conformance issues in your code, thereby improving your code's correctness and portability. This option disables permissive behaviors and sets the `/Zc` compiler options for strict conformance. In the integrated development environment (IDE), this option also underlines nonconforming code.

/std:clatest

The `/std:clatest` option enables all currently implemented compiler and standard library features proposed for C23. There is no `/std:c23` at the time of writing, but once one becomes available, you can use it to build C23 code.

Debugging

I've been programming professionally for 42 years. Once or maybe twice during that time, I've written a program that compiled and ran correctly on the first try. For all the other times, there is debugging.

Let's debug a faulty program. The program shown in [Listing 11-6](#) is an early version of the `vstrcat` function. We reviewed a finished version of this program in [Chapter 7](#), but this version is not yet ready to deploy.

```
#include <stdarg.h>
#include <string.h>
#include <stdio.h>
#include <stddef.h>

#define name_size 20U

char *vstrcat(char *buff, size_t buff_length, ...) {
    char *ret = buff;
    va_list list;
    va_start(list, buff_length);
```

```

    const char *part = nullptr;
    size_t offset = 0;
    while ((part = va_arg(list, const char *))) {
        buff = (char *)memcpy(buff, part, '\0', buff_length - offset)
- 1;
        if (buff == nullptr) {
            ret[0] = '\0';
            break;
        }
        offset = buff - ret;
    }

    va_end(list);
    return ret;
}

int main() {
    char name[name_size] = "";
    char first[] = "Robert";
    char middle[] = "C.";
    char last[] = "Seacord";

    puts(
        vstrcat(
            name, sizeof(name), first, " ",
            middle, " ", last, nullptr
        )
    );
}

```

Listing 11-6: Printing an error

When we run this program as shown, it outputs my name as expected:

Robert C. Seacord

However, we also want to ensure that this program, which uses a fixed-size array for name, properly handles the case where the full name is larger than the name array. To test this, we can change the size of the array to a too-small value:

```
#define name_size 10U
```

Now, when we run the program, we learn we have a problem but not much more than that:

```
$ ./bug
Segmentation fault
```

Instead of adding print statements, we'll take the plunge and debug this program using Visual Studio Code on Linux. Just running this program in the debugger, as shown in [Figure 11-1](#), provides us with some information that we didn't previously have.




Figure 11-1: Debugging a program in Visual Studio Code

We can see from the CALL STACK pane that we are crashing in the `__memmove_avx_unaligned_erms` function in `libc`.

```
libc.so.6!__memmove_avx_unaligned_erms()
(\x86_64\multiarch\memmove-vec-unaligned-erms.S:314)
vstrcat(char * buff, size_t buff_length) (\home\rsc\bug.c:17)
main() (\home\rsc\bug.c:32)
```

We can also see that the segmentation fault is occurring on the line with the call to `memcpy`. There isn't much else going on in this line, so it's reasonable to surmise that this function is a `memcpy` helper function. It's seldom the case that the bug is in the implementation of the library function, so we'll assume for now that we're passing an invalid set of arguments.

Before looking at the arguments, let's review the description of the `memcpy`

function from the C23 standard:

```
#include <string.h>
```

```
void *memccpy(void * restrict s1, const void * restrict s2, int c, size_t n);
```

The `memccpy` function copies characters from the object pointed to by `s2` into the object pointed to by `s1`, stopping after the first occurrence of character `c` (converted to an unsigned `char`) is copied or after `n` characters are copied, whichever comes first. If copying takes place between objects that overlap, the behavior is undefined.

From the Variables pane in the debugger, we can see that the part we are adding looks correct:

part: 0x7fffffffddcd6 "Seacord"

The `ret` alias to the start of `ret` also has an expected value:

ret: 0x7fffffffddcde "Robert C. "

The value stored in `buff`, however, seems odd, as it has the same value as an EOF (−1):

buff: 0xffffffffffffffff <error: Cannot access memory at address 0xffffffffffffffff>

The `buff` parameter is a character pointer that is assigned the return value from `memccpy`. So once again, let's check the standard to see what this function returns:

The `memccpy` function returns a pointer to the character after the copy of `c` in `s1`, or a null pointer if `c` was not found in the first `n` characters of `s2`.

According to the C standard, this function can return only a null pointer or a pointer to a character in `s1` (`buff`, in this program). The storage for `buff` begins at 0x7fffffffddcde and extends for only 10 bytes, so neither of these explain the 0xffffffffffffffff value, so the mystery deepens.

It's time to examine the behavior of the `vstrcat` function more closely. We'll set a breakpoint on line 12 near the beginning of the function and start debugging. The buttons along the left of the title bar allow you to continue, step over, step into, step out, restart, and stop debugging. Starting from line 12, we can single-step through the program by clicking the Step Over button. The `vstrcat` function loops several times, so you'll have to step through a few iterations of the loop, watching the values in the VARIABLES pane. If you do

this carefully, you'll eventually see that `buff` is set to `0xffffffffffffffff` on line 18 following the call to the `memcpy` function, as shown in [Figure 11-2](#). This isn't detected by the null pointer test, and the segmentation fault occurs on the next iteration.

It was here that I had my eureka moment. The `memcpy` function returns a null pointer to indicate that `'\0'` was not found in the first `buff_length - offset` characters of `part`. However, we are subtracting 1 from the value returned by `memcpy` so that `buff` points to the first occurrence of `'\0'` rather than just after it. This works when the character is found, but when it isn't found, we subtract 1 from a null pointer, which is technically undefined behavior in C. On this implementation, the null pointer is represented by a 0 value. Subtracting 1 from 0 wraps around and produces the `0xffffffffffffffff` value for `buff` before we can test it. Consequently, the error condition is not detected, and the subsequent call to `memcpy` results in the segmentation fault.




Figure 11-2: An interesting program state

Now that we have discovered the root cause, the bug can be repaired by moving the minus-one subtraction after the null pointer check, which results in the final version of the program shown in [Chapter 7](#).

Unit Testing

Testing increases your confidence that your code is defect free. *Unit tests* are small programs that exercise your code. *Unit testing* is a process that validates that each unit of the software performs as designed. A *unit* is the smallest testable part of any software; in C, this is typically an individual function or data abstraction.

You can write simple tests that resemble normal application code (see [Listing 11-7](#), for example), but it can be beneficial to use a *unit-testing framework*. Several unit-testing frameworks are available, including Google Test, CUnit, CppUnit, Unity, and others. We'll examine the most popular of these, based on a recent survey of the C development ecosystem by JetBrains (<https://www.jetbrains.com/lp/devecosystem-2023/c/>): Google Test.

Google Test works for Linux, Windows, and macOS. Tests are written in C++, so you get to learn another (related) programming language for testing purposes. CUnit and Unity are good alternatives if you want to restrict your testing to pure C.

In Google Test, you write assertions to verify the tested code's behavior. Google Test assertions, which are function-like macros, are the real language of the tests. If a test crashes or has a failed assertion, it fails; otherwise, it succeeds. An assertion's result can be success, nonfatal failure, or fatal failure. If a fatal failure occurs, the current function is aborted; otherwise, the program continues normally.

We'll use Google Test on Ubuntu Linux. To install it, follow the directions from the Google Test GitHub page at <https://github.com/google/googletest/tree/main/googletest>.

Once Google Test is installed, we'll set up a unit test for the `get_error` function shown in [Listing 11-7](#). This function returns an error message string corresponding to the error number passed in as an argument. You'll need to include the headers that declare the `errno_t` type and the `strerrorlen_s` and `strerror_s` functions. Save it in a file named `error.c` so that the build instructions described later in this section will work properly.

`error.c`

```
char *get_error(errno_t errnum) {
    rsize_t size = strerrorlen_s(errnum) + 1;
    char* msg = malloc(size);
    if (msg != nullptr) {
        errno_t status = strerror_s(msg, size, errnum);
        if (status != 0) {
```

```
        strncpy_s(msg, size, "unknown error", size - 1);
    }
}
return msg;
}
```

Listing 11-7: The `get_error` function

This function calls both the `strerrorlen_s` and `strerror_s` functions defined in the normative but optional Annex K, “Bounds-checking interfaces” (described in [Chapter 7](#)).

Unfortunately, neither GCC nor Clang implements Annex K, so instead we’ll use the Safeclib implementation developed by Reini Urban and available from GitHub (<https://github.com/rurban/safeclib>).

You can install `libsafec-dev` on Ubuntu with the following command:

```
% sudo apt install libsafec-dev
```

[Listing 11-8](#) contains a unit test suite for the `get_error` function named `GetErrorTest`. A *test suite* is a set of test cases to be executed in a specific test cycle. The `GetErrorTest` suite consists of two test cases: `KnownError` and `UnknownError`. A *test case* is a set of preconditions, inputs, actions (where applicable), expected results, and postconditions, developed based on test conditions (<https://glossary.istqb.org>). Save this code in a file named `tests.cc`.

`tests.cc`

```
#include <gtest/gtest.h>
#include <errno.h>
#define errno_t int

// implemented in a C source file
❶ extern "C" char* get_error(errno_t errnum);

namespace {
❷ TEST(GetErrorTest, KnownError) {
    EXPECT_STREQ(get_error(ENOMEM), "Cannot allocate memory");
    EXPECT_STREQ(get_error(ENOTSOCK), "Socket operation on non-so
cket");
    EXPECT_STREQ(get_error(EPIPE), "Broken pipe");
}
```

```

    TEST(GetErrorTest, UnknownError) {
        EXPECT_STREQ(get_error(-1), "Unknown error -1");
    }
} // namespace

int main(int argc, char** argv) {
    ::testing::InitGoogleTest(&argc, argv);
    return RUN_ALL_TESTS();
}

```

Listing 11-8: Unit tests for the `get_error` function

Most of the C++ code is boilerplate and can be copied without modification, including, for example, the `main` function, which invokes the function-like macro `RUN_ALL_TESTS` to execute your tests. The two parts that aren't boilerplate are the extern "C" declaration ❶ and the tests ❷. The extern "C" declaration changes the linkage requirements so that the C++ compiler linker doesn't mangle the function name, as it is wont to do. You need to add a similar declaration for each function being tested, or you can simply include the C header file within an extern "C" block as follows:

```

extern "C" {
    #include "api_to_test.h"
}

```

An extern "C" declaration is necessary only when compiling with C but linking with C++.

Both test cases are specified using the `TEST` macro, which takes two arguments. The first argument is the name of the test suite, and the second argument is the name of the test case.

Insert Google Test assertions, along with any additional C++ statements you wish to include, in the function body. In [Listing 11-8](#), we used the `EXPECT_STREQ` assertion, which verifies that two strings have the same content. The `strerror_s` function returns a locale-specific message string, which can vary between implementations.

We used the assertion on several error numbers to verify that the function is returning the correct string for each error number. The `EXPECT_STREQ` assertion is a nonfatal assertion because testing can continue even when this specific assertion fails. This is typically preferable to fatal assertions, as it lets you detect and fix multiple bugs in a single run-edit-compile cycle. If it's not possible to

continue testing after an initial failure (because a subsequent operation relies on a previous result, for example), you can use the fatal `ASSERT_STREQ` assertion.

[Listing 11-9](#) shows a simple *CMakeLists.txt* file that can be used to build the tests. This file assumes that the C functions we're testing can be found in the *error.c* file and that the implementations of the Annex K functions are provided by the *safec* library.

```
cmake_minimum_required(VERSION 3.21)
cmake_policy(SET CMP0135 NEW)
project(chapter-11)

# GoogleTest requires at least C++14
set(CMAKE_CXX_STANDARD 14)
set(CMAKE_CXX_STANDARD_REQUIRED ON)
set(CMAKE_C_STANDARD 23)

include(FetchContent)
FetchContent_Declare(
  googletest
  URL https://github.com/google/googletest/archive/03597a01ee50ed
33e9dfd640b249b4be3799d395.zip
)

FetchContent_MakeAvailable(googletest)

include(ExternalProject)
ExternalProject_Add(
  libsafec
  BUILD_IN_SOURCE 1
  URL https://github.com/rurban/safeclib/releases/download/v3.7.1
/safeclib-3.7.1.tar.gz
  CONFIGURE_COMMAND autoreconf --install
  COMMAND ./configure --prefix=${CMAKE_BINARY_DIR}/libsafec
)
ExternalProject_Get_Property(libsafec install_dir)
include_directories(${install_dir}/src/libsafec/include)
link_directories(${install_dir}/src/libsafec/src/.libs/)

enable_testing()

add_library(error error.c)
add_dependencies(error libsafec)
add_executable(tests tests.cc)

target_link_libraries(
```

```
    tests
    error
    safec
    GTest::gtest_main
)

include(GoogleTest)
gtest_discover_tests(tests)
```

Listing 11-9: The CMakeLists.txt file

If you prefer to install `libsafec-dev` using the `apt install` command, remove the lines specific to installing `libsafec`.

Build and run the tests using the following sequence of commands:

```
$ cmake -S . -B build
$ cmake --build build
$ ./build/tests
```

The test case tests for several error numbers from `<errno.h>`. How many of these should be tested depends on what you're trying to accomplish. Ideally, the tests should be comprehensive, which would mean adding an assertion for every error number in `<errno.h>`. This can become tiresome, however; once you have established that your code is working, you're mostly just testing that the underlying C standard library functions you're using are implemented correctly. Instead, we could test the error numbers we're likely going to retrieve, but doing so can again become tiresome because we'd have to identify all the functions called in the program and which error codes they may return. We opted to implement a few spot checks for several randomly selected error numbers from different locations in the list.

[Listing 11-10](#) shows the result of running this test.

```
$ ./build/tests
[=====] Running 2 tests from 1 test suite.
[-----] Global test environment set-up.
[-----] 2 tests from GetErrorTest
[RUN      ] GetErrorTest.KnownError
[         OK] GetErrorTest.KnownError (0 ms)
[RUN      ] GetErrorTest.UnknownError
/home/rcs/tests.cc:19: Failure
Expected equality of these values:
    get_error(-1)
```



```

    Which is: "Unknown error -1"
    "unknown error"
[   FAILED   ] GetErrorTest.UnknownError (0 ms)
[-----] 2 tests from GetErrorTest (0 ms total)

[-----] Global test environment tear-down
[=====] 2 tests from 1 test suite ran. (0 ms total)
[   PASSED   ] 1 test.
[   FAILED   ] 1 test, listed below:
[   FAILED   ] GetErrorTest.UnknownError

1 FAILED TEST

```

Listing 11-10: An unsuccessful test run

From the test output, you can see that two tests were executed from one test suite. The KnownError test case passed, and the UnknownError test case failed. The UnknownError test failed because the following assertion returned false:

```
EXPECT_STREQ(get_error(-1), "unknown error");
```

The test assumed that the error path in the `get_error` function would execute and return the string "unknown error"). Instead, the `strerror_s` function returned the string "Unknown error -1". Examining the source code for the `strerror_s` function (at https://github.com/rurban/safeclib/blob/master/src/str/strerror_s.c), we can see that the function does return error codes. Consequently, it's clear that the function doesn't treat an unknown error number as an error. Checking the C standard, we see that "strerror_s shall map any value of type `int` to a message," so the `strerror_s` function is implemented correctly, but our assumptions about how it behaved were incorrect.

There is a defect in the implementation of the `get_error` function in that it indicates "unknown error" when the `strerror_s` function fails, but according to the standard:

The `strerror_s` function returns zero if the length of the desired string was less than `maxsize` and there was no runtime-constraint violation. Otherwise, the `strerror_s` function returns a nonzero value.

Consequently, if the `strerror_s` function returns a nonzero value, a serious error has occurred that's bad enough to reconsider the design of this function. Instead of returning a string on an error condition, it should probably return a null pointer or otherwise handle the error in a manner consistent with the overall

error handling strategy for your system. [Listing 11-11](#) updates the function to return a null pointer value.

```
char *get_error(errno_t errnum) {
    rsize_t size = strerrorlen_s(errnum) + 1;
    char* msg = malloc(size);
    if (msg != nullptr) {
        errno_t status = strerror_s(msg, size, errnum);
        if (status != 0) return nullptr;
    }
    return msg;
}
```

Listing 11-11: The `get_error` function

We need to repair the test to check for the correct string returned by `get_error(-1)`:

```
EXPECT_STREQ(get_error(-1), "Unknown error -1");
```

After making this change, rebuilding, and rerunning the tests, we can see that both test cases succeeded as shown in [Listing 11-12](#).

```
$ ./build/tests
[=====] Running 2 tests from 1 test suite.
[-----] Global test environment set-up.
[-----] 2 tests from GetErrorTest
[RUN      ] GetErrorTest.KnownError
[         OK] GetErrorTest.KnownError (0 ms)
[RUN      ] GetErrorTest.UnknownError
[         OK] GetErrorTest.UnknownError (0 ms)
[-----] 2 tests from GetErrorTest (0 ms total)

[-----] Global test environment tear-down
[=====] 2 tests from 1 test suite ran. (0 ms total)
[ PASSED  ] 2 tests.
```

Listing 11-12: A successful test run

In addition to discovering a design error, we also discovered that our tests are incomplete, as we failed to test the error case. We should add further tests to ensure that error cases are handled correctly. Adding these tests is left as an

exercise for the reader.

Static Analysis

Static analysis includes any process for assessing code without executing it (ISO/IEC TS 17961:2013) to provide information about possible software defects.

Static analysis has practical limitations, as the correctness of software is computationally undecidable. For example, the halting theorem of computer science states that there are programs whose exact control flow cannot be determined statically. As a result, any property dependent on control flow—such as halting—may not be decidable for some programs. Consequently, static analysis may fail to report flaws or may report flaws where they don't exist.

A failure to report a real flaw in the code is known as a *false negative*. False negatives are serious analysis errors, as they may leave you with a false sense of security. Most tools err on the side of caution and, as a result, generate false positives. A *false positive* is a test result that incorrectly indicates that a flaw is present. Tools might report some high-risk flaws and miss other flaws as an unintended consequence of trying not to overwhelm the user with false positives. False positives can also occur when the code is too complex to completely analyze. The use of function pointers and libraries can make false positives more likely.

Ideally, tools are both complete and sound in their analysis. An analyzer is considered *sound* if it cannot give a false-negative result. An analyzer is considered *complete* if it cannot issue false positives. The possibilities for a given rule are outlined in [Figure 11-3](#).

		False Positive	
		Yes	No
False Negative	Yes	Incomplete and unsound	Complete and unsound
	No	Incomplete and sound	Complete and sound

Figure 11-3: Completeness and soundness

Compilers perform limited analysis, providing diagnostics about highly

localized issues in code that don't require much reasoning. For example, when comparing a signed value to an unsigned value, the compiler may issue a diagnostic about a type mismatch because it doesn't require additional information to identify the error. As mentioned earlier in this chapter, there are numerous compiler flags, such as `/w4` for Visual C++ and `-Wall` for GCC and Clang, that control compiler diagnostics.

Compilers generally provide high-quality diagnostics, and you shouldn't ignore them. Always try to understand the reason for the warning and rewrite the code to eliminate the error, rather than simply quieting warnings by adding type casts or making arbitrary changes until the warning goes away. See the CERT C rule MSC00-C, "Compile cleanly at high warning levels," for more information on this topic.

Once you've addressed compiler warnings in your code, you can use a separate static analyzer to identify additional flaws. Static analyzers will diagnose more complex defects by evaluating the expressions in your program, performing in-depth control and data flow analysis, and reasoning about the possible ranges of values and control flow paths taken.

Having a tool locate and identify specific errors in your program is much, much easier than hours of testing and debugging, and it's much less costly than deploying defective code. A wide variety of free and commercial static analysis tools are available. For example, Visual C++ has incorporated a static analyzer that you can invoke with the `/analyze` flag. Visual C++ analysis allows you to specify which rule sets (such as recommended, security, or internationalization) you would like to run or whether to run them all. For more information on Visual C++'s static code analysis, see Microsoft's website at <https://learn.microsoft.com/en-us/visualstudio/code-quality>. Similarly, Clang has incorporated a static analyzer that can be run as a stand-alone tool or within Xcode (<https://clang-analyzer.llvm.org>). Beginning with version 10, GCC has introduced static analysis that's enabled through the `-fanalyzer` option. Commercial tools also exist, such as CodeQL from GitHub, TrustInSoft Analyzer, SonarQube from SonarSource, Coverity from Synopsys, LDRA Testbed, Helix QAC from Perforce, and others.

Many static analysis tools have nonoverlapping capabilities, so it may make sense to use more than one.

Dynamic Analysis

Dynamic analysis is the process of evaluating a system or component during

execution. It's also referred to as *runtime analysis*, among other similar names.

A common approach to dynamic analysis is to *instrument* the code—for example, by enabling compile-time flags that inject extra instructions into the executable—and then run the instrumented executable. The debug memory allocation library `dmalloc` described in [Chapter 6](#) takes a similar approach. The `dmalloc` library provides replacement memory management routines with runtime-configurable debugging facilities. You can control the behavior of these routines by using a command line utility (also called `dmalloc`) to detect memory leaks and to discover and report defects such as writing outside the bounds of an object and using a pointer after it's been freed.

The advantage of dynamic analysis is that it has a low false-positive rate, so if one of these tools flags a problem, fix it!

A drawback of dynamic analysis is that it requires sufficient code coverage. If a defective code path is not exercised during the testing process, the defect won't be found. Another drawback is that the instrumentation may change other aspects of the program in undesirable ways, such as adding performance overhead or increasing the binary size. Unlike other dynamic analysis tools, the `FORTIFY_SOURCE` macro mentioned earlier in this chapter provides lightweight support for detecting buffer overflows so that it can be enabled in a production build with no noticeable impacts on performance.

AddressSanitizer

AddressSanitizer (ASan, <https://github.com/google/sanitizers/wiki/AddressSanitizer>) is an example of an effective dynamic analysis tool that is available (for free) for several compilers. Several related sanitizers exist, including ThreadSanitizer, MemorySanitizer, Hardware-Assisted AddressSanitizer, and UndefinedBehaviorSanitizer. Many other dynamic analysis tools are available, both commercial and free. For more information on sanitizers, see <https://github.com/google/sanitizers>. I'll demonstrate the value of these tools by discussing AddressSanitizer in some detail.

ASan is a dynamic memory error detector for C and C++ programs. It's incorporated into LLVM version 3.1 and GCC version 4.8, as well as later versions of these compilers. ASan is also available starting with Visual Studio 2019.

This dynamic analysis tool can find a variety of memory errors, including the following:

- Use after free (dangling pointer dereference)
- Heap, stack, and global buffer overflow
- Use after return
- Use after scope
- Initialization order bugs
- Memory leaks

To demonstrate ASan's usefulness, we'll start by replacing the `get_error` function from [Listing 11-7](#) with the `print_error` function shown in [Listing 11-13](#).

error.c

```
errno_t print_error(errno_t errnum) {
    rsize_t size = strerrorlen_s(errnum) + 1;
    char* msg = malloc(size);
    if (msg == nullptr) return ENOMEM;
    errno_t status = strerror_s(msg, size, errnum);
    if (status != 0) return EINVAL;
    fputs(msg, stderr);
    return EOK;
}
```

Listing 11-13: The `print_error` function

We'll also replace the unit test suite for the `get_error` function with the unit test suite for the `print_error` function shown in [Listing 11-14](#).

tests.cc

```
TEST(PrintTests, ZeroReturn) {
    EXPECT_EQ(print_error(ENOMEM), 0);
    EXPECT_EQ(print_error(ENOTSOCK), 0);
    EXPECT_EQ(print_error(EPIPE), 0);
}
```

Listing 11-14: The `PrintTests` test suite

This Google Test code defines the `PrintTests` test suite that contains a single test case, `ZeroReturn`. This test case uses the nonfatal `EXPECT_EQ` assertion to test for a return value of 0 from several calls to the `print_error` function to print some randomly selected error numbers. Next, we need to build

and run this code on Ubuntu Linux.

Running the Tests

Running the revised tests from [Listing 11-14](#) produces the positive results shown in [Listing 11-15](#).

```
$ ./build/tests
[=====] Running 1 test from 1 test suite.
[-----] Global test environment set-up.
[-----] 1 test from PrintTests
[RUN      ] PrintTests.ZeroReturn
[         OK] PrintTests.ZeroReturn (0 ms)
[-----] 1 test from PrintTests (0 ms total)

[-----] Global test environment tear-down
[=====] 1 test from 1 test suite ran. (0 ms total)
[ PASSED ] 1 test.
```

Listing 11-15: A test run of the `PrintTests` test suite

An inexperienced tester may look at these results and mistakenly think, “Hey, this code is working!” However, you should take additional steps to improve your confidence that your code is free from defects. Now that we have a working test harness in place, it’s time to instrument the code.

Instrumenting the Code

You can instrument your code by using ASan to compile and link your program with the `-fsanitize=address` flag. [Table 11-3](#) shows some compiler flags that are commonly used with ASan.

Table 11-3: Compiler and Linker Flags Commonly Used with AddressSanitizer

Flag	Purpose
<code>-fsanitize=address</code>	Enable AddressSanitizer (must be passed to both the compiler and the linker)
<code>-g3</code>	Get symbolic debugging information
<code>-fno-omit-frame-pointer</code>	Leave frame pointers to get more informative stack traces in error messages
<code>-fsanitize-blacklist=path</code>	Pass a blacklist file
<code>-fno-common</code>	Do not treat global variables as common variables (allows ASan to instrument them)

Select the compiler and linker flags you want to use from [Table 11-3](#) and

add them to your *CMakeLists.txt* file using the `add_compile_options` and `add_link_options` commands:

```
add_compile_options(-g3 -fno-omit-frame-pointer -fno-common -fsanitize=address)
add_link_options(-fsanitize=address)
```

Do not enable sanitization in the build phase because the inserted runtime instrumentation can cause false positives.

As previously mentioned, AddressSanitizer works with Clang, GCC, and Visual C++. (See <https://devblogs.microsoft.com/cppblog/addresssanitizer-asan-for-windows-with-msvc> for more information on ASan support.)

Depending on which version of which compiler you're using, you may also need to define the following environment variables:

```
ASAN_OPTIONS=symbolize=1
ASAN_SYMBOLIZER_PATH=/path/to/llvm_build/bin/llvm-symbolizer
```

Try rebuilding and rerunning your tests with these environmental variables set.

Running the Instrumented Tests

The unit test suite you wrote using Google Test should continue to pass but will also exercise your code, allowing AddressSanitizer to detect additional problems. You should now see the additional output in [Listing 11-16](#) from running `./build/tests`.

```
==22489==ERROR: LeakSanitizer: detected memory leaks

Direct leak of 31 byte(s) in 1 object(s) allocated from:
    #0 0x7f2bcf9f58ff in __interceptor_malloc
    ../../../../src/libsanitizer/asan/asan_malloc_linux.cpp:69
    #1 0x557d3105f6da in print_error /home/rcc/test/error.c:21
    #2 0x557d3105d314 in TestBody /home/rcc/test/tests.cc:28

// --snip--
```

Listing 11-16: An instrumented test run of `PrintTests`

[Listing 11-16](#) shows only the first finding of several that are produced. Most

of this stack trace is redacted because it is from the test infrastructure itself and is uninteresting because it doesn't help locate the defects.

AddressSanitizer's LeakSanitizer component has "detected memory leaks" and informs us that this is a direct leak of 31 bytes from one object. The stack trace identifies the filename and line number related to the diagnostic:

```
#1 0x557d3105f6da in print_error /home/rcs/test/error.c:21
```

This line of code contains the call to `malloc` in the `print_error` function:

```
errno_t print_error(errno_t errnum) {
    rsize_t size = strerrorlen_s(errnum) + 1;
    char* msg = malloc(size);
    // --snip--
}
```

This is an obvious error; the return value from `malloc` is assigned to an automatic variable defined within the scope of the `print_error` function and never freed. We lose the opportunity to free this allocated memory after the function returns, and the lifetime of the object holding the pointer to the allocated memory ends. To fix this problem, add a call to `free(msg)` after the allocated storage is no longer required but before the function returns. Rerun the tests and repair any additional defects until you're satisfied with the quality of your program.

EXERCISES

1. Use static analysis to evaluate the defective code from [Listing 11-6](#). Did the static analysis provide any additional findings?
2. Add further tests to exercise the error handling paths in both the `get_error` ([Listing 11-8](#)) and `print_error` ([Listing 11-14](#)) functions.
3. Evaluate the remaining results from the `./tests` test instrumented with AddressSanitizer ([Listing 11-15](#)). Eliminate the remaining true-positive errors detected.
4. Instrument the `./tests` program using other sanitizers available at <https://github.com/google/sanitizers> and address any issues you find.
5. Use these and similar testing, debugging, and analysis techniques on your real-world code.
6. Use profile-guided optimization to optimize the prime factorization program from [Chapter 10](#). Refer to your compiler's documentation for details.

Summary

In this chapter, you learned about static and runtime assertions and were introduced to some of the more important and recommended compiler flags for GCC, Clang, and Visual C++. You also learned how to debug, test, and analyze your code by using both static and dynamic analysis.

These are the important last lessons in this book, because you'll find you spend a considerable amount of time as a professional C programmer debugging and analyzing your code. I posted the following on social media a while back, and it summarizes my (and other C programmers) relationship with the C programming language:

- Language I dislike: C
- Language I begrudgingly respect: C
- Language I think is overrated: C
- Language I think is underrated: C
- Language I like: C

Future Directions

With C23 complete, the committee can turn its attention to the next revision of the C programming language, C2Y. This will likely be published in 2029. While that may seem like a long time, it's roughly half the time previous C standard editions required.

The C committee has already approved a new charter to document our principles (Seacord et al. 2024). While the committee is dedicated to maintaining the traditional spirit of C, there will be a renewed focus on security and safety. For C2Y, we'll likely improve automatic type inference, expand `constexpr` support, and potentially adopt lambdas and other features from C++. We're also working on a novel `defer` feature for error handling and resource management. The C floating-point group will continue its work to update to IEEE 754:2019. A technical specification for a provenance-aware memory object model for C (ISO/IEC CD TS 6010:2024) should be published soon and hopefully incorporated into C2Y.

APPENDIX

THE FIFTH EDITION OF THE C STANDARD (C23)

with Aaron Ballman



The latest (fifth) edition of the C standard (ISO/IEC 9899:2024) is nicknamed C23.

C23 maintains the *spirit of C*, while adding new features and functions to improve the safety, security, and capabilities of the language.

Attributes

The `[[attributes]]` syntax was added to C23 to specify additional information for various source constructs such as types, objects, identifiers, or blocks (Ballman 2019). Prior to C23, similar features were provided in an implementation-defined (nonportable) manner:

```
__declspec(deprecated)
__attribute__((warn_unused_result))
int func(const char *str) __attribute__((nonnull(1)));
```

Starting with C23, attributes can be specified as follows:

```
[[deprecated, nodiscard]]
int func(
    const char *str [[gnu::nonnull]]
);
```

Like C++, syntactic location determines apportionment. Attributes include

deprecated, fallthrough, maybe_unused, nodiscard, unsequenced, and reproducible. The attribute syntax supports both standard attributes and vendor-specific attributes. The `__has_c_attribute` conditional inclusion operator can be used for feature testing.

Keywords

The C language is often ridiculed for having ugly keywords. C typically defines new keywords using reserved identifiers that begin with an underscore character (`_`) followed by a capital letter.

C23 introduced more natural spellings for these keywords (Gustedt 2022). In Table A-1, C11 keywords using this convention are shown on the left, and the more natural spellings introduced in C23 are shown on the right.

Table A-1: Keyword Spellings

Value	Type
<code>_Bool</code>	<code>bool</code>
<code>_Static_assert</code>	<code>static_assert</code>
<code>_Thread_local</code>	<code>thread_local</code>
<code>_Alignof</code>	<code>alignof</code>
<code>_Alignas</code>	<code>alignas</code>

Another update is the introduction of the `nullptr` constant. The well-worn `NULL` macro has a pointer type or maybe an integer type. It will implicitly convert to any scalar type, so it's not particularly type safe. The `nullptr` constant has type `nullptr_t` and will implicitly convert only to a pointer type, `void`, or `bool`.

Integer Constant Expressions

Integer constant expressions are not a portable construct; vendors can extend them. For example, the array in `func` may or may not be a variable-length array (VLA):

```
void func() {
    static const int size = 12;
    int array[size]; // might be a VLA
}
```

C23 adds `constexpr` variables (which imply the `const` qualifier) when you really want something to be a constant (Gilding and Gustedt 2022a):

```
void func() {  
    static constexpr int Size = 12;  
    int Array[Size]; // never a VLA  
}
```

C23 doesn't support constexpr functions yet, only objects. Structure members cannot be constexpr.

Enumeration Types

C enumeration types seem normal through C17 but have some strange behaviors. For example, the underlying integer type is implementation defined and could be either a signed or unsigned integer type. C23 now allows the programmer to specify the underlying type for enumerations (Meneide and Pygott 2022):

```
enum E : unsigned short {  
    Valid = 0, // has type unsigned short  
    NotValid = 0x1FFFF // error, too big  
};  
  
// can forward declare with fixed type  
enum F : int;
```

You can also declare enumeration constants larger than int:

```
// has underlying type unsigned long  
enum G {  
    BiggerThanInt = 0xFFFF'FFFF'0000L,  
};
```

Type Inference

C23 enhanced the auto type specifier for single object definitions using type inference (Gilding and Gustedt 2022b). It's basically the same idea as in C++, but auto cannot appear in function signatures:

```
const auto l = 0L; // l is const long  
auto huh = "test"; // huh is char *, not char[5] or const char *  
void func();  
auto f = func; // f is void (*)()  
auto x = (struct S){ // x is struct S
```

```
    1, 2, 3.0
};
#define swap(a, b) \
    do {auto t = (a); (a) = (b); (b) = t;} \
    while (0)
```

typeof Operators

C23 adds support for `typeof` and `typeof_unqual` operators. These are like `decltype` in C++ and are used to specify a type based on another type or the type of an expression. The `typeof` operator retains qualifiers, while the `typeof_unqual` strips qualifiers, including `_Atomic`.

K&R C Functions

K&R C allowed functions to be declared without prototypes:

```
int f();
int f(a, b) int a, b; {return 0;}
```

K&R C functions were deprecated 35 years ago and are finally being removed from the standard. All functions now have prototypes. An empty parameter list used to mean “takes any number of arguments” and now means “takes zero arguments,” the same as C++. It is possible to emulate “accepts zero or more args” via a variadic function signature: `int f(...)`; which is now possible because `va_start` no longer requires passing the parameter before the `...` to the call.

Preprocessor

New features have been added to C23 to improve preprocessing. The `#elifdef` directive complements `#ifdef` and also has an `#elifndef` form. The `#warning` directive complements `#error` but does not halt translation. The `__has_include` operator tests for the existence of a header file, and the `__has_c_attribute` operator tests for the existence of a standard or vendor attribute.

The `#embed` directive embeds external data directly into the source code via the preprocessor:

```
unsigned char buffer[] = {
    #embed "/dev/urandom" limit(32) // embeds 32 chars from /dev/uran
    dom
```

```

};
struct FileObject {
    unsigned int MagicNumber;
    unsigned _BitInt(8) RGBA[4];
    struct Point {
        unsigned int x, y;
    } UpperLeft, LowerRight;
} Obj = {
#ifdef __has_embed(SomeFile.bin) == __STDC__EMBED_FOUND__
// embeds contents of file as struct
// initialization elements
#embed "SomeFile.bin"
#endif
};

```

Integer Types and Representations

Starting in C23, two's complement is the only allowed integer representation (Bastien and Gustedt 2019). Signed integer overflow remains undefined behavior. The `int8_t`, `int16_t`, `int32_t`, and `int64_t` types are now portably available everywhere. The `[u]intmax_t` types are no longer maximal and are only required to represent long long values, not extended or bit-precise integer values.

C23 also introduces bit-precise integer types (Blower et al. 2020). These are signed and unsigned types that allow you to specify the bit-width. These integers do not undergo integer promotions, so they remain the size you requested. Bit-width includes the sign bit, so `_BitInt(2)` is the smallest signed bit-precise integer. `BITINT_MAXWIDTH` specifies the maximum width of a bit-precise integer. It must be at least `ULLONG_WIDTH` but can be much larger (Clang supports $> 2M$ bits).

In C17, adding two nibbles required some bit twiddling:

```

unsigned int add(
    unsigned int L, unsigned int R)
{
    unsigned int LC = L & 0xF;
    unsigned int RC = R & 0xF;
    unsigned int Res = LC + RC;
    return Res & 0xF;
}

```

This is much simpler with `_BitInt`:

```
unsigned _BitInt(4) add(
    unsigned _BitInt(4) L,
    unsigned _BitInt(4) R)
{
    return L + R;
}
```

C23 also added binary literals. The integer literals 0b00101010101, 0x155, 341, and 0525 all express the same value. You can also now use digit separators for improved readability, for example: 0b0000'1111'0000'1100, 0xF'0C, 3'852, and 07'414.

C23 finally has checked integer operations that will detect overflow and wraparound in addition, subtraction, and multiplication operations (Svoboda 2021):

```
#include <stdckdint.h> // new header

bool ckd_add(Type1 *Result, Type2 L, Type3 R);
bool ckd_sub(Type1 *Result, Type2 L, Type3 R);
bool ckd_mul(Type1 *Result, Type2 L, Type3 R);
```

Division is not supported, and it only works with integer types other than plain char, bool, or bit-precise integers. Type1, Type2, and Type3 can be different types. These functions return false if the mathematical result of the operation can be represented by Type1; otherwise, they return true. These functions make it easier to comply with the CERT C Coding Standard and MISRA C guidelines, but it is still awkward to compose operations.

unreachable Function-Like Macro

The unreachable function-like macro is provided in <stddef.h>. It expands to a void expression; reaching the expression during execution is undefined behavior. This allows you to give hints to the optimizer about flow control that is impossible to reach (Gustedt 2021).

As with anything you tell the optimizer to assume, use it with caution, because the optimizer will believe you even if you're wrong. The following is a typical example of how unreachable might be used in practice:

```
#include <stdlib.h>
enum Color {Red, Green, Blue};
```

```
int func(enum Color C) {
    switch (C) {
        case Red: return do_red();
        case Green: return do_green();
        case Blue: return do_blue();
    }
    unreachable(); // unhandled value
}
```

Bit and Byte Utilities

C23 introduces a collection of bit and byte utilities in the `<stdbit.h>` header (Meneide 2023). These include functions to:

- Count the number of 1s or 0s in a bit pattern
- Count the number of leading or trailing 1s or 0s
- Find the first leading or trailing 1 or 0
- Test whether a single bit is set
- Determine the smallest number of bits required to represent a value
- Determine the next smallest or largest power of two based on a value

For example, the following code can be used to count the number of consecutive 0 bits in a value, starting from the most significant bit:

```
#include <stdbit.h>
void func(uint32_t V) {
    int N = stdc_leading_zeros(V);
    // use the leading zero count N
}
```

Prior to C23, this operation is considerably more involved:

```
void func(uint32_t V) {
    int N = 32;
    unsigned R;
    R = V >> 16;
    if (R != 0) {N -= 16; V = R;}
    R = V >> 8;
    if (R != 0) {N -= 8; V = R;}
    R = V >> 4;
    if (R != 0) {N -= 4; V = R;}
    R = V >> 2;
```

```
if (R != 0) {N -= 2; V = R;}
R = V >> 1;
if (R != 0) N -= 2;
else      N -= V;
// use the leading zero count N
}
```

IEEE Floating-Point Support

C23 updates IEEE floating-point support by integrating TS 18661-1, 2, and 3 (ISO/IEC TS 18661-1 2014, ISO/IEC TS 18661-2 2015, ISO/IEC TS 18661-3 2015). Annex F now has parity with the IEEE standard for floating-point arithmetic (IEEE 754-2019). Annex F also applies to decimal floats:

`_Decimal32`, `_Decimal64`, and `_Decimal128`. You cannot mix decimal operations with binary, complex, or imaginary floats, however. Annex H (previously the language-independent arithmetic annex) supports interchange, extended floating types, and nonarithmetic interchange formats. It allows for binary16, graphics processing unit (GPU) data, binary, or decimal representations.

Math library changes support `<math.h>` operations on `_DecimalN`, `_FloatN`, and `_FloatNx` types. Special variants of exponents, logarithms, powers, and π -based trig functions; improved functions for min/max, total ordering, and testing for numerical properties; and functions allowing fine-grained control of conversions between floating-point values and integers or strings have been added.

The `memset_explicit` function has been added for when you really need to clear memory. It's the same as `memset`, except the optimizer cannot remove a call to it. The `strdup` and `strndup` functions have been adopted from POSIX.

REFERENCES

- American National Standards Institute (ANSI). 1986. “Information Systems—Coded Character Sets—7-Bit American National Standard Code for Information Interchange (7-Bit ASCII).” ANSI X3.4-1986.
- Ballman, Aaron. 2019. “Attributes in C.” ISO/IEC, WG14 N2335 (C International Standardization Working Group). <http://www.open-std.org/jtc1/sc22/wg14/www/docs/n2335.pdf>.
- Bastien, J.F., and Jens Gustedt. 2019. “Two’s Complement Sign Representation for C2x.” ISO/IEC JTC1/SC22/WG14 (C International Standardization Working Group). <https://www.open-std.org/jtc1/sc22/wg14/www/docs/n2412.pdf>.
- Blower, Melanie, Tommy Hoffner, and Erich Keane. 2020. “Adding Fundamental Type for N-bit Integers.” Document no. N2534. <http://www.open-std.org/jtc1/sc22/wg14/www/docs/n2534.pdf>.
- Boute, Raymond T. 1992. “The Euclidean Definition of the Functions div and mod.” *ACM Transactions on Programming Language and Systems* 14, no. 2 (April): 127–144. <http://dx.doi.org/10.1145/128861.128862>.
- Dijkstra, Edsger. 1968. “Go To Statement Considered Harmful.” *Communications of the ACM* 11, no. 3. <https://homepages.cwi.nl/~storm/teaching/reader/Dijkstra68.pdf>.
- Gilding, Alex, and Jens Gustedt. 2022a. “The constexpr Specifier for Object Definitions.” ISO/IEC JTC1/SC22/WG14 (C International Standardization Working Group). <https://www.open-std.org/jtc1/sc22/wg14/www/docs/n3018.htm>.
- . 2022b. “Type Inference for Object Definitions.” ISO/IEC JTC1/SC22/WG14 (C International Standardization Working Group). <https://www.open-std.org/jtc1/sc22/wg14/www/docs/n3007.htm>.
- Gustedt, Jens. 2021. “Add Annotations for Unreachable Control Flow.” ISO/IEC

- JTC1/SC22/WG14 (C International Standardization Working Group). <https://www.open-std.org/jtc1/sc22/wg14/www/docs/n2826.pdf>.
- . 2022. “Revise Spelling of Keywords.” ISO/IEC JTC1/SC22/WG14 (C International Standardization Working Group). <https://www.open-std.org/jtc1/sc22/wg14/www/docs/n2934.pdf>.
- Gustedt, Jens, and JeanHeyd Meneide. 2022. “Introduce the nullptr Constant.” ISO/IEC JTC1/SC22/WG14 (C International Standardization Working Group). <https://www.open-std.org/jtc1/sc22/wg14/www/docs/n3042.htm>.
- Hopcroft, John E., and Jeffrey D. Ullman. 1979. *Introduction to Automata Theory, Languages, and Computation*. Reading, MA: Addison-Wesley.
- IEEE. 2019. “IEEE Standard for Floating-Point Arithmetic.” IEEE Std 754-2008. <https://ieeexplore.ieee.org/document/4610935>.
- IEEE and The Open Group. 2018. “Standard for Information Technology—Portable Operating System Interface (POSIX), Base Specifications,” Issue 7. IEEE Std 1003.1.
- International Organization for Standardization and International Electrotechnical Commission. 2014. *Systems and software Quality Requirements and Evaluation (SQuaRE)—System and software quality models*. ISO/IEC 25000:2014. Geneva, Switzerland: International Organization for Standardization.
- ISO/IEC. 1990. “Programming Languages—C,” 1st ed. ISO/IEC 9899:1990.
- . 1999. “Programming Languages—C,” 2nd ed. ISO/IEC 9899:1999.
- . 2007. “Information Technology—Programming Languages, Their Environments and System Software Interfaces—Extensions to the C Library—Part 1: Bounds-Checking Interfaces.” ISO/IEC TR 24731-1:2007.
- . 2011. “Programming Languages—C,” 3rd ed. ISO/IEC 9899:2011.
- . 2013. “Information Technology—Programming Languages, Their Environments and System Software Interfaces—C Secure Coding Rules.” ISO/IEC TS 17961:2013.
- . 2014. “Floating-Point Extensions for C—Part 1: Binary Floating-Point Arithmetic.” ISO/IEC TS 18661-1:2014.
- . 2015. “Floating-Point Extensions for C—Part 2: Decimal Floating-Point Arithmetic.” ISO/IEC TS 18661-2:2015.
- . 2015. “Floating-Point Extensions for C—Part 3: Interchange and

- Extended Types.” ISO/IEC TS 18661-3:2015.
- . 2018. “Programming Languages—C,” 4th ed. ISO/IEC 9899:2018.
- . 2024. “Programming Languages—C,” 5th ed. ISO/IEC 9899:2024.
- . 2024. “Information Technology—Programming Languages—C—A Provenance-Aware Memory Object Model for C.” Draft Technical Specification. ISO/IEC CD TS 6010:2024.
- ISO/IEC/IEEE. 2011. “Information Technology—Microprocessor Systems—Floating-Point Arithmetic.” ISO/IEC/IEEE 60559:2011. <https://www.iso.org/obp/ui/#iso:std:57469:en>.
- Kernighan, Brian W., and Dennis M. Ritchie. 1988. *The C Programming Language*, 2nd ed. Upper Saddle River, NJ: Prentice Hall.
- Knuth, Donald. 1997. “Information Structures.” In *Fundamental Algorithms*, 3rd ed., 438–442. Vol. 1 of *The Art of Computer Programming*. Boston: Addison-Wesley.
- Kuhn, Markus. 1999. “UTF-8 and Unicode FAQ for Unix/Linux.” June 4, 1999. <https://www.cl.cam.ac.uk/~mgk25/unicode.html>.
- Lamport, Leslie. 1979. “How to Make a Multiprocessor Computer That Correctly Executes Multiprocess Programs.” *IEEE Transactions on Computers* C-28 9 (September): 690–691.
- Lewin, Michael. 2012. “All About XOR.” *Overload Journal* 109 (June). <https://accu.org/index.php/journals/1915>.
- Meneide, JeanHeyd. 2023. “More Modern Bit Utilities.” ISO/IEC JTC1/SC22/WG14 (C International Standardization Working Group). Last updated February 7, 2023. <https://www.open-std.org/jtc1/sc22/wg14/www/docs/n3104.htm>.
- Meneide, JeanHeyd, and Clive Pygott. 2022. “Enhancements to Enumerations.” ISO/IEC JTC1/SC22/WG14 (C International Standardization Working Group). Last updated July 19, 2022. <https://www.open-std.org/jtc1/sc22/wg14/www/docs/n3030.htm>.
- Ritchie, Dennis. 1993. “The Development of the C Language.” In *The Second ACM SIGPLAN History of Programming Languages Conference (HOPL-II)*, April 20–23, 1993, Cambridge, Massachusetts, 201–208. New York: Association for Computing Machinery.
- Saks, Dan. 2002. “Tag vs. Type Names.” October 1, 2002.

<https://www.embedded.com/tag-vs-type-names/>.

Schoof, René. 2008. “Four Primality Testing Algorithms.” Submitted January 24, 2008. arXiv preprint arXiv:0801.3840. <https://doi.org/10.48550/arXiv.0801.3840>.

Seacord, Robert C. 2013. *Secure Coding in C and C++*, 2nd ed. Boston: Addison-Wesley Professional.

———. 2014. *The CERT® C Coding Standard: 98 Rules for Developing Safe, Reliable, and Secure Systems*, 2nd ed. Boston: Addison-Wesley Professional.

———. 2017. “Uninitialized Reads.” *Communications of the ACM* 60, no. 4 (March): 40–44. <https://doi.org/10.1145/3024920>.

———. 2019. “Bounds-Checking Interfaces: Field Experience and Future Directions.” February 3, 2019. <https://www.open-std.org/jtc1/sc22/wg14/www/docs/n2336.pdf>.

Seacord, Robert C., et al. 2024. *The C Standard Charter*. Last updated June 12, 2024. <https://www.open-std.org/jtc1/sc22/wg14/www/docs/n3280.htm>.

Seacord, Robert C., Martin Uecker, Jens Gustedt, and Philipp Klaus Krause. 2024. “Accessing Byte Arrays, v4.” <https://www.open-std.org/jtc1/sc22/wg14/www/docs/n3254.pdf>.

Svoboda, David. 2021. “Checked N-Bit Integers?” WG N2867. Last updated November 28, 2021. <https://open-std.org/JTC1/SC22/WG14/www/docs/n2867.pdf>.

TIOBE. 2024. TIOBE Index. <https://www.tiobe.com/tiobe-index/>.

The Unicode Consortium. 2023. *The Unicode Standard*, version 15.1.0. Mountain View, CA: The Unicode Consortium. <https://www.unicode.org/versions/Unicode15.1.0/>.

Weimer, Florian. 2018. “Recommended Compiler and Linker Flags for GCC.” *Red Hat Developer*, March 21, 2018. <https://developers.redhat.com/blog/2018/03/21/compiler-and-linker-flags-gcc>.

INDEX

A

- abort function, [162](#), [177](#), [232](#)
- abort_handler_s function, [164–165](#)
- ABS macro, [54–56](#)
- addition (+) operator, [25](#)
- additive operators, [83–84](#)
- address, [14](#)
- AddressSanitizer, [132](#), [253–257](#)
 - compiler flags used with, [255](#)
- address space layout randomization (ASLR), [239](#)
- Advanced Encryption Standard (AES), [56](#)
- aligned_alloc function, [120](#)
- alignment requirements, [41–42](#)
- alignof operator, [92](#)
- alloca function, [128–129](#)
- allocated storage duration, [35](#), [116](#)
- /analyze flag, [241](#)
- Annex K bounds-checking interfaces
 - gets_s function, [161–162](#)
 - runtime constraints, [163–164](#)
 - strcpy_s function, [162–163](#)
- application programming interface (API), [168](#), [216](#)
- argument checking, [156–157](#)
- arithmetic types, [19–22](#), [47](#)
 - conversion, [64–72](#)
 - enumeration, [21](#)
 - floating-point, [21–22](#), [59–64](#)
 - integer, [48–59](#)
 - operators, [83–85](#)
 - pointer, [94–96](#)
- arm_missile function, [19](#)
- ARRAY_SIZE identifier, [203](#)
- array types, [25–27](#)
- ASCII, [138](#)
- assertions
 - runtime, [232–234](#)
 - static, [230–232](#)
- assert macro, [232](#)

`assignInterestRate` function, [103–104](#)
assignment (=) operator, [74](#)
associativity, [78–80](#)
`ATOMIC_VAR_INIT` macro, [207](#)
attributes, [44–45](#), [259–260](#)
automatic storage duration, [15](#)
auto storage-class specifier, [38–39](#)

B

backslash (\), [143](#)
basic execution character set, [19](#)
basic multilingual plane (BMP), [139](#), [146](#)
big-endian ordering, [191](#)
binary constants, [58](#)
binary resources, embedding, [209–210](#)
binary streams, [172](#)
 reading from and writing to, [188–191](#)
bin directory, [225–227](#)
bit and byte utilities, [264–265](#)
`_BitInt` type, [56–57](#)
bit-precise constants, [58](#)
bit-precise integers, [20](#), [56–59](#)
bitwise operators
 bitwise AND (&), [87](#)
 bitwise exclusive OR (^), [88](#)
 bitwise inclusive OR (|), [88–89](#)
 complement, [85–86](#)
 shift, [86–87](#)
block, [14](#). *See also* [compound statements](#)
block scope, [34](#)
Boeing 787, [51](#)
Boolean type, [18–19](#)
break statement, [111–112](#)
buffering, [169–170](#)
buffer overflows, [106](#), [123](#), [152](#), [159](#), [160](#), [164–165](#), [188](#), [239](#), [240](#)
byte-oriented stream, [171](#)
bytes, [154](#)

C

C, [xxv–xxvi](#)
 history and development of, [xxiii–xxv](#)
 getting started with, [1–11](#)
call-by-value language, [16](#)
`calloc` function, [120–121](#)
`cardinal_points` enumeration, [21](#)
cast operators, [90–91](#)
`cc` command, [4](#)
central processing units (CPUs), [41](#), [85](#)
`char16_t` type, [142](#)

`char32_t` type, [142](#)

characters

- ASCII, [138](#)

- code points, [138](#)

- constants, [142–143](#)

- conversion of, [146–149](#)

- data types, [140–142](#)

- escape sequences, [143–144](#)

- execution character set, [140](#)

- Linux, [144–145](#)

- literals, [142–143](#)

- narrow, [140–142](#), [146](#)

- reading and writing, [177–180](#)

- source character set, [140](#)

- Unicode, [138–140](#)

- wide, [19](#), [140–142](#), [146](#)

character string literal, [150–152](#)

`char` type, [19](#), [140–141](#)

Clang, [44–45](#), [93](#), [102](#), [104](#), [129](#), [131](#), [134](#), [145](#), [196](#), [198](#), [203](#), [206](#), [226](#), [252](#), [255](#)

- flags, [235–240](#)

- installation, [8](#)

- predefined macros list, [210](#)

`clearerr` function, [169](#)

`clear_stdin` function, [231](#)

`close` function, [177](#)

code

- building, [225–227](#)

- page, [138](#)

- points, [138](#), [154](#)

- reuse, [215](#)

- units, [139](#), [154](#)

coercion, [65](#)

cohesion, [214–215](#)

collection, [216](#)

`collection_type` identifier, [217](#)

comma (,) operator, [94](#)

`CommandLineToArgvW` function, [145](#)

common extensions, [11](#)

compilation process, preprocessor, [196](#)

compilers, [7–8](#)

compiler settings and flags, [234–235](#)

- GCC and Clang, [235–240](#)

- Visual C++, [240–241](#)

compiling and running a program, [3–5](#)

complement operators, [85–86](#)

complex types, [22](#)

componentization, [213–218](#)

- code reuse, [215](#)

- cohesion, [214–215](#)

- coupling, [214–215](#)

- data abstraction, [215–216](#)
- opaque types, [217–218](#)
- compound assignment operators, [93–94](#)
- compound statements, [14](#), [98–99](#)
- conditional inclusion, [198–202](#)
 - generating diagnostics, [200–201](#)
 - header guards, [201–202](#)
- conditional (?:) operator, [91–92](#)
- considered behavior, [10](#)
- constants
 - character, [142–143](#)
 - floating, [64](#)
 - integer, [57–59](#)
 - binary, [58](#)
 - decimal, [57](#)
 - hexadecimal, [57–58](#)
 - octal, [57](#)
- constexpr storage-class specifier, [37–38](#)
- const qualifier, [32](#)
- continue statement, [111](#)
- control flow, [97–113](#). *See also* [jump statements](#); [selection statements](#)
 - compound statement, [14](#), [98–99](#)
 - expression statements, [97–98](#)
 - iteration statements, [105–108](#)
- control flow guard (CFG), [240–241](#)
- controlling expression, [99–105](#)
- conversion
 - arithmetic, [64–72](#)
 - characters, [146–149](#)
 - floating-type demotions, [71](#)
 - floating-type to integer-type and vice versa, [71](#)
 - implicit, [69–70](#)
 - integer, [65–67](#), [70–71](#)
 - safe conversion, [70–71](#)
 - specifiers, [187](#)
 - usual arithmetic conversions, [67–69](#)
- convert_arg function, [222](#)
- convert_cmd_line_args function, [222](#)
- copy_process function, [110](#)
- coupling, [214–215](#)
- C standard library, [147–149](#)

D

- dangling pointers, dealing with, [125–126](#)
- data abstraction, [215–216](#)
- debugging, [241–245](#)
- decimal constants, [57](#)
- decimal floating types, [22](#)
- declarations, [74](#)

- declarations, integer, [49](#)
- decrement (--) operator, [77](#)
- defined operator, [199](#)
- #define preprocessing directive, [202–203](#)
- defragmentation, [117](#)
- derived types
 - array, [25–27](#)
 - function, [22–23](#)
 - pointer, [23–25](#)
 - structure, [27–28](#)
 - union, [28–29](#)
- diagnostics, generating, [200–201](#)
- dmalloc (debug memory allocation), [132–134](#)
 - library, [252](#)
- double type, [59–62](#)
- do...while loop, [231](#)
- do...while statement, [106–107](#)
- dup_string function, [233](#)
- dynamically allocated memory, [115](#)
- dynamic analysis, [252–253](#)
- dynamic library, [218–219](#)
 - debugging, [132–135](#)
 - flexible array members, [127–128](#)
 - memory management, [117–126](#)
 - safety-critical systems, [134–135](#)
 - storage duration, [116–117](#)
 - using, [117](#)

E

- editors, [6–7](#)
- 8-bit representation, [51](#)
- else statement, [199](#)
- #embed preprocessor directive, [209](#)
- endianness, [28](#), [191–193](#)
- #endif preprocessing directive, [199](#)
- enumerations, [21](#), [261](#)
- environment variables, [164](#)
- equal (==) operator, [70](#)
- #error preprocessing directive, [201](#)
- escape sequences, [143–144](#)
- Euclidean division, [84](#)
- evaluation, [75–76](#)
 - indeterminately sequenced, [81](#)
 - order of, [80–82](#)
 - unsequenced, [81](#)
- executables, [218–219](#)
- execution character set, [140](#)
- EXIT_SUCCESS macro, [2–3](#)
- EXPECT_STREQ assertion, [247](#)

- explicit undefined behaviors, [10](#)
- expressions, [73](#)
 - evaluations, [75–76](#)
 - function invocation, [76–77](#)
 - simple assignment, [74–75](#)
 - statement, [97–98](#)
- Extended ASCII, [138](#)
- extended grapheme cluster, [154–155](#)
- extended integer types, [20](#)
- extern "C" declaration, [247](#)
- extern specifier, [37](#)
- extern storage class specifier, [220](#)

F

- fclose function, [176–177](#), [181](#), [188](#)
- feof function, [169](#)
- ferror function, [169](#)
- fgetc function, [178](#)
- fgetpos function, [181](#)
- field-programmable gate arrays (FPGAs), [56](#)
- file access modes, [174](#)
- file descriptor, [174](#)
- file inclusion, preprocessor, [197–198](#)
- FILE object, [168](#), [174](#), [177](#)
- file pointers, [168](#)
- file scope, [34](#)
- files, I/O
 - closing, [176–177](#)
 - creating, [172–176](#)
 - opening, [172–176](#)
- file status flags, [175](#)
- flag argument, [55](#)
- flags, [235–240](#)
- flexible array members, [127–128](#)
- floating-point
 - arithmetic, [21–22](#), [62](#)
 - C model, [60–62](#)
 - constants, [64](#)
 - encoding, [59–60](#)
 - types, [59–60](#)
 - values, [62–64](#)
- float type, [59–60](#)
- flooring division, [84](#)
- flushing, [170](#), [180](#)
- fopen function, [172–174](#), [181](#)
- format string, [185](#)
- formatted output, [5](#)
- formatted text streams, reading, [184–188](#)
- for statement, [107–108](#)

- [_FORTIFY_SOURCE macro, 239](#)
- [fpclassify macro, 63–64](#)
- [FPGAs, 56](#)
- [-fpic flag, 239](#)
- [-fpie flag, 239](#)
- [fread function, 188–191](#)
- [free_aligned_sized function, 124–126](#)
- [free function, 123–224](#)
- [free_sized function, 124](#)
- [freestanding environment, 2](#)
- [fseek function, 173, 180–181](#)
- [fsetpos function, 173, 181](#)
- [-fstack-protector-strong option, 240](#)
- [functions, 2, 14](#)
 - [code reuse, 215](#)
 - [declarator, 22](#)
 - [definition, 23](#)
 - [designator, 76–77](#)
 - [invocation, 76–77](#)
 - [objects and, 14](#)
 - [prototype, 23](#)
 - [return values, 4–5](#)
 - [scope, 34](#)
 - [type, 22–23](#)
 - [variadic, 76](#)
- [fwrite function, 188–191](#)

G

- [GCC. See \[GNU Compiler Collection\]\(#\)](#)
- [generic selection expression, 207](#)
- [getc function, 178](#)
- [getchar function, 178](#)
- [get_error function, 246](#)
- [get_file_size function, 181](#)
- [get_password function, 159](#)
- [gets function, 160–161](#)
- [gets_s function, 161–162](#)
- [-glevel flag, 236–237](#)
- [GNU Compiler Collection \(GCC\), \[xxv–xxvi\]\(#\), \[8\]\(#\), \[21\]\(#\), \[93\]\(#\), \[102–104\]\(#\), \[129\]\(#\), \[131\]\(#\), \[133\]\(#\), \[134\]\(#\), \[144\]\(#\), \[196\]\(#\), \[198\]\(#\), \[203\]\(#\), \[206\]\(#\), \[210\]\(#\), \[226\]\(#\), \[252\]\(#\), \[253\]\(#\)](#)
 - [compiler and linker flags, 235–240](#)
- [Google Test, 245–246](#)
- [goto statement, 109–110](#)
- [/guard:cf flag, 240–241](#)

H

- [__has_c_attribute operator, 262](#)
- [__has_include operator, 262](#)
- [__has_include preprocessor operator, 198](#)

- header files, [2](#), [216](#)
- header guards, [201–202](#)
- heap manager, [116–117](#)
- hexadecimal constants, [57–58](#)
- hidden scope, [34](#)

I

- identifier, [14](#)
- IDEs, [6–7](#)
- IEEE floating-point support, [265](#)
- if...else ladder, [102](#)
- if statement, [99–102](#), [199](#)
- ignore_handler_s function, [164–165](#)
- implementation-defined behaviors, [9](#)
- implicit conversion, [65](#), [69–70](#)
- implicit undefined behavior, [10](#)
- #include preprocessor directive, [2](#)
- incomplete array type, [127](#)
- increment (++) operator, [77](#)
- indeterminately sequenced evaluation, [81](#)
- indirection (*) operator, [16–17](#), [24–25](#)
- infinite loop, [106](#)
- initializer, [74](#)
- inner scope, [34](#)
- input/output (I/O), [2](#), [33](#), [146](#), [167–168](#)
- integers
 - bit-precise, [20](#), [56–57](#)
 - constant expressions, [260–261](#)
 - constants, [57–59](#)
 - conversion rank, [65–66](#)
 - conversions, [70–71](#)
 - declarations, [49](#)
 - overflow, [54–56](#)
 - padding, [48](#)
 - precision, [48](#)
 - promotions, [66–67](#)
 - ranges, [48](#)
 - signed, [52–56](#)
 - unsigned, [49–52](#)
 - width, [48](#)
- integrated development environments (IDEs), [6–7](#)
- intentional behavior, [10](#)
- Internet Protocol (IP), [191](#)
- int type, [141](#)
- I/O. See [input/output](#)
- is_prime function, [224](#)

J

- jump statements

break, [111–112](#)
continue, [111](#)
goto, [109–110](#)
return, [112–113](#)

K

K&R C functions, [262](#)
keywords, [260](#)
KnownError test case, [249](#)
Knuth, Donald, [116](#), [214](#)

L

labels, [34](#)
libiconv function, [149](#)
libraries, [218](#)
lifetime, determining, [15](#)
lines, reading and writing, [177–180](#)
linkage, [219–221](#)
link phase, [218](#)
Linux, [144–145](#)
literals
 character, [142–143](#)
 string, [150–152](#)
little-endian ordering, [191](#)
locale, [140](#)
locale-specific behavior, [11](#)
locator value, [74](#)
logical AND (&&) operator, [89–90](#)
logical negation (!) operator, [83](#)
logical OR (|) operator, [89](#)
long double type, [59–62](#)

M

macro
 definitions, [202–211](#)
 embedded binary resources, [209–210](#)
 generic selection expression as, [208](#)
 predefined, [210–211](#)
 replacement, [205–207](#)
 type-generic, [207–209](#)
 with automatic type inference, [209](#)
 undefining, [204–205](#)
 unsafe expansion, [206](#)
main entry point, [145–146](#)
main function, [2](#), [17](#), [40](#), [225](#)
malloc function, [118–120](#)
matrix, [26](#)
max function, [76](#)

- mbrtoc16 function, [148](#)
- mbsrtowcs function, [147](#)
- mbstowcs function, [147](#)
- mbtowc function, [147](#)
- memcpy function, [157–159](#), [243–244](#)
- memcpy function, [157](#)
- __memmove_avx_unaligned_erms function, [243](#)
- memory
 - allocating without declaring type, [118–119](#)
 - leaks, [117](#)
 - avoiding, [121–122](#)
 - management of, [117–126](#)
 - manager, [116–117](#)
 - reading uninitialized memory, [119–120](#)
 - states of, [126–127](#)
- memset_explicit function, [159–160](#)
- memset function, [159–160](#)
- memset_s function, [159–160](#)
- Miller-Rabin primality test, [224](#)
- MultiByteToWideChar function, [149](#)
- multiplicative operators, [84](#)

N

- NAME macro, [205](#)
- namespace, [30](#)
- narrow characters, [140–142](#), [146](#)
- narrow string, [149](#)
- NDEBUG macro, [233](#)
- nesting, [34](#)
- not-a-number (NaN), [63](#)
- NULL macro, [24](#)
- null pointer, calling realloc with, [122](#)
- nullptr pointer, [126](#)
- num_args parameter, [223](#)

O

- objects, [13–14](#). *See also* [types](#)
 - storage
 - class, [36–39](#)
 - duration, [35–36](#)
- octal constants, [57](#)
- o flag, [235–236](#)
- opaque types, [217–218](#)
- open file description, [174](#)
- open function, [174–176](#)
- operating system (OS), [116](#)
- operators, [73](#)
 - alignof, [92](#)
 - arithmetic, [83–85](#)

- associativity, [78–80](#)
- bitwise, [85–89](#)
- cast, [90–91](#)
- comma (,), [94](#)
- compound assignment, [93–94](#)
- conditional (? :), [91–92](#)
- decrement (--), [77](#)
- increment (++), [77](#)
- logical, [89–90](#)
- order of evaluation, [80–82](#)
- postfix, [77](#)
- precedence, [78–80](#)
- prefix, [77](#)
- relational, [93](#)
- sizeof, [82–83](#)
- order of operations, [78](#)
- original equipment manufacturer (OEM), [145–146](#)
- outer scope, [34](#)
- overflow, integer, [54–56](#)

P

- padding, [48](#)
- parameters, [16](#)
- pass-by-value language, [16](#)
- pedantic flag, [238](#)
- pedantic mode, [11](#)
- /permissive- flag, [241](#)
- pie flag, [239](#)
- planes, [138–139](#)
- ++1 operation, [77](#)
- pointer arithmetic, [94–96](#)
- pointers, [14](#), [23–25](#)
- portability, [9–11](#)
 - common extensions, [11](#)
 - implementation-defined behaviors, [9](#)
 - locale-specific behavior, [11](#)
 - undefined behavior, [10–11](#)
 - unspecified behaviors, [10](#)
- Portable Operating System Interface (POSIX), [24](#), [164–165](#), [183–184](#), [189](#)
- postfix operators, [77](#)
- precedence, operator, [78–80](#)
- precision, [48](#)
- predefined macros, [210–211](#)
- predefined streams, [170–171](#)
- predicate. *See* [assertions](#)
- prefix operators, [77](#)
- preprocessing directives, [196](#)
- preprocessor, [195](#)
 - compilation process, [196](#)

- conditional inclusion, [198–202](#)
- file inclusion, [197–198](#)
- macro definitions, [202–211](#)
- translation phases, [196](#)
- prime number, [221](#)
- prnterr function, [129](#), [130](#)
- print_error function, [130](#), [253–254](#)
- printf function, [5](#)
- print_help function, [221](#)
- program structure
 - building code, [225–227](#)
 - componentization, [213–218](#)
 - executables, [218–219](#)
 - linkage, [219–221](#)
 - simple program, [221–225](#)
- promotions, integer, [66–67](#)
- pthread.h header, [200](#)
- public interface, [215–216](#)
- putc function, [178](#)
- puts function, [2](#), [5](#), [178](#)

Q

- qualified types, [31–34](#)

R

- random-access memory (RAM), [184](#)
- real floating types, [22](#)
- reallocarray function, [117–118](#), [123](#)
- realloc function, [121–122](#)
- rec.signame, [186](#), [188](#), [190](#)
- referenced type, [23](#)
- register storage-class specifier, [38](#)
- relational operators, [93](#)
- representable value, [48](#)
- restrict-qualified pointer, [33–34](#)
- return statement, [112–113](#)
- return values, function, [4–5](#)
- rewind function, [173](#), [182–183](#)
- rmdir function, [184](#)
- RUN_ALL_TESTS macro, [247](#)
- runtime analysis, [252–253](#)
- runtime assertions, [232–234](#)
- runtime constraints, [163–164](#)
- rvalue (right operand), [74](#)

S

- safe conversion, [70–71](#)
- safety-critical systems, [134–135](#)

- scope, [34–35](#)
 - block, [34](#)
 - file, [34](#)
 - function, [34](#)
 - function prototype, [34](#)
- /sd1 flag, [241](#)
- Secure Hash Algorithm (SHA), [56](#)
- selection statements
 - if, [99–102](#)
 - switch, [102–104](#)
- sequence points, [81–82](#)
- set_constraint_handler_s function, [163](#)
- setlocale function, [148](#)
- 7-bit ASCII, [138](#)
- shadowed scope, [34](#)
- shared flag, [239](#)
- shift operations, [86–87](#)
- show_classification function, [63](#)
- side effects, [55](#), [76](#)
- signed char type, [19](#)
- signed integers, [20](#)
- sign extension, [70](#)
 - integer overflow, [54–56](#)
 - representation, [52–54](#)
- simple assignment, [74–75](#)
- simple program, structuring, [221–225](#)
- sin function, [207](#)
- single quote ('), [143](#)
- sizeof operator, [82–83](#), [131](#), [154](#)
- sizeof(size++) operand, [131](#)
- small types, [66–67](#)
- software development kit (SDK), [146](#)
- source character set, [140](#)
- source files, [216](#)
- spaghetti code, [109](#)
- standard error stream (stderr), [171](#)
- standard input stream (stdin), [170](#)
- standard output stream (stdout), [170](#)
- states, memory, [126–127](#)
- static
 - analysis, [251–252](#)
 - assertion, [230–232](#)
 - keyword, [220](#)
 - library, [218](#)
- static storage-class specifier, [36–37](#)
- __STDC_ENDIAN_BIG__ macro, [192](#)
- __STDC_ENDIAN_LITTLE__ macro, [191](#)
- __STDC_ENDIAN_NATIVE__ macro, [192](#)
- /std:clatest flag, [241](#)
- std= flag, [238](#)

- storage
 - class, [36–39](#)
 - duration, [35–36](#)
 - heap manager, [116–117](#)
 - memory manager, [116–117](#)
 - using dynamically allocated memory, [117](#)
 - other forms of, [128–132](#)
- strcpy function, [155–156](#)
- strcpy_s function, [162–163](#)
- streams, I/O
 - binary, [172](#)
 - buffering, [169–170](#)
 - error and end-of-file indicators, [168–169](#)
 - orientation, [171](#)
 - predefined, [170–171](#)
 - text, [172](#)
- strerrorlen_s function, [246](#)
- strerror_s function, [246](#), [249](#)
- strictly conforming programs, [9](#)
- string-handling functions, [152–165](#)
- stringizing, [205](#)
- strings, [137–138](#), [149–152](#)
- strlen function, [154–155](#)
- strndup function, [164](#)
- structure member (.) operator, [27](#)
- structure pointer (->) operator, [27](#)
- structure type (struct), [27–28](#)
- subnormal numbers, [62](#)
- subscript ([]) operator, [25](#)
- substatement, [99–102](#)
- supplementary characters, [139](#)
- surrogates, [139](#)
- swap function, [16–17](#)
- switch statement, [102–104](#)

T

- tags, [29–31](#)
- tempfile, [170–171](#)
- temporary files, [184](#)
- test suite, [246](#)
- text stream, [172](#)
- thread_local storage-class specifier, [37](#)
- threads.h header, [200](#)
- thread storage duration, [35](#), [37](#)
- time-of-check to time-of-use (ToCToU), [33](#)
- token pasting, [206](#)
- translation phases, [196](#)
- translation unit, [196](#)
- Transmission Control Protocol (TCP), [191](#)

- truncating division, [84](#)
- type, [14](#)
- typedef storage-class specifier, [38](#)
- type-generic macros, [207–209](#)
- type inference, [261](#)
- typeof operators, [39–40](#), [262–264](#)
 - integer types and representation, [263–264](#)
 - K&R C functions, [262](#)
 - preprocessor, [262](#)
- typeof_unqual operator, [39–40](#)
- types
 - object
 - arithmetic, [19–22](#)
 - Boolean, [18–19](#)
 - character, [19](#)
 - void, [22](#)
 - definitions, [26–27](#)
 - derived, [22–29](#)
 - array, [25–27](#)
 - function, [22–23](#)
 - pointer, [23–25](#)
 - structure, [27–28](#)
 - union, [28–29](#)
 - qualifiers, [31–34](#)
 - const, [32](#)
 - restrict, [33–34](#)
 - volatile, [32–33](#)
 - variably modified, [42–44](#)

U

- Ubuntu Linux, [246](#)
- UINT_MAX expression, [50–52](#)
- unary & (address-of) operator, [17](#)
- unary + and – operators, [83](#)
- undefined behavior, portability and, [10–11](#)
- Unicode scalar value, [139](#)
- Unicode Standard (Unicode), [138–140](#)
- Unicode transformation formats (UTFs), [139](#)
- uninitialized memory, [119–120](#)
- union types, [28–29](#)
- unit testing, [245–251](#)
- Universal Serial Bus (USB) ports, [168](#)
- unreachable function-like macro, [264](#)
- unsequenced evaluations, [81](#)
- unsigned char type, [19](#)
- unsigned integers, [20](#)
 - representation, [49–50](#)
 - wraparound, [50–52](#)
- unsigned-preserving approach, [67](#)

unspecified behaviors, [10](#)
Urban, Reini, [246](#)
User Datagram Protocol (UDP), [191](#)
usual arithmetic conversions, [67–69](#)

V

value computation, [75–76](#)
value-preserving approach, [67](#)
valueReturnedIfTrue operand, [91–92](#)
values, swapping, [15–18](#)
variable-length arrays (VLAs), [129–132](#), [260](#)
variables, [14](#)
 declaring, [14–18](#)
variably modified types (VMTs), [42–44](#)
Visual C++, [21](#), [91](#), [139](#), [142](#), [145](#), [165](#), [196](#), [203](#), [240–241](#), [252](#), [255](#)
Visual Studio Code (VS Code), [6–7](#), [242](#)
Visual Studio IDE, Microsoft, [6](#), [8](#)
void type, [22](#)
volatile-qualified type, [32–33](#)
vstrcat function, [241](#)

W

-Wall flag, [237–238](#)
wchar_t type, [141–142](#)
-Wconversion flag, [238](#)
wcslen function, [154](#)
wcsrtombs function, [147](#)
wcstombs function, [147](#)
-Werror flag, [238](#)
-Wextra flag, [237–238](#)
while statement, [105–106](#)
wide characters, [19](#), [140–142](#), [146](#)
wide-oriented stream, [171](#)
wide string, [150](#)
width, integer, [48](#)
-WI flag, [239](#)
Win32 conversion APIs, [149](#)
Windows, [145–146](#)
-wl, -z, noexecstack linker option, [239–240](#)
wmain entry point, [145–146](#)
word, [41](#)
wraparound, [50–52](#)